

6.036/6.862: Introduction to Machine Learning

Lecture: starts Tuesdays 9:35am (Boston time zone)

Course website: introml.odl.mit.edu

Who's talking? Prof. Tamara Broderick

Questions? discourse.odl.mit.edu (“Lecture 2” category)

Materials: Will all be available at course website

Last Time

- I. Machine learning setup
- II. Linear classifiers
- III. Learning algorithms

Today's Plan

- I. Perceptron algorithm
- II. Harder and easier linear classification
- III. Perceptron theorem

Recall: Classifiers

Recall: Classifiers

- A linear classifier:

$$h(x; \theta, \theta_0)$$

Recall: Classifiers

- A linear classifier:

$$h(x; \theta, \theta_0) = \text{sign}(\theta^\top x + \theta_0)$$

Recall: Classifiers

- A linear classifier:

$$h(x; \theta, \theta_0) = \text{sign}(\theta^\top x + \theta_0)$$

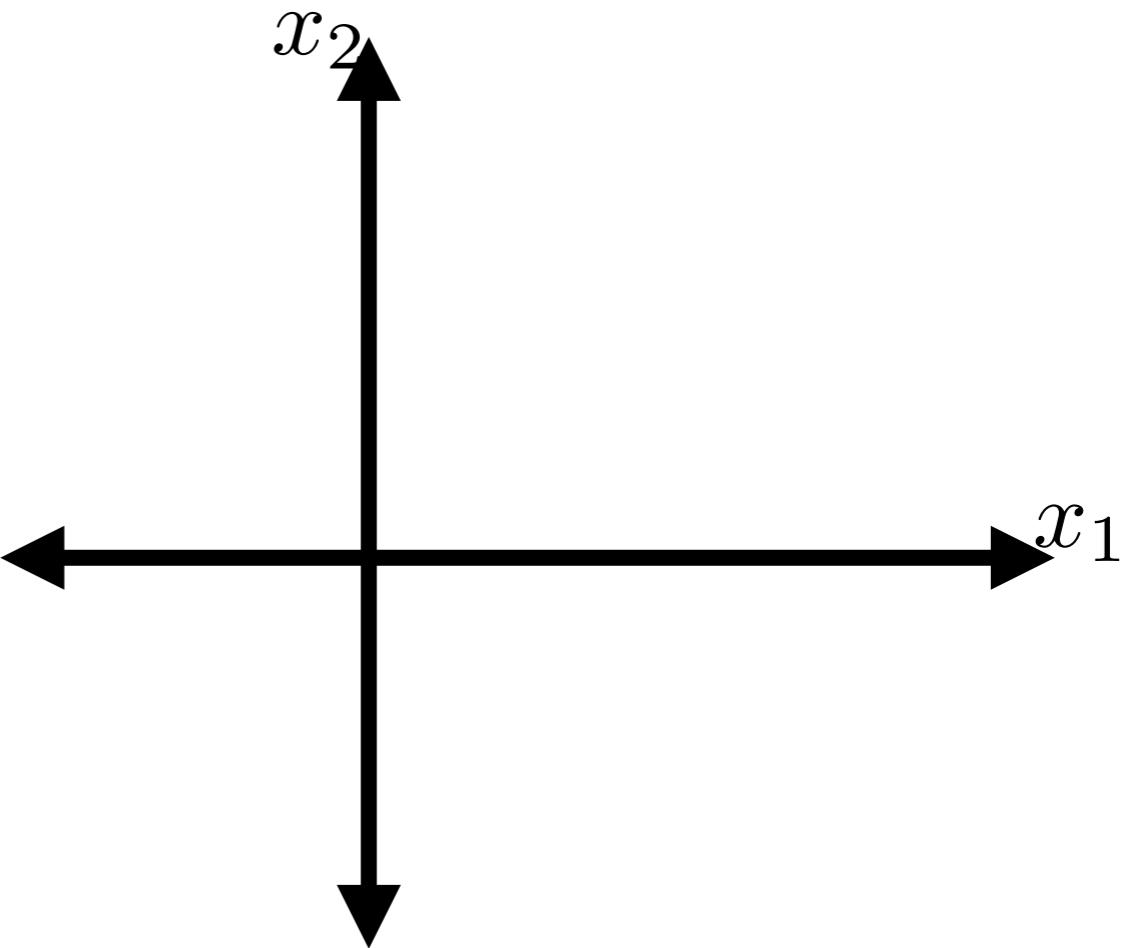
$$= \begin{cases} +1 & \text{if } \theta^\top x + \theta_0 > 0 \\ -1 & \text{if } \theta^\top x + \theta_0 \leq 0 \end{cases}$$

Recall: Classifiers

- A linear classifier:

$$h(x; \theta, \theta_0) = \text{sign}(\theta^\top x + \theta_0)$$

$$= \begin{cases} +1 & \text{if } \theta^\top x + \theta_0 > 0 \\ -1 & \text{if } \theta^\top x + \theta_0 \leq 0 \end{cases}$$

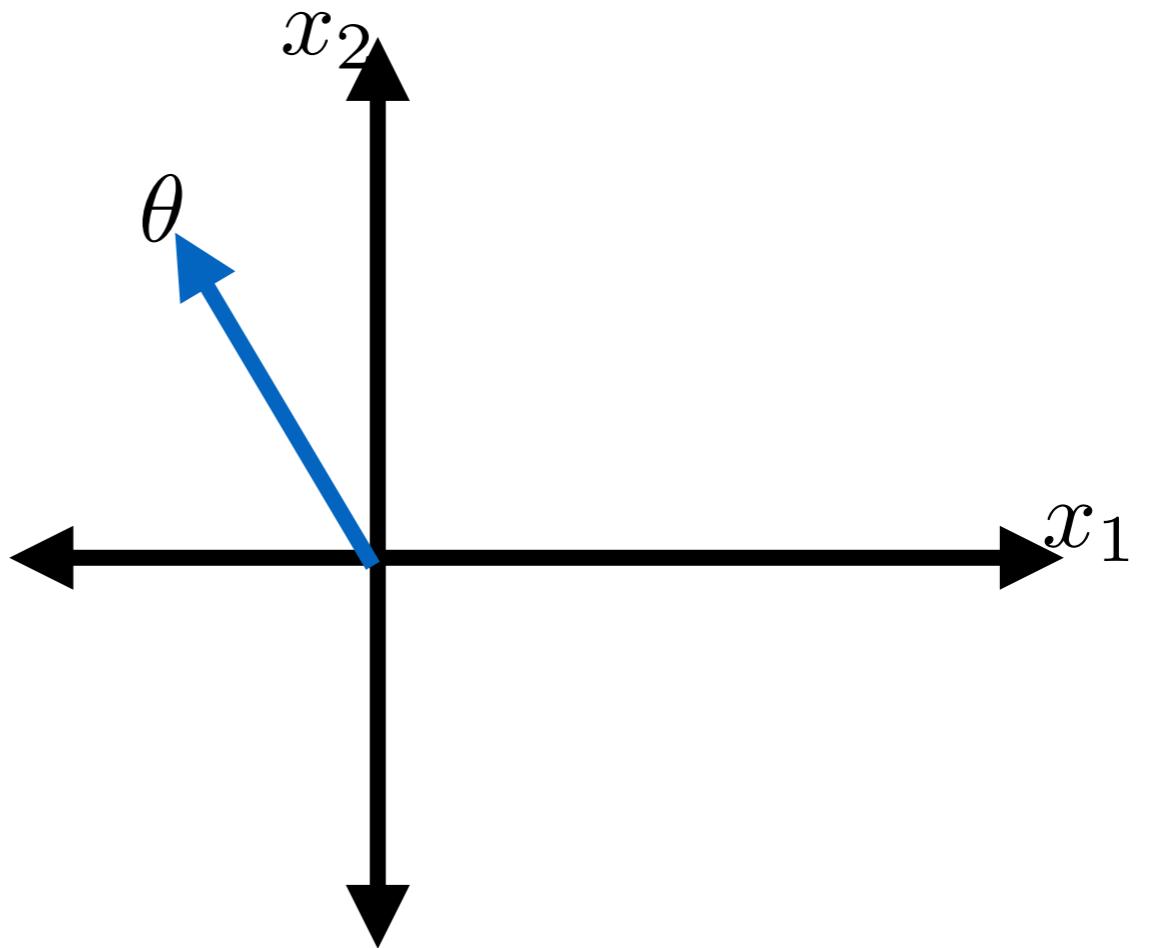


Recall: Classifiers

- A linear classifier:

$$h(x; \theta, \theta_0) = \text{sign}(\theta^\top x + \theta_0)$$

$$= \begin{cases} +1 & \text{if } \theta^\top x + \theta_0 > 0 \\ -1 & \text{if } \theta^\top x + \theta_0 \leq 0 \end{cases}$$

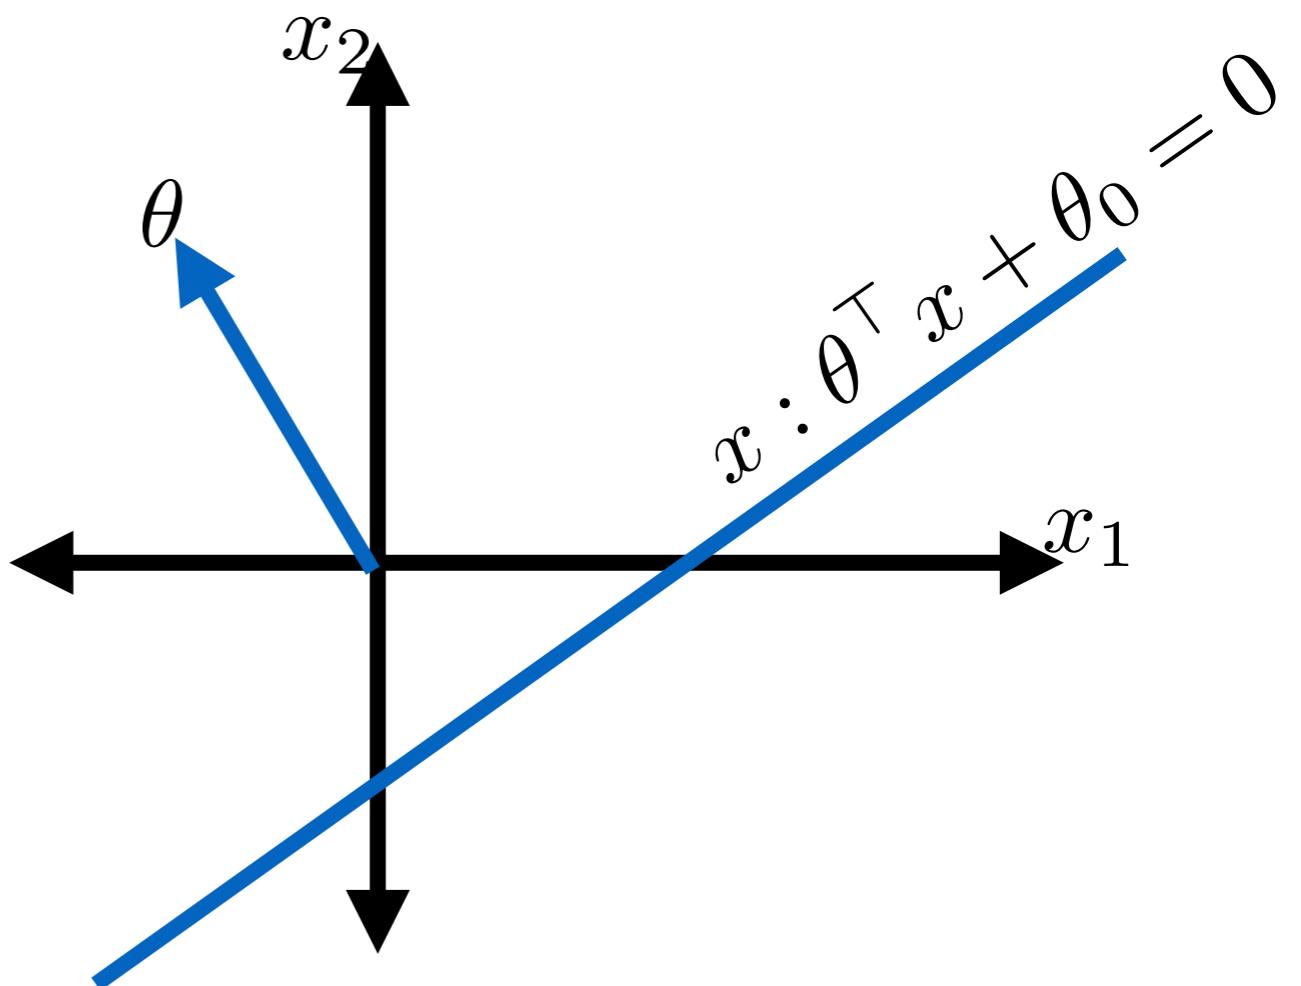


Recall: Classifiers

- A linear classifier:

$$h(x; \theta, \theta_0) = \text{sign}(\theta^\top x + \theta_0)$$

$$= \begin{cases} +1 & \text{if } \theta^\top x + \theta_0 > 0 \\ -1 & \text{if } \theta^\top x + \theta_0 \leq 0 \end{cases}$$

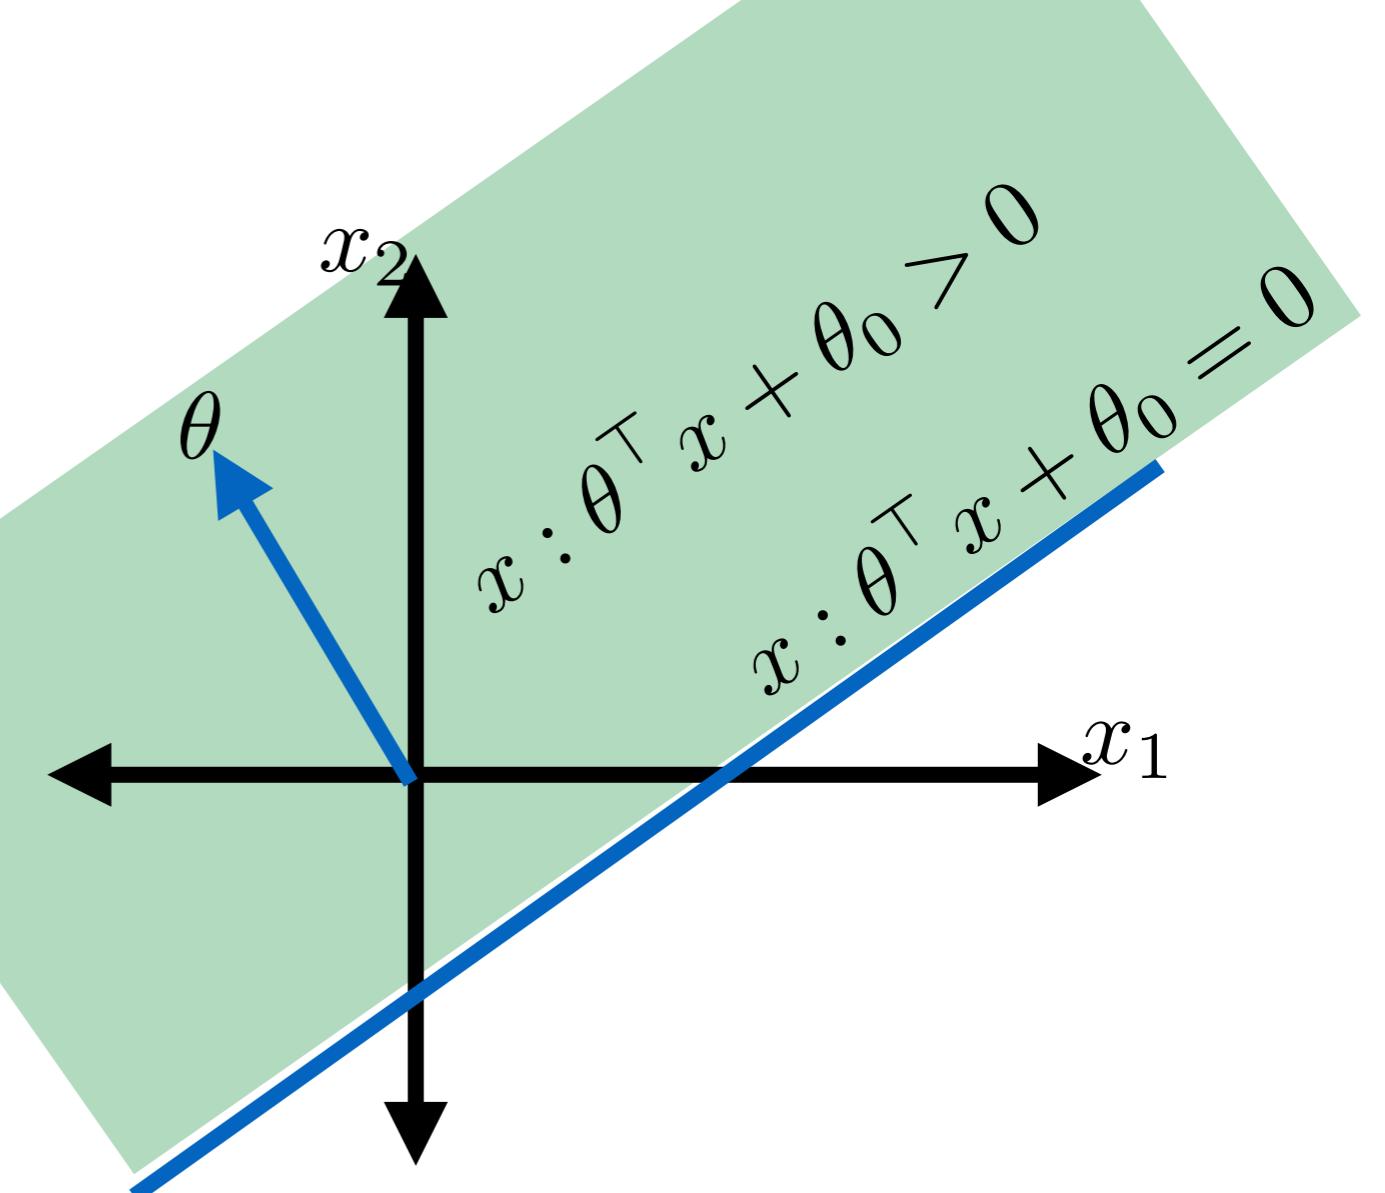


Recall: Classifiers

- A linear classifier:

$$h(x; \theta, \theta_0) = \text{sign}(\theta^\top x + \theta_0)$$

$$= \begin{cases} +1 & \text{if } \theta^\top x + \theta_0 > 0 \\ -1 & \text{if } \theta^\top x + \theta_0 \leq 0 \end{cases}$$

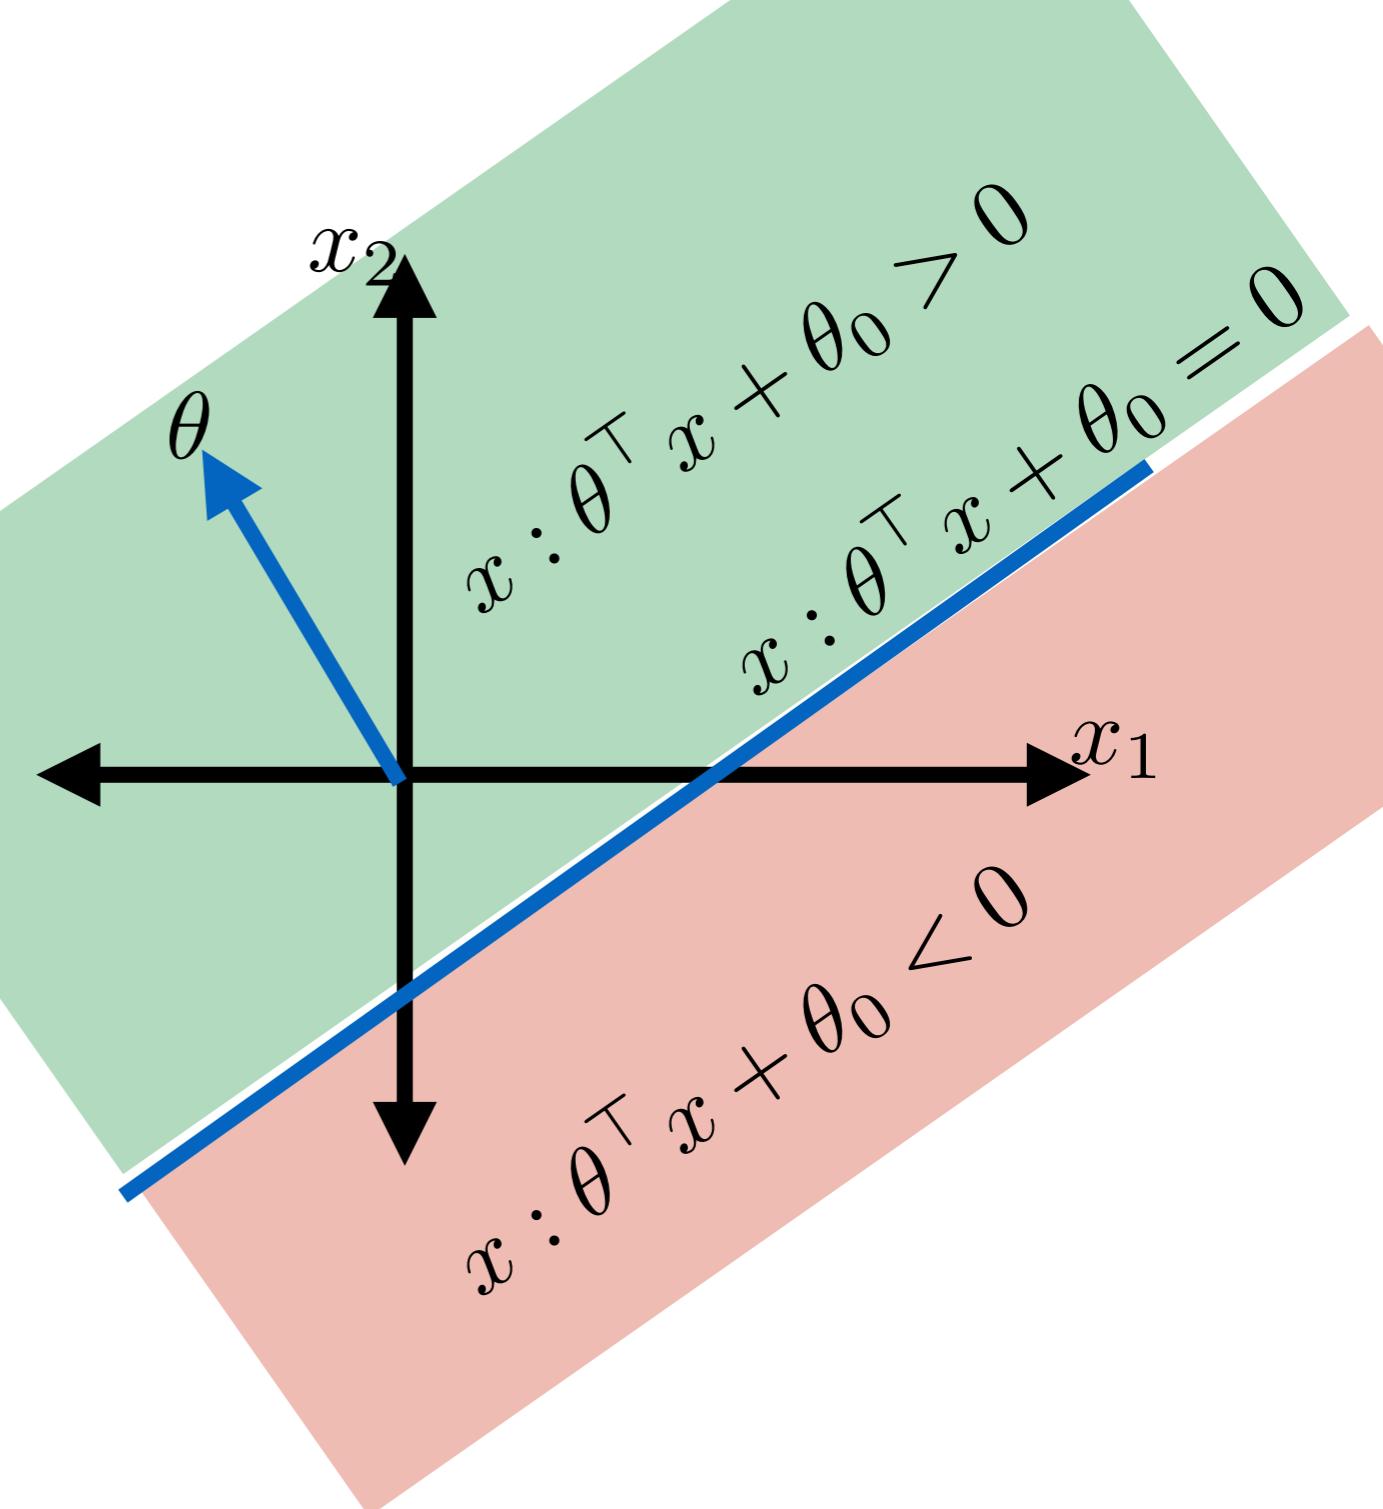


Recall: Classifiers

- A linear classifier:

$$h(x; \theta, \theta_0) = \text{sign}(\theta^\top x + \theta_0)$$

$$= \begin{cases} +1 & \text{if } \theta^\top x + \theta_0 > 0 \\ -1 & \text{if } \theta^\top x + \theta_0 \leq 0 \end{cases}$$

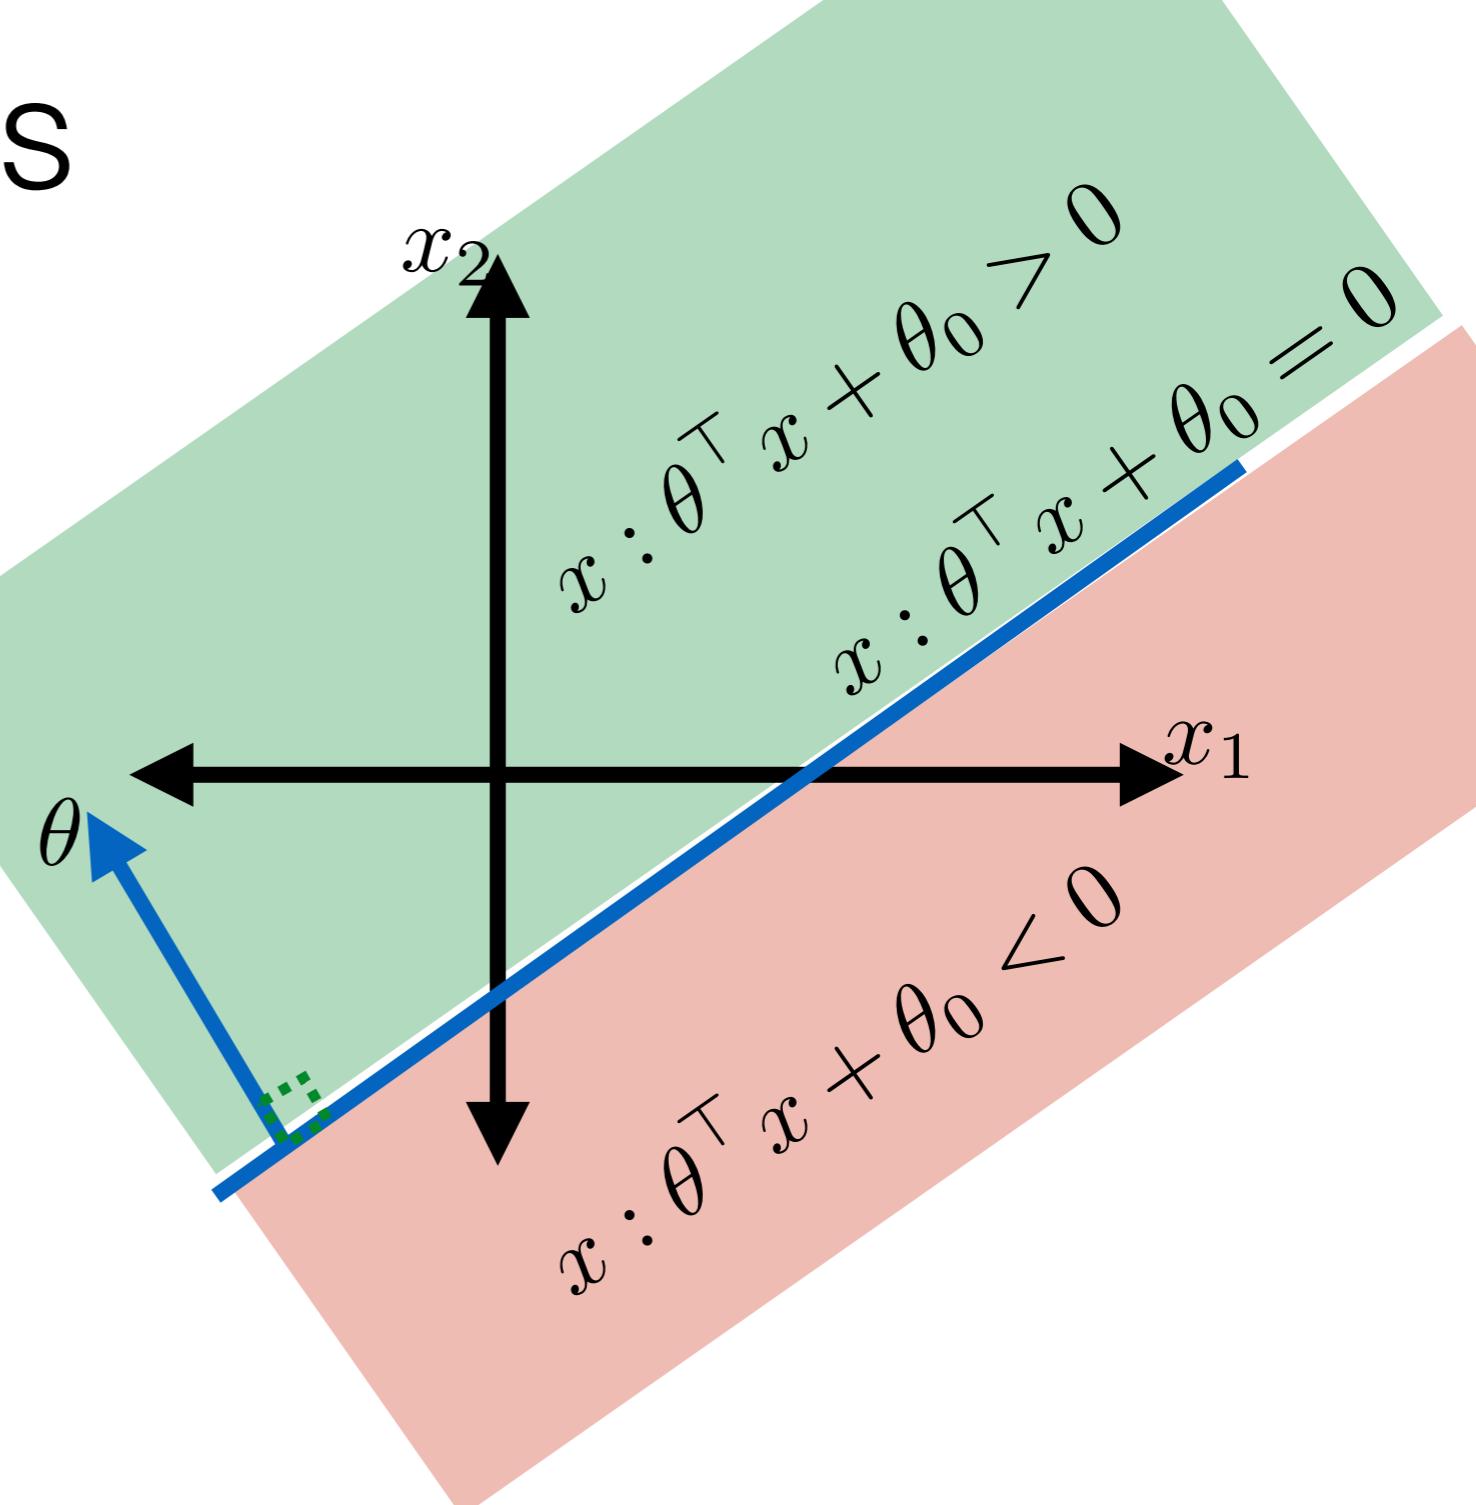


Recall: Classifiers

- A linear classifier:

$$h(x; \theta, \theta_0) = \text{sign}(\theta^\top x + \theta_0)$$

$$= \begin{cases} +1 & \text{if } \theta^\top x + \theta_0 > 0 \\ -1 & \text{if } \theta^\top x + \theta_0 \leq 0 \end{cases}$$

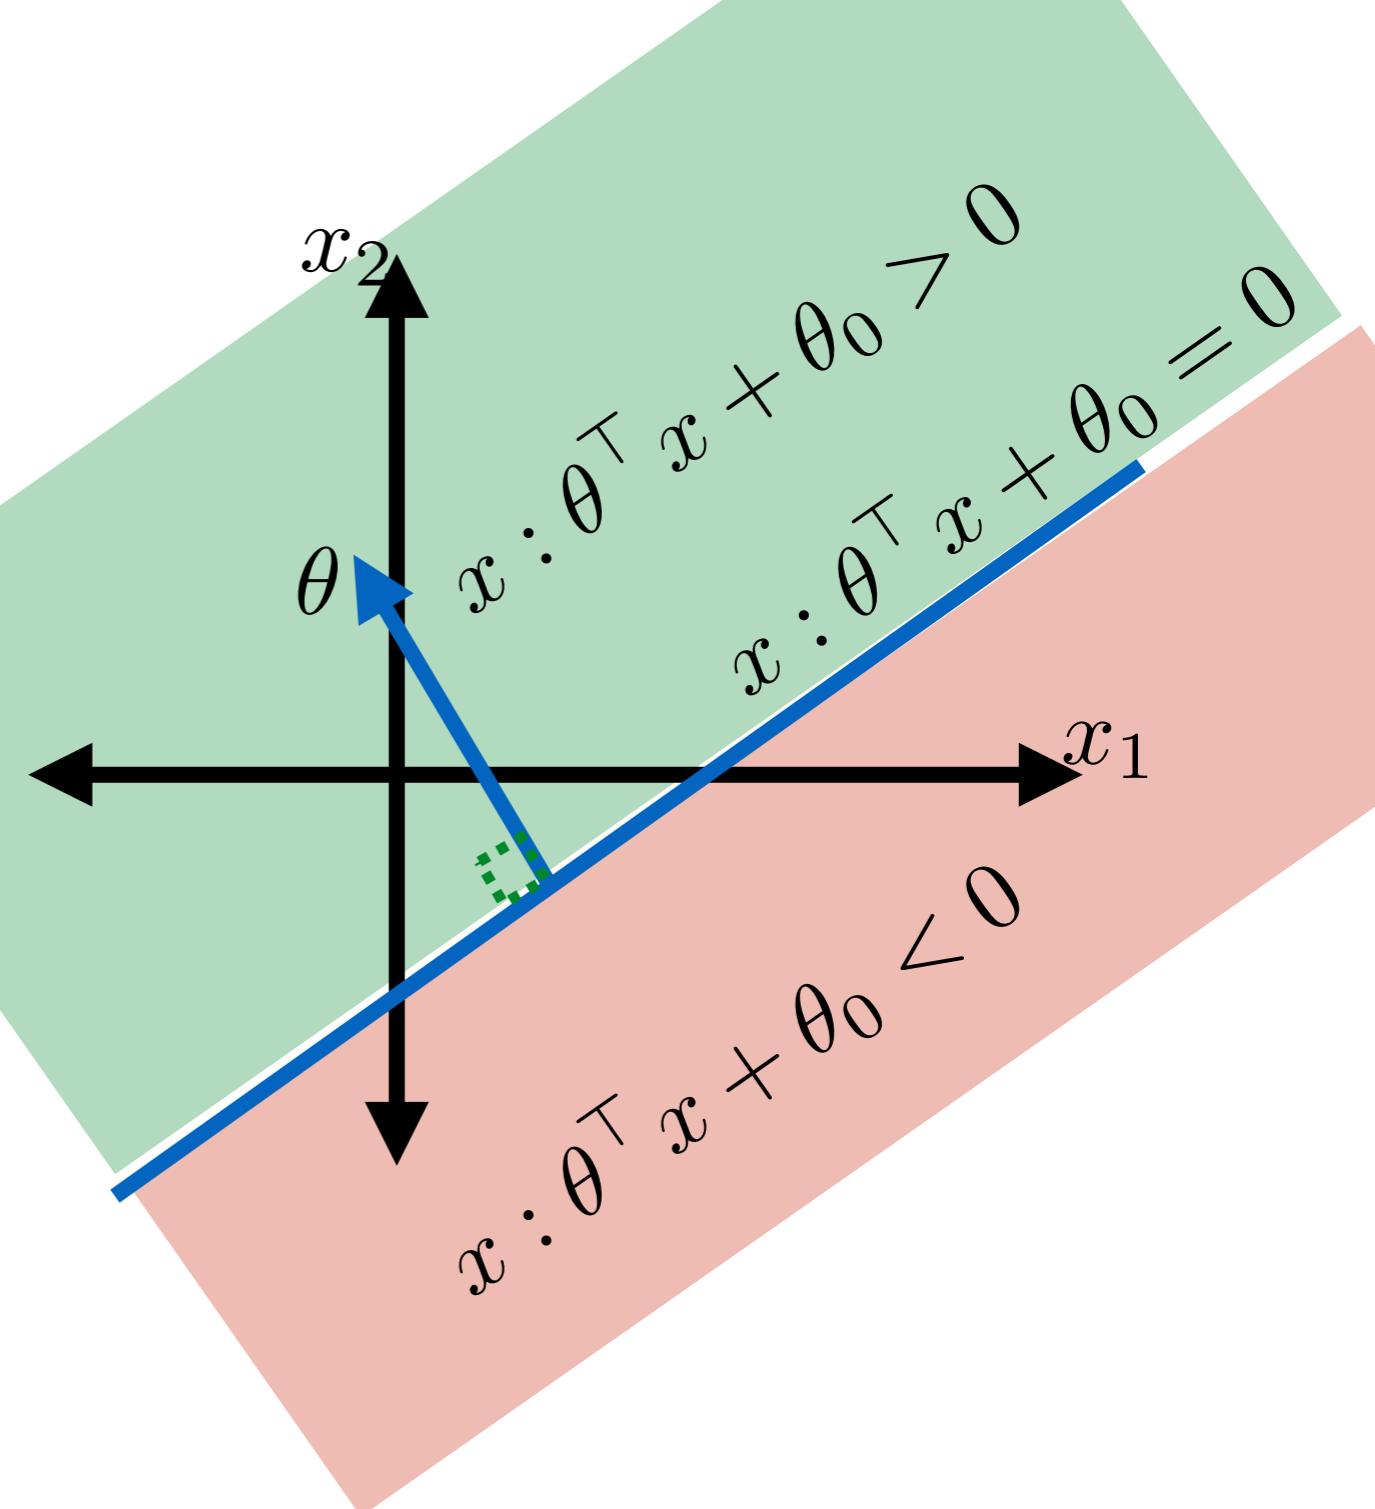


Recall: Classifiers

- A linear classifier:

$$h(x; \theta, \theta_0) = \text{sign}(\theta^\top x + \theta_0)$$

$$= \begin{cases} +1 & \text{if } \theta^\top x + \theta_0 > 0 \\ -1 & \text{if } \theta^\top x + \theta_0 \leq 0 \end{cases}$$

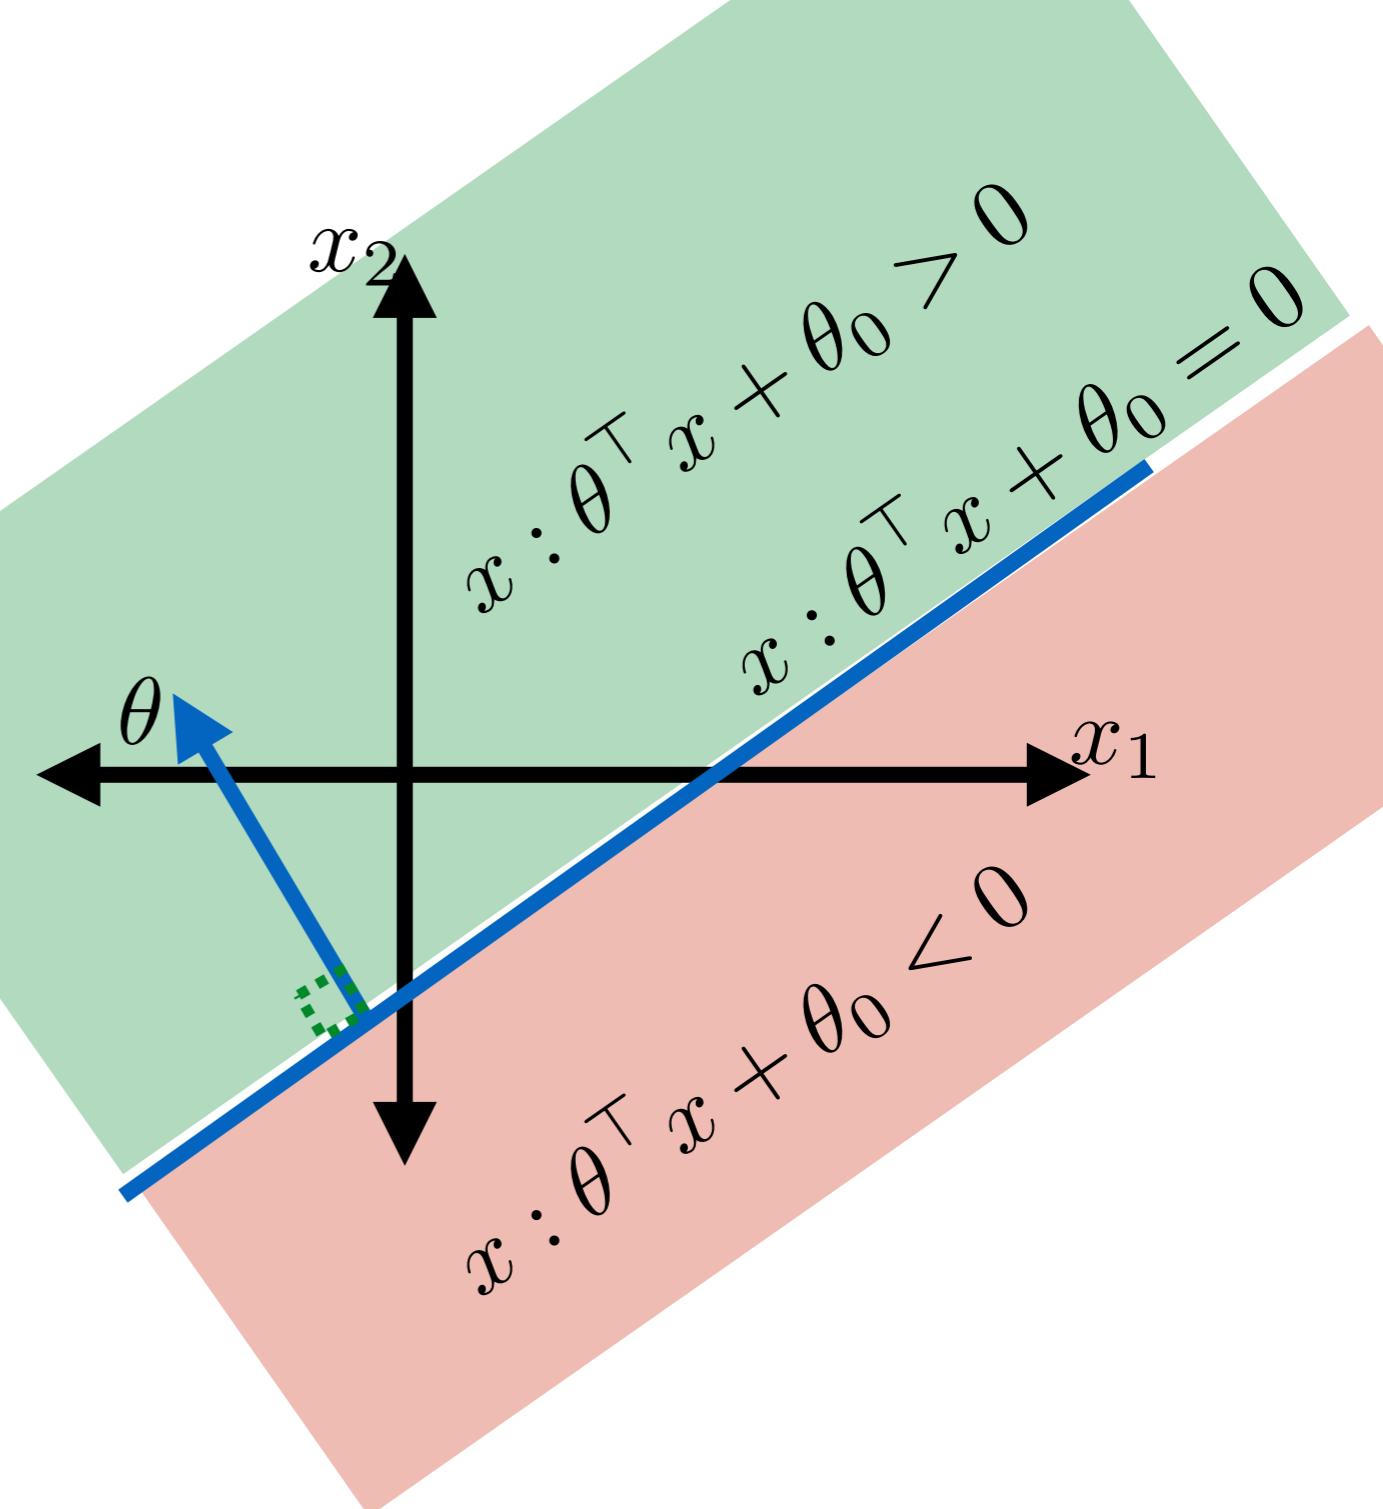


Recall: Classifiers

- A linear classifier:

$$h(x; \theta, \theta_0) = \text{sign}(\theta^\top x + \theta_0)$$

$$= \begin{cases} +1 & \text{if } \theta^\top x + \theta_0 > 0 \\ -1 & \text{if } \theta^\top x + \theta_0 \leq 0 \end{cases}$$

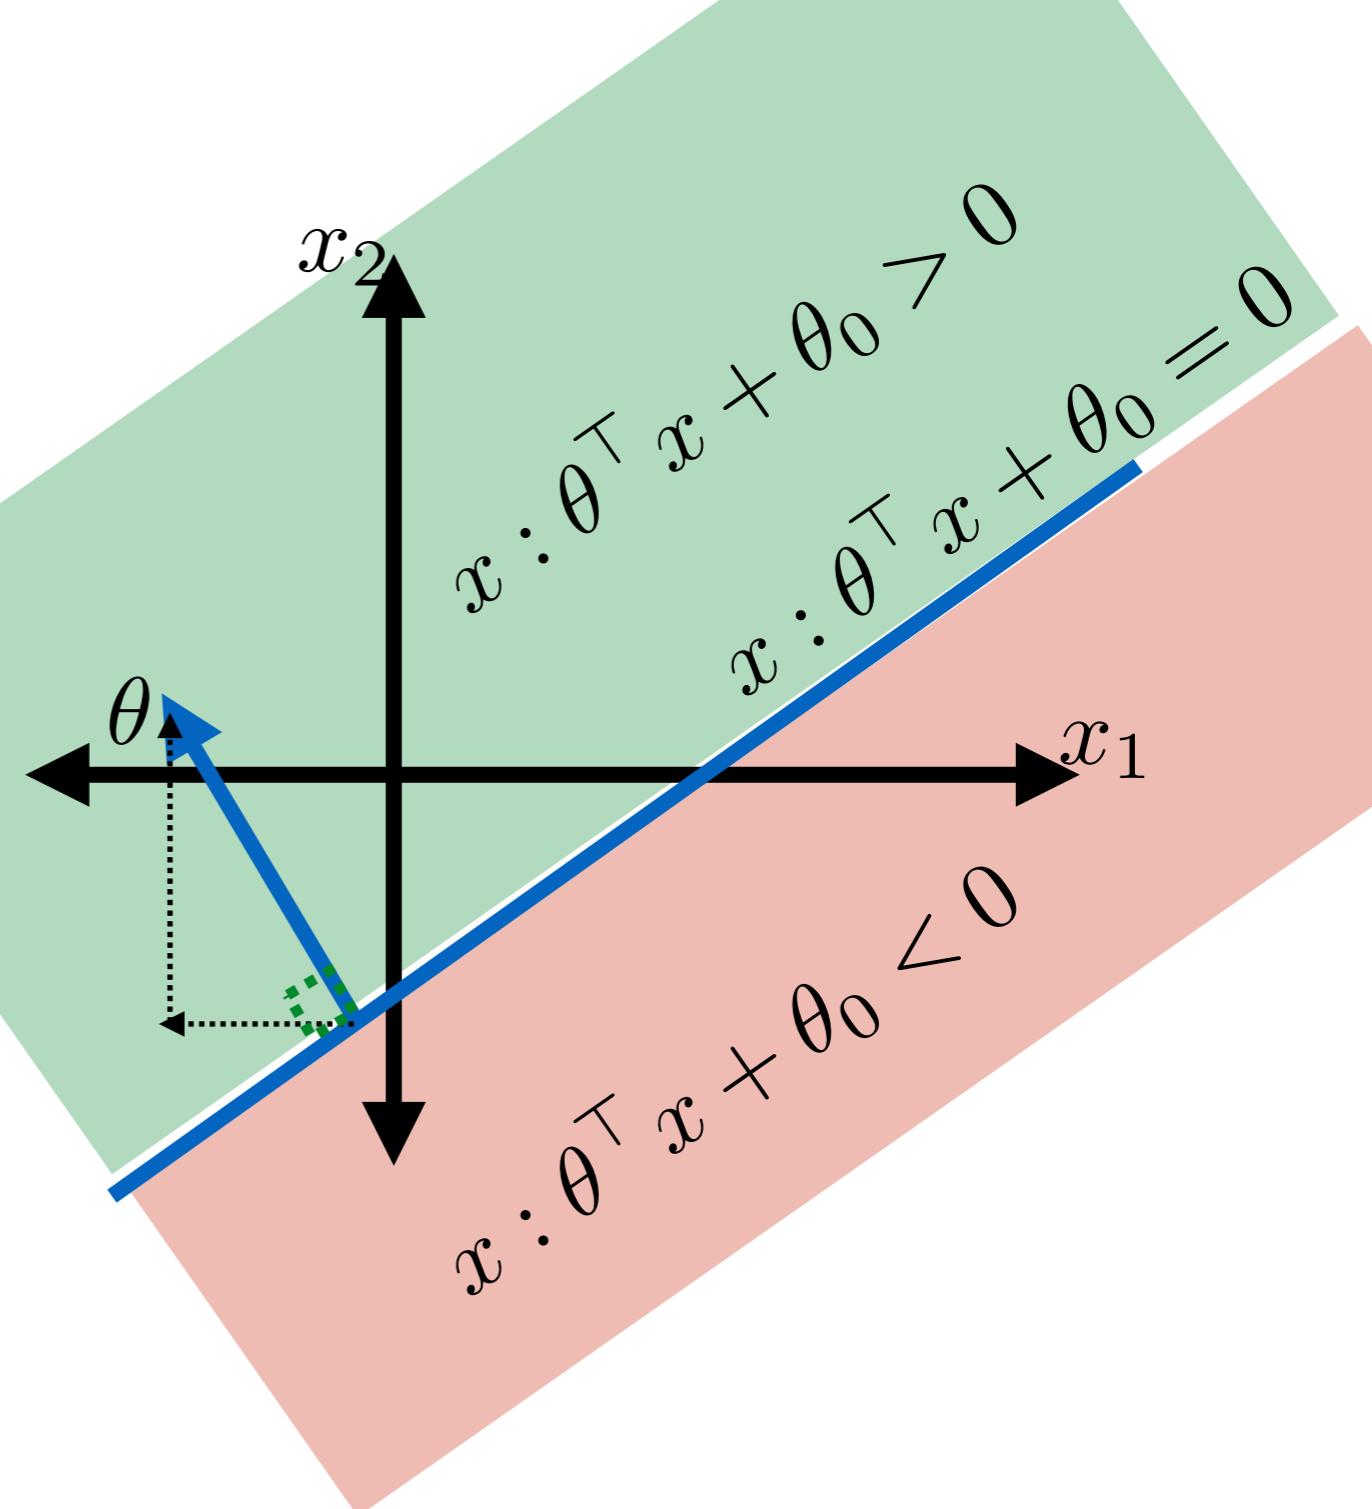


Recall: Classifiers

- A linear classifier:

$$h(x; \theta, \theta_0) = \text{sign}(\theta^\top x + \theta_0)$$

$$= \begin{cases} +1 & \text{if } \theta^\top x + \theta_0 > 0 \\ -1 & \text{if } \theta^\top x + \theta_0 \leq 0 \end{cases}$$

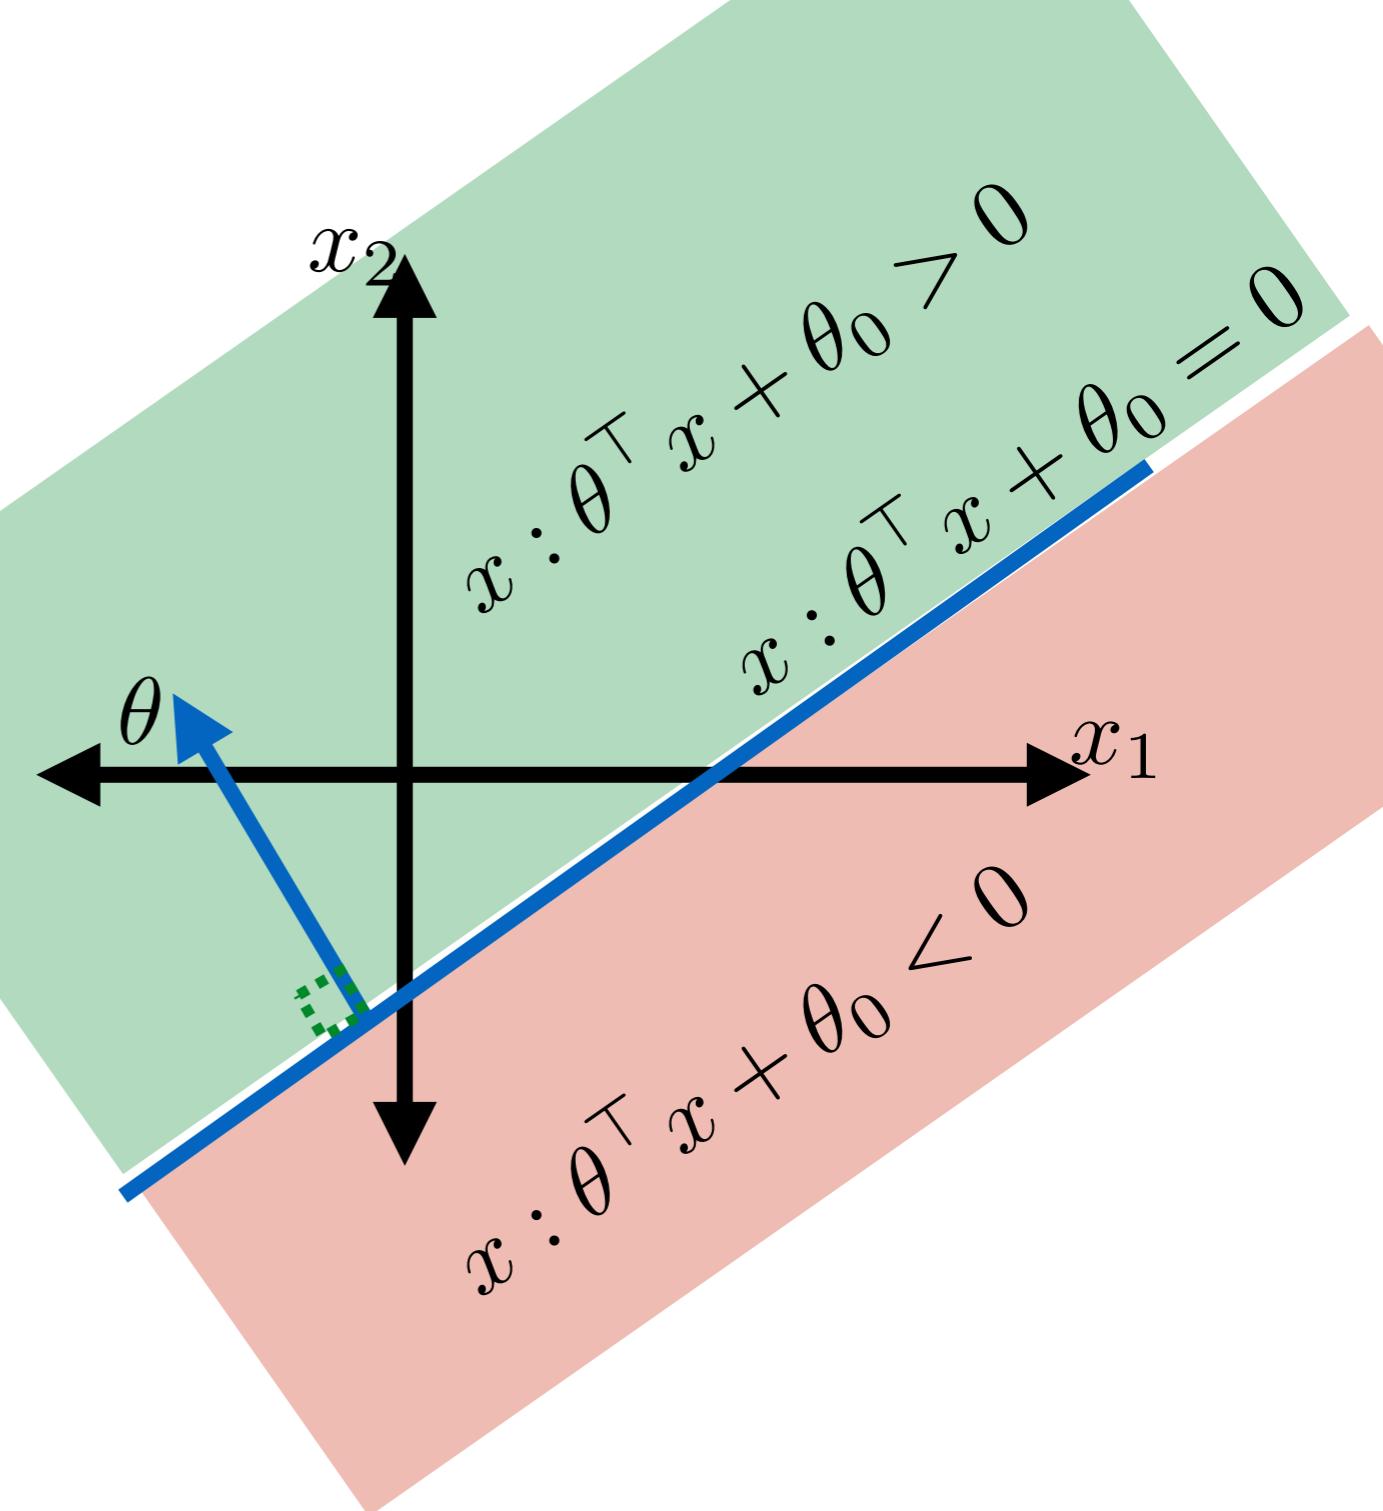


Recall: Classifiers

- A linear classifier:

$$h(x; \theta, \theta_0) = \text{sign}(\theta^\top x + \theta_0)$$

$$= \begin{cases} +1 & \text{if } \theta^\top x + \theta_0 > 0 \\ -1 & \text{if } \theta^\top x + \theta_0 \leq 0 \end{cases}$$

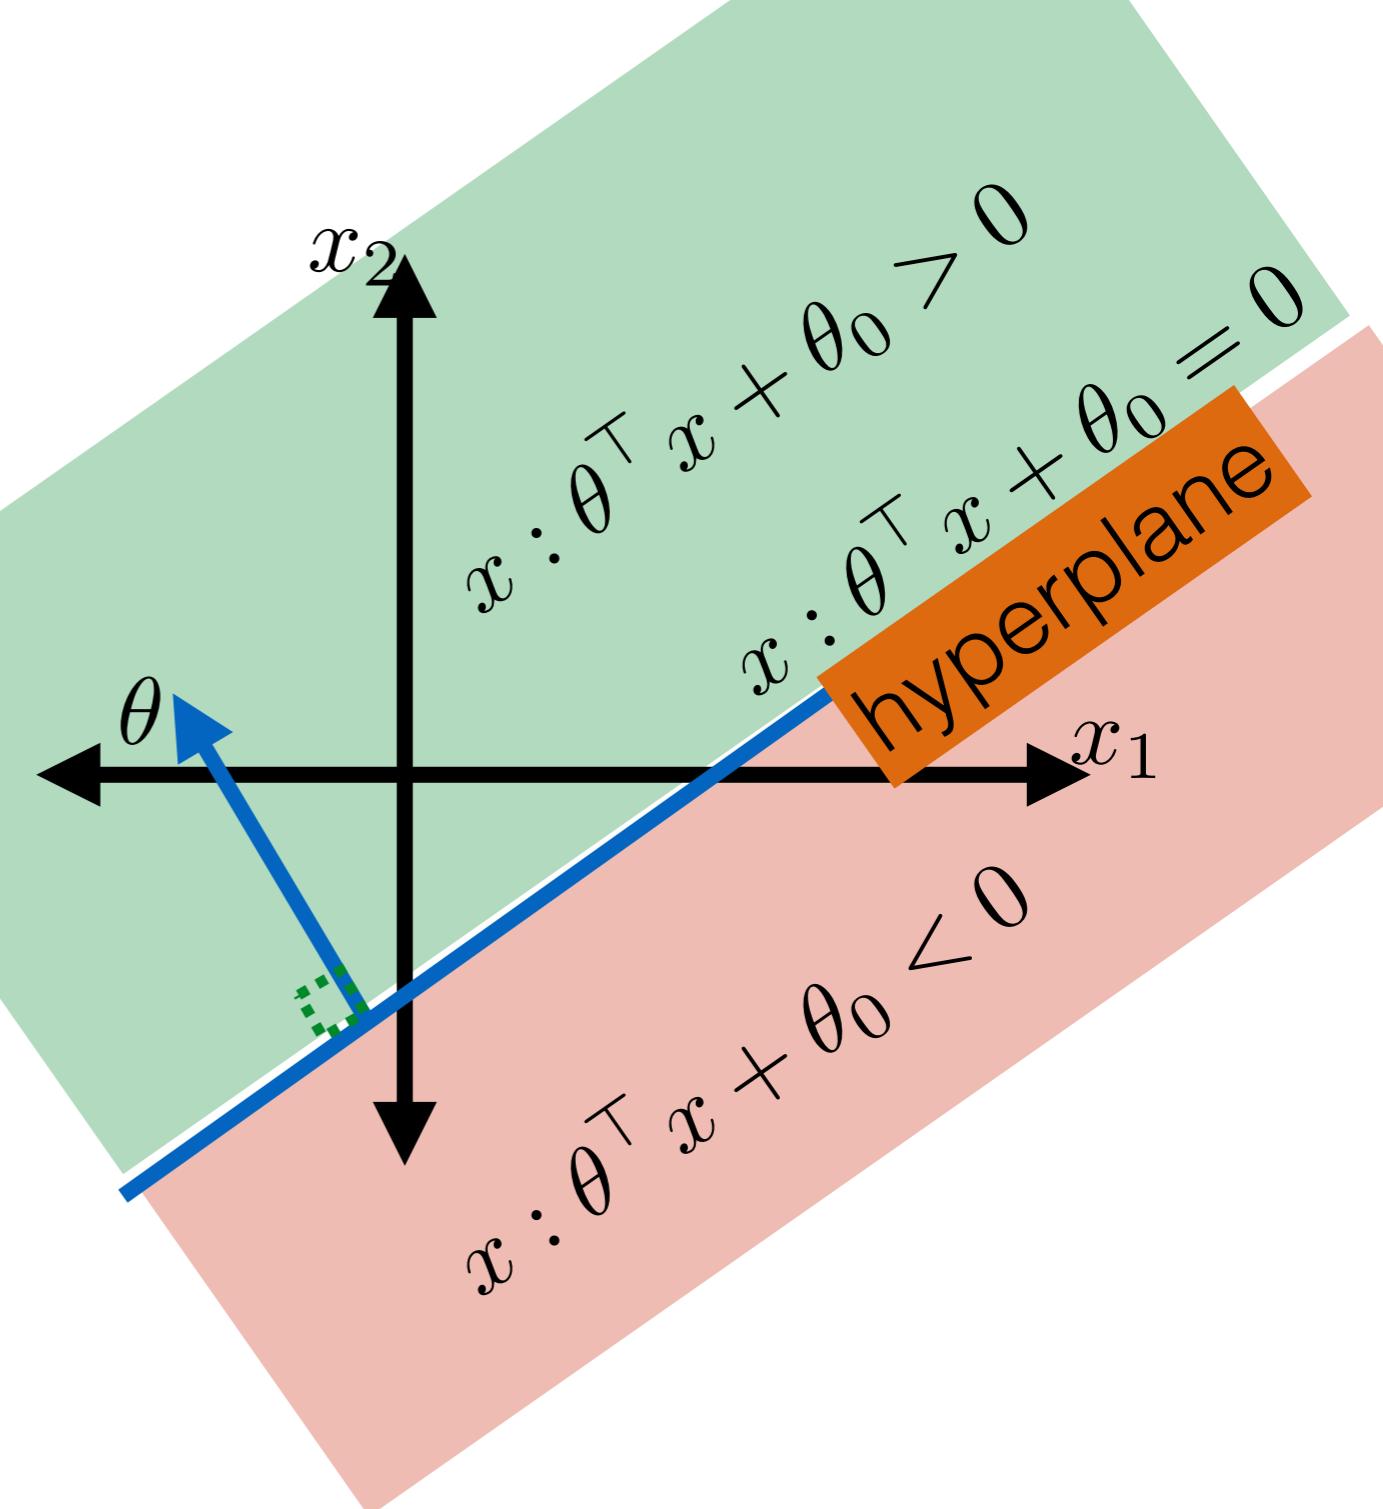


Recall: Classifiers

- A linear classifier:

$$h(x; \theta, \theta_0) = \text{sign}(\theta^\top x + \theta_0)$$

$$= \begin{cases} +1 & \text{if } \theta^\top x + \theta_0 > 0 \\ -1 & \text{if } \theta^\top x + \theta_0 \leq 0 \end{cases}$$

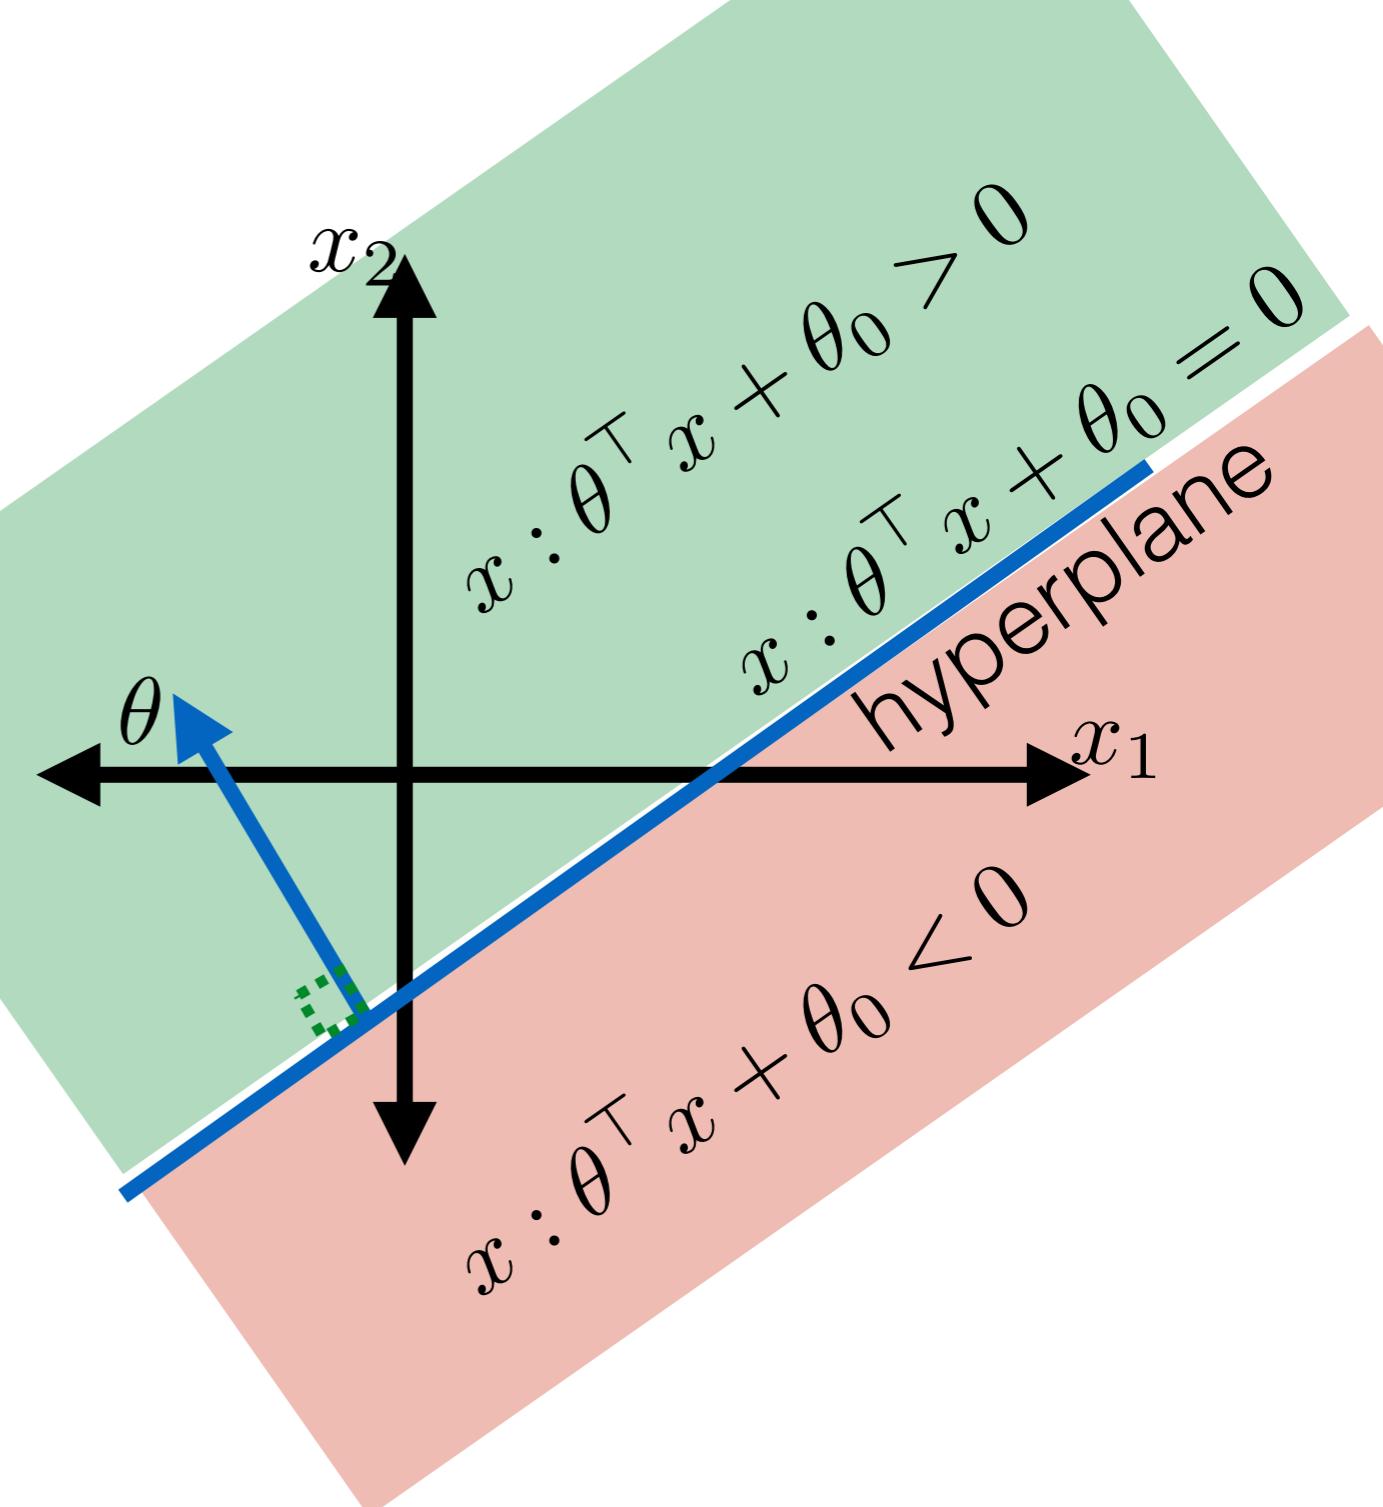


Recall: Classifiers

- A linear classifier:

$$h(x; \theta, \theta_0) = \text{sign}(\theta^\top x + \theta_0)$$

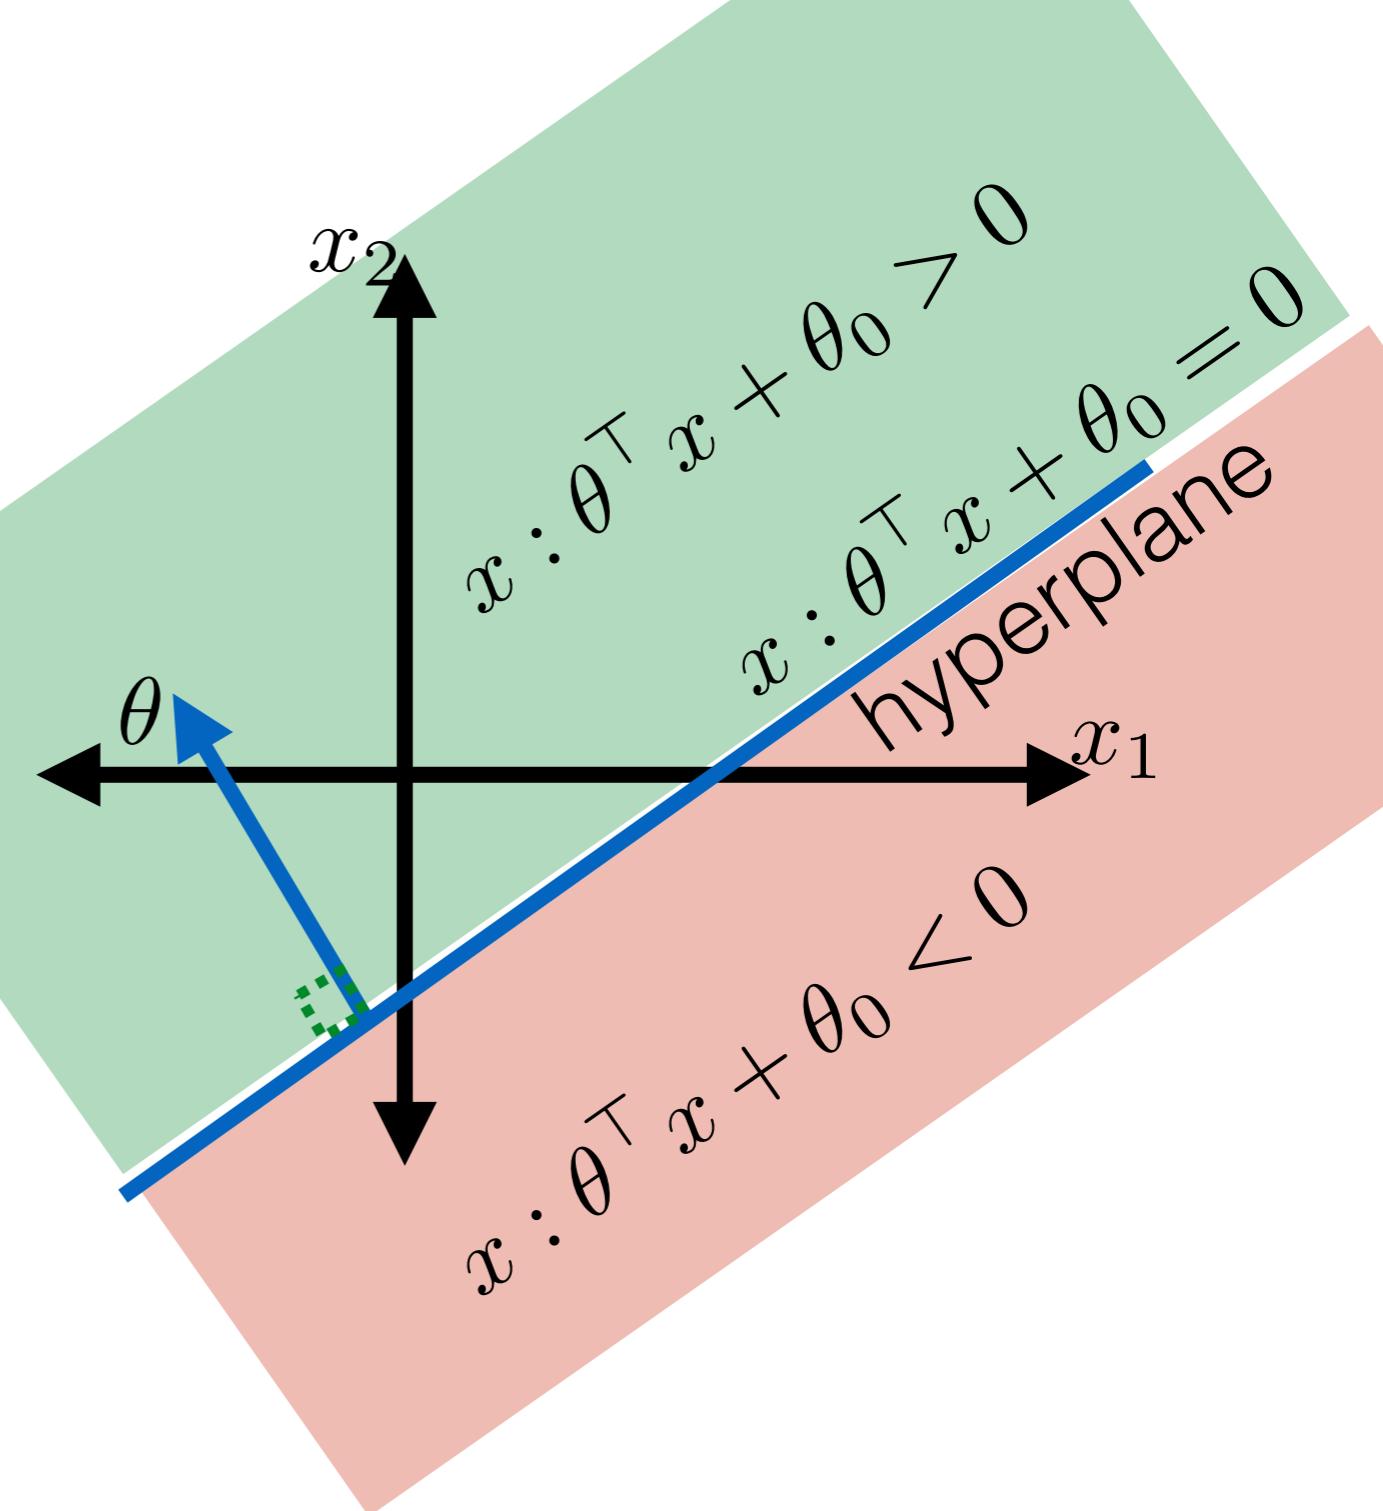
$$= \begin{cases} +1 & \text{if } \theta^\top x + \theta_0 > 0 \\ -1 & \text{if } \theta^\top x + \theta_0 \leq 0 \end{cases}$$



Recall: Classifiers

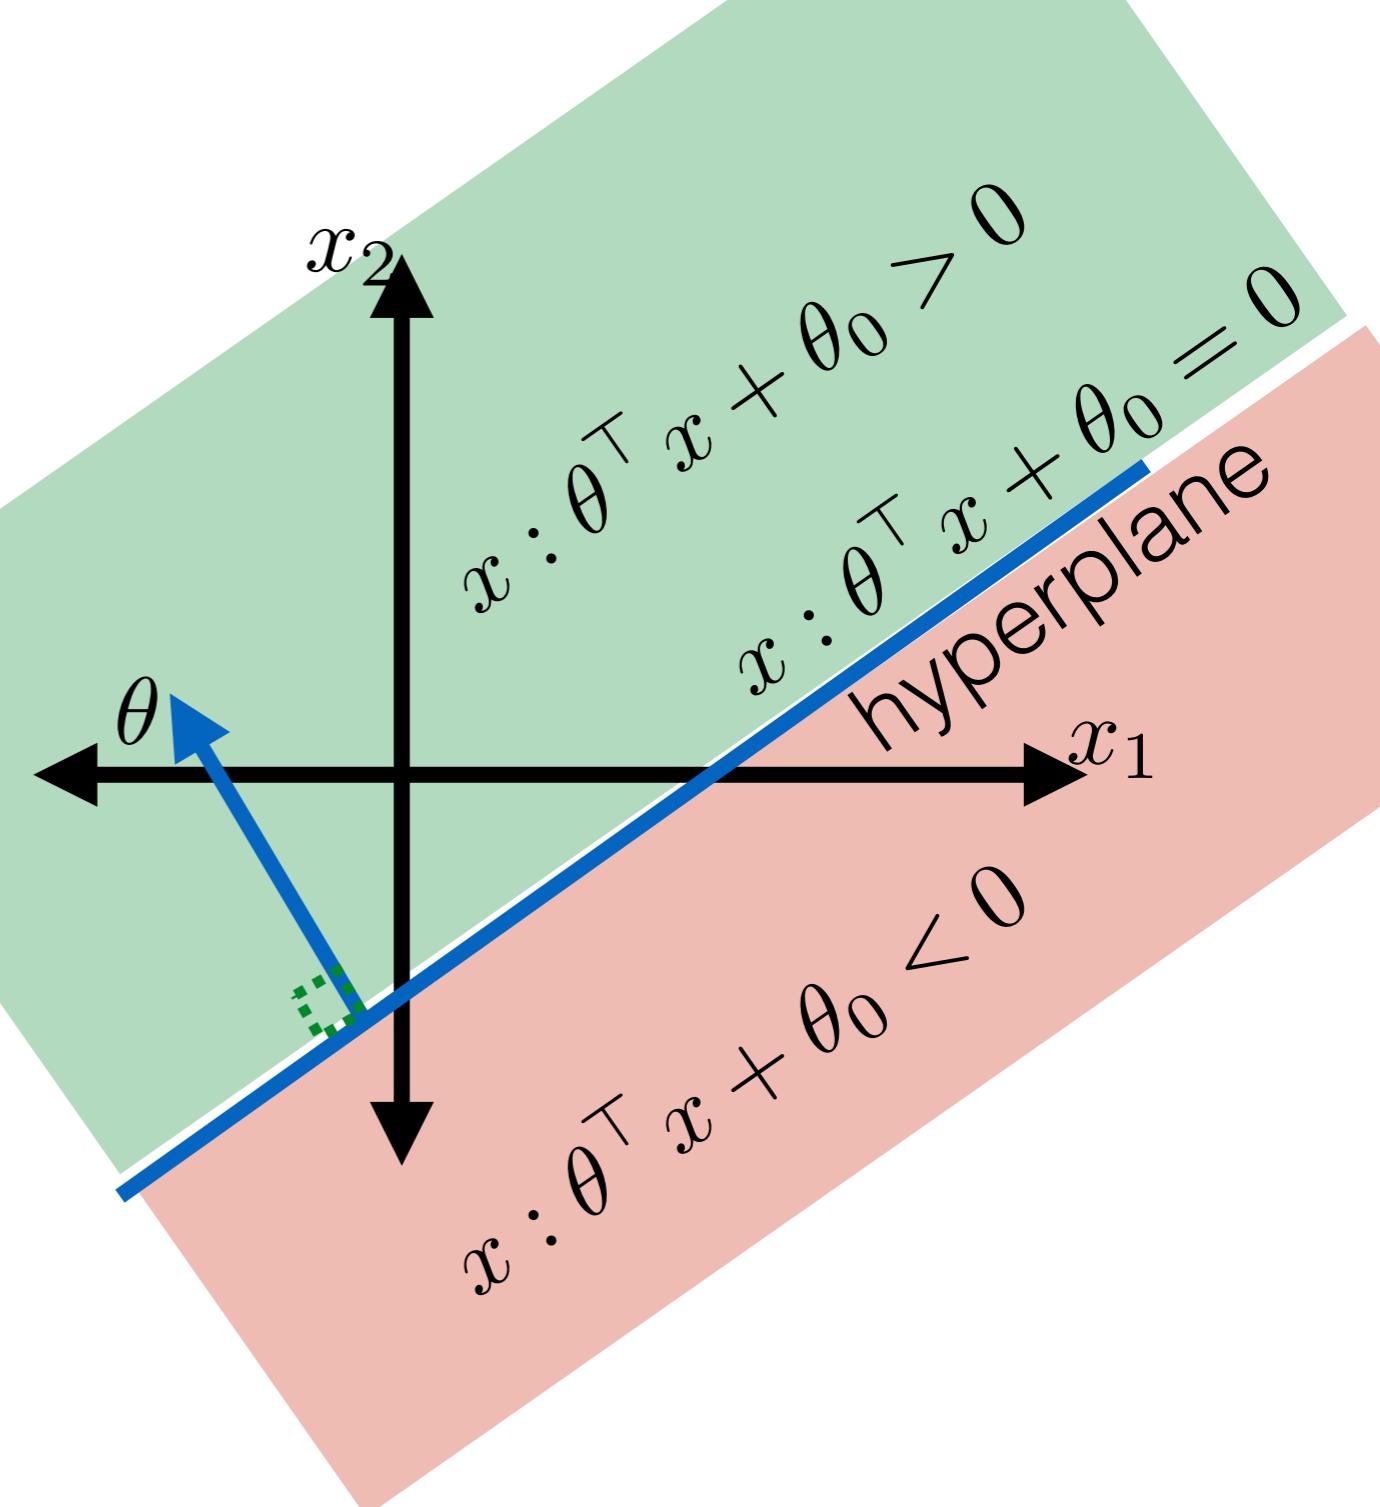
- A linear classifier:
$$h(x; \theta, \theta_0) = \text{sign}(\theta^\top x + \theta_0)$$

$$= \begin{cases} +1 & \text{if } \theta^\top x + \theta_0 > 0 \\ -1 & \text{if } \theta^\top x + \theta_0 \leq 0 \end{cases}$$
- Hypothesis class \mathcal{H} of all linear classifiers



Recall: Classifiers

- A linear classifier:
$$h(x; \theta, \theta_0) = \text{sign}(\theta^\top x + \theta_0)$$
$$= \begin{cases} +1 & \text{if } \theta^\top x + \theta_0 > 0 \\ -1 & \text{if } \theta^\top x + \theta_0 \leq 0 \end{cases}$$
- Hypothesis class \mathcal{H} of all linear classifiers
- 0-1 Loss
$$L(g, a) = \begin{cases} 0 & \text{if } g = a \\ 1 & \text{else} \end{cases}$$



Recall: Classifiers

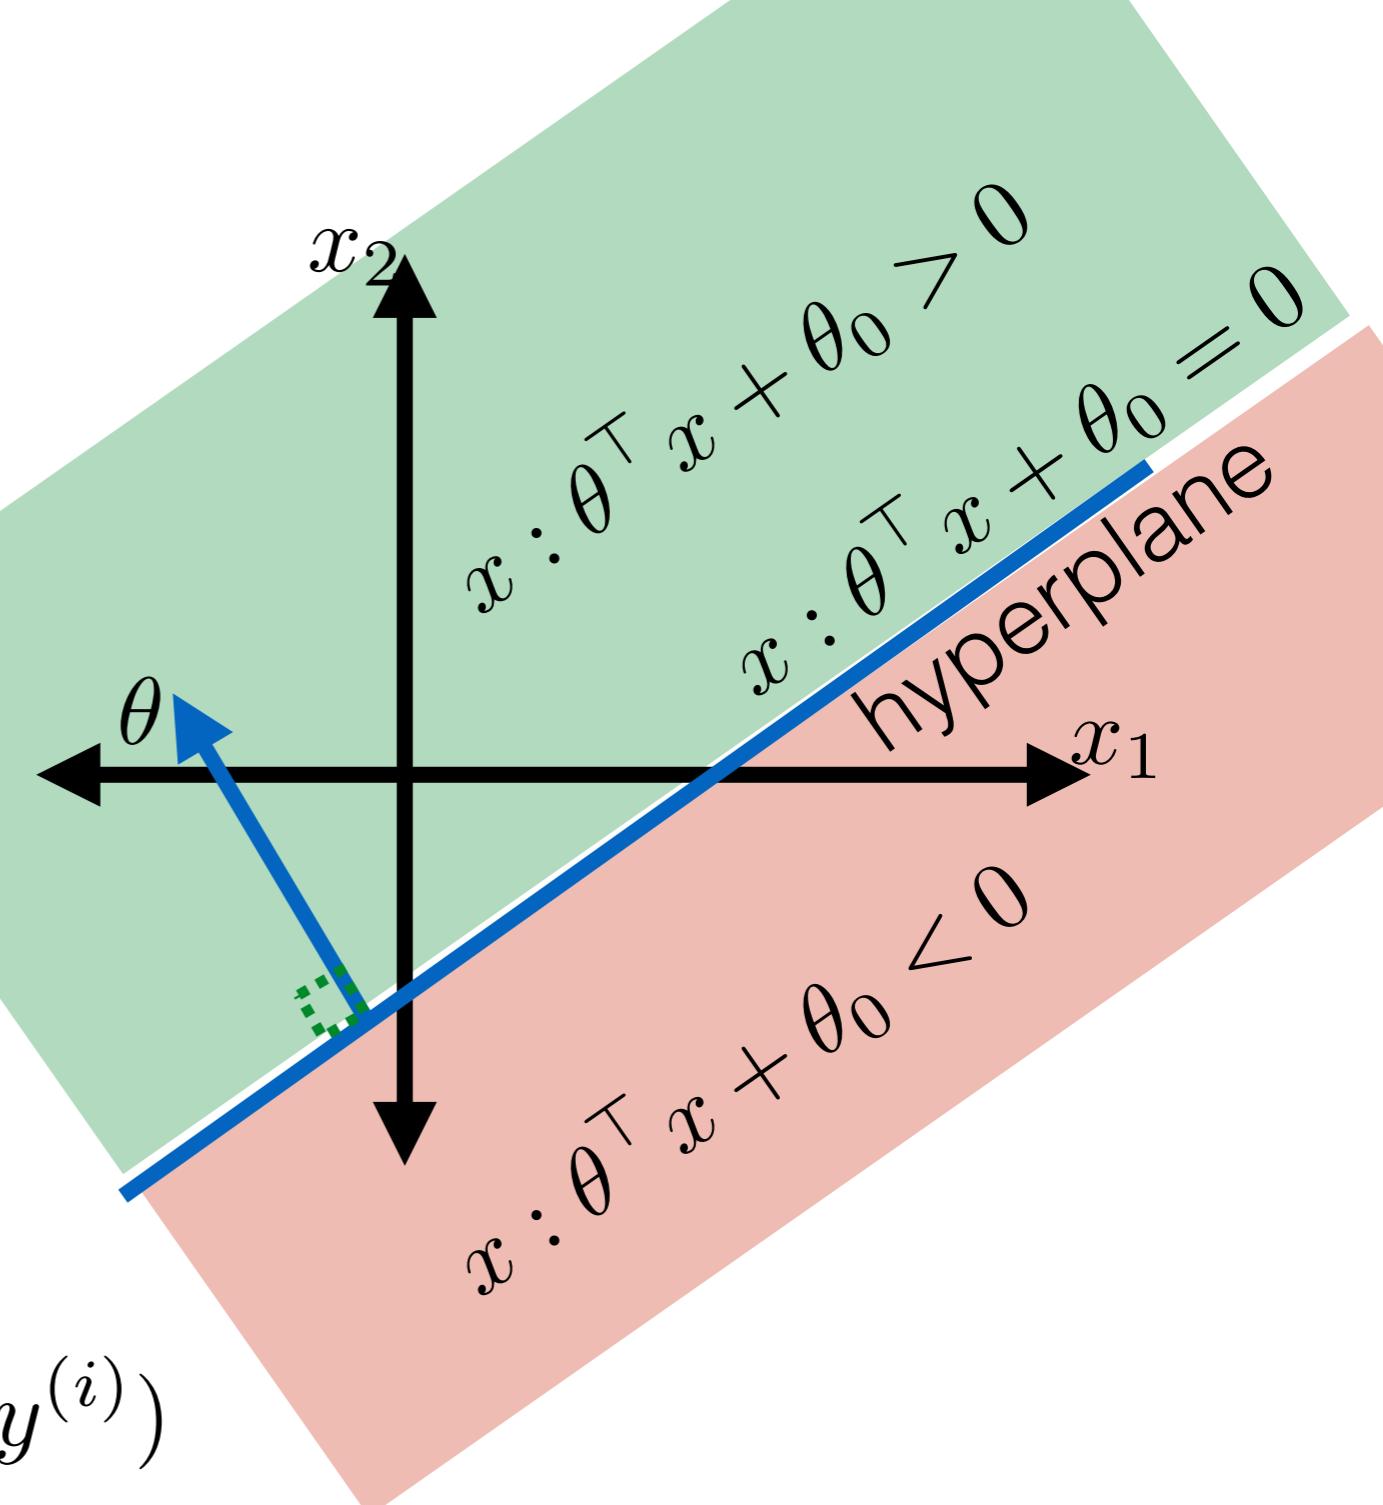
- A linear classifier:
$$h(x; \theta, \theta_0) = \text{sign}(\theta^\top x + \theta_0)$$

$$= \begin{cases} +1 & \text{if } \theta^\top x + \theta_0 > 0 \\ -1 & \text{if } \theta^\top x + \theta_0 \leq 0 \end{cases}$$

- Hypothesis class \mathcal{H} of all linear classifiers

- 0-1 Loss
$$L(g, a) = \begin{cases} 0 & \text{if } g = a \\ 1 & \text{else} \end{cases}$$

- Training error
$$\mathcal{E}_n(h) = \frac{1}{n} \sum_{i=1}^n L(h(x^{(i)}), y^{(i)})$$



Recall: Classifiers

- A linear classifier:
$$h(x; \theta, \theta_0) = \text{sign}(\theta^\top x + \theta_0)$$

$$= \begin{cases} +1 & \text{if } \theta^\top x + \theta_0 > 0 \\ -1 & \text{if } \theta^\top x + \theta_0 \leq 0 \end{cases}$$

- Hypothesis class \mathcal{H} of all linear classifiers

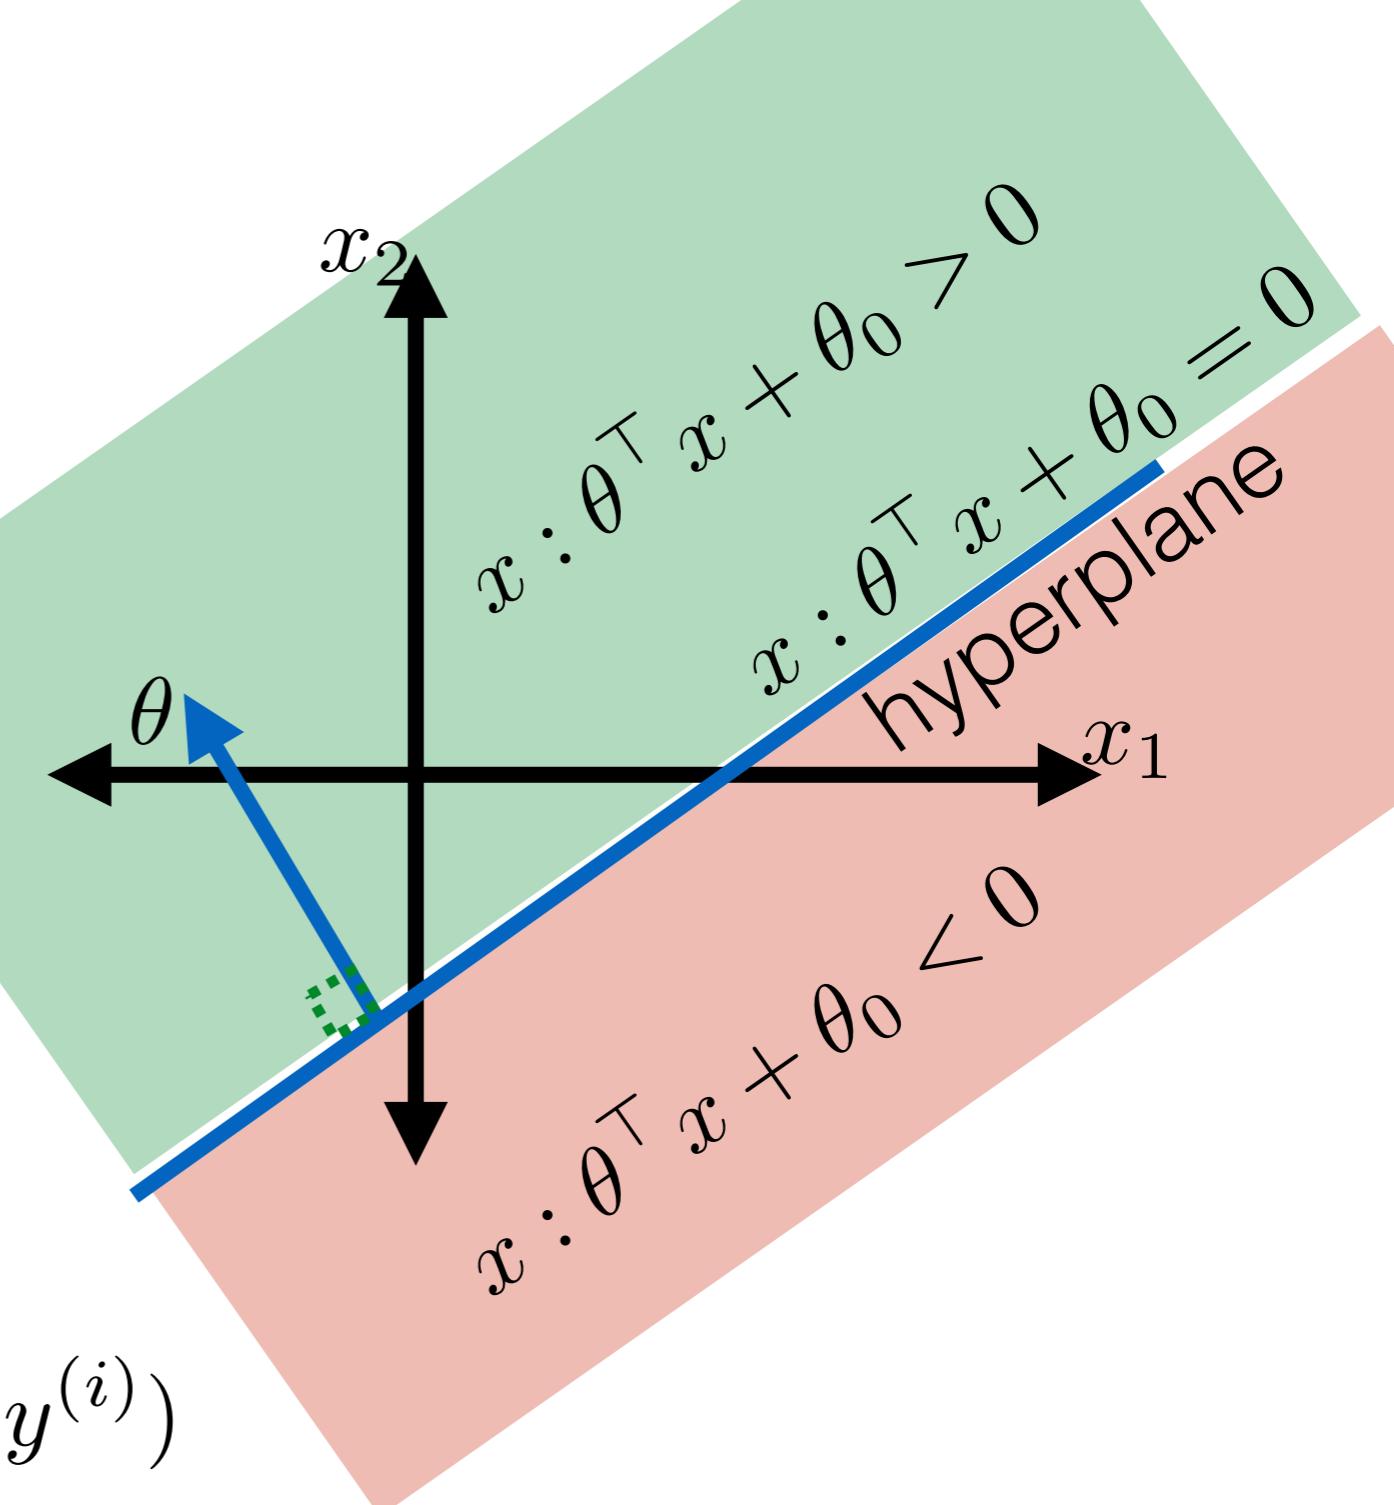
- 0-1 Loss

$$L(g, a) = \begin{cases} 0 & \text{if } g = a \\ 1 & \text{else} \end{cases}$$

- Training error

$$\mathcal{E}_n(h) = \frac{1}{n} \sum_{i=1}^n L(h(x^{(i)}), y^{(i)})$$

- Example learning algorithm (given hypotheses $h^{(j)}$)



Recall: Classifiers

- A linear classifier:
$$h(x; \theta, \theta_0) = \text{sign}(\theta^\top x + \theta_0)$$

$$= \begin{cases} +1 & \text{if } \theta^\top x + \theta_0 > 0 \\ -1 & \text{if } \theta^\top x + \theta_0 \leq 0 \end{cases}$$

- Hypothesis class \mathcal{H} of all linear classifiers

- 0-1 Loss

$$L(g, a) = \begin{cases} 0 & \text{if } g = a \\ 1 & \text{else} \end{cases}$$

- Training error

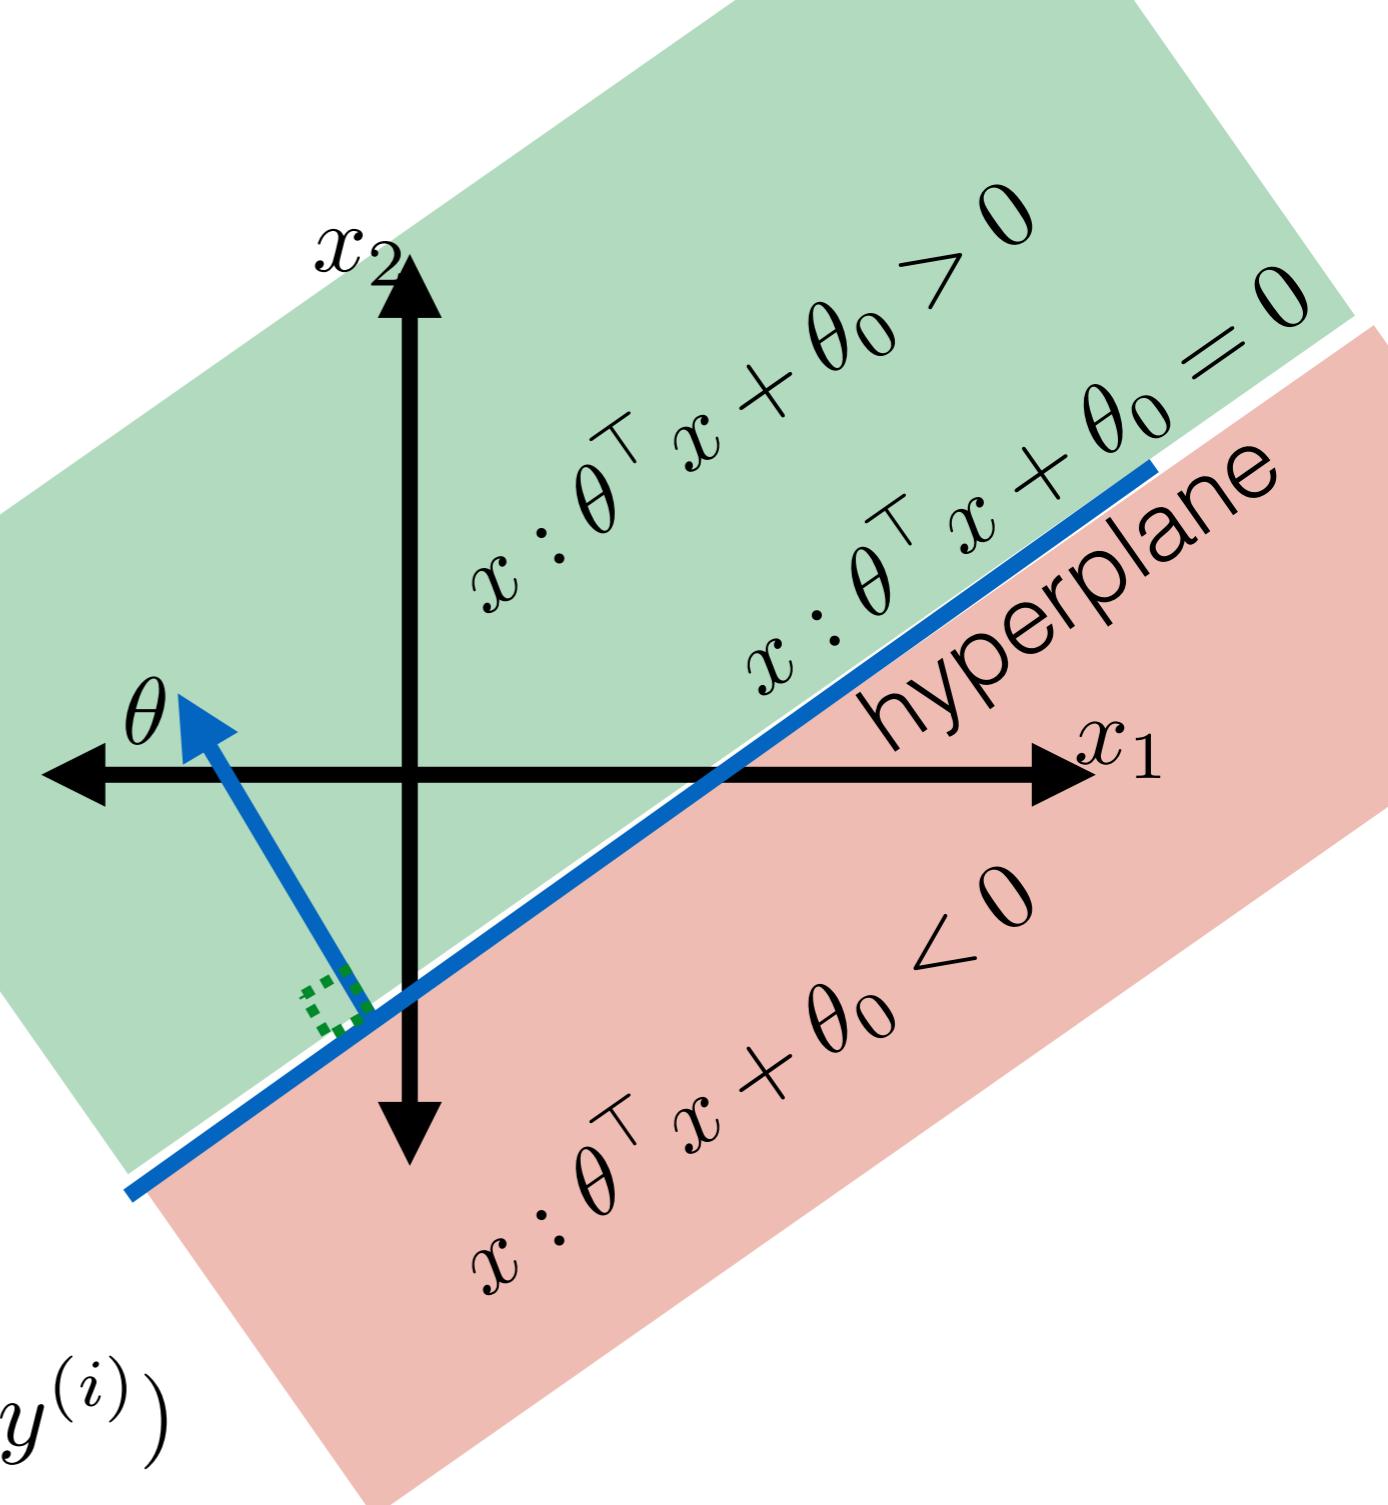
$$\mathcal{E}_n(h) = \frac{1}{n} \sum_{i=1}^n L(h(x^{(i)}), y^{(i)})$$

- Example learning algorithm (given hypotheses $h^{(j)}$)

Ex_learning_alg(\mathcal{D}_n ; k)

Set $j^* = \operatorname{argmin}_{j \in \{1, \dots, k\}} \mathcal{E}_n(h^{(j)})$

Return $h^{(j^*)}$



Recall: Classifiers

- A linear classifier:
$$h(x; \theta, \theta_0) = \text{sign}(\theta^\top x + \theta_0)$$

$$= \begin{cases} +1 & \text{if } \theta^\top x + \theta_0 > 0 \\ -1 & \text{if } \theta^\top x + \theta_0 \leq 0 \end{cases}$$

- Hypothesis class \mathcal{H} of all linear classifiers

- 0-1 Loss

$$L(g, a) = \begin{cases} 0 & \text{if } g = a \\ 1 & \text{else} \end{cases}$$

- Training error

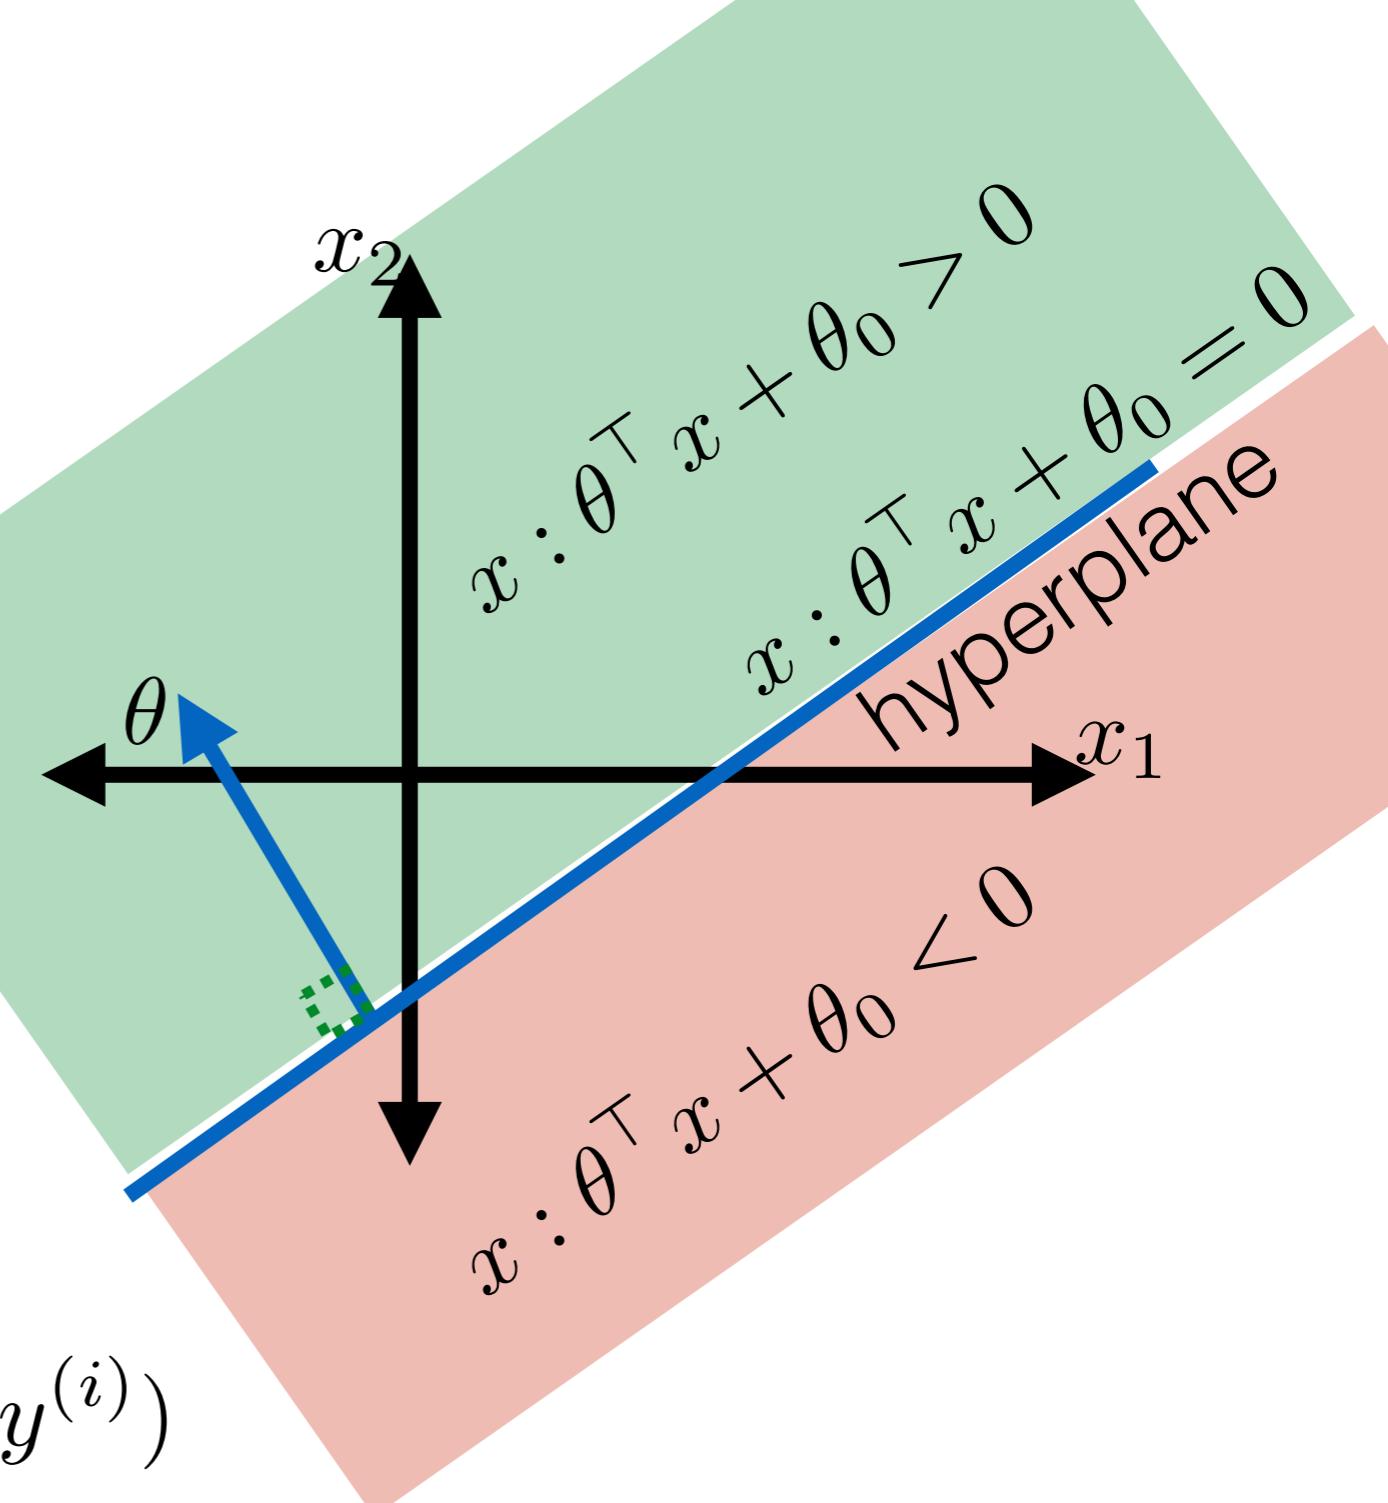
$$\mathcal{E}_n(h) = \frac{1}{n} \sum_{i=1}^n L(h(x^{(i)}), y^{(i)})$$

- Example learning algorithm (given hypotheses $h^{(j)}$)

Ex_learning_alg(\mathcal{D}_n ; k)

Set $j^* = \operatorname{argmin}_{j \in \{1, \dots, k\}} \mathcal{E}_n(h^{(j)})$

Return $h^{(j^*)}$



[demo]

Perceptron Algorithm

Perceptron Algorithm

Perceptron

Perceptron Algorithm

Perceptron(\mathcal{D}_n ; τ)

Perceptron Algorithm

Perceptron(\mathcal{D}_n ; τ)

Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$

Initialize $\theta_0 = 0$

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$ [How many 0s?]

Initialize $\theta_0 = 0$

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$ [How many 0s?]

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$ [How many 0s?]

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

Perceptron Algorithm

Perceptron(\mathcal{D}_n ; τ)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^T$ [How many 0s?]

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

Perceptron Algorithm

Perceptron(\mathcal{D}_n ; τ)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$ [How many 0s?]

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$ [How many 0s?]

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$ [How many 0s?]

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

Perceptron Algorithm

Perceptron(\mathcal{D}_n ; τ)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$ [How many 0s?]

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$ [How many 0s?]

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$ [How many 0s?]

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

Perceptron Algorithm

Perceptron (\mathcal{D}_n ; τ)

Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$ [How many 0s?]

Initialize $\theta_0 = 0$

for $t = 1$ to τ [i.e. True if either:

changed = False

for i = 1 to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

[How many 0s?]

[i.e. True if either:

A. point is not on the line
 & prediction is wrong

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

[How many 0s?]

[i.e. True if either:

 A. point is not on the line
 & prediction is wrong

 B. point is on the line

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

[How many 0s?]

[i.e. True if either:

 A. point is not on the line
 & prediction is wrong

 B. point is on the line

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

[How many 0s?]

[i.e. True if either:

- A. point is not on the line & prediction is wrong
- B. point is on the line
- C. initial step]

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

[How many 0s?]

[i.e. True if either:

- A. point is not on the line & prediction is wrong
- B. point is on the line
- C. initial step]

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

 Set $\theta = \theta + y^{(i)}x^{(i)}$

 Set $\theta_0 = \theta_0 + y^{(i)}$

[How many 0s?]

[i.e. True if either:

- A. point is not on the line & prediction is wrong
- B. point is on the line
- C. initial step]

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

 Set $\theta = \theta + y^{(i)}x^{(i)}$

 Set $\theta_0 = \theta_0 + y^{(i)}$

 changed = True

[How many 0s?]

[i.e. True if either:

- A. point is not on the line & prediction is wrong
- B. point is on the line
- C. initial step]

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

 Set $\theta = \theta + y^{(i)}x^{(i)}$

 Set $\theta_0 = \theta_0 + y^{(i)}$

 changed = True

if not changed

break

[How many 0s?]

[i.e. True if either:

- A. point is not on the line & prediction is wrong
- B. point is on the line
- C. initial step]

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

 Set $\theta = \theta + y^{(i)}x^{(i)}$

 Set $\theta_0 = \theta_0 + y^{(i)}$

 changed = True

if not changed

break

Return θ, θ_0

[How many 0s?]

[i.e. True if either:

- A. point is not on the line & prediction is wrong
- B. point is on the line
- C. initial step]

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

 Set $\theta = \theta + y^{(i)}x^{(i)}$

 Set $\theta_0 = \theta_0 + y^{(i)}$

 changed = True

if not changed

break

Return θ, θ_0

[How many 0s?]

[i.e. True if either:

- A. point is not on the line & prediction is wrong
- B. point is on the line
- C. initial step]

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^T$

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

 Set $\theta = \theta + y^{(i)}x^{(i)}$

 Set $\theta_0 = \theta_0 + y^{(i)}$

 changed = True

if not changed

break

Return θ, θ_0

[How many 0s?]

[i.e. True if either:

- A. point is not on the line & prediction is wrong
- B. point is on the line
- C. initial step]

What does an update do?

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

 Set $\theta = \theta + y^{(i)}x^{(i)}$

 Set $\theta_0 = \theta_0 + y^{(i)}$

 changed = True

if not changed

break

Return θ, θ_0

[How many 0s?]

[i.e. True if either:

- A. point is not on the line & prediction is wrong
- B. point is on the line
- C. initial step]

What does an update do?

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

 Set $\theta = \theta + y^{(i)}x^{(i)}$

 Set $\theta_0 = \theta_0 + y^{(i)}$

 changed = True

if not changed

break

Return θ, θ_0

[How many 0s?]

[i.e. True if either:

- A. point is not on the line & prediction is wrong
- B. point is on the line
- C. initial step]

What does an update do?

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^T$

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

 Set $\theta = \theta + y^{(i)}x^{(i)}$

 Set $\theta_0 = \theta_0 + y^{(i)}$

 changed = True

if not changed

break

Return θ, θ_0

[How many 0s?]

[i.e. True if either:

- A. point is not on the line & prediction is wrong
- B. point is on the line
- C. initial step]

What does an update do?

$$y^{(i)} \left(\theta_{\text{updated}}^\top x^{(i)} + \theta_{0,\text{updated}} \right)$$

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

 Set $\theta = \theta + y^{(i)}x^{(i)}$

 Set $\theta_0 = \theta_0 + y^{(i)}$

 changed = True

if not changed

break

Return θ, θ_0

[How many 0s?]

[i.e. True if either:

- A. point is not on the line & prediction is wrong
- B. point is on the line
- C. initial step]

What does an update do?

$$y^{(i)} \left((\theta + y^{(i)}x^{(i)})^\top x^{(i)} + (\theta_0 + y^{(i)}) \right)$$

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

 Set $\theta = \theta + y^{(i)}x^{(i)}$

 Set $\theta_0 = \theta_0 + y^{(i)}$

 changed = True

if not changed

break

Return θ, θ_0

[How many 0s?]

[i.e. True if either:

- A. point is not on the line & prediction is wrong
- B. point is on the line
- C. initial step]

What does an update do?

$$\begin{aligned} & y^{(i)} \left((\theta + y^{(i)}x^{(i)})^\top x^{(i)} + (\theta_0 + y^{(i)}) \right) \\ &= y^{(i)}(\theta^\top x^{(i)} + \theta_0) + (y^{(i)})^2(x^{(i)\top} x^{(i)} + 1) \end{aligned}$$

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

 Set $\theta = \theta + y^{(i)}x^{(i)}$

 Set $\theta_0 = \theta_0 + y^{(i)}$

 changed = True

if not changed

break

Return θ, θ_0

[How many 0s?]

[i.e. True if either:

- A. point is not on the line & prediction is wrong
- B. point is on the line
- C. initial step]

What does an update do?

$$\begin{aligned} & y^{(i)} \left((\theta + y^{(i)}x^{(i)})^\top x^{(i)} + (\theta_0 + y^{(i)}) \right) \\ &= y^{(i)}(\theta^\top x^{(i)} + \theta_0) + (y^{(i)})^2(x^{(i)\top} x^{(i)} + 1) \end{aligned}$$

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

 Set $\theta = \theta + y^{(i)}x^{(i)}$

 Set $\theta_0 = \theta_0 + y^{(i)}$

 changed = True

if not changed

break

Return θ, θ_0

[How many 0s?]

[i.e. True if either:

- A. point is not on the line & prediction is wrong
- B. point is on the line
- C. initial step]

What does an update do?

$$\begin{aligned} & y^{(i)} \left((\theta + y^{(i)}x^{(i)})^\top x^{(i)} + (\theta_0 + y^{(i)}) \right) \\ &= y^{(i)}(\theta^\top x^{(i)} + \theta_0) + (y^{(i)})^2(x^{(i)\top} x^{(i)} + 1) \end{aligned}$$

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

 Set $\theta = \theta + y^{(i)}x^{(i)}$

 Set $\theta_0 = \theta_0 + y^{(i)}$

 changed = True

if not changed

break

Return θ, θ_0

[How many 0s?]

[i.e. True if either:

- A. point is not on the line & prediction is wrong
- B. point is on the line
- C. initial step]

What does an update do?

$$\begin{aligned} & y^{(i)} \left((\theta + y^{(i)}x^{(i)})^\top x^{(i)} + (\theta_0 + y^{(i)}) \right) \\ &= y^{(i)}(\theta^\top x^{(i)} + \theta_0) + (y^{(i)})^2(x^{(i)\top} x^{(i)} + 1) \end{aligned}$$

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

 Set $\theta = \theta + y^{(i)}x^{(i)}$

 Set $\theta_0 = \theta_0 + y^{(i)}$

 changed = True

if not changed

break

Return θ, θ_0

[How many 0s?]

[i.e. True if either:

- A. point is not on the line & prediction is wrong
- B. point is on the line
- C. initial step]

What does an update do?

$$\begin{aligned} & y^{(i)} \left((\theta + y^{(i)}x^{(i)})^\top x^{(i)} + (\theta_0 + y^{(i)}) \right) \\ &= y^{(i)}(\theta^\top x^{(i)} + \theta_0) + (y^{(i)})^2(x^{(i)\top} x^{(i)} + 1) \end{aligned}$$

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

 Set $\theta = \theta + y^{(i)}x^{(i)}$

 Set $\theta_0 = \theta_0 + y^{(i)}$

 changed = True

if not changed

break

Return θ, θ_0

[How many 0s?]

[i.e. True if either:

- A. point is not on the line & prediction is wrong
- B. point is on the line
- C. initial step]

What does an update do?

$$\begin{aligned} & y^{(i)} \left((\theta + y^{(i)}x^{(i)})^\top x^{(i)} + (\theta_0 + y^{(i)}) \right) \\ &= y^{(i)}(\theta^\top x^{(i)} + \theta_0) + (y^{(i)})^2(x^{(i)\top} x^{(i)} + 1) \end{aligned}$$

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

 Set $\theta = \theta + y^{(i)}x^{(i)}$

 Set $\theta_0 = \theta_0 + y^{(i)}$

 changed = True

if not changed

break

Return θ, θ_0

[How many 0s?]

[i.e. True if either:

- A. point is not on the line & prediction is wrong
- B. point is on the line
- C. initial step]

What does an update do?

$$y^{(i)} \left((\theta + y^{(i)}x^{(i)})^\top x^{(i)} + (\theta_0 + y^{(i)}) \right)$$
$$= y^{(i)}(\theta^\top x^{(i)} + \theta_0) + (y^{(i)})^2(x^{(i)\top} x^{(i)} + 1)$$

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

 Set $\theta = \theta + y^{(i)}x^{(i)}$

 Set $\theta_0 = \theta_0 + y^{(i)}$

 changed = True

if not changed

break

Return θ, θ_0

[How many 0s?]

[i.e. True if either:

- A. point is not on the line & prediction is wrong
- B. point is on the line
- C. initial step]

What does an update do?

$$\begin{aligned} & y^{(i)} \left((\theta + y^{(i)}x^{(i)})^\top x^{(i)} + (\theta_0 + y^{(i)}) \right) \\ &= y^{(i)}(\theta^\top x^{(i)} + \theta_0) + (y^{(i)})^2(x^{(i)\top} x^{(i)} + 1) \end{aligned}$$

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

 Set $\theta = \theta + y^{(i)}x^{(i)}$

 Set $\theta_0 = \theta_0 + y^{(i)}$

 changed = True

if not changed

break

Return θ, θ_0

[How many 0s?]

[i.e. True if either:

- A. point is not on the line & prediction is wrong
- B. point is on the line
- C. initial step]

What does an update do?

$$\begin{aligned} & y^{(i)} \left((\theta + y^{(i)}x^{(i)})^\top x^{(i)} + (\theta_0 + y^{(i)}) \right) \\ &= y^{(i)}(\theta^\top x^{(i)} + \theta_0) + (y^{(i)})^2(x^{(i)\top} x^{(i)} + 1) \end{aligned}$$

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

 Set $\theta = \theta + y^{(i)}x^{(i)}$

 Set $\theta_0 = \theta_0 + y^{(i)}$

 changed = True

if not changed

break

Return θ, θ_0

[How many 0s?]

[i.e. True if either:

- A. point is not on the line & prediction is wrong
- B. point is on the line
- C. initial step]

What does an update do?

$$\begin{aligned} & y^{(i)} \left((\theta + y^{(i)}x^{(i)})^\top x^{(i)} + (\theta_0 + y^{(i)}) \right) \\ &= y^{(i)}(\theta^\top x^{(i)} + \theta_0) + (y^{(i)})^2(x^{(i)\top} x^{(i)} + 1) \end{aligned}$$

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

 Set $\theta = \theta + y^{(i)}x^{(i)}$

 Set $\theta_0 = \theta_0 + y^{(i)}$

 changed = True

if not changed

break

Return θ, θ_0

[How many 0s?]

[i.e. True if either:

- A. point is not on the line & prediction is wrong
- B. point is on the line
- C. initial step]

What does an update do?

$$\begin{aligned} & y^{(i)} \left((\theta + y^{(i)}x^{(i)})^\top x^{(i)} + (\theta_0 + y^{(i)}) \right) \\ &= y^{(i)}(\theta^\top x^{(i)} + \theta_0) + (y^{(i)})^2(x^{(i)\top} x^{(i)} + 1) \end{aligned}$$

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

 Set $\theta = \theta + y^{(i)}x^{(i)}$

 Set $\theta_0 = \theta_0 + y^{(i)}$

 changed = True

if not changed

break

Return θ, θ_0

[How many 0s?]

[i.e. True if either:

- A. point is not on the line & prediction is wrong
- B. point is on the line
- C. initial step]

What does an update do?

$$\begin{aligned} & y^{(i)} \left((\theta + y^{(i)}x^{(i)})^\top x^{(i)} + (\theta_0 + y^{(i)}) \right) \\ &= y^{(i)}(\theta^\top x^{(i)} + \theta_0) + (y^{(i)})^2(x^{(i)\top} x^{(i)} + 1) \end{aligned}$$

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

 Set $\theta = \theta + y^{(i)}x^{(i)}$

 Set $\theta_0 = \theta_0 + y^{(i)}$

 changed = True

if not changed

break

Return θ, θ_0

[How many 0s?]

[i.e. True if either:

- A. point is not on the line & prediction is wrong
- B. point is on the line
- C. initial step]

What does an update do?

$$\begin{aligned} & y^{(i)} \left((\theta + y^{(i)}x^{(i)})^\top x^{(i)} + (\theta_0 + y^{(i)}) \right) \\ &= y^{(i)}(\theta^\top x^{(i)} + \theta_0) + (y^{(i)})^2(x^{(i)\top} x^{(i)} + 1) \end{aligned}$$

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

 Set $\theta = \theta + y^{(i)}x^{(i)}$

 Set $\theta_0 = \theta_0 + y^{(i)}$

 changed = True

if not changed

break

Return θ, θ_0

[How many 0s?]

[i.e. True if either:

- A. point is not on the line & prediction is wrong
- B. point is on the line
- C. initial step]

What does an update do?

$$\begin{aligned} & y^{(i)} \left((\theta + y^{(i)}x^{(i)})^\top x^{(i)} + (\theta_0 + y^{(i)}) \right) \\ &= y^{(i)}(\theta^\top x^{(i)} + \theta_0) + (y^{(i)})^2(x^{(i)\top} x^{(i)} + 1) \end{aligned}$$

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

 Set $\theta = \theta + y^{(i)}x^{(i)}$

 Set $\theta_0 = \theta_0 + y^{(i)}$

 changed = True

if not changed

break

Return θ, θ_0

[How many 0s?]

[i.e. True if either:

- A. point is not on the line & prediction is wrong
- B. point is on the line
- C. initial step]

What does an update do?

$$\begin{aligned} & y^{(i)} \left((\theta + y^{(i)}x^{(i)})^\top x^{(i)} + (\theta_0 + y^{(i)}) \right) \\ &= y^{(i)}(\theta^\top x^{(i)} + \theta_0) + (y^{(i)})^2(x^{(i)\top} x^{(i)} + 1) \\ &= y^{(i)}(\theta^\top x^{(i)} + \theta_0) + (\|x^{(i)}\|^2 + 1) \end{aligned}$$

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

 Set $\theta = \theta + y^{(i)}x^{(i)}$

 Set $\theta_0 = \theta_0 + y^{(i)}$

 changed = True

if not changed

break

Return θ, θ_0

[How many 0s?]

[i.e. True if either:

- A. point is not on the line & prediction is wrong
- B. point is on the line
- C. initial step]

What does an update do?

$$\begin{aligned} & y^{(i)} \left((\theta + y^{(i)}x^{(i)})^\top x^{(i)} + (\theta_0 + y^{(i)}) \right) \\ &= y^{(i)}(\theta^\top x^{(i)} + \theta_0) + (y^{(i)})^2(x^{(i)\top} x^{(i)} + 1) \\ &= y^{(i)}(\theta^\top x^{(i)} + \theta_0) + (\|x^{(i)}\|^2 + 1) \end{aligned}$$

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

 Set $\theta = \theta + y^{(i)}x^{(i)}$

 Set $\theta_0 = \theta_0 + y^{(i)}$

 changed = True

if not changed

break

Return θ, θ_0

[How many 0s?]

[i.e. True if either:

- A. point is not on the line & prediction is wrong
- B. point is on the line
- C. initial step]

What does an update do?

$$\begin{aligned} & y^{(i)} \left((\theta + y^{(i)}x^{(i)})^\top x^{(i)} + (\theta_0 + y^{(i)}) \right) \\ &= y^{(i)}(\theta^\top x^{(i)} + \theta_0) + (y^{(i)})^2(x^{(i)\top} x^{(i)} + 1) \\ &= y^{(i)}(\theta^\top x^{(i)} + \theta_0) + (\|x^{(i)}\|^2 + 1) \end{aligned}$$

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

 Set $\theta = \theta + y^{(i)}x^{(i)}$

 Set $\theta_0 = \theta_0 + y^{(i)}$

 changed = True

if not changed

break

Return θ, θ_0

[How many 0s?]

[i.e. True if either:

- A. point is not on the line & prediction is wrong
- B. point is on the line
- C. initial step]

What does an update do?

$$\begin{aligned} & y^{(i)} \left((\theta + y^{(i)}x^{(i)})^\top x^{(i)} + (\theta_0 + y^{(i)}) \right) \\ &= y^{(i)}(\theta^\top x^{(i)} + \theta_0) + (y^{(i)})^2(x^{(i)\top} x^{(i)} + 1) \\ &= y^{(i)}(\theta^\top x^{(i)} + \theta_0) + (\|x^{(i)}\|^2 + 1) \end{aligned}$$

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

 Set $\theta = \theta + y^{(i)}x^{(i)}$

 Set $\theta_0 = \theta_0 + y^{(i)}$

 changed = True

if not changed

break

Return θ, θ_0

[How many 0s?]

[i.e. True if either:

- A. point is not on the line & prediction is wrong
- B. point is on the line
- C. initial step]

What does an update do?

$$\begin{aligned} & y^{(i)} \left((\theta + y^{(i)}x^{(i)})^\top x^{(i)} + (\theta_0 + y^{(i)}) \right) \\ &= y^{(i)}(\theta^\top x^{(i)} + \theta_0) + (y^{(i)})^2(x^{(i)\top} x^{(i)} + 1) \\ &= y^{(i)}(\theta^\top x^{(i)} + \theta_0) + (\|x^{(i)}\|^2 + 1) \end{aligned}$$

Perceptron Algorithm

Perceptron($\mathcal{D}_n ; \tau$)

 Initialize $\theta = [0 \ 0 \ \dots \ 0]^\top$

 Initialize $\theta_0 = 0$

for $t = 1$ to τ

 changed = False

for $i = 1$ to n

if $y^{(i)}(\theta^\top x^{(i)} + \theta_0) \leq 0$

 Set $\theta = \theta + y^{(i)}x^{(i)}$

 Set $\theta_0 = \theta_0 + y^{(i)}$

 changed = True

if not changed

break

Return θ, θ_0

[demo]

[How many 0s?]

[i.e. True if either:

- A. point is not on the line & prediction is wrong
- B. point is on the line
- C. initial step]

What does an update do?

$$y^{(i)} \left((\theta + y^{(i)}x^{(i)})^\top x^{(i)} + (\theta_0 + y^{(i)}) \right)$$

$$= y^{(i)}(\theta^\top x^{(i)} + \theta_0) + (y^{(i)})^2(x^{(i)\top} x^{(i)} + 1)$$

$$= y^{(i)}(\theta^\top x^{(i)} + \theta_0) + (\|x^{(i)}\|^2 + 1)$$

Let's Talk About Classifier Quality

Let's Talk About Classifier Quality

- *Definition:* A training set \mathcal{D}_n is **linearly separable** if there exist θ, θ_0 such that, for every point index $i \in \{1, \dots, n\}$, we have

$$y^{(i)}(\theta^\top x^{(i)} + \theta_0) > 0$$

Let's Talk About Classifier Quality

- *Definition:* A training set \mathcal{D}_n is **linearly separable** if there exist θ, θ_0 such that, for every point index $i \in \{1, \dots, n\}$, we have

$$y^{(i)}(\theta^\top x^{(i)} + \theta_0) > 0$$

Let's Talk About Classifier Quality

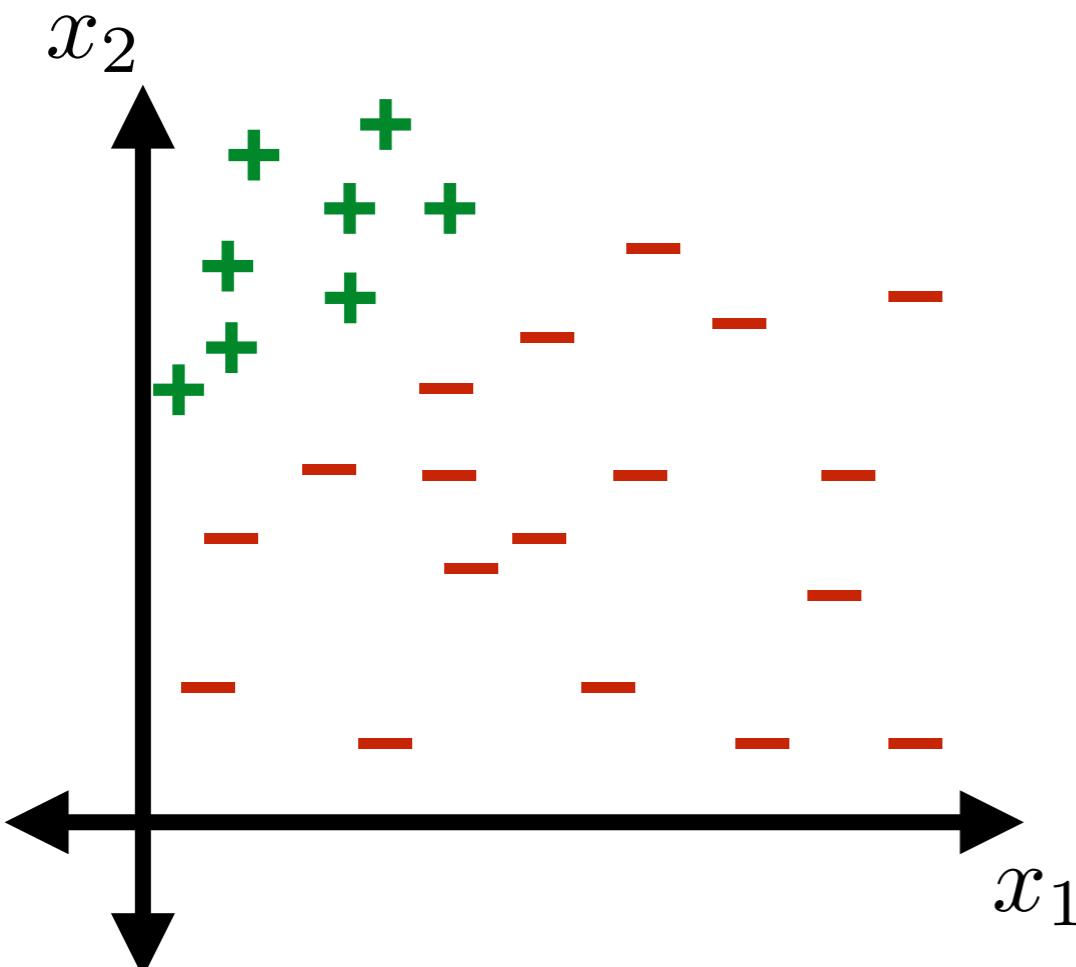
- *Definition:* A training set \mathcal{D}_n is **linearly separable** if there exist θ, θ_0 such that, for every point index $i \in \{1, \dots, n\}$, we have

$$y^{(i)}(\theta^\top x^{(i)} + \theta_0) > 0$$

Let's Talk About Classifier Quality

- *Definition:* A training set \mathcal{D}_n is **linearly separable** if there exist θ, θ_0 such that, for every point index $i \in \{1, \dots, n\}$, we have

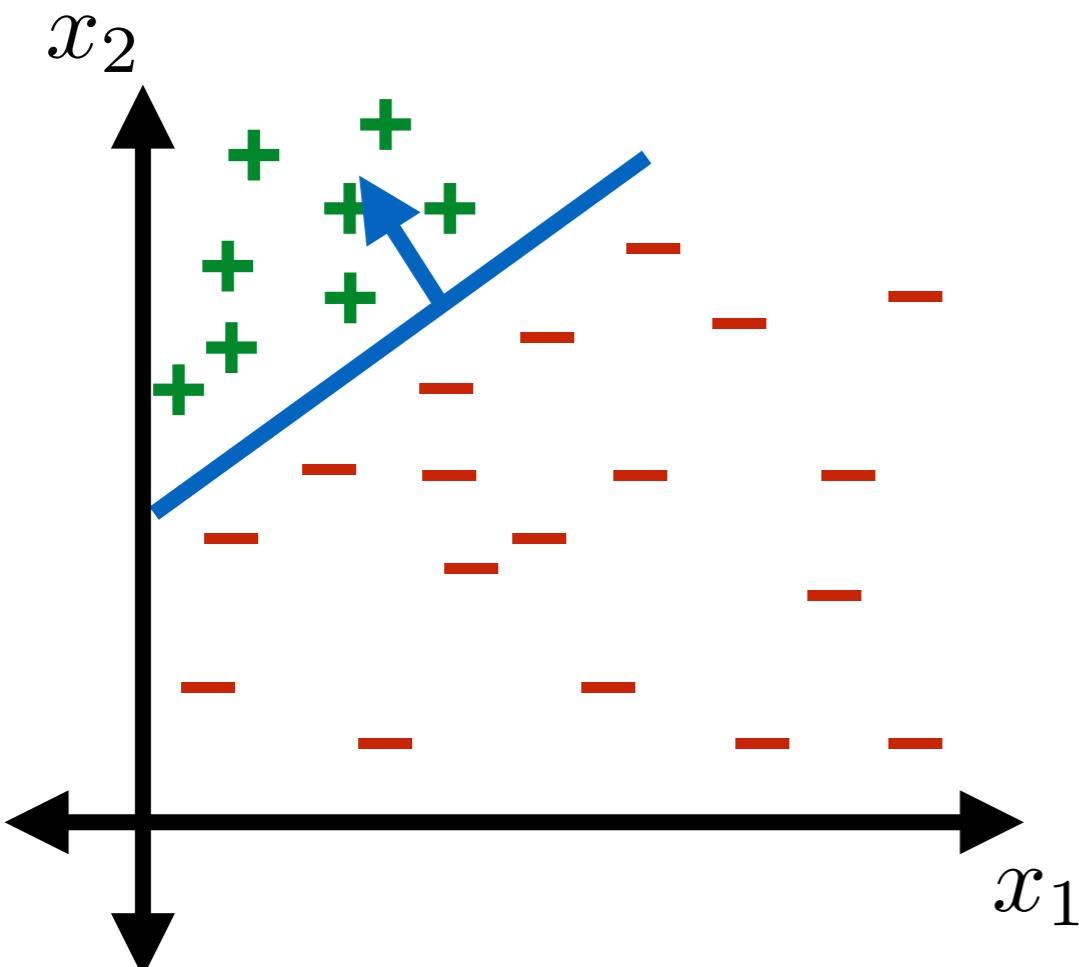
$$y^{(i)}(\theta^\top x^{(i)} + \theta_0) > 0$$



Let's Talk About Classifier Quality

- *Definition:* A training set \mathcal{D}_n is **linearly separable** if there exist θ, θ_0 such that, for every point index $i \in \{1, \dots, n\}$, we have

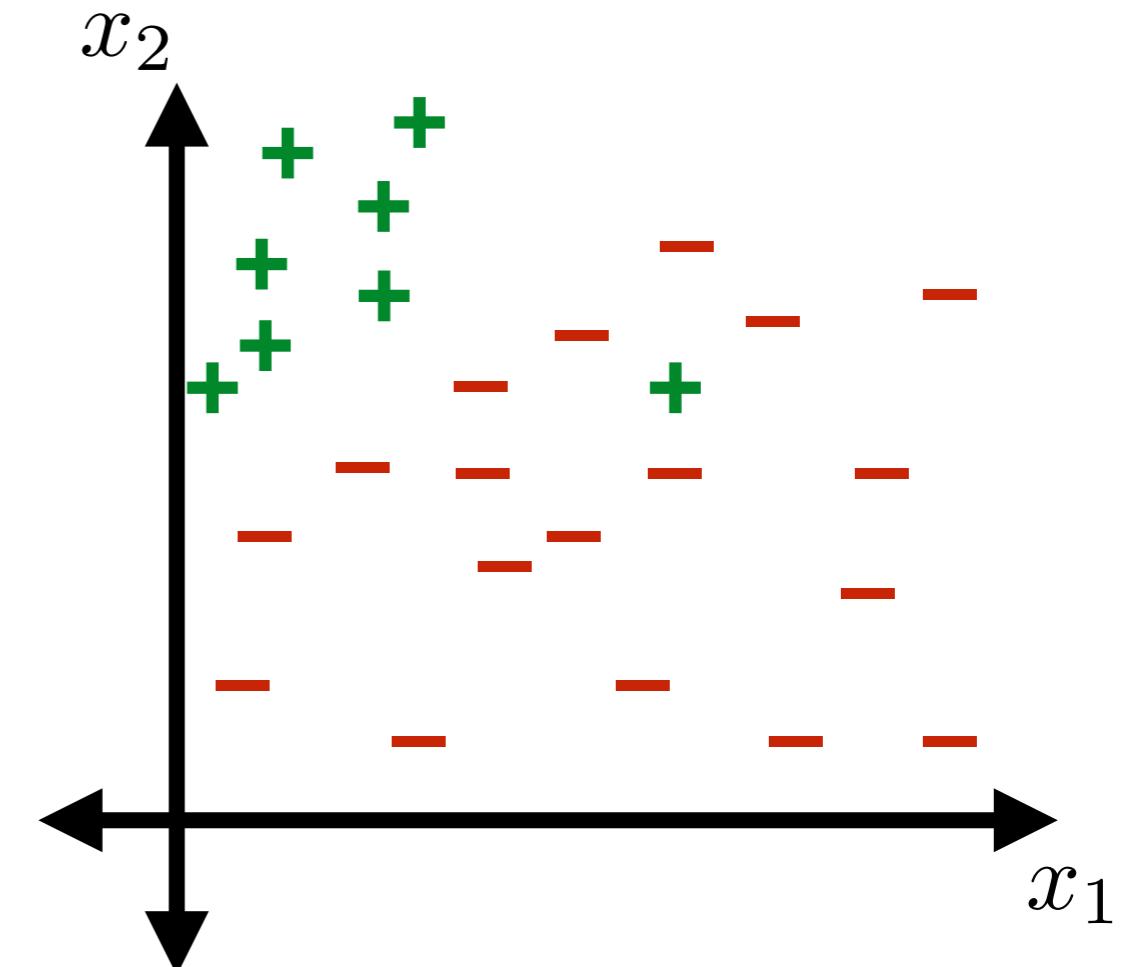
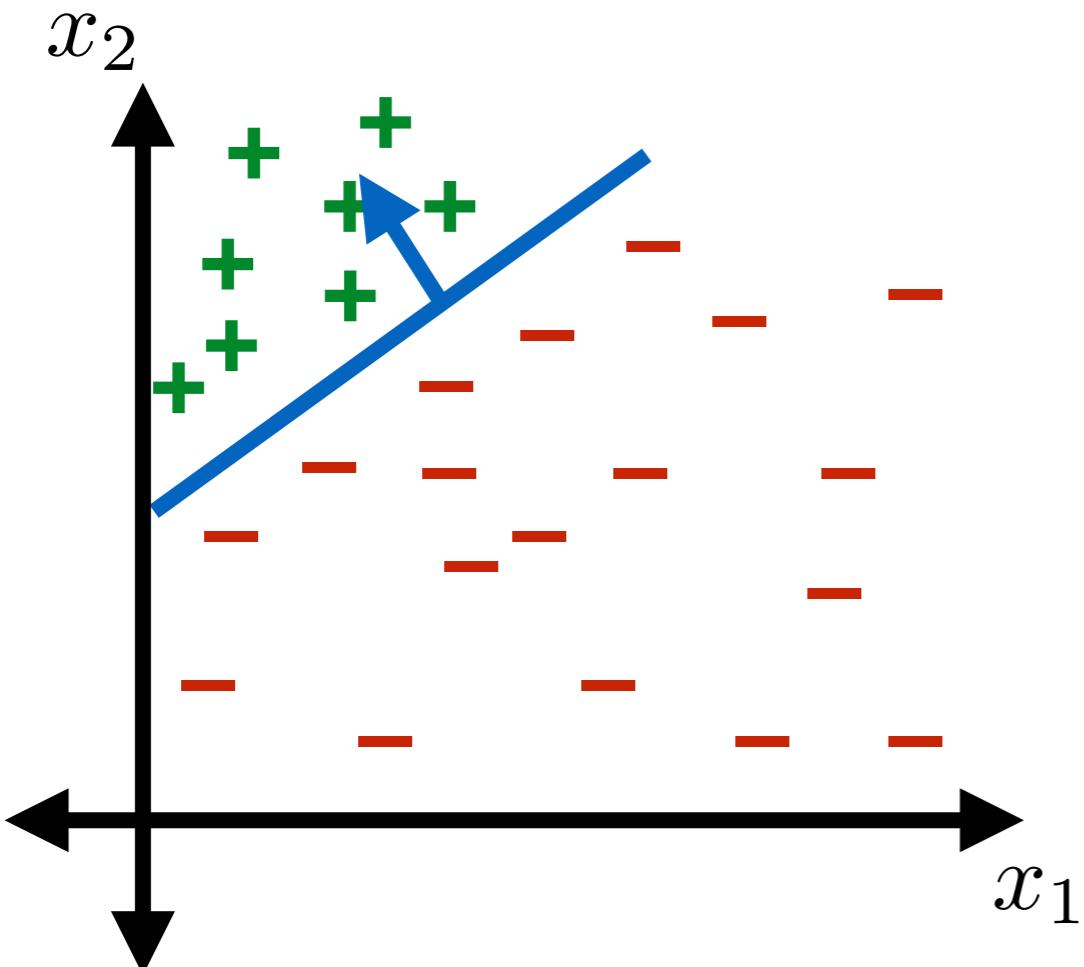
$$y^{(i)}(\theta^\top x^{(i)} + \theta_0) > 0$$



Let's Talk About Classifier Quality

- *Definition:* A training set \mathcal{D}_n is **linearly separable** if there exist θ, θ_0 such that, for every point index $i \in \{1, \dots, n\}$, we have

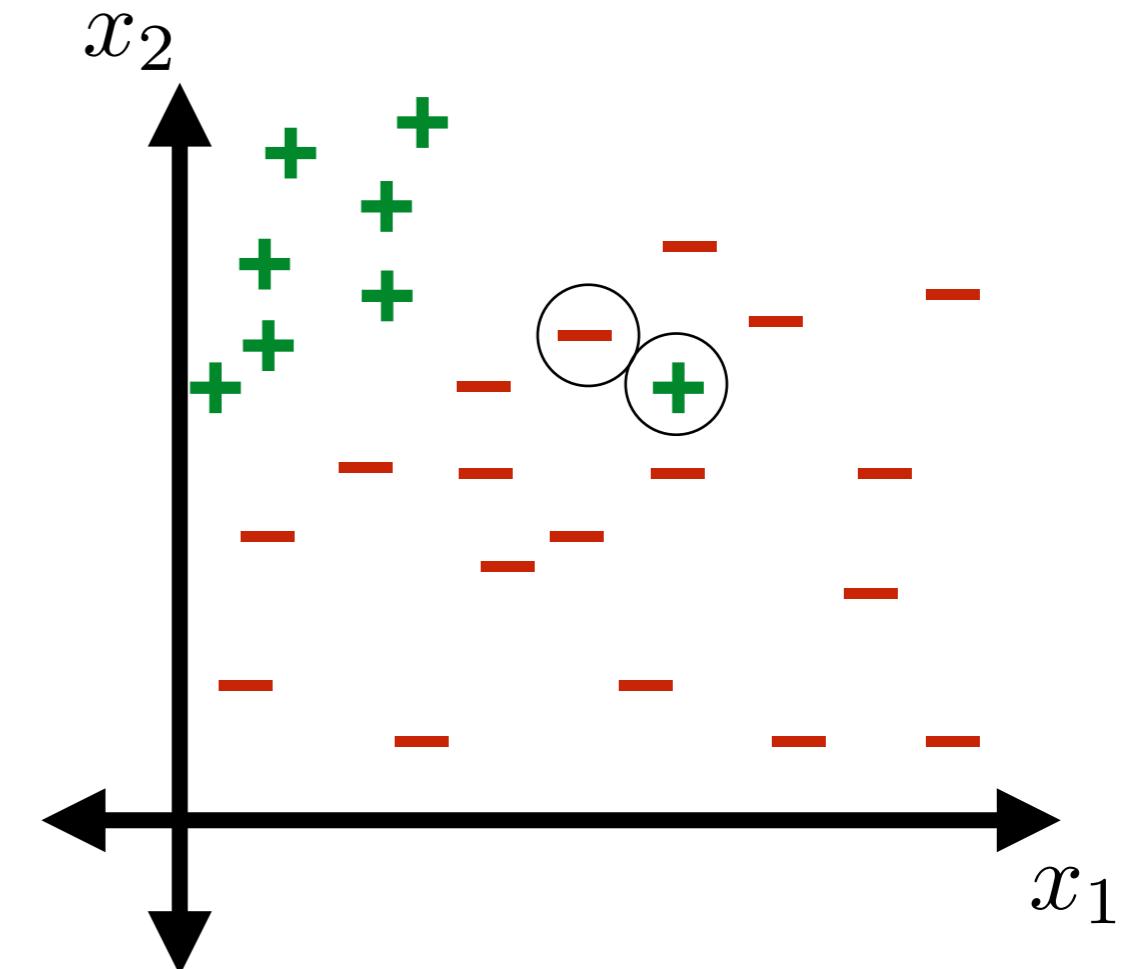
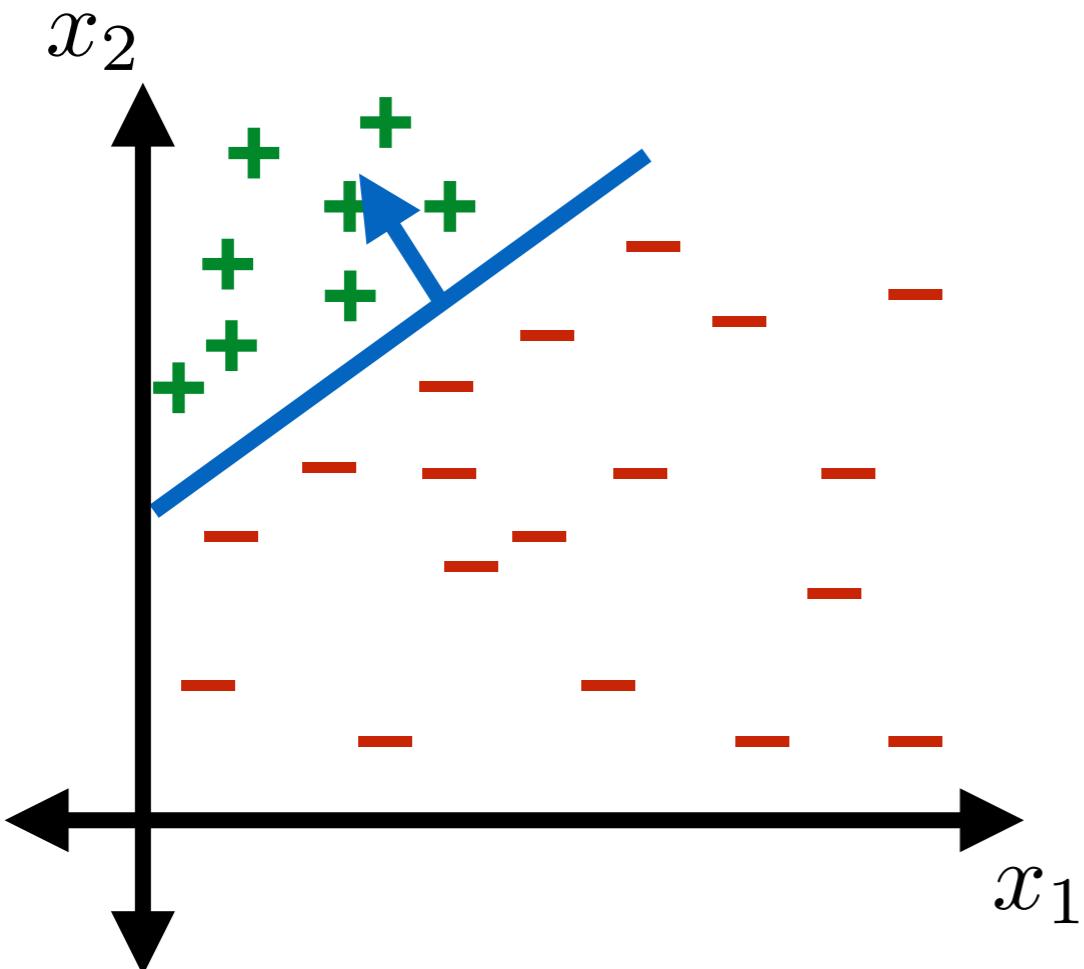
$$y^{(i)}(\theta^\top x^{(i)} + \theta_0) > 0$$



Let's Talk About Classifier Quality

- *Definition:* A training set \mathcal{D}_n is **linearly separable** if there exist θ, θ_0 such that, for every point index $i \in \{1, \dots, n\}$, we have

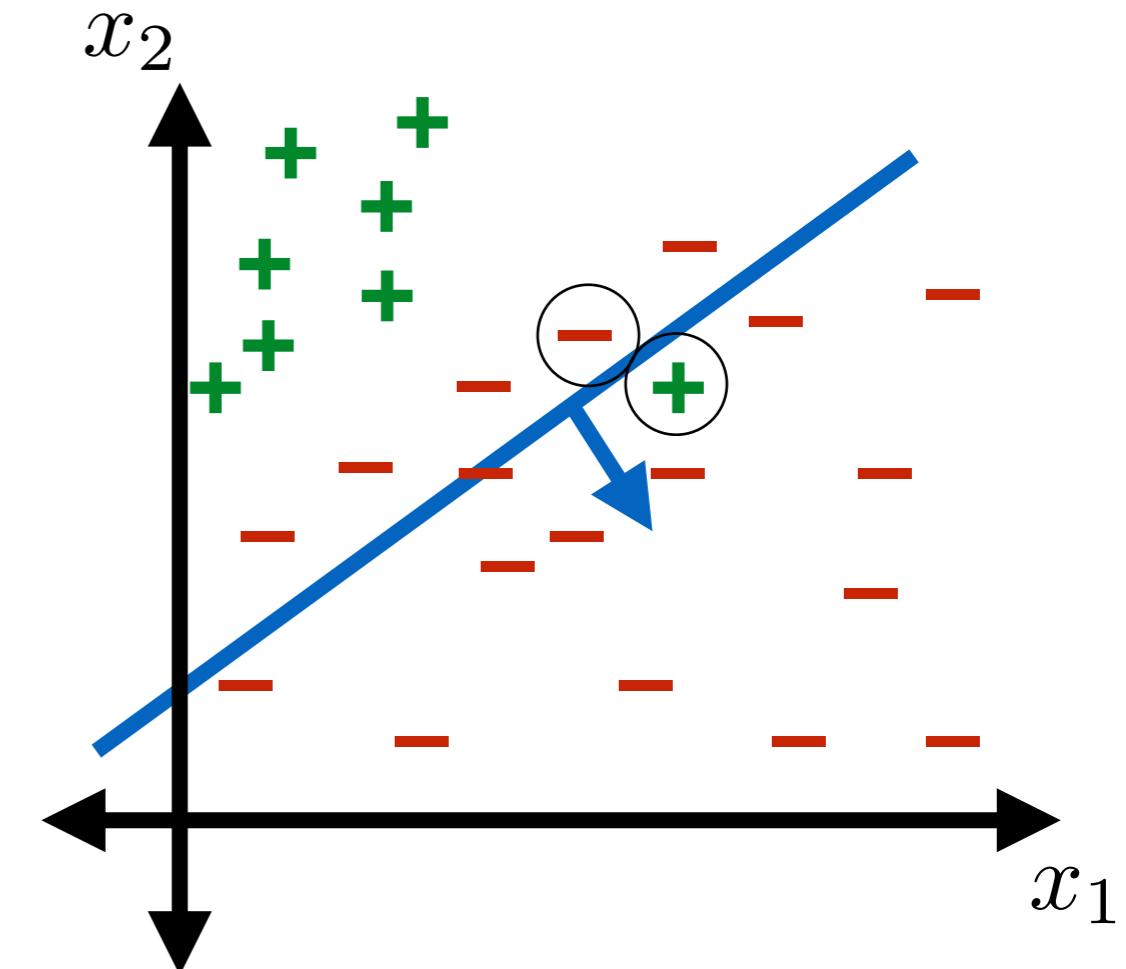
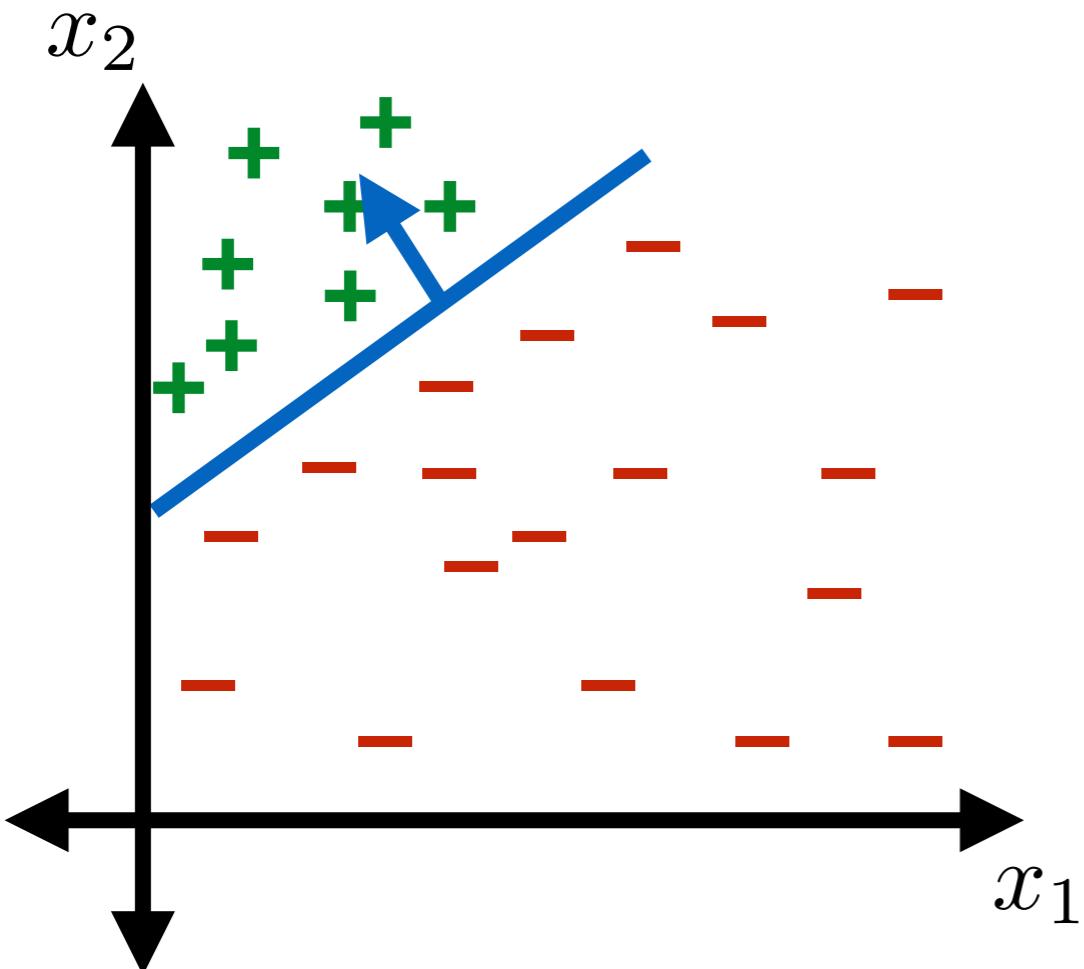
$$y^{(i)}(\theta^\top x^{(i)} + \theta_0) > 0$$



Let's Talk About Classifier Quality

- *Definition:* A training set \mathcal{D}_n is **linearly separable** if there exist θ, θ_0 such that, for every point index $i \in \{1, \dots, n\}$, we have

$$y^{(i)}(\theta^\top x^{(i)} + \theta_0) > 0$$

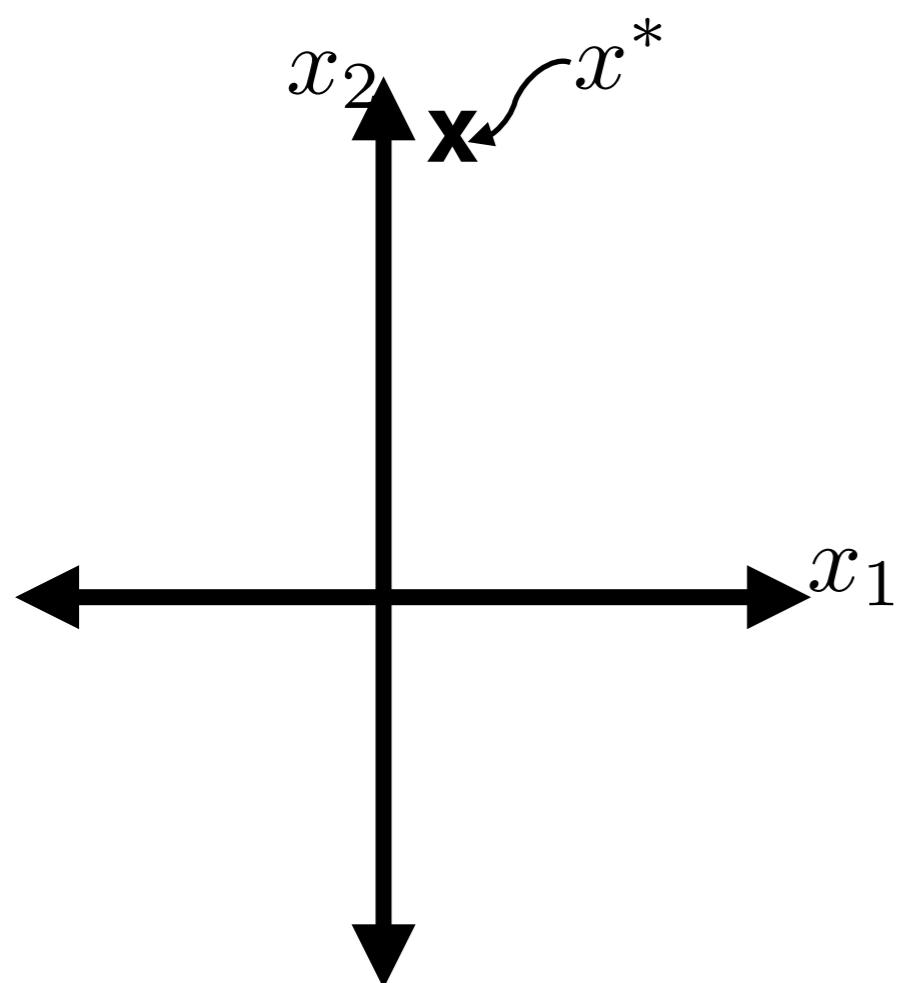


Let's Talk About Classifier Quality

Math facts!

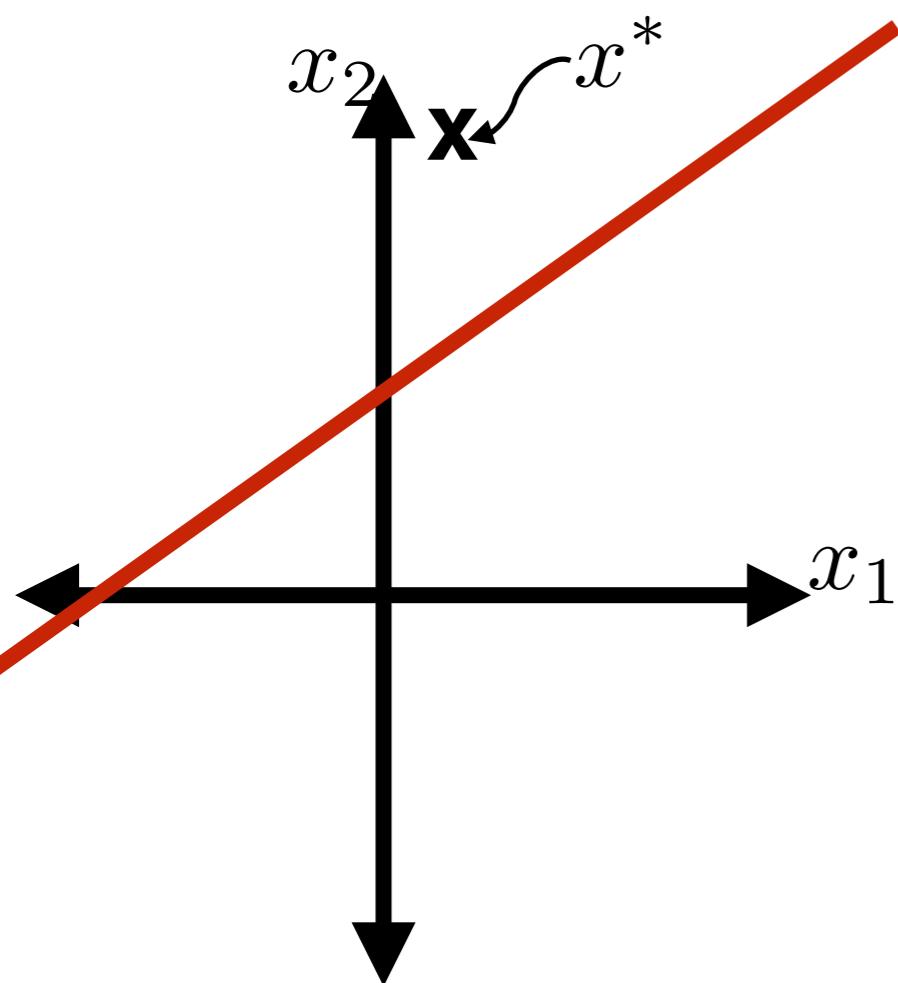
Let's Talk About Classifier Quality

Math facts!



Let's Talk About Classifier Quality

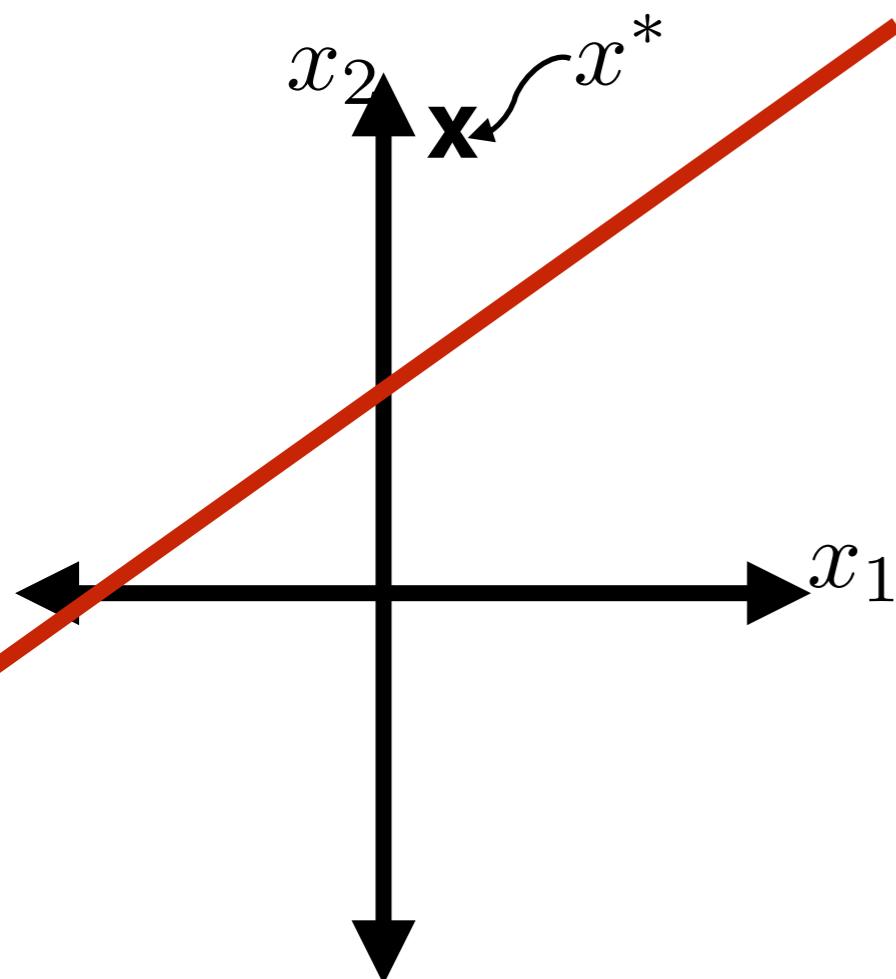
Math facts!



Let's Talk About Classifier Quality

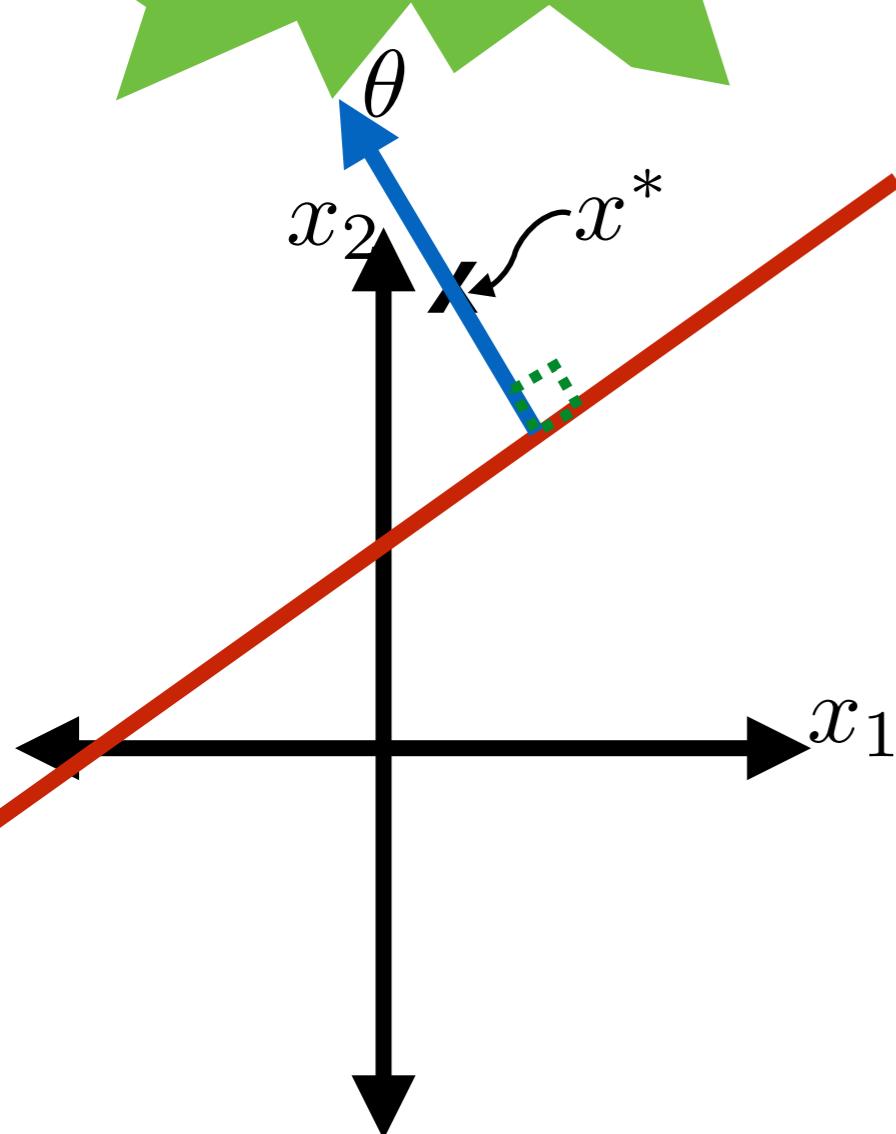
Math facts!

- The signed distance from a hyperplane defined by θ, θ_0 to a point x^* is:



Let's Talk About Classifier Quality

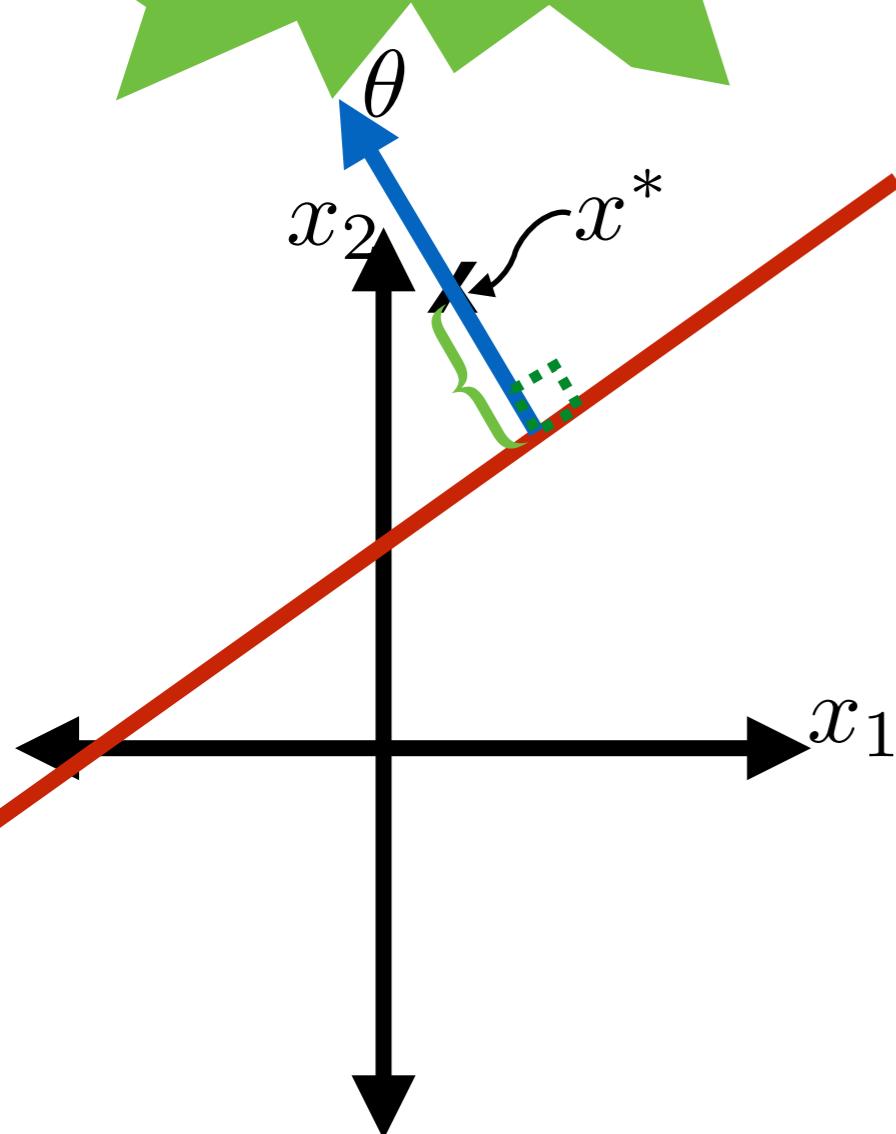
Math facts!



- The signed distance from a hyperplane defined by θ, θ_0 to a point x^* is:

Let's Talk About Classifier Quality

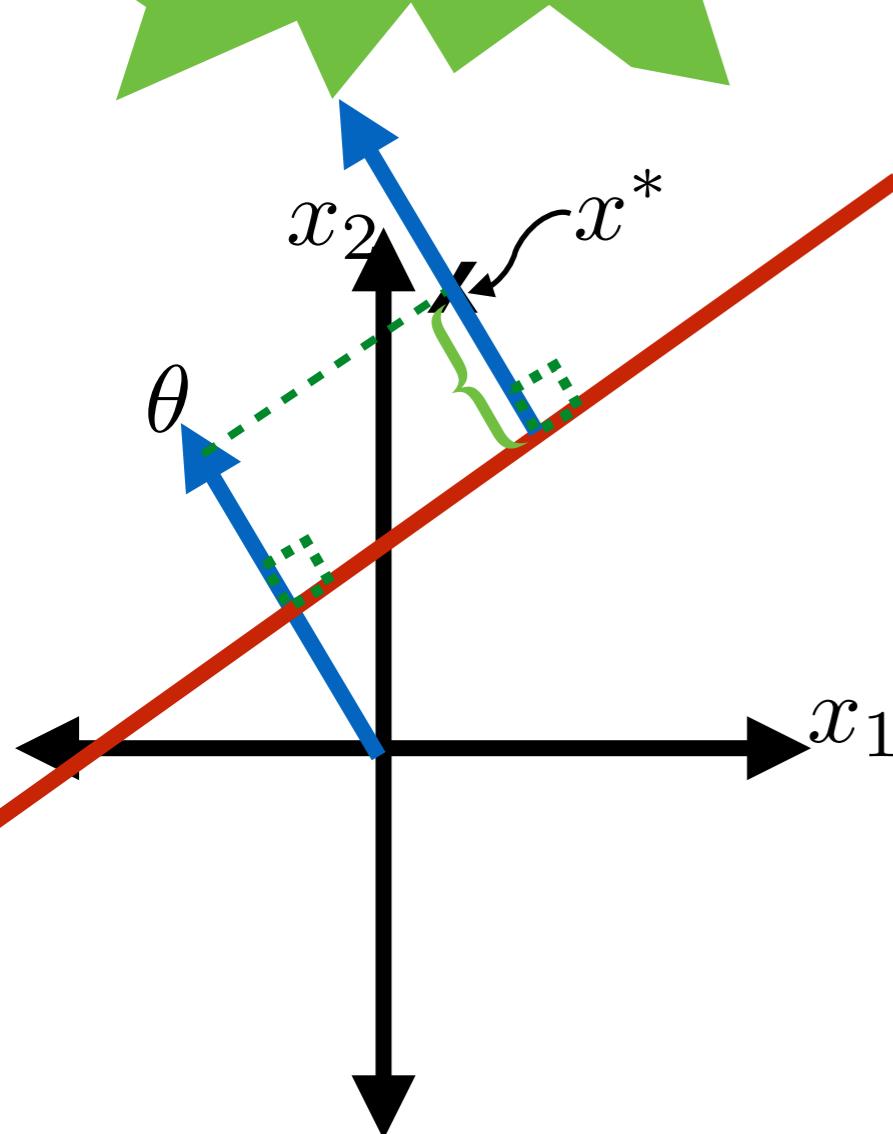
Math facts!



- The signed distance from a hyperplane defined by θ, θ_0 to a point x^* is:

Let's Talk About Classifier Quality

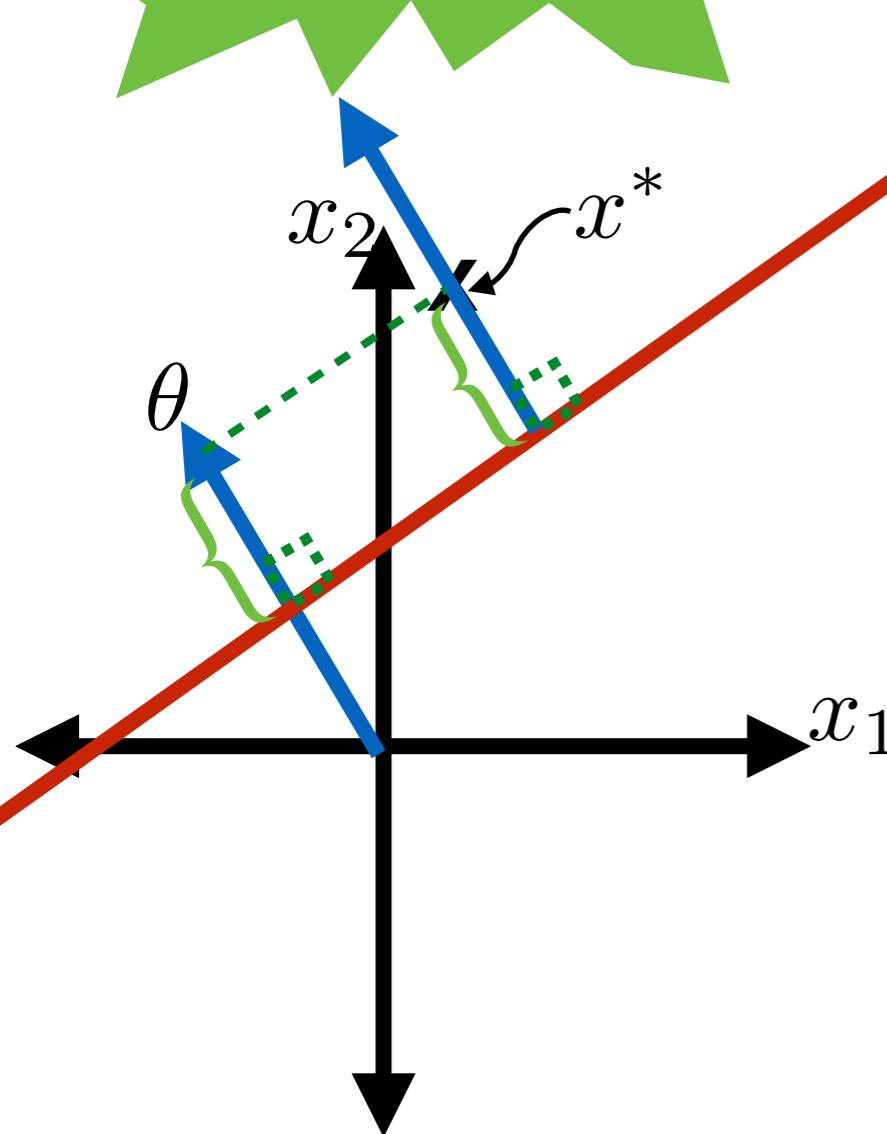
Math facts!



- The signed distance from a hyperplane defined by θ, θ_0 to a point x^* is:

Let's Talk About Classifier Quality

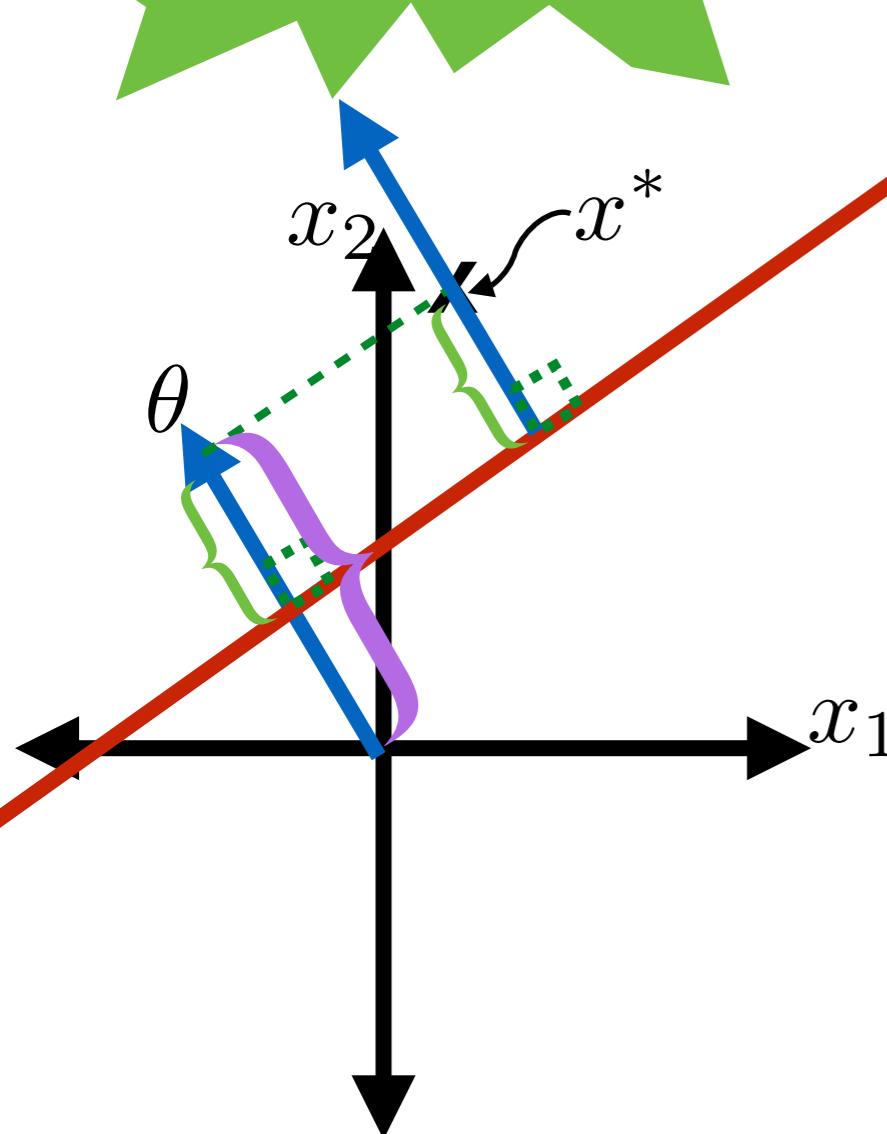
Math facts!



- The signed distance from a hyperplane defined by θ, θ_0 to a point x^* is:

Let's Talk About Classifier Quality

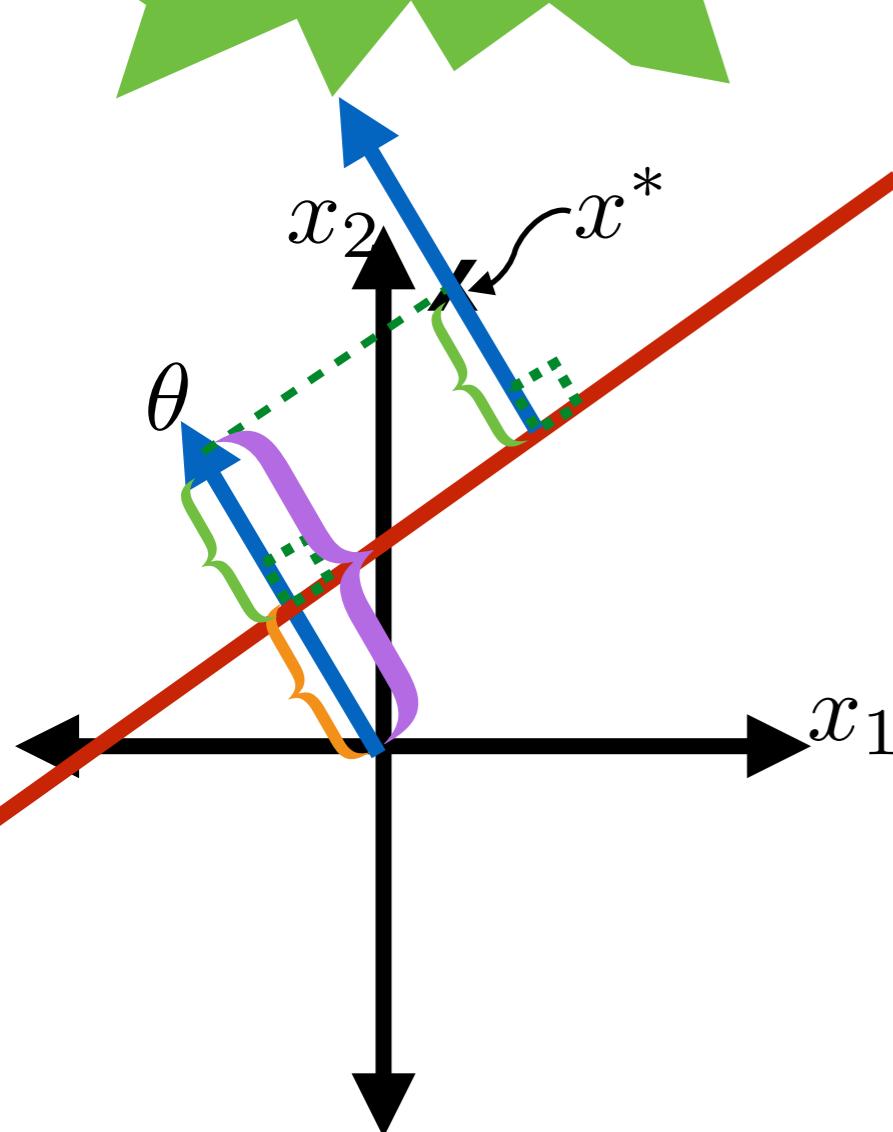
Math facts!



- The signed distance from a hyperplane defined by θ, θ_0 to a point x^* is:

Let's Talk About Classifier Quality

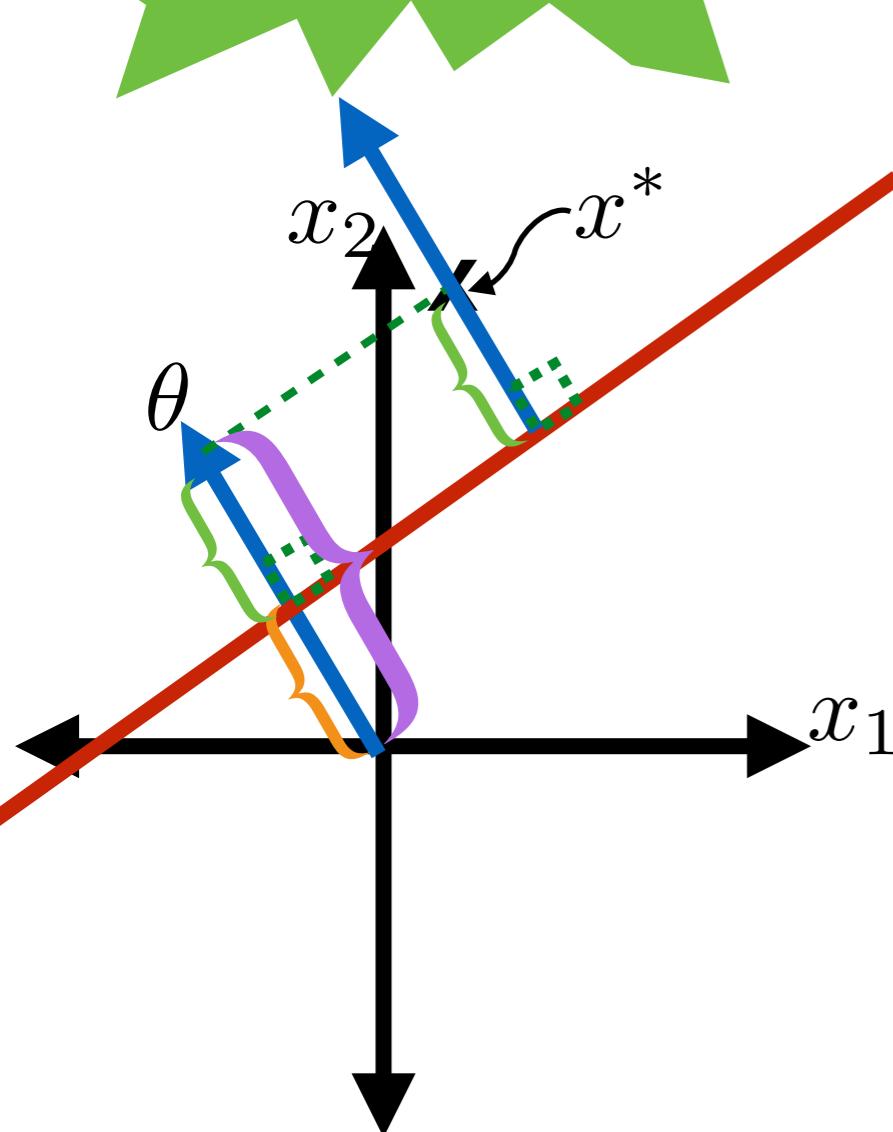
Math facts!



- The signed distance from a hyperplane defined by θ, θ_0 to a point x^* is:

Let's Talk About Classifier Quality

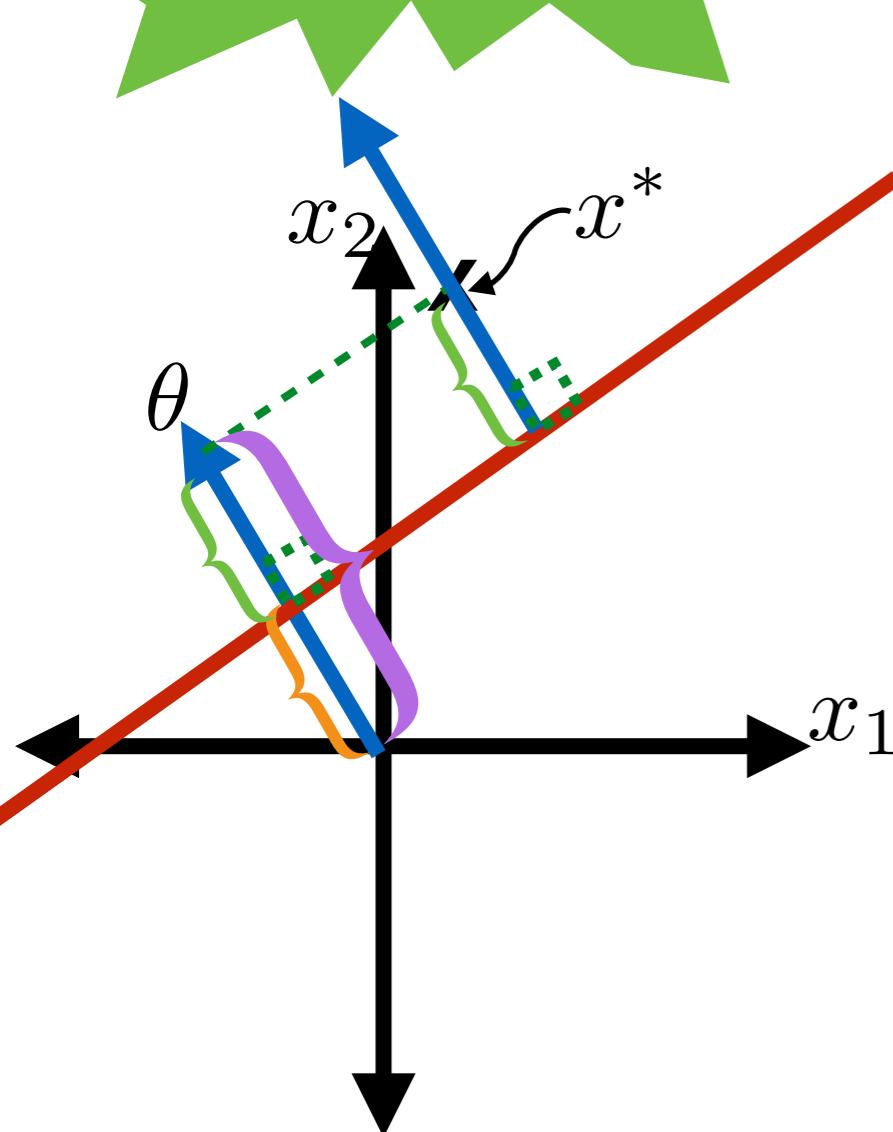
Math facts!



- The signed distance from a hyperplane defined by θ, θ_0 to a point x^* is:
= projection of x^* on θ

Let's Talk About Classifier Quality

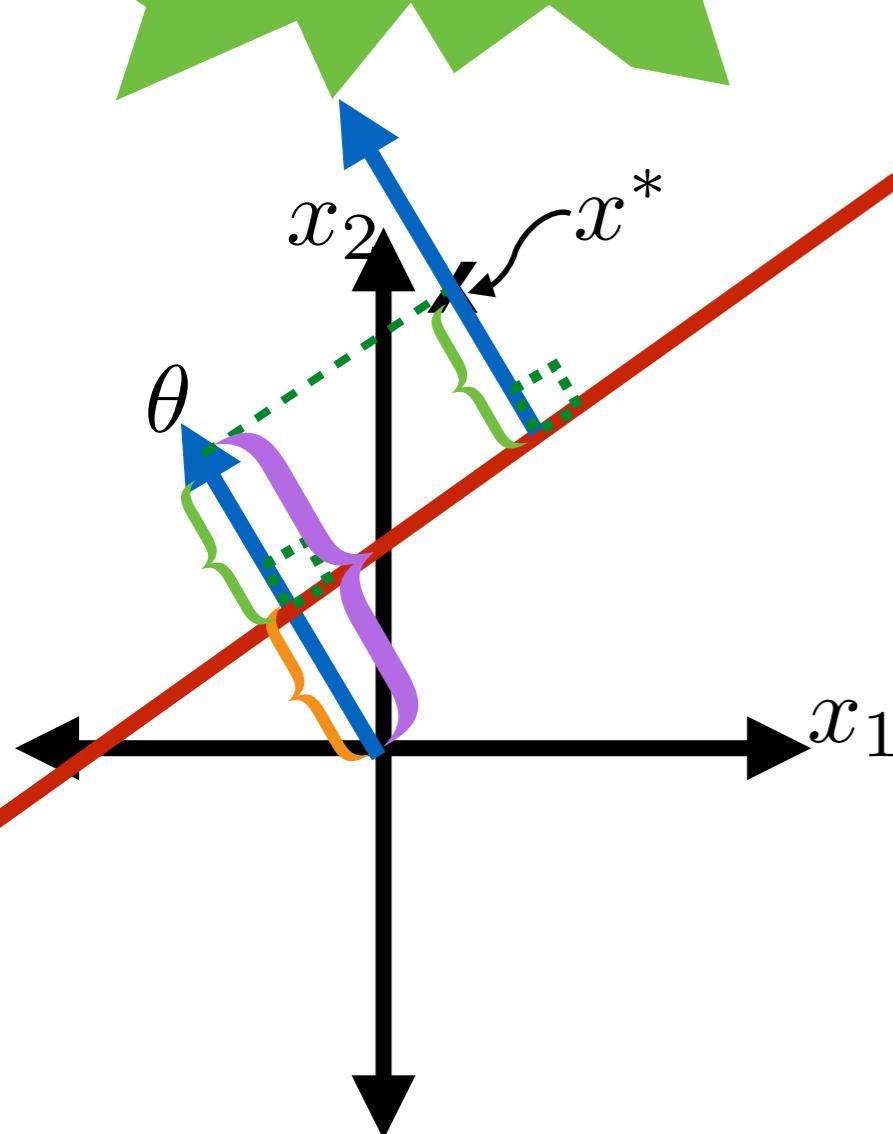
Math facts!



- The signed distance from a hyperplane defined by θ, θ_0 to a point x^* is:
 - = projection of x^* on θ
 - signed distance of line to origin

Let's Talk About Classifier Quality

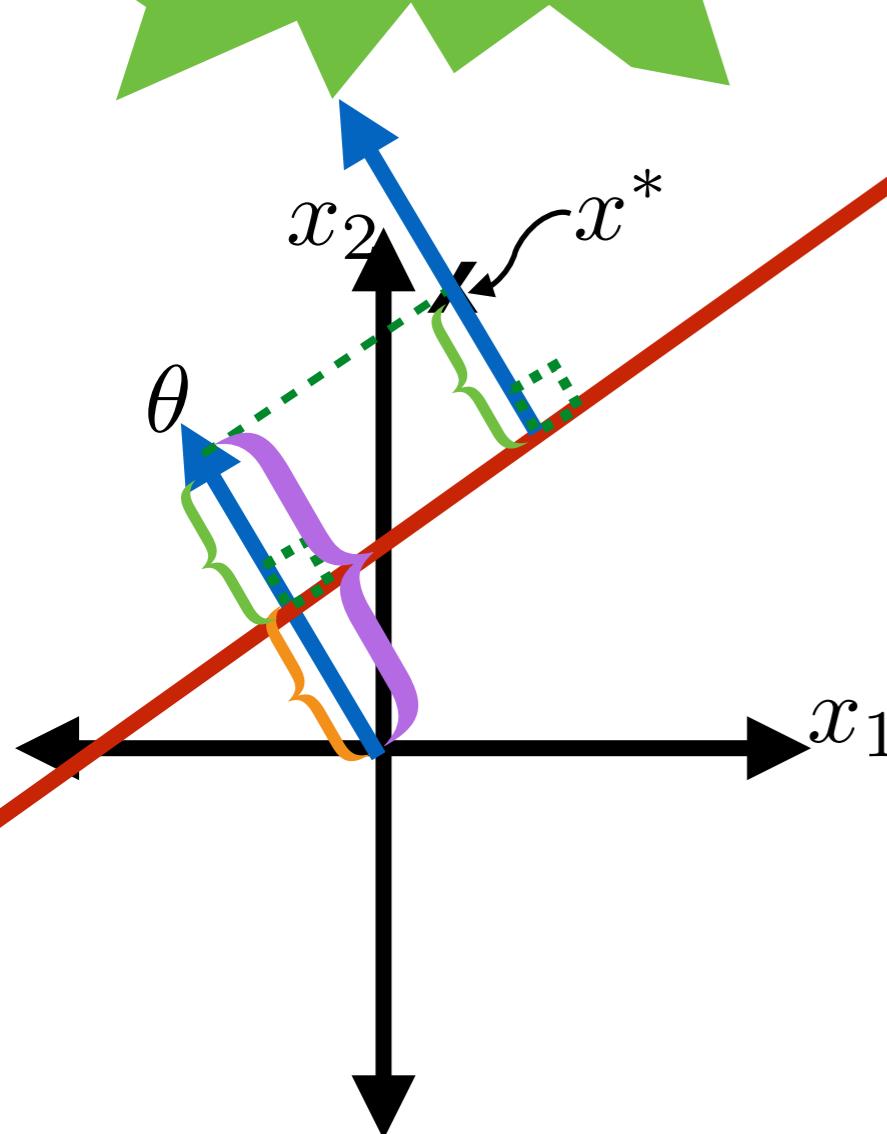
Math facts!



- The signed distance from a hyperplane defined by θ, θ_0 to a point x^* is:
 - = projection of x^* on θ
 - signed distance of line to origin
- $$= \frac{\theta^\top x^*}{\|\theta\|}$$

Let's Talk About Classifier Quality

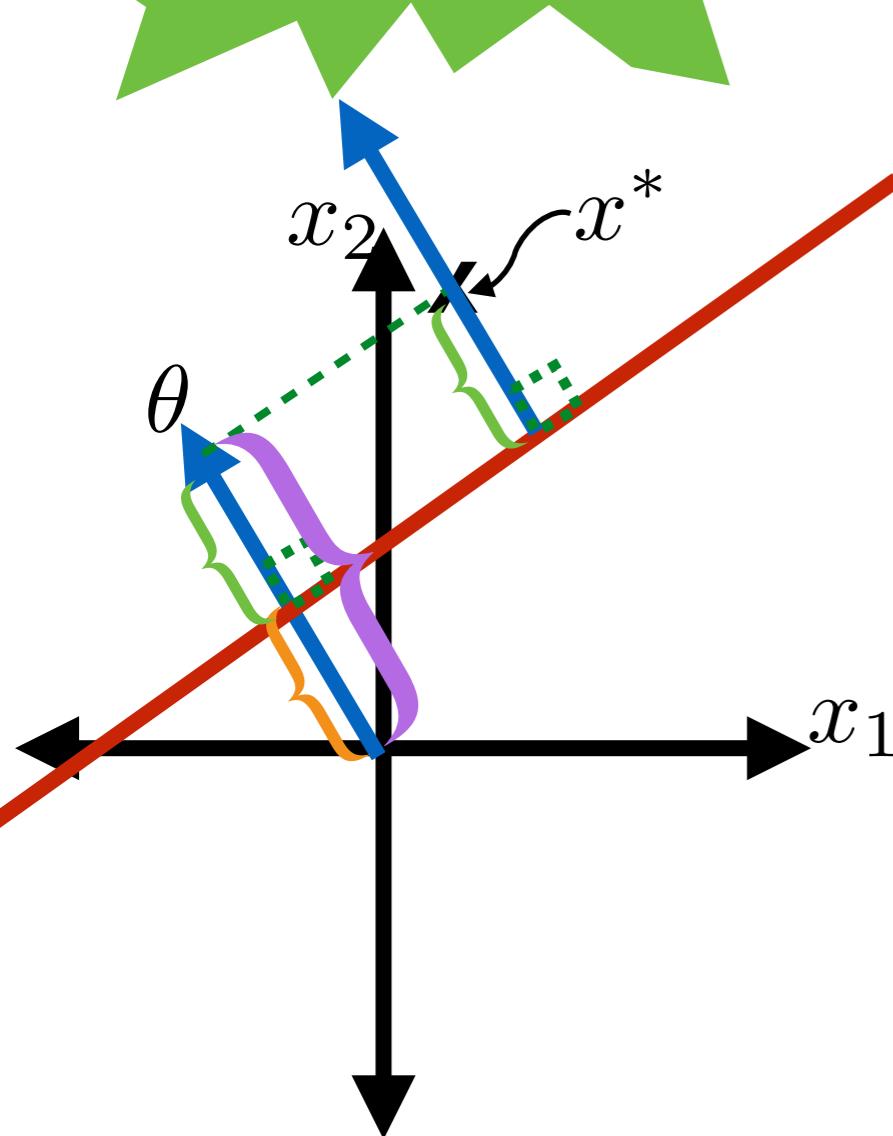
Math facts!



- The signed distance from a hyperplane defined by θ, θ_0 to a point x^* is:
 - = projection of x^* on θ
 - signed distance of line to origin
- $$= \frac{\theta^\top x^*}{\|\theta\|} - \frac{-\theta_0}{\|\theta\|}$$

Let's Talk About Classifier Quality

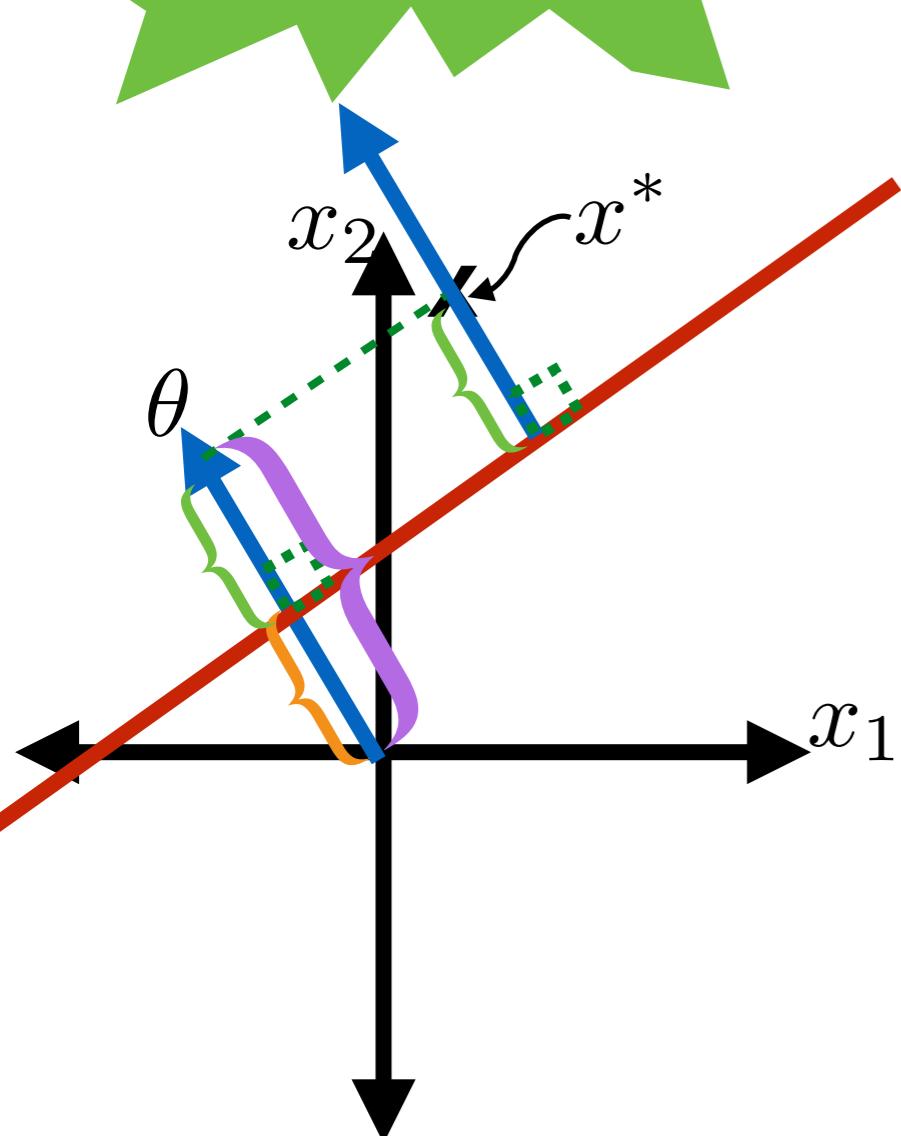
Math facts!



- The signed distance from a hyperplane defined by θ, θ_0 to a point x^* is:
 - = projection of x^* on θ
 - signed distance of line to origin
- $$= \frac{\theta^\top x^*}{\|\theta\|} - \frac{-\theta_0}{\|\theta\|} = \frac{\theta^\top x^* + \theta_0}{\|\theta\|}$$

Let's Talk About Classifier Quality

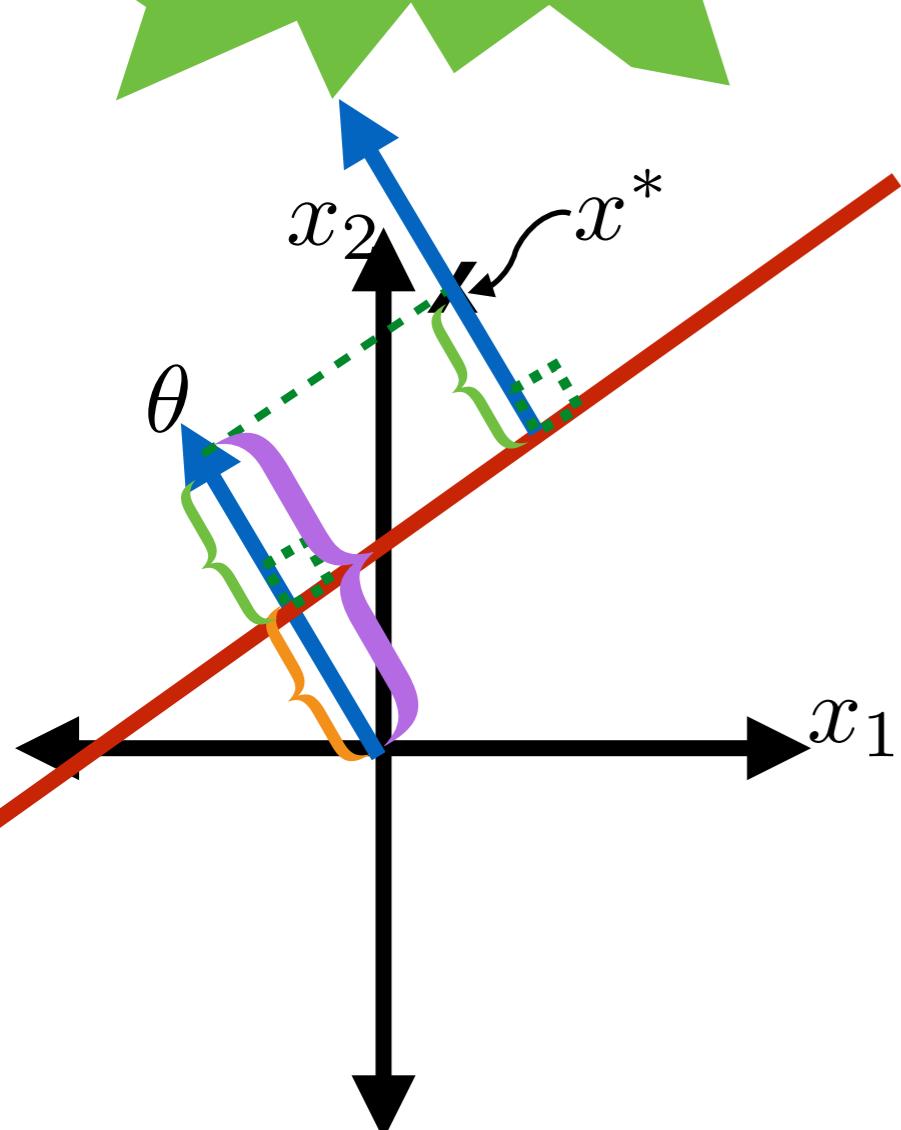
Math facts!



- The signed distance from a hyperplane defined by θ, θ_0 to a point x^* is:
 - = projection of x^* on θ
 - signed distance of line to origin
$$= \frac{\theta^\top x^*}{\|\theta\|} - \frac{-\theta_0}{\|\theta\|} = \frac{\theta^\top x^* + \theta_0}{\|\theta\|}$$
- *Definition:* The **margin of the labelled point** (x^*, y^*) with respect to the hyperplane defined by θ, θ_0 is:

Let's Talk About Classifier Quality

Math facts!



- The signed distance from a hyperplane defined by θ, θ_0 to a point x^* is:
 - = projection of x^* on θ
 - signed distance of line to origin
$$= \frac{\theta^\top x^*}{\|\theta\|} - \frac{-\theta_0}{\|\theta\|} = \frac{\theta^\top x^* + \theta_0}{\|\theta\|}$$
- *Definition:* The **margin of the labelled point** (x^*, y^*) with respect to the hyperplane defined by θ, θ_0 is:
$$y^* \left(\frac{\theta^\top x^* + \theta_0}{\|\theta\|} \right)$$

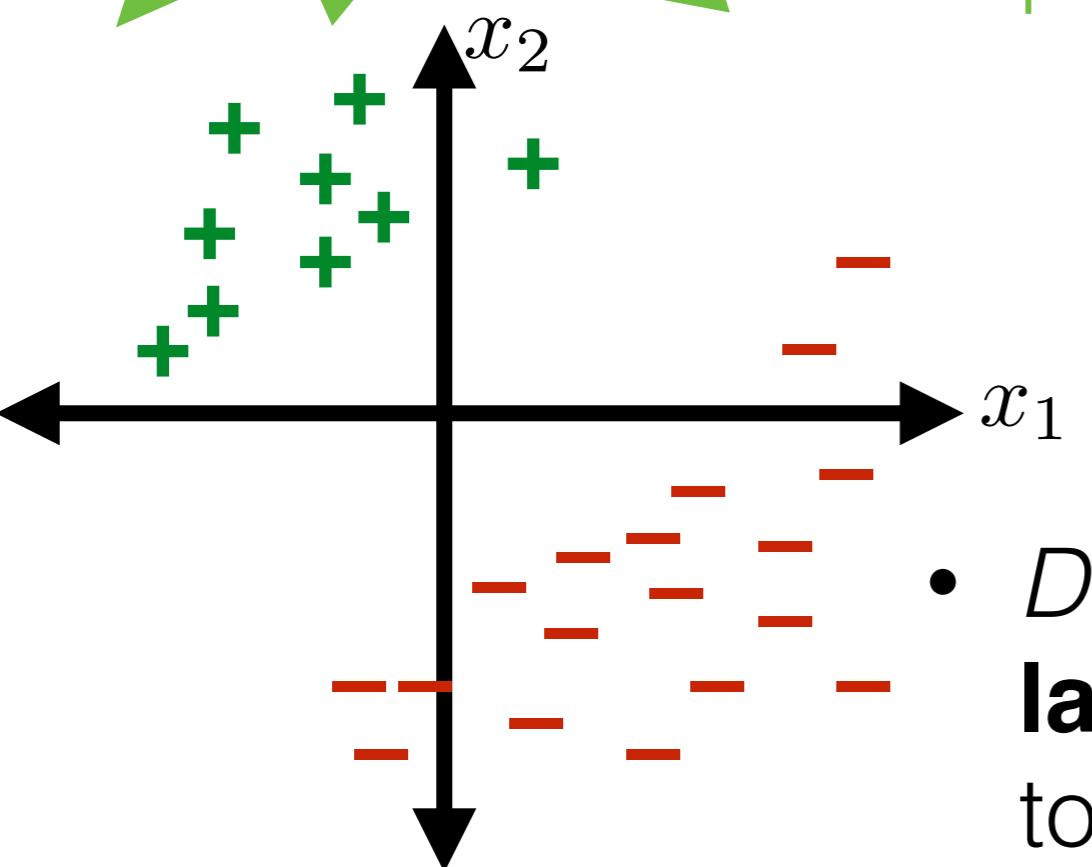
Let's Talk About Classifier Quality

Math facts!

- The signed distance from a hyperplane defined by θ, θ_0 to a point x^* is:
 - = projection of x^* on θ
 - signed distance of line to origin
$$= \frac{\theta^\top x^* - \theta_0}{\|\theta\|} = \frac{\theta^\top x^* + \theta_0}{\|\theta\|}$$
- *Definition:* The **margin of the labelled point** (x^*, y^*) with respect to the hyperplane defined by θ, θ_0 is:
$$y^* \left(\frac{\theta^\top x^* + \theta_0}{\|\theta\|} \right)$$

Let's Talk About Classifier Quality

Math facts!



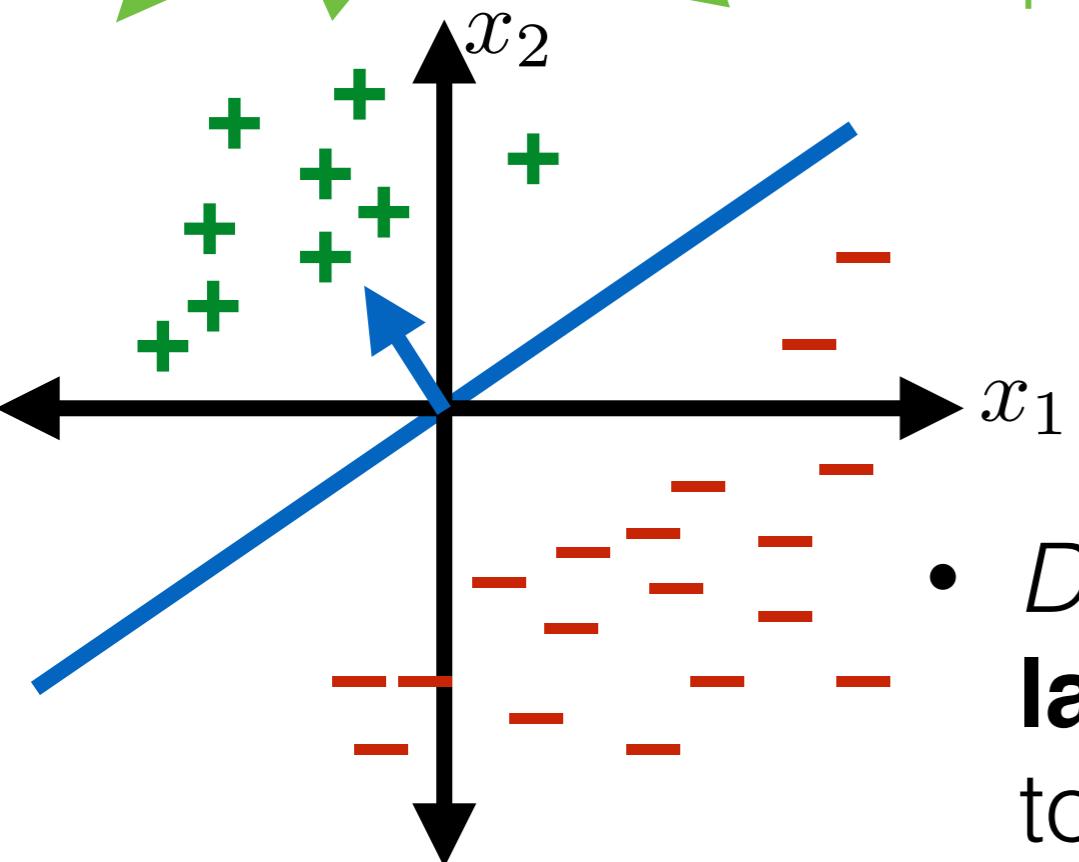
- The signed distance from a hyperplane defined by θ, θ_0 to a point x^* is:
 - = projection of x^* on θ
 - signed distance of line to origin
$$= \frac{\theta^\top x^* - \theta_0}{\|\theta\|} = \frac{\theta^\top x^* + \theta_0}{\|\theta\|}$$

- Definition:* The **margin of the labelled point** (x^*, y^*) with respect to the hyperplane defined by θ, θ_0 is:

$$y^* \left(\frac{\theta^\top x^* + \theta_0}{\|\theta\|} \right)$$

Let's Talk About Classifier Quality

Math facts!



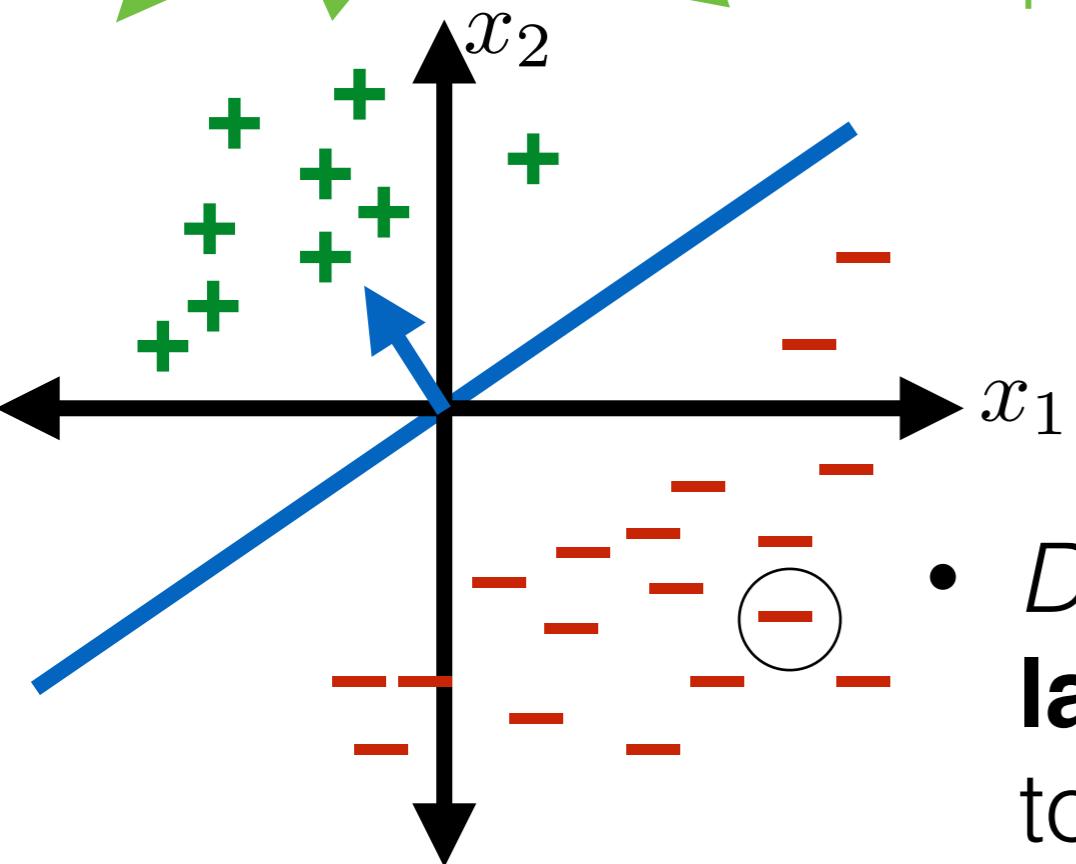
- The signed distance from a hyperplane defined by θ, θ_0 to a point x^* is:
 - = projection of x^* on θ
 - signed distance of line to origin
$$= \frac{\theta^\top x^* - \theta_0}{\|\theta\|} = \frac{\theta^\top x^* + \theta_0}{\|\theta\|}$$

- Definition:* The **margin of the labelled point** (x^*, y^*) with respect to the hyperplane defined by θ, θ_0 is:

$$y^* \left(\frac{\theta^\top x^* + \theta_0}{\|\theta\|} \right)$$

Let's Talk About Classifier Quality

Math facts!



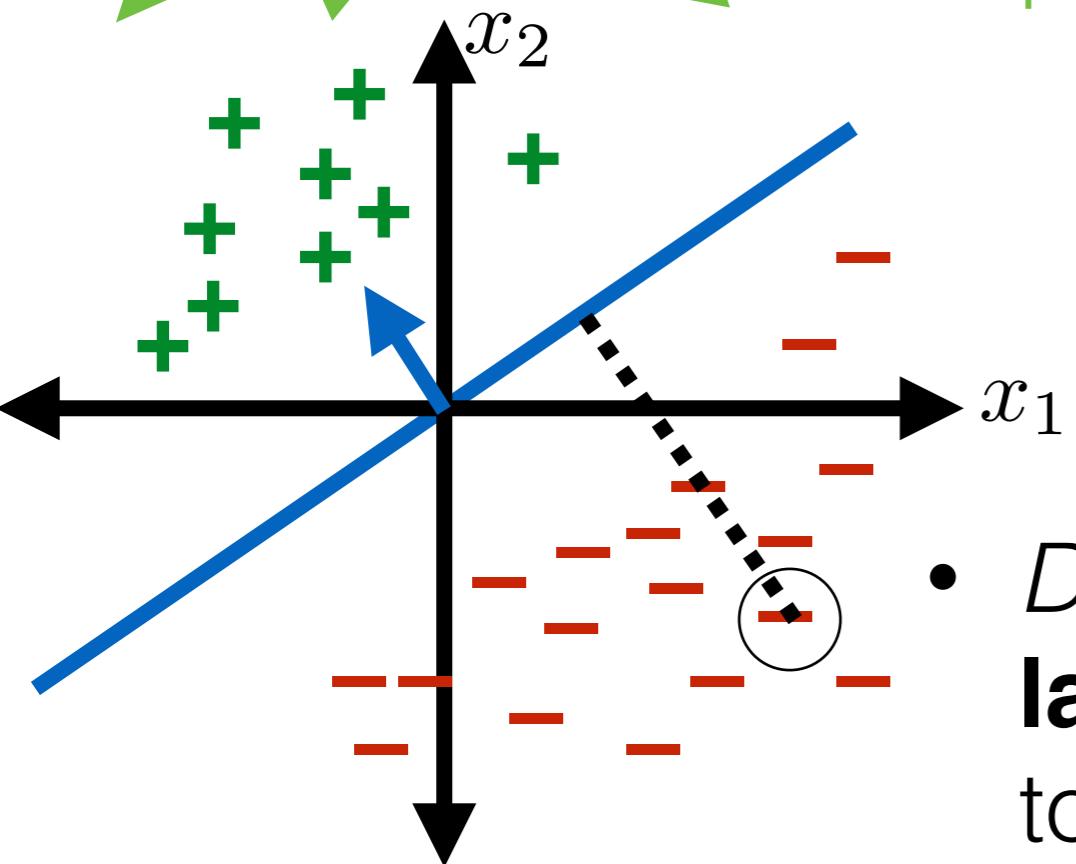
- The signed distance from a hyperplane defined by θ, θ_0 to a point x^* is:
 - = projection of x^* on θ
 - signed distance of line to origin
$$= \frac{\theta^\top x^* - \theta_0}{\|\theta\|} = \frac{\theta^\top x^* + \theta_0}{\|\theta\|}$$

- Definition:* The **margin of the labelled point** (x^*, y^*) with respect to the hyperplane defined by θ, θ_0 is:

$$y^* \left(\frac{\theta^\top x^* + \theta_0}{\|\theta\|} \right)$$

Let's Talk About Classifier Quality

Math facts!



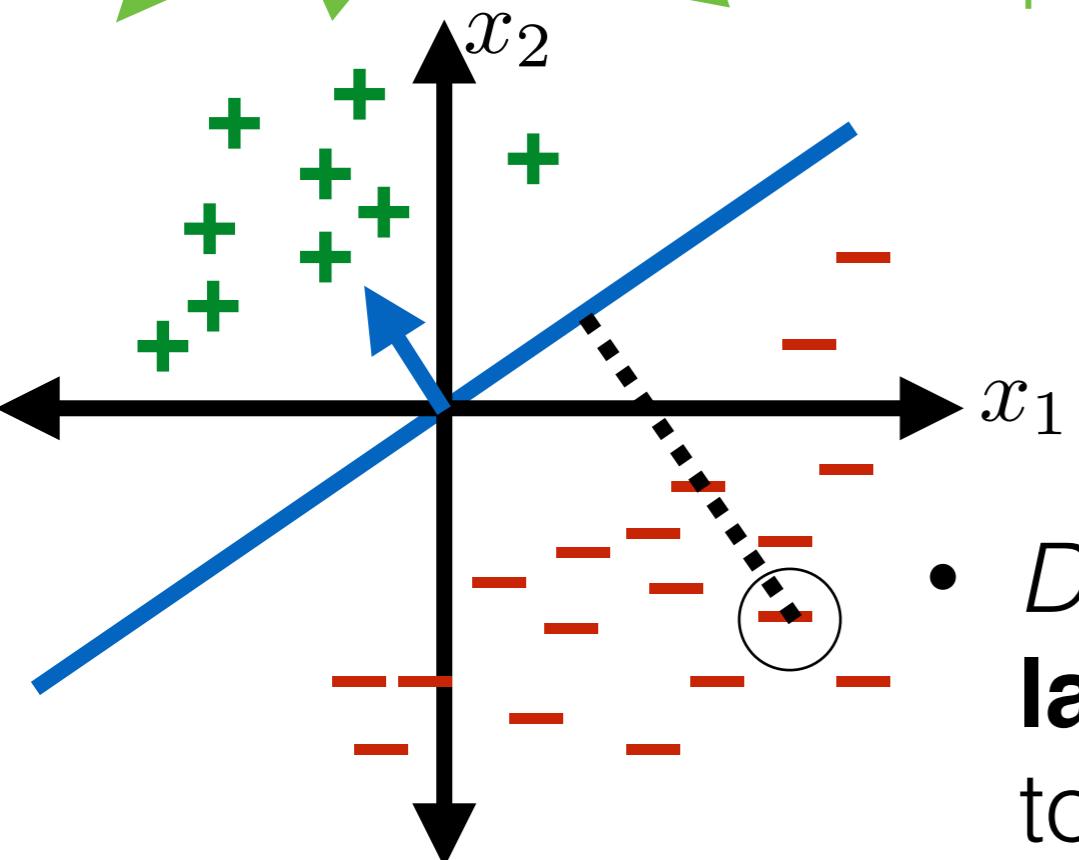
- The signed distance from a hyperplane defined by θ, θ_0 to a point x^* is:
 - = projection of x^* on θ
 - signed distance of line to origin
$$= \frac{\theta^\top x^* - \theta_0}{\|\theta\|} = \frac{\theta^\top x^* + \theta_0}{\|\theta\|}$$

- Definition:* The **margin of the labelled point** (x^*, y^*) with respect to the hyperplane defined by θ, θ_0 is:

$$y^* \left(\frac{\theta^\top x^* + \theta_0}{\|\theta\|} \right)$$

Let's Talk About Classifier Quality

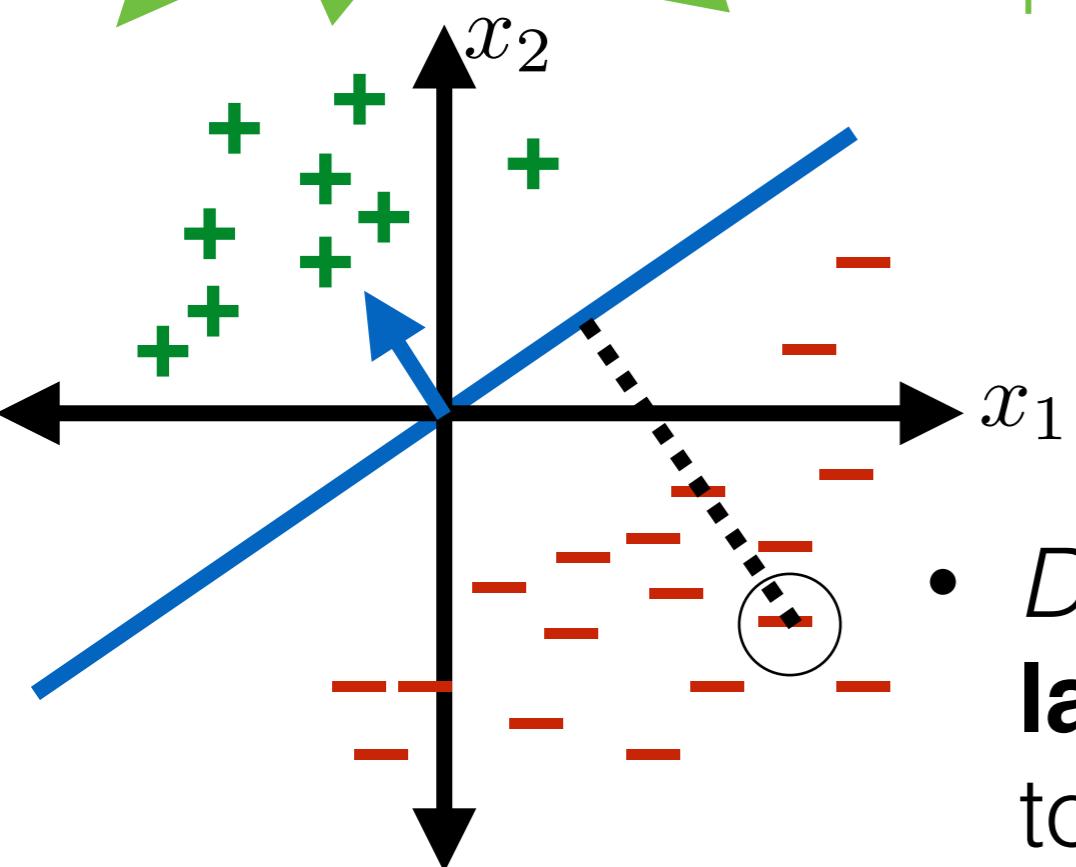
Math facts!



- The signed distance from a hyperplane defined by θ, θ_0 to a point x^* is:
 - = projection of x^* on θ
 - signed distance of line to origin
$$= \frac{\theta^\top x^* - \theta_0}{\|\theta\|} = \frac{\theta^\top x^* + \theta_0}{\|\theta\|}$$
- Definition:* The **margin of the labelled point** (x^*, y^*) with respect to the hyperplane defined by θ, θ_0 is:
$$y^* \left(\frac{\theta^\top x^* + \theta_0}{\|\theta\|} \right)$$
- Definition:* The **margin of the training set** \mathcal{D}_n with respect to the hyperplane defined by θ, θ_0 is:
5

Let's Talk About Classifier Quality

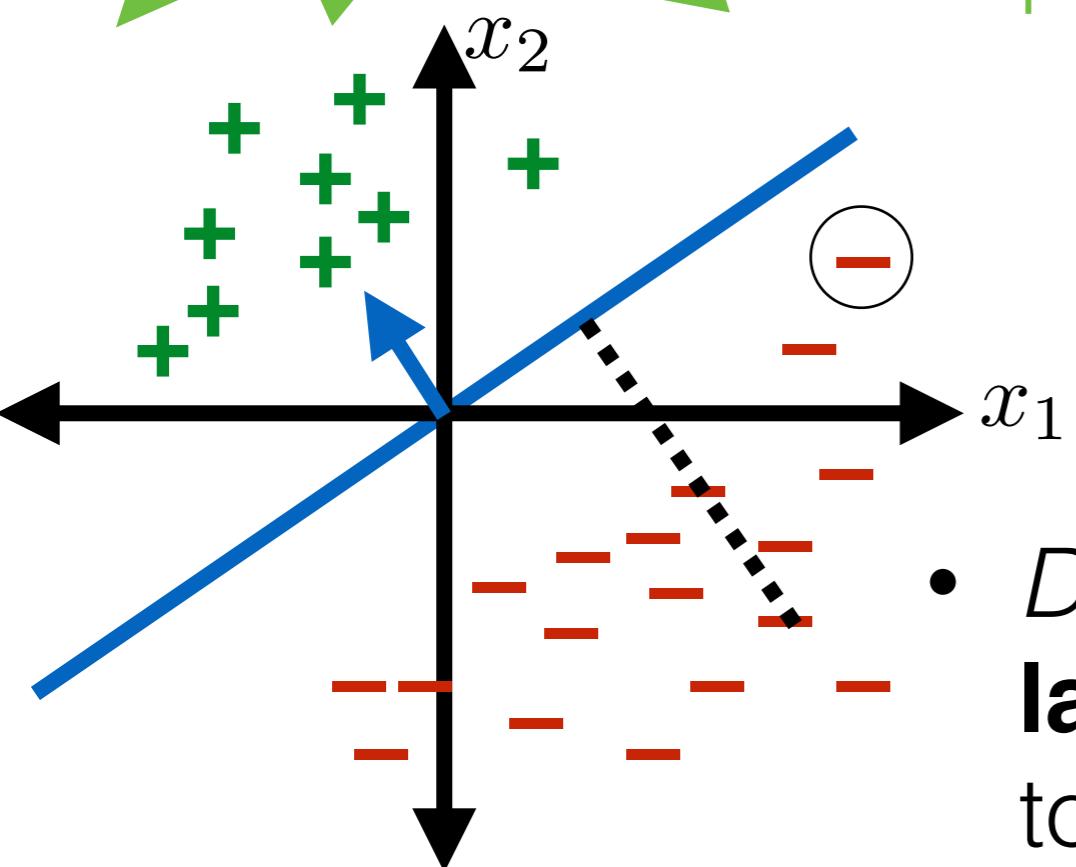
Math facts!



- The signed distance from a hyperplane defined by θ, θ_0 to a point x^* is:
= projection of x^* on θ
– signed distance of line to origin
$$= \frac{\theta^\top x^*}{\|\theta\|} - \frac{-\theta_0}{\|\theta\|} = \frac{\theta^\top x^* + \theta_0}{\|\theta\|}$$
- Definition:* The **margin of the labelled point** (x^*, y^*) with respect to the hyperplane defined by θ, θ_0 is:
$$y^* \left(\frac{\theta^\top x^* + \theta_0}{\|\theta\|} \right)$$
- Definition:* The **margin of the training set** \mathcal{D}_n with respect to the hyperplane defined by θ, θ_0 is:
$$\min_{i \in \{1, \dots, n\}} y^{(i)} \left(\frac{\theta^\top x^{(i)} + \theta_0}{\|\theta\|} \right)$$

Let's Talk About Classifier Quality

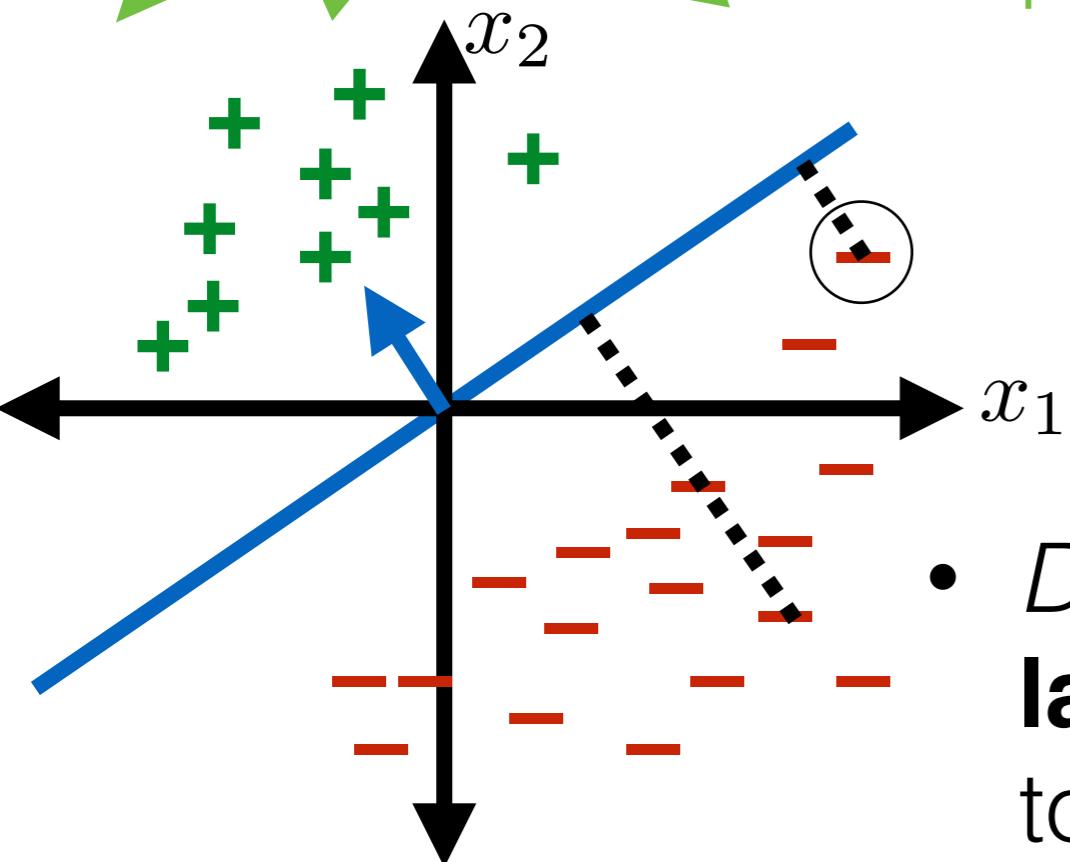
Math facts!



- The signed distance from a hyperplane defined by θ, θ_0 to a point x^* is:
= projection of x^* on θ
– signed distance of line to origin
$$= \frac{\theta^\top x^* - \theta_0}{\|\theta\|} = \frac{\theta^\top x^* + \theta_0}{\|\theta\|}$$
- Definition:* The **margin of the labelled point** (x^*, y^*) with respect to the hyperplane defined by θ, θ_0 is:
$$y^* \left(\frac{\theta^\top x^* + \theta_0}{\|\theta\|} \right)$$
- Definition:* The **margin of the training set** \mathcal{D}_n with respect to the hyperplane defined by θ, θ_0 is:
$$\min_{i \in \{1, \dots, n\}} y^{(i)} \left(\frac{\theta^\top x^{(i)} + \theta_0}{\|\theta\|} \right)$$

Let's Talk About Classifier Quality

Math facts!



- The signed distance from a hyperplane defined by θ, θ_0 to a point x^* is:
 - = projection of x^* on θ
 - signed distance of line to origin
$$= \frac{\theta^\top x^* - \theta_0}{\|\theta\|} = \frac{\theta^\top x^* + \theta_0}{\|\theta\|}$$

- Definition:* The **margin of the labelled point** (x^*, y^*) with respect to the hyperplane defined by θ, θ_0 is:

$$y^* \left(\frac{\theta^\top x^* + \theta_0}{\|\theta\|} \right)$$

- Definition:* The **margin of the training set** \mathcal{D}_n with respect to the hyperplane defined by θ, θ_0 is:

$$\min_{i \in \{1, \dots, n\}} y^{(i)} \left(\frac{\theta^\top x^{(i)} + \theta_0}{\|\theta\|} \right)$$

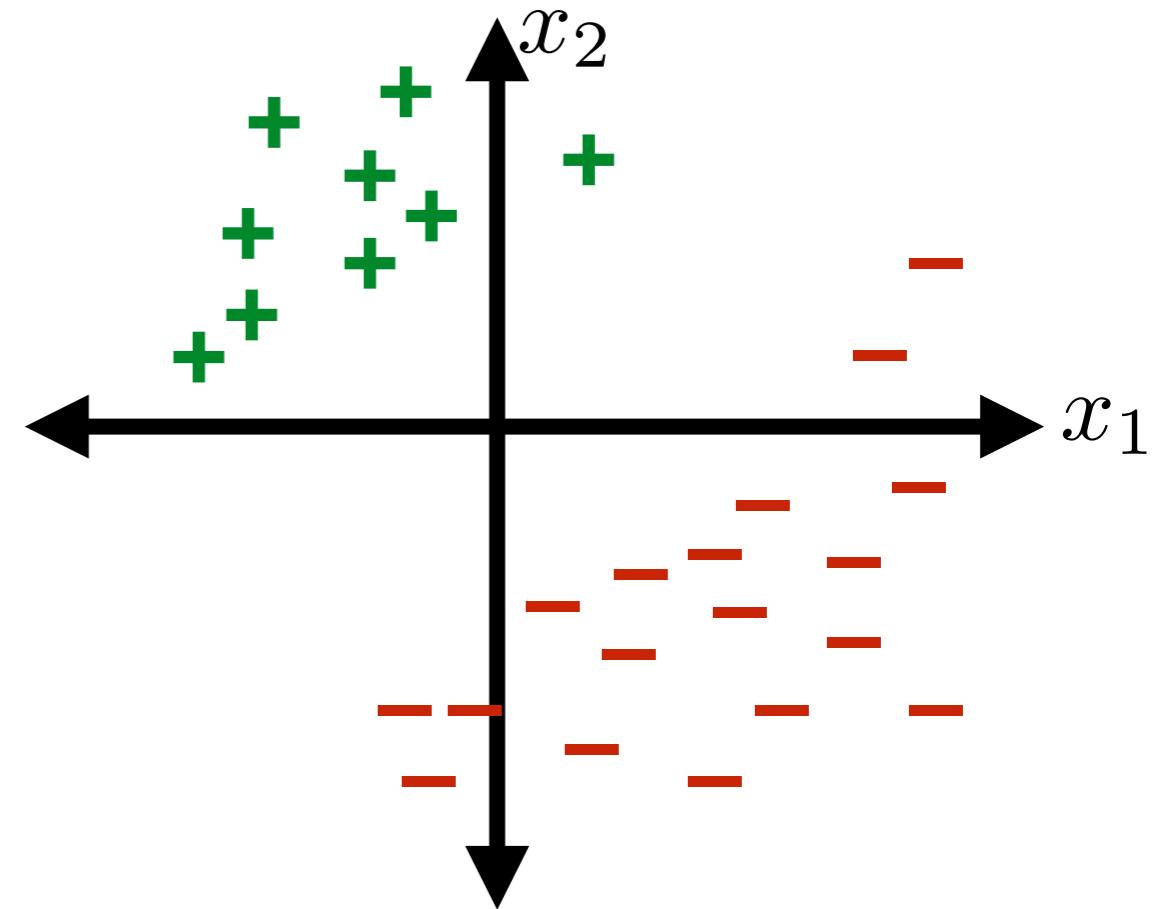
Theorem: Perceptron Performance

Theorem: Perceptron Performance

- **Assumptions:**

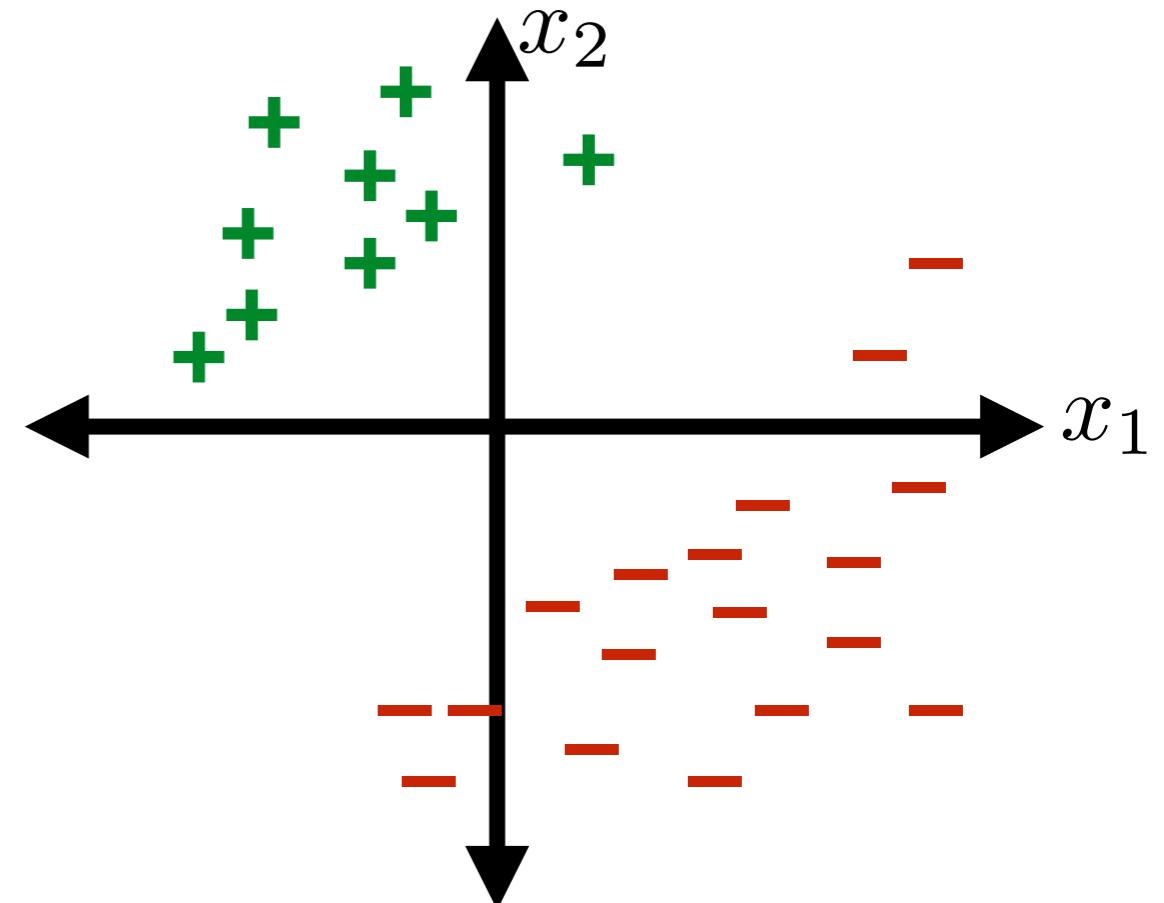
Theorem: Perceptron Performance

- **Assumptions:**



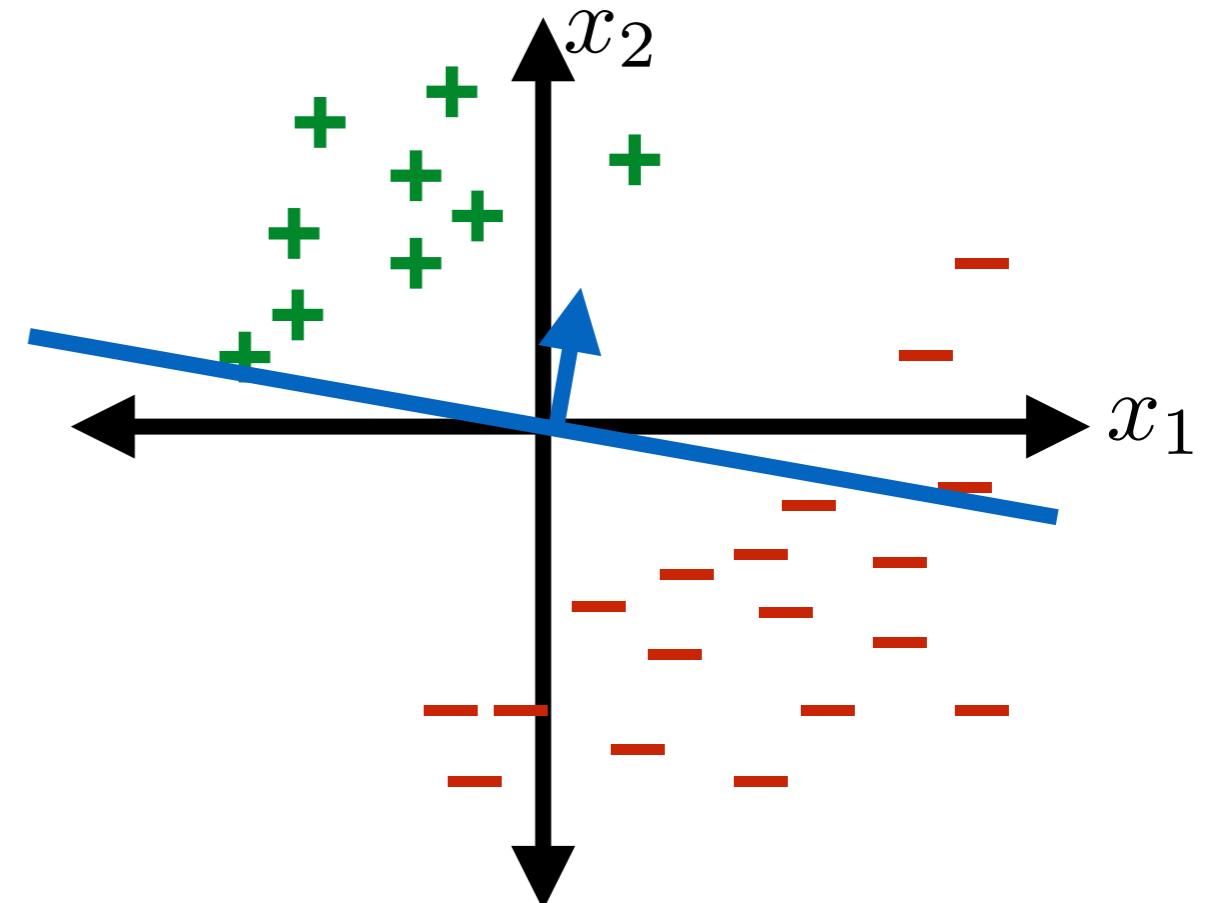
Theorem: Perceptron Performance

- **Assumptions:**
 - A. Our hypothesis class = classifiers with separating hyperplanes that pass through the origin (i.e. $\theta_0 = 0$)



Theorem: Perceptron Performance

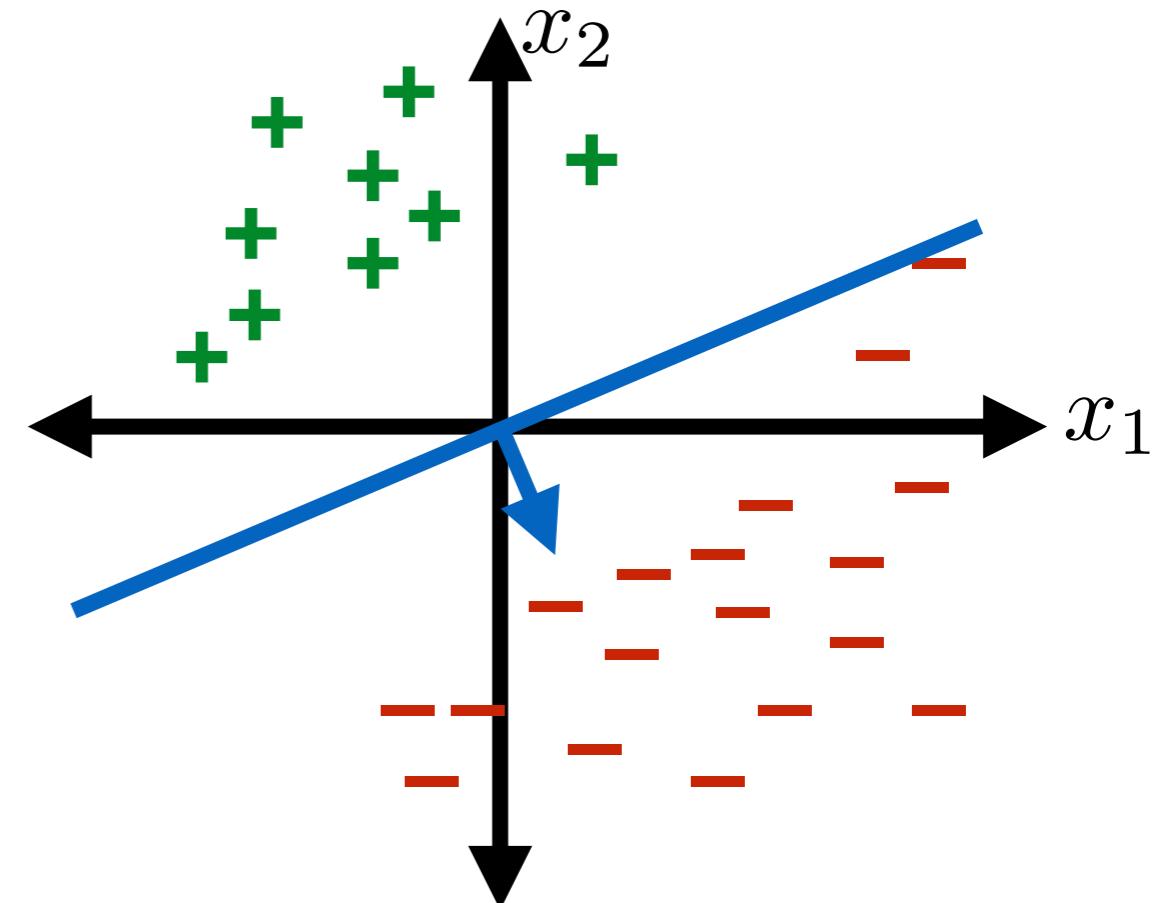
- **Assumptions:**
 - A. Our hypothesis class = classifiers with separating hyperplanes that pass through the origin (i.e. $\theta_0 = 0$)



Theorem: Perceptron Performance

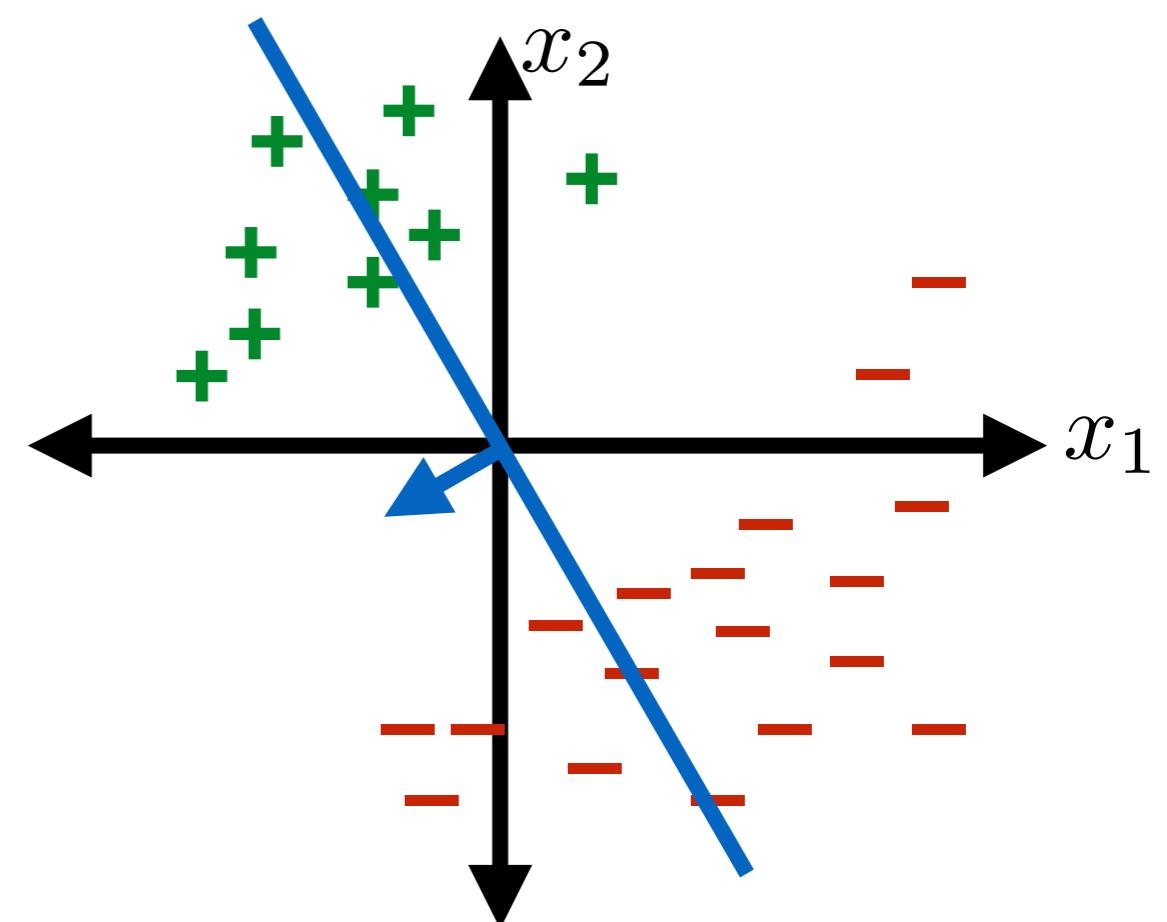
- **Assumptions:**

- A. Our hypothesis class = classifiers with separating hyperplanes that pass through the origin (i.e. $\theta_0 = 0$)



Theorem: Perceptron Performance

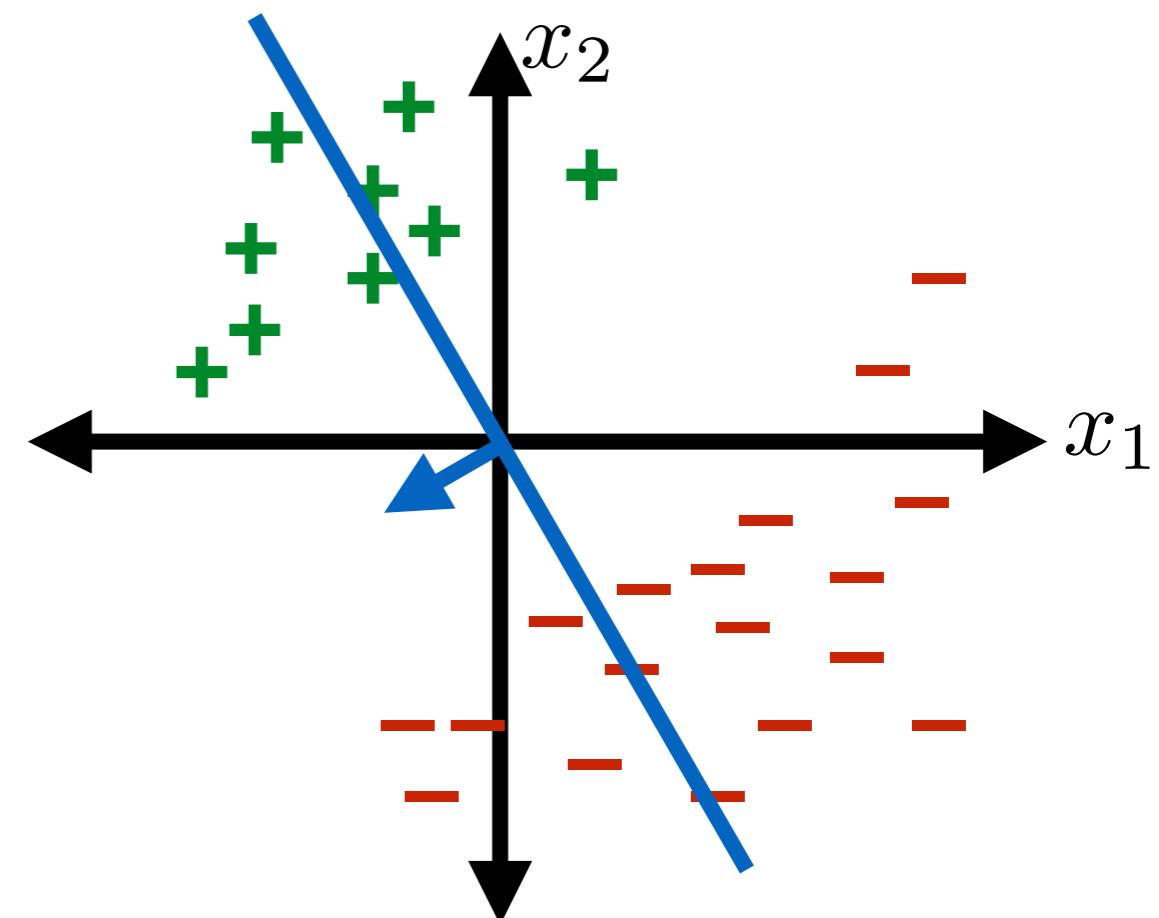
- **Assumptions:**
 - A. Our hypothesis class = classifiers with separating hyperplanes that pass through the origin (i.e. $\theta_0 = 0$)



Theorem: Perceptron Performance

- **Assumptions:**

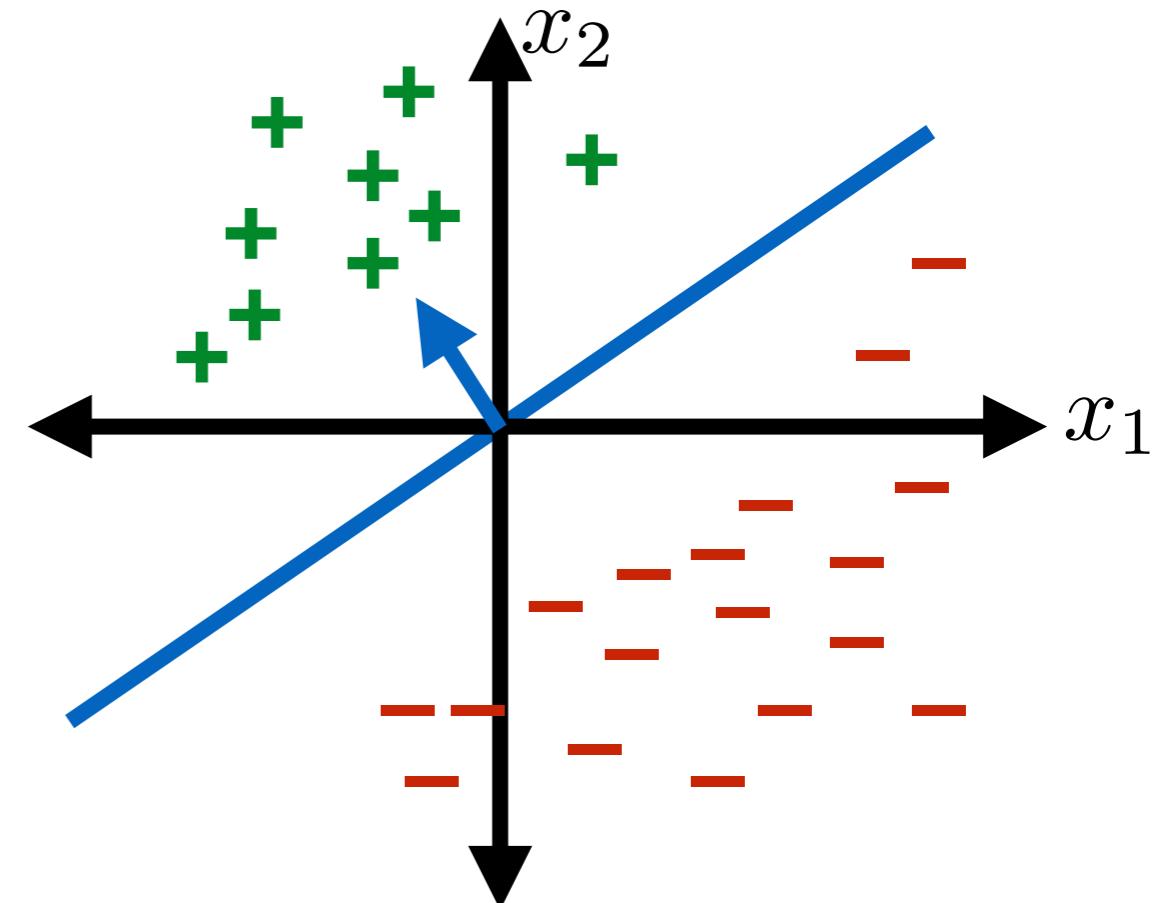
- A. Our hypothesis class = classifiers with separating hyperplanes that pass through the origin (i.e. $\theta_0 = 0$)
- B. There exist θ^* and γ such that $\gamma > 0$ and, for every $i \in \{1, \dots, n\}$, we have $y^{(i)} \left(\frac{\theta^{*\top} x^{(i)}}{\|\theta\|} \right) > \gamma$



Theorem: Perceptron Performance

- **Assumptions:**

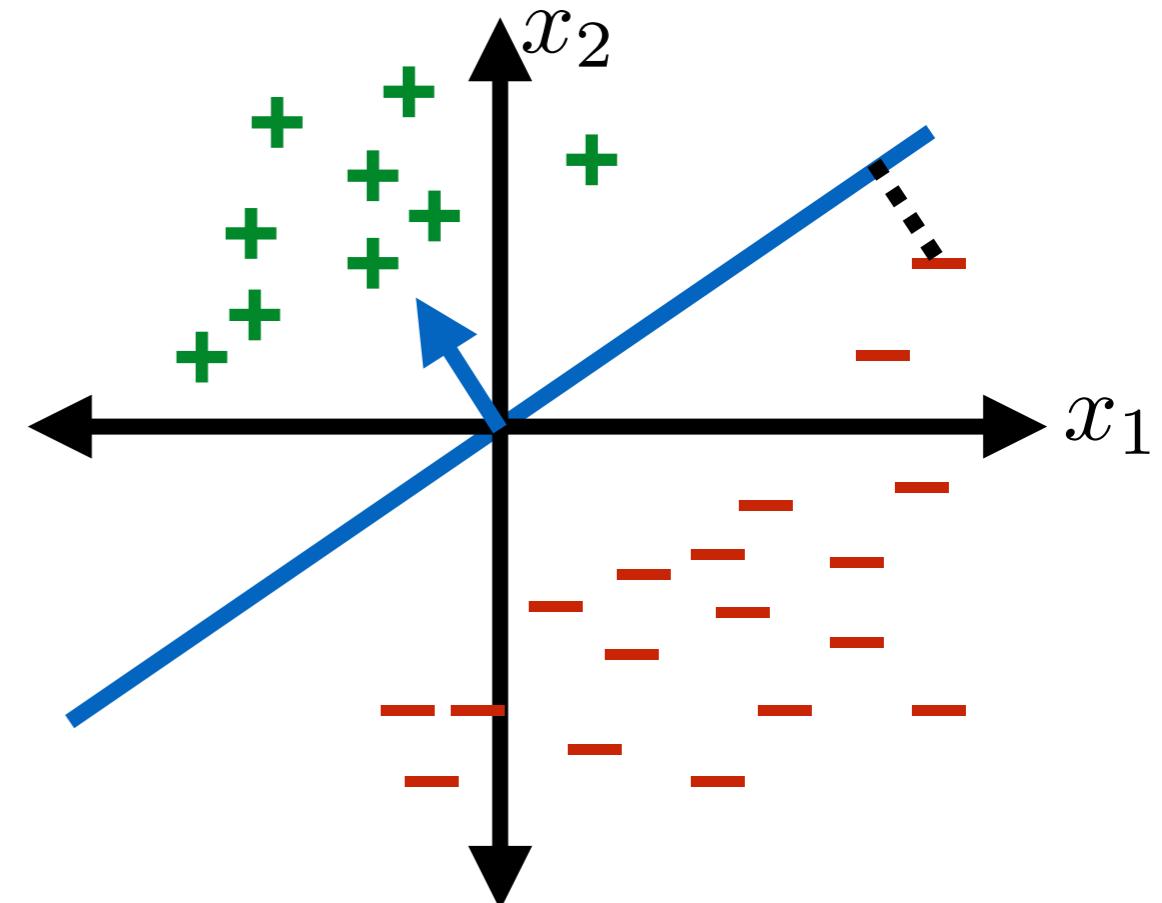
- A. Our hypothesis class = classifiers with separating hyperplanes that pass through the origin (i.e. $\theta_0 = 0$)
- B. There exist θ^* and γ such that $\gamma > 0$ and, for every $i \in \{1, \dots, n\}$, we have $y^{(i)} \left(\frac{\theta^{*\top} x^{(i)}}{\|\theta\|} \right) > \gamma$



Theorem: Perceptron Performance

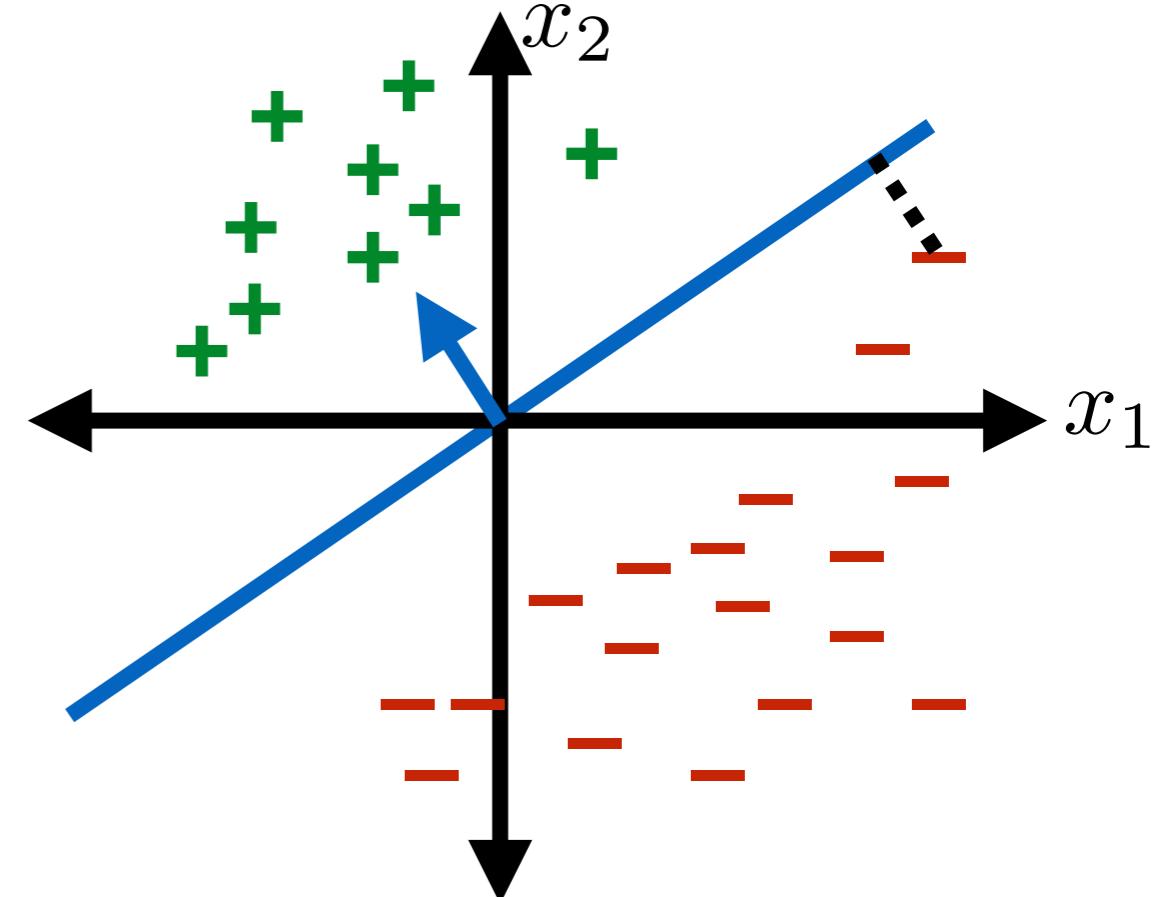
- **Assumptions:**

- A. Our hypothesis class = classifiers with separating hyperplanes that pass through the origin (i.e. $\theta_0 = 0$)
- B. There exist θ^* and γ such that $\gamma > 0$ and, for every $i \in \{1, \dots, n\}$, we have $y^{(i)} \left(\frac{\theta^{*\top} x^{(i)}}{\|\theta\|} \right) > \gamma$



Theorem: Perceptron Performance

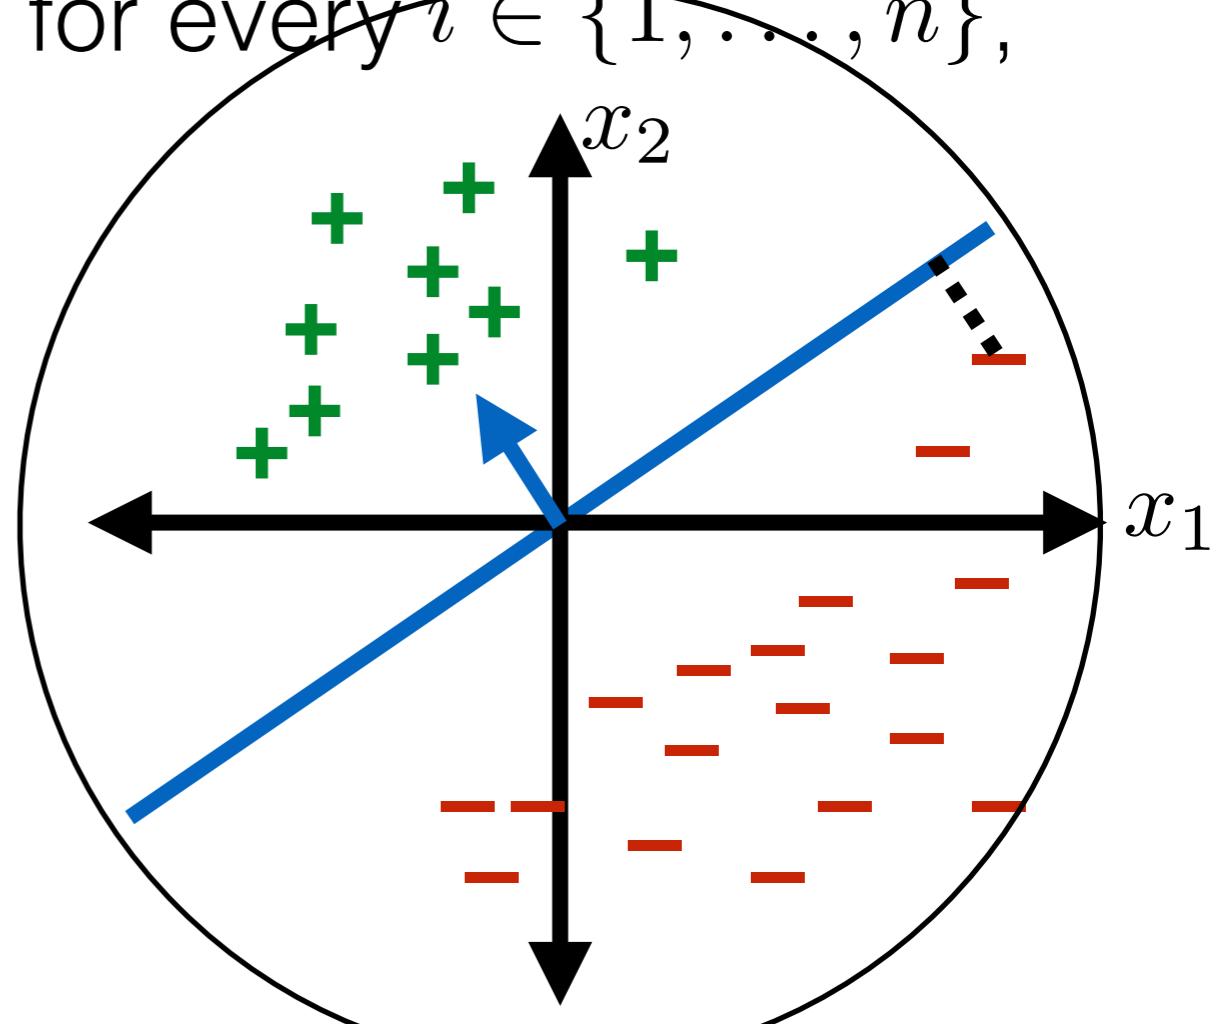
- **Assumptions:**
 - A. Our hypothesis class = classifiers with separating hyperplanes that pass through the origin (i.e. $\theta_0 = 0$)
 - B. There exist θ^* and γ such that $\gamma > 0$ and, for every $i \in \{1, \dots, n\}$, we have $y^{(i)} \left(\frac{\theta^{*\top} x^{(i)}}{\|\theta\|} \right) > \gamma$
 - C. There exists R such that, for every $i \in \{1, \dots, n\}$, we have $\|x^{(i)}\| \leq R$



Theorem: Perceptron Performance

- **Assumptions:**

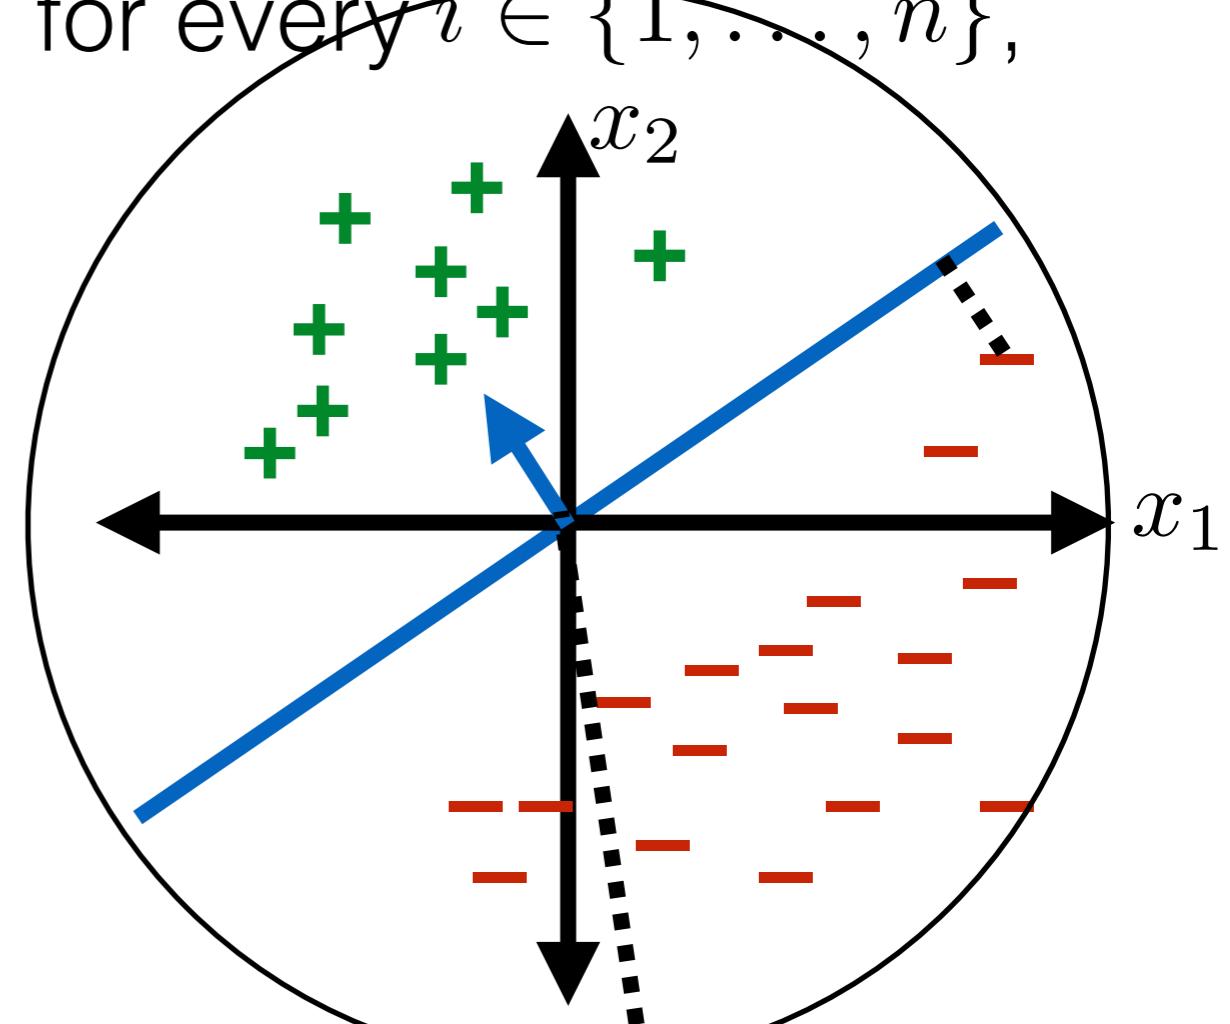
- A. Our hypothesis class = classifiers with separating hyperplanes that pass through the origin (i.e. $\theta_0 = 0$)
- B. There exist θ^* and γ such that $\gamma > 0$ and, for every $i \in \{1, \dots, n\}$, we have $y^{(i)} \left(\frac{\theta^{*\top} x^{(i)}}{\|\theta\|} \right) > \gamma$
- C. There exists R such that, for every $i \in \{1, \dots, n\}$, we have $\|x^{(i)}\| \leq R$



Theorem: Perceptron Performance

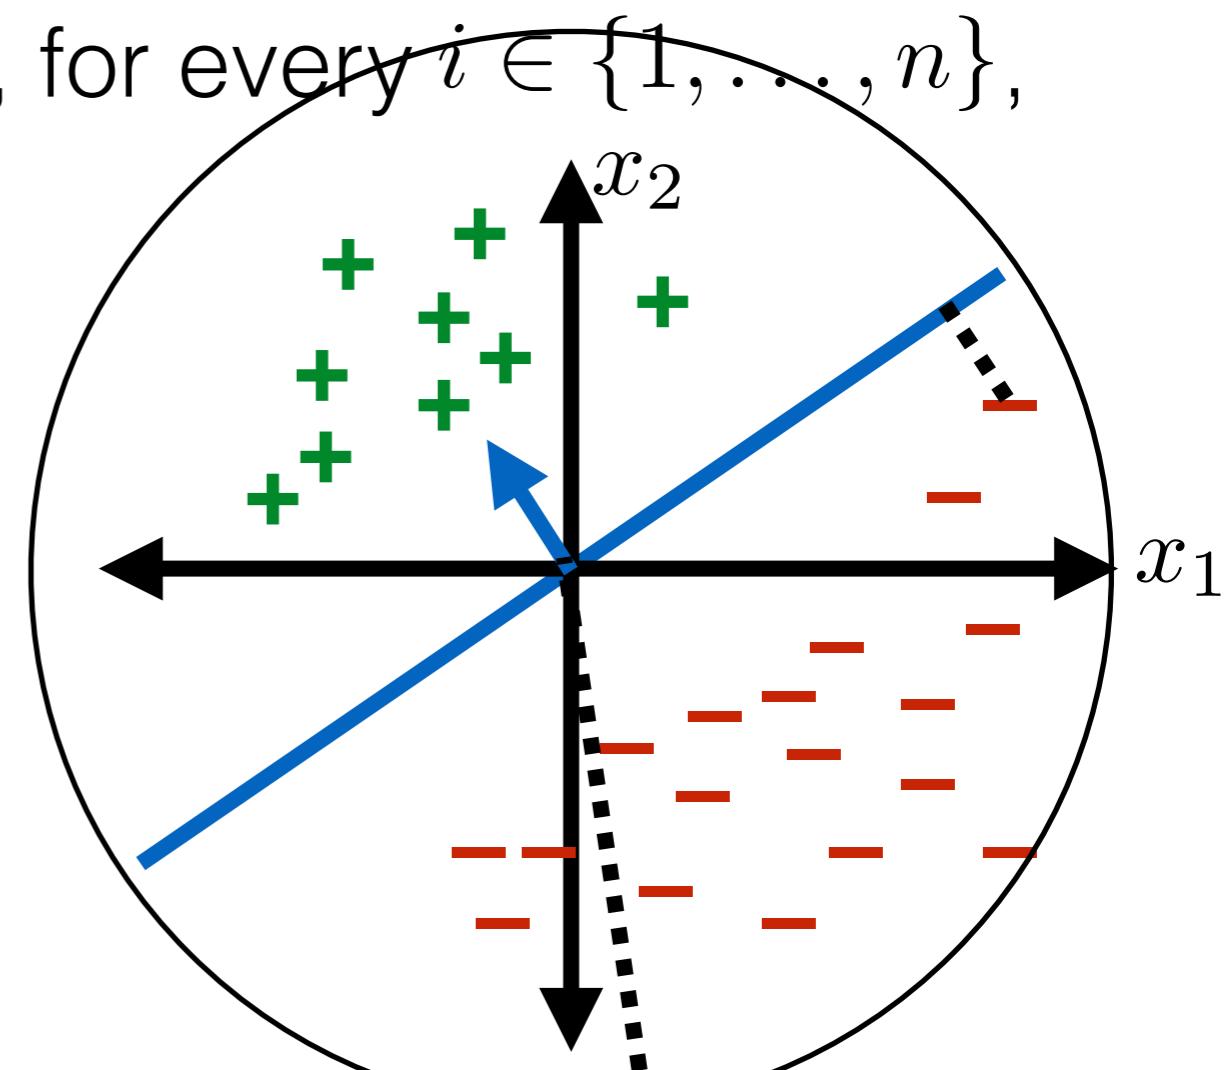
- **Assumptions:**

- A. Our hypothesis class = classifiers with separating hyperplanes that pass through the origin (i.e. $\theta_0 = 0$)
- B. There exist θ^* and γ such that $\gamma > 0$ and, for every $i \in \{1, \dots, n\}$, we have $y^{(i)} \left(\frac{\theta^{*\top} x^{(i)}}{\|\theta\|} \right) > \gamma$
- C. There exists R such that, for every $i \in \{1, \dots, n\}$, we have $\|x^{(i)}\| \leq R$



Theorem: Perceptron Performance

- **Assumptions:**
 - A. Our hypothesis class = classifiers with separating hyperplanes that pass through the origin (i.e. $\theta_0 = 0$)
 - B. There exist θ^* and γ such that $\gamma > 0$ and, for every $i \in \{1, \dots, n\}$, we have $y^{(i)} \left(\frac{\theta^{*\top} x^{(i)}}{\|\theta\|} \right) > \gamma$
 - C. There exists R such that, for every $i \in \{1, \dots, n\}$, we have $\|x^{(i)}\| \leq R$
- **Conclusion:** Then the perceptron algorithm will make at most $(R/\gamma)^2$ updates to θ . Once it goes through a pass of i without changes, the training error of its hypothesis will be 0.



Why classifiers through the origin?

Why classifiers through the origin?

- If we're clever, we don't lose any flexibility

Why classifiers through the origin?

- If we're clever, we don't lose any flexibility
 - Classifier with offset

Why classifiers through the origin?

- If we're clever, we don't lose any flexibility
 - Classifier with offset

$$x \in \mathbb{R}^d, \theta \in \mathbb{R}^d, \theta_0 \in \mathbb{R}$$

$$x : \theta^\top x + \theta_0 = 0$$

Why classifiers through the origin?

- If we're clever, we don't lose any flexibility
 - Classifier with offset

$$x \in \mathbb{R}^d, \theta \in \mathbb{R}^d, \theta_0 \in \mathbb{R}$$

$$x : \theta^\top x + \theta_0 = 0$$

- Classifier without offset

Why classifiers through the origin?

- If we're clever, we don't lose any flexibility
 - Classifier with offset

$$x \in \mathbb{R}^d, \theta \in \mathbb{R}^d, \theta_0 \in \mathbb{R}$$

$$x : \theta^\top x + \theta_0 = 0$$

- Classifier without offset

$$x_{\text{new}} \in \mathbb{R}^{d+1}, \theta_{\text{new}} \in \mathbb{R}^{d+1}$$

$$x_{\text{new}} = [x_1, x_2, \dots, x_d, 1]^\top, \theta_{\text{new}} = [\theta_1, \theta_2, \dots, \theta_d, \theta_0]^\top$$

Why classifiers through the origin?

- If we're clever, we don't lose any flexibility
 - Classifier with offset

$$x \in \mathbb{R}^d, \theta \in \mathbb{R}^d, \theta_0 \in \mathbb{R}$$

$$x : \theta^\top x + \theta_0 = 0$$

- Classifier without offset

$$x_{\text{new}} \in \mathbb{R}^{d+1}, \theta_{\text{new}} \in \mathbb{R}^{d+1}$$

$$x_{\text{new}} = [x_1, x_2, \dots, x_d, 1]^\top, \theta_{\text{new}} = [\theta_1, \theta_2, \dots, \theta_d, \theta_0]^\top$$

$$x_{\text{new}, 1:d} : \theta_{\text{new}}^\top x_{\text{new}} = 0$$

Why classifiers through the origin?

- If we're clever, we don't lose any flexibility
 - Classifier with offset

$$x \in \mathbb{R}^d, \theta \in \mathbb{R}^d, \theta_0 \in \mathbb{R}$$

$$x : \theta^\top x + \theta_0 = 0$$

- Classifier without offset

$$x_{\text{new}} \in \mathbb{R}^{d+1}, \theta_{\text{new}} \in \mathbb{R}^{d+1}$$

$$x_{\text{new}} = [x_1, x_2, \dots, x_d, 1]^\top, \theta_{\text{new}} = [\theta_1, \theta_2, \dots, \theta_d, \theta_0]^\top$$

$$x_{\text{new}, 1:d} : \theta_{\text{new}}^\top x_{\text{new}} = 0$$

Why classifiers through the origin?

- If we're clever, we don't lose any flexibility
 - Classifier with offset

$$x \in \mathbb{R}^d, \theta \in \mathbb{R}^d, \theta_0 \in \mathbb{R}$$
$$x : \theta^\top x + \theta_0 \leq 0$$

- Classifier without offset

$$x_{\text{new}} \in \mathbb{R}^{d+1}, \theta_{\text{new}} \in \mathbb{R}^{d+1}$$

$$x_{\text{new}} = [x_1, x_2, \dots, x_d, 1]^\top, \theta_{\text{new}} = [\theta_1, \theta_2, \dots, \theta_d, \theta_0]^\top$$
$$x_{\text{new}, 1:d} : \theta_{\text{new}}^\top x_{\text{new}} \leq 0$$

Why classifiers through the origin?

- If we're clever, we don't lose any flexibility
 - Classifier with offset

$$x \in \mathbb{R}^d, \theta \in \mathbb{R}^d, \theta_0 \in \mathbb{R}$$

$$x : \theta^\top x + \theta_0 \begin{matrix} < \\ > \end{matrix} 0$$

- Classifier without offset

$$x_{\text{new}} \in \mathbb{R}^{d+1}, \theta_{\text{new}} \in \mathbb{R}^{d+1}$$

$$x_{\text{new}} = [x_1, x_2, \dots, x_d, 1]^\top, \theta_{\text{new}} = [\theta_1, \theta_2, \dots, \theta_d, \theta_0]^\top$$

$$x_{\text{new}, 1:d} : \theta_{\text{new}}^\top x_{\text{new}} \begin{matrix} < \\ > \end{matrix} 0$$

Why classifiers through the origin?

- If we're clever, we don't lose any flexibility
 - Classifier with offset

$$x \in \mathbb{R}^d, \theta \in \mathbb{R}^d, \theta_0 \in \mathbb{R}$$

$$x : \theta^\top x + \theta_0 \stackrel{<}{\stackrel{>}{=}} 0$$

- Classifier without offset

$$x_{\text{new}} \in \mathbb{R}^{d+1}, \theta_{\text{new}} \in \mathbb{R}^{d+1}$$

$$x_{\text{new}} = [x_1, x_2, \dots, x_d, 1]^\top, \theta_{\text{new}} = [\theta_1, \theta_2, \dots, \theta_d, \theta_0]^\top$$

$$x_{\text{new}, 1:d} : \theta_{\text{new}}^\top x_{\text{new}} \stackrel{<}{\stackrel{>}{=}} 0$$

Why classifiers through the origin?

- If we're clever, we don't lose any flexibility
 - Classifier with offset

$$x \in \mathbb{R}^d, \theta \in \mathbb{R}^d, \theta_0 \in \mathbb{R}$$

$$x : \theta^\top x + \theta_0 \stackrel{<}{\stackrel{>}{=}} 0$$

- Classifier without offset

$$x_{\text{new}} \in \mathbb{R}^{d+1}, \theta_{\text{new}} \in \mathbb{R}^{d+1}$$

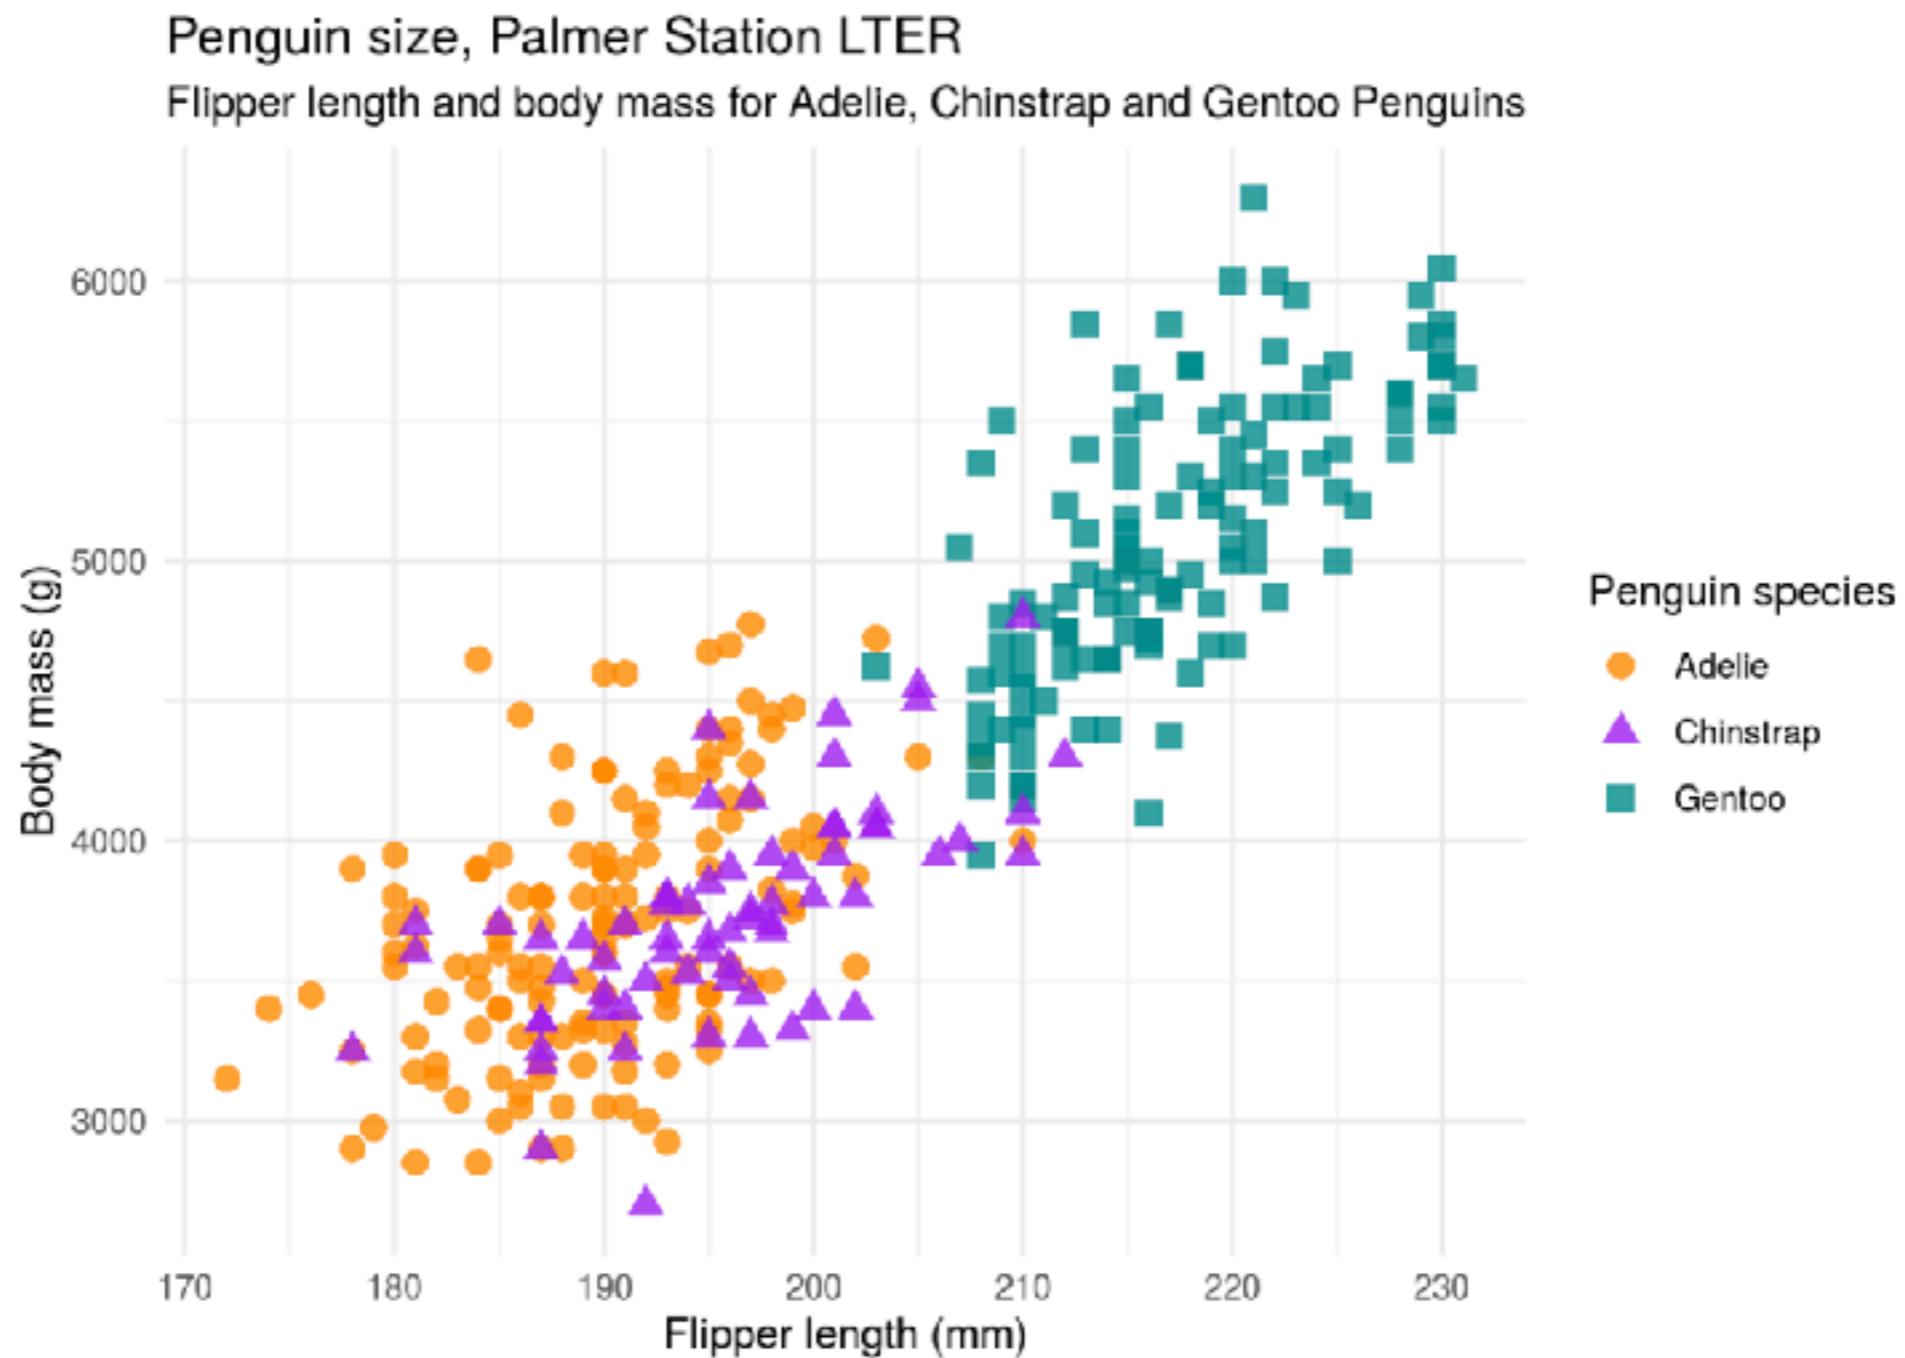
$$x_{\text{new}} = [x_1, x_2, \dots, x_d, 1]^\top, \theta_{\text{new}} = [\theta_1, \theta_2, \dots, \theta_d, \theta_0]^\top$$

$$x_{\text{new}, 1:d} : \theta_{\text{new}}^\top x_{\text{new}} \stackrel{<}{\stackrel{>}{=}} 0$$

- Can first convert to “expanded” feature space, then apply theorem

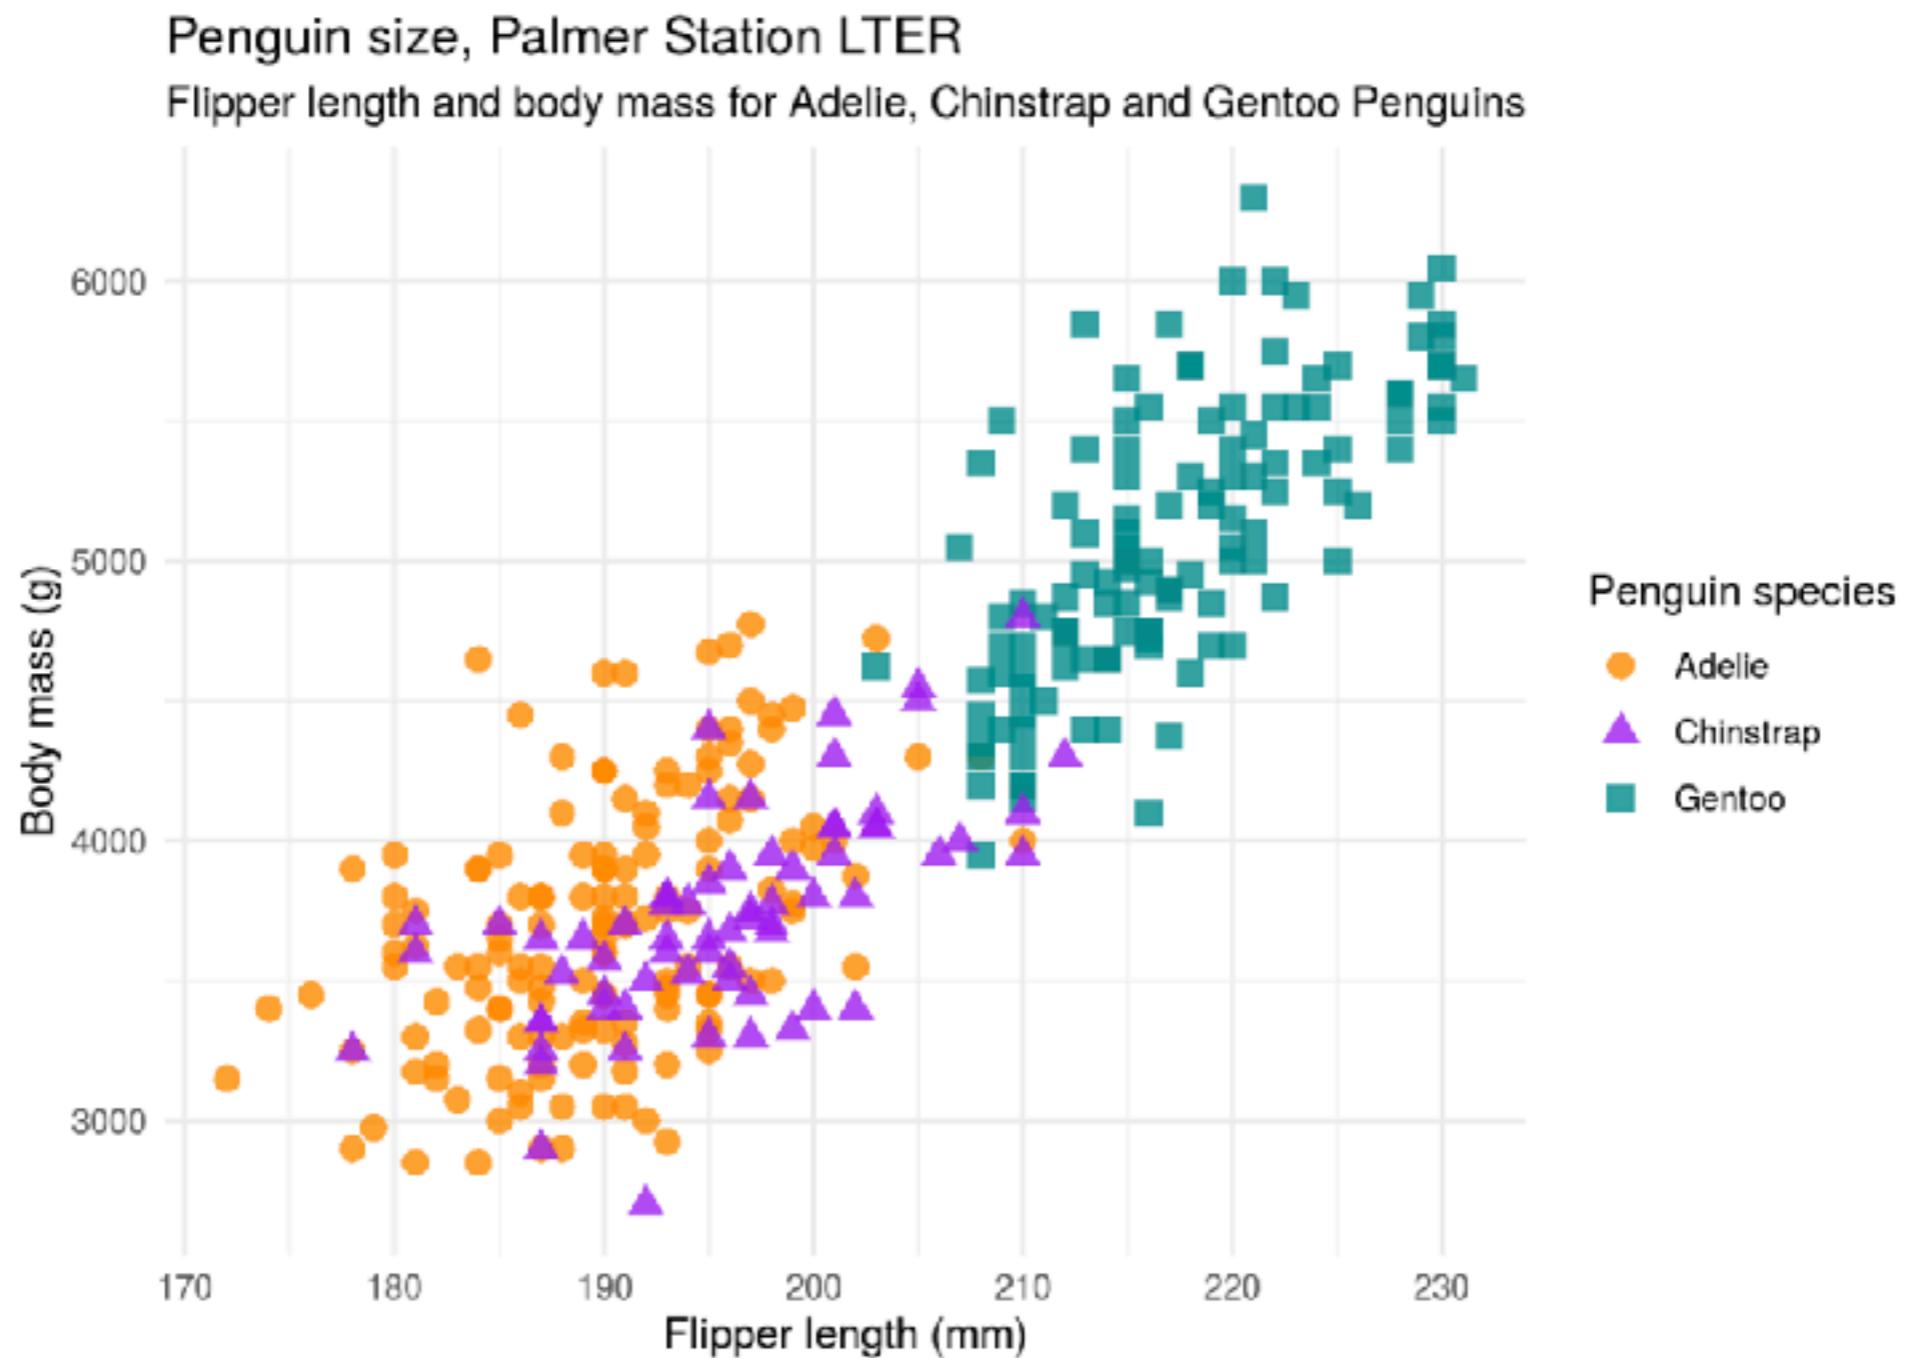
Problem: data not linearly separable

- Typical real data sets aren't linearly separable



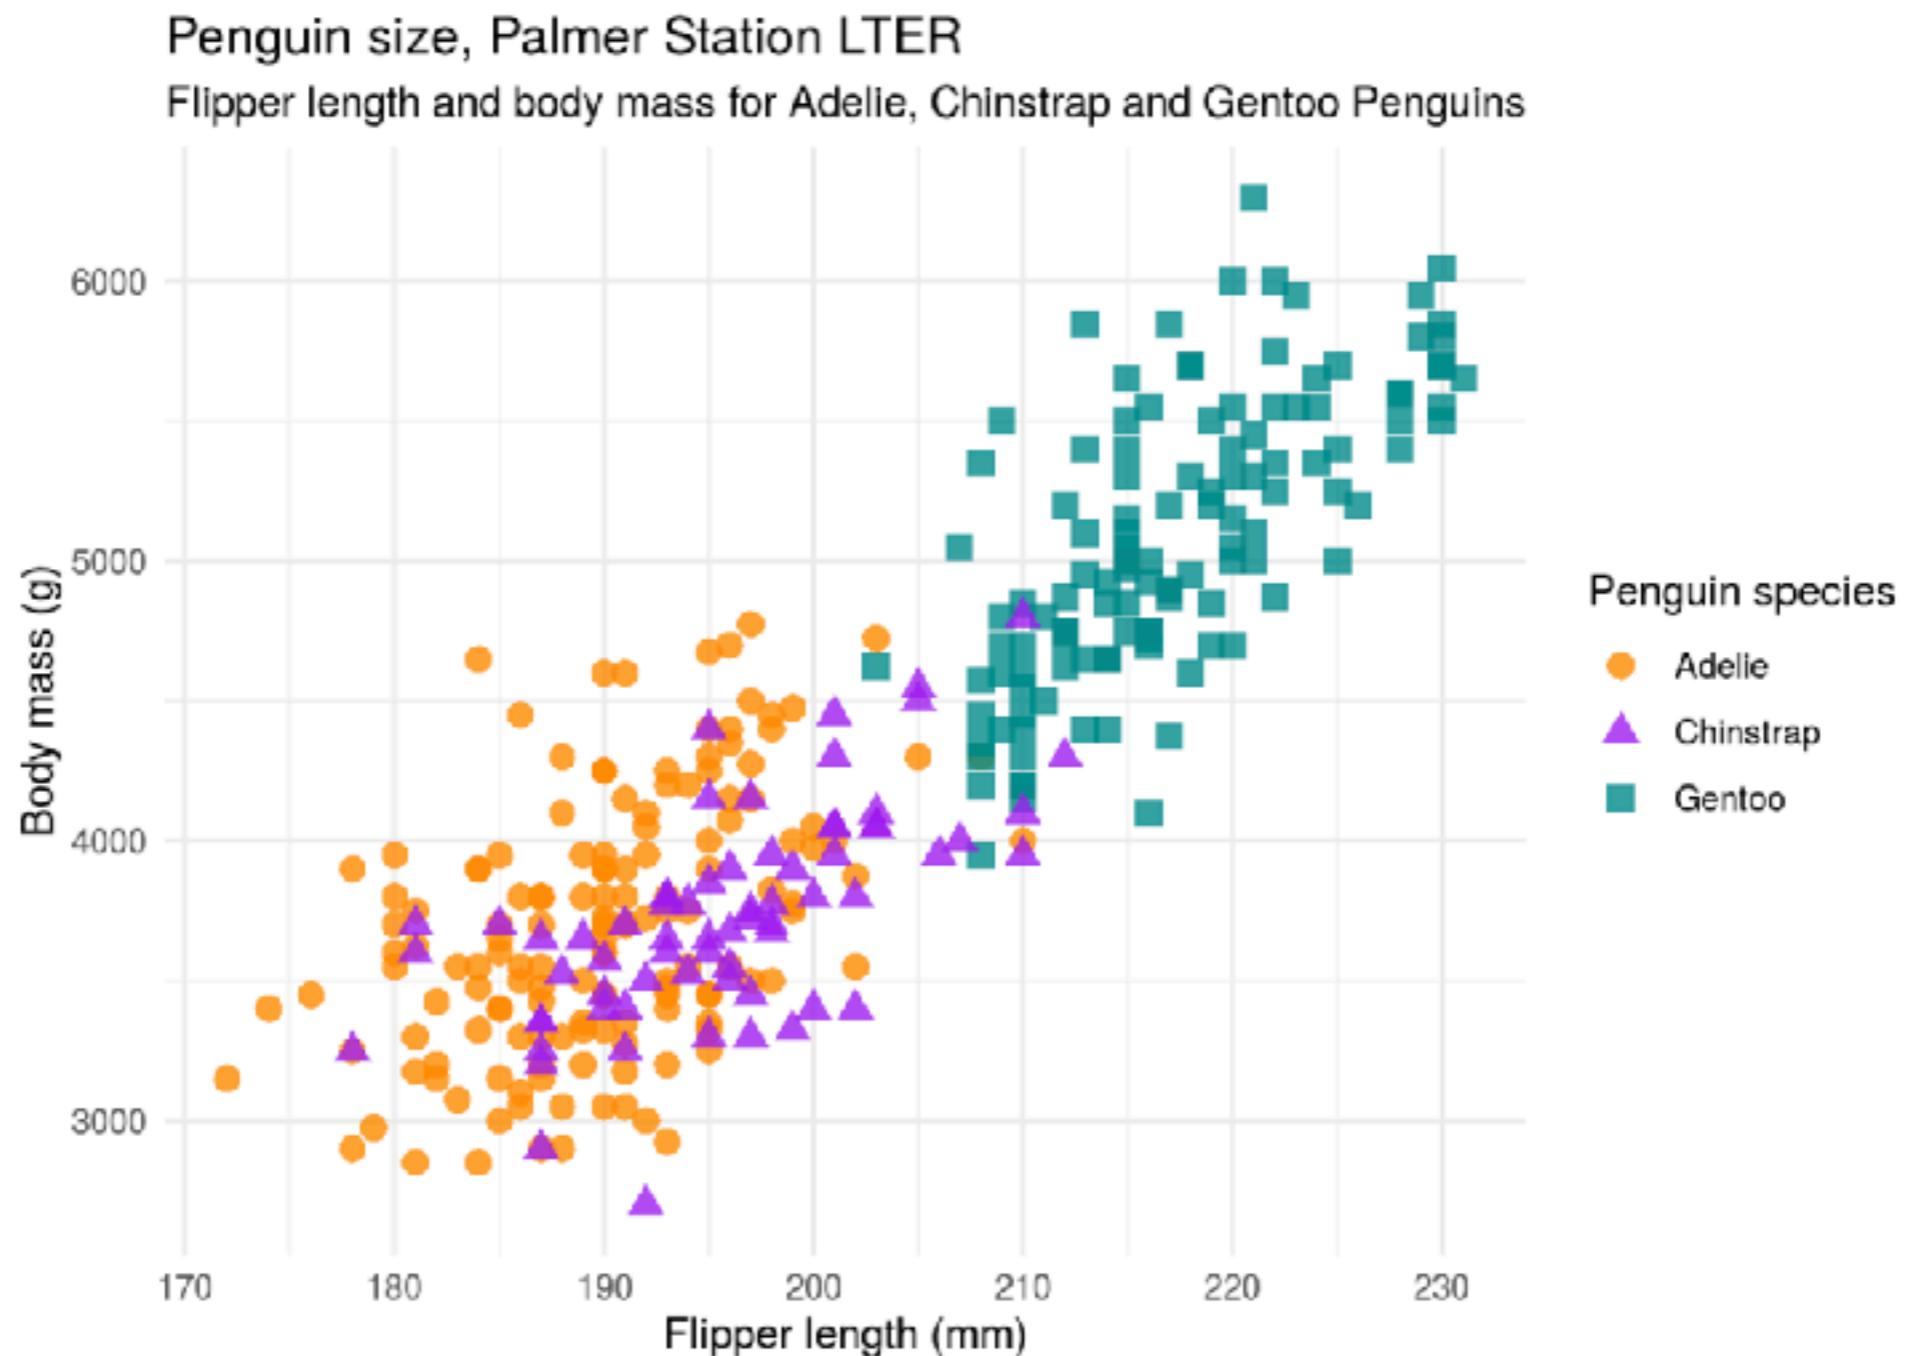
Problem: data not linearly separable

- Typical real data sets aren't linearly separable [demo]



Problem: data not linearly separable

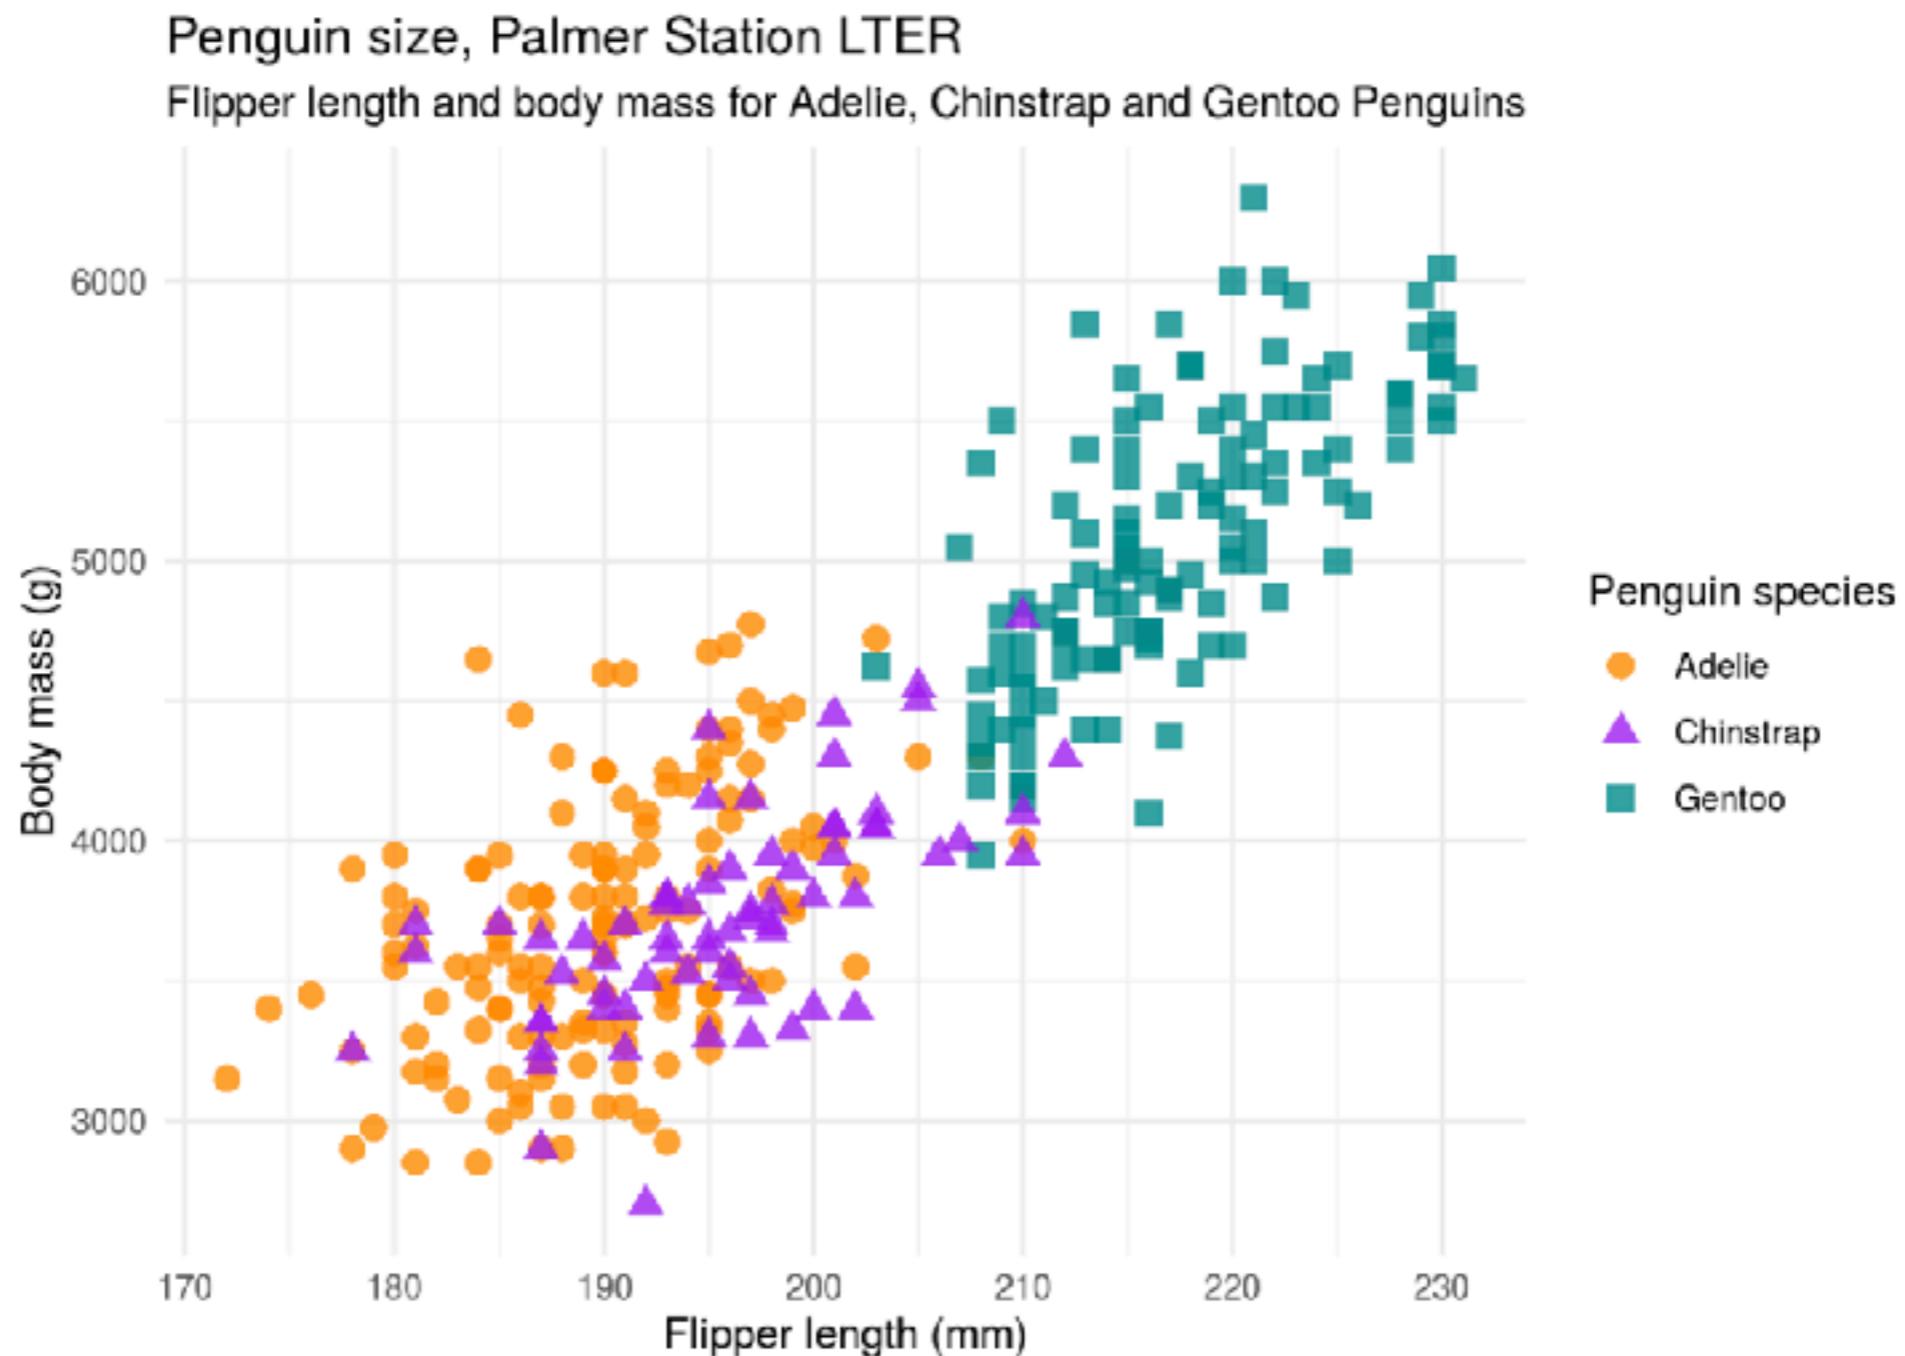
- Typical real data sets aren't linearly separable [demo]



- What can we do?

Problem: data not linearly separable

- Typical real data sets aren't linearly separable [demo]



- What can we do? See upcoming lectures!

Machine Learning Tasks

Machine Learning Tasks

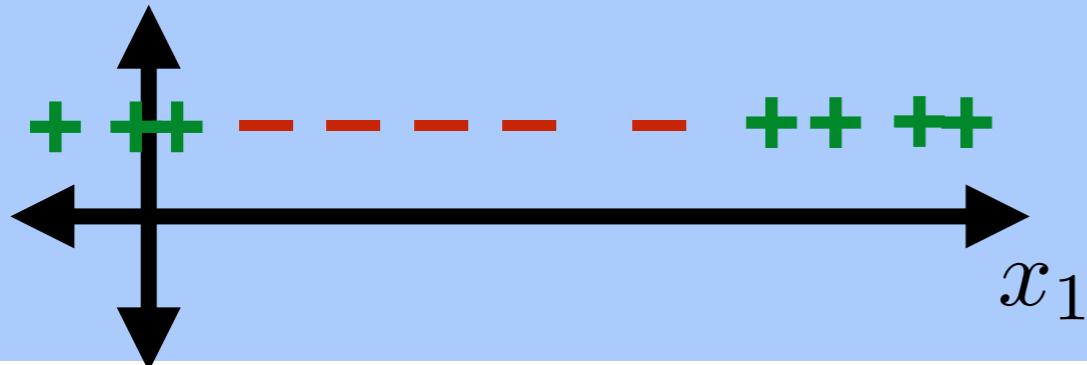
- **Binary/two-class classification**

Machine Learning Tasks

- **Binary/two-class classification:**
Learn a mapping: $\mathbb{R}^d \rightarrow \{-1, +1\}$

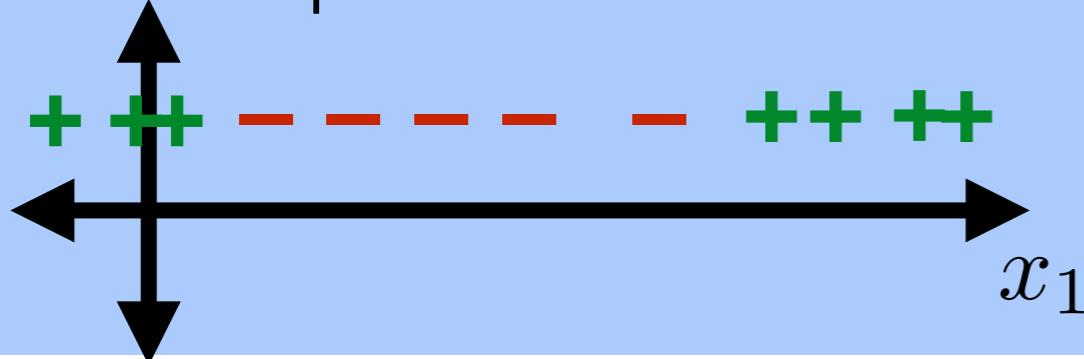
Machine Learning Tasks

- **Binary/two-class classification:**
Learn a mapping: $\mathbb{R}^d \rightarrow \{-1, +1\}$



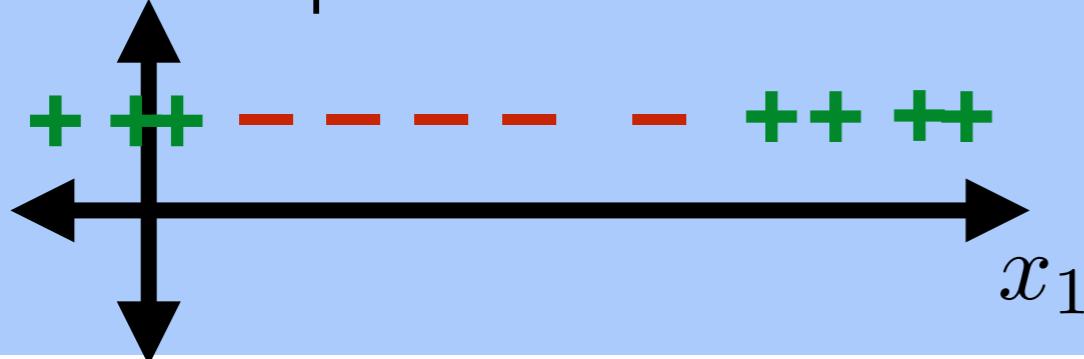
Machine Learning Tasks

- **Binary/two-class classification:**
Learn a mapping: $\mathbb{R}^d \rightarrow \{-1, +1\}$
 - Example: **linear classification**



Machine Learning Tasks

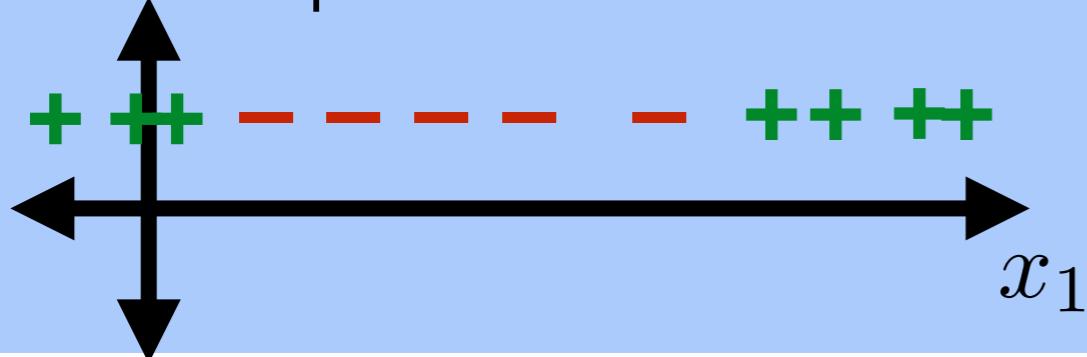
- **Binary/two-class classification:**
Learn a mapping: $\mathbb{R}^d \rightarrow \{-1, +1\}$
 - Example: **linear classification**



- **Multi-class classification:**

Machine Learning Tasks

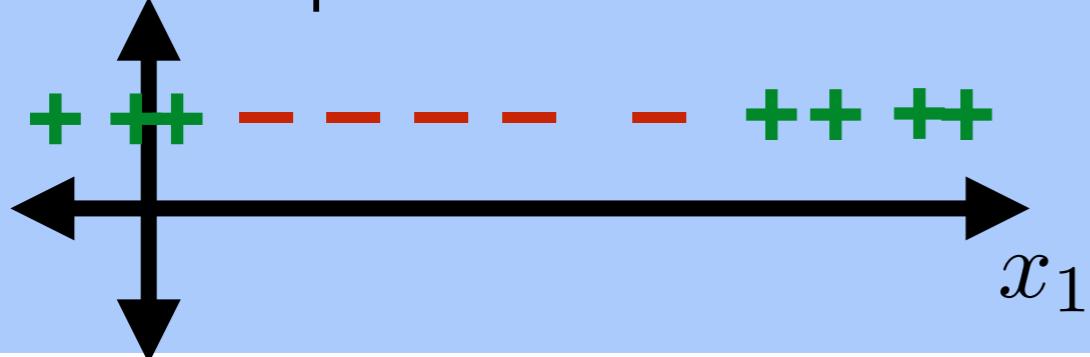
- **Binary/two-class classification:**
Learn a mapping: $\mathbb{R}^d \rightarrow \{-1, +1\}$
 - Example: **linear classification**



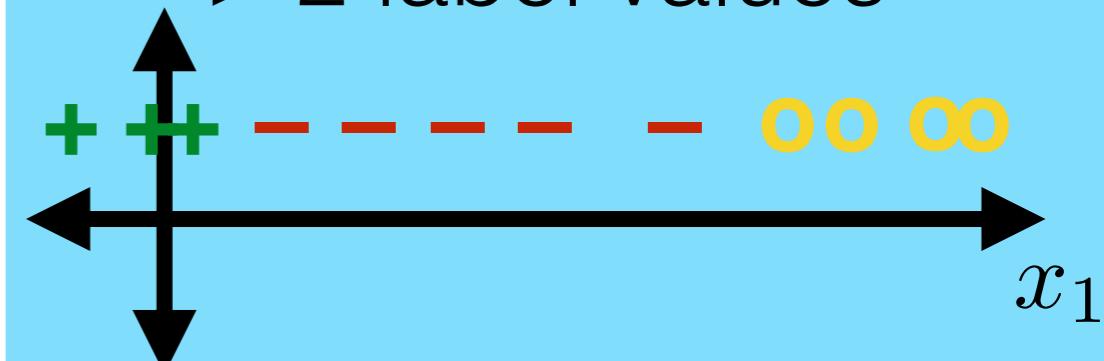
- **Multi-class classification:**
> 2 label values

Machine Learning Tasks

- **Binary/two-class classification:**
Learn a mapping: $\mathbb{R}^d \rightarrow \{-1, +1\}$
 - Example: **linear classification**

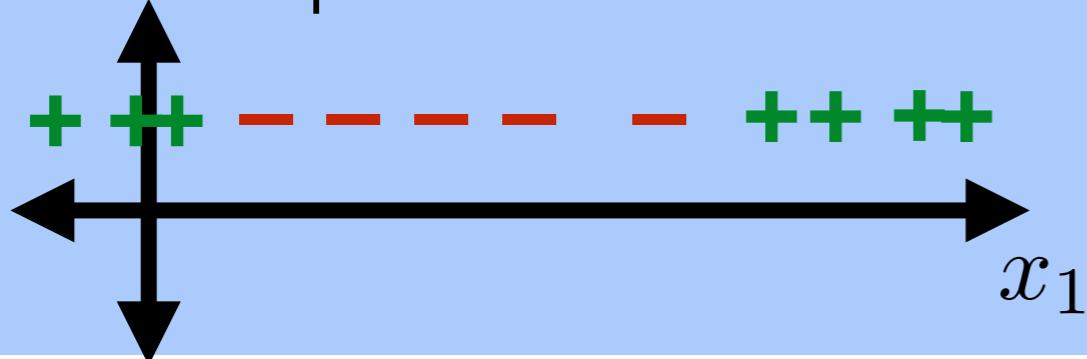


- **Multi-class classification:**
> 2 label values

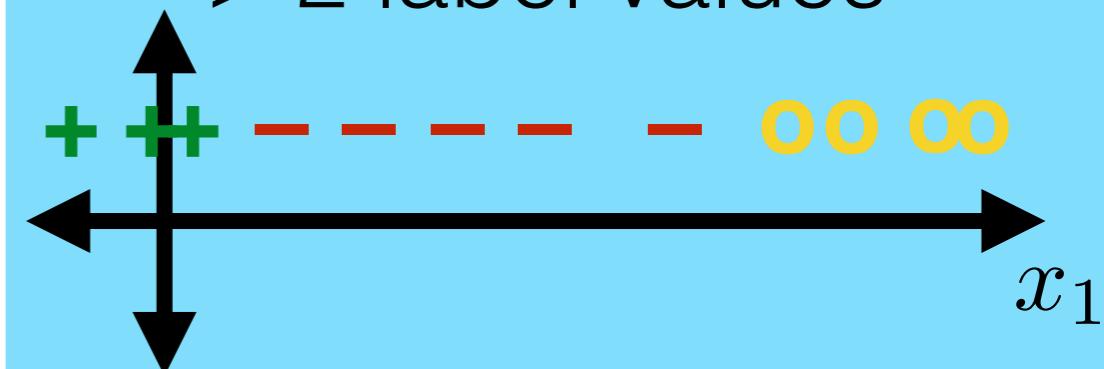


Machine Learning Tasks

- **Binary/two-class classification:**
Learn a mapping: $\mathbb{R}^d \rightarrow \{-1, +1\}$
 - Example: **linear classification**

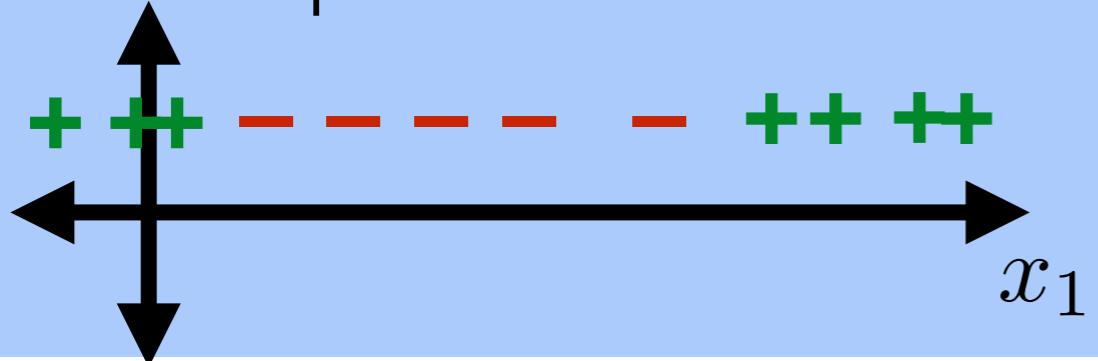


- **Multi-class classification:**
 > 2 label values



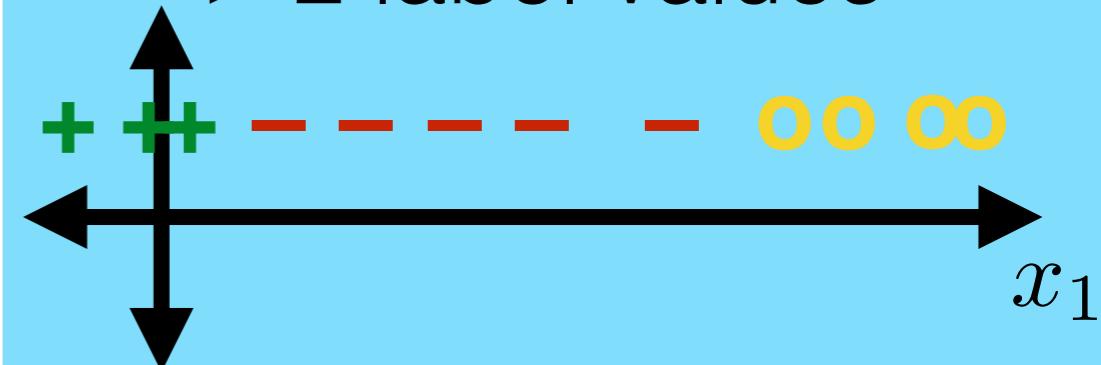
Machine Learning Tasks

- **Binary/two-class classification:**
Learn a mapping: $\mathbb{R}^d \rightarrow \{-1, +1\}$
 - Example: **linear classification**



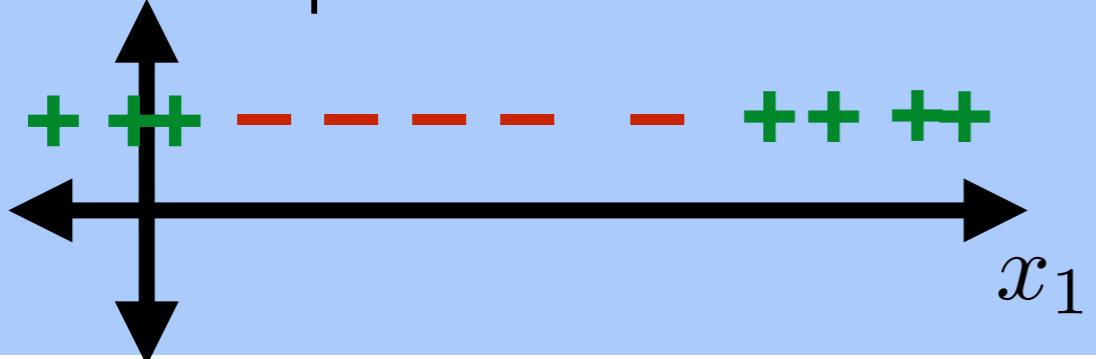
- **Classification**

- **Multi-class classification:**
 > 2 label values



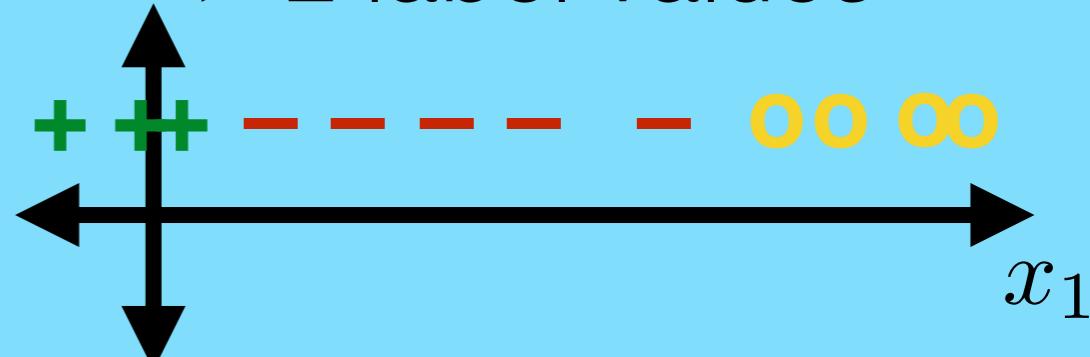
Machine Learning Tasks

- **Binary/two-class classification:**
Learn a mapping: $\mathbb{R}^d \rightarrow \{-1, +1\}$
 - Example: **linear classification**

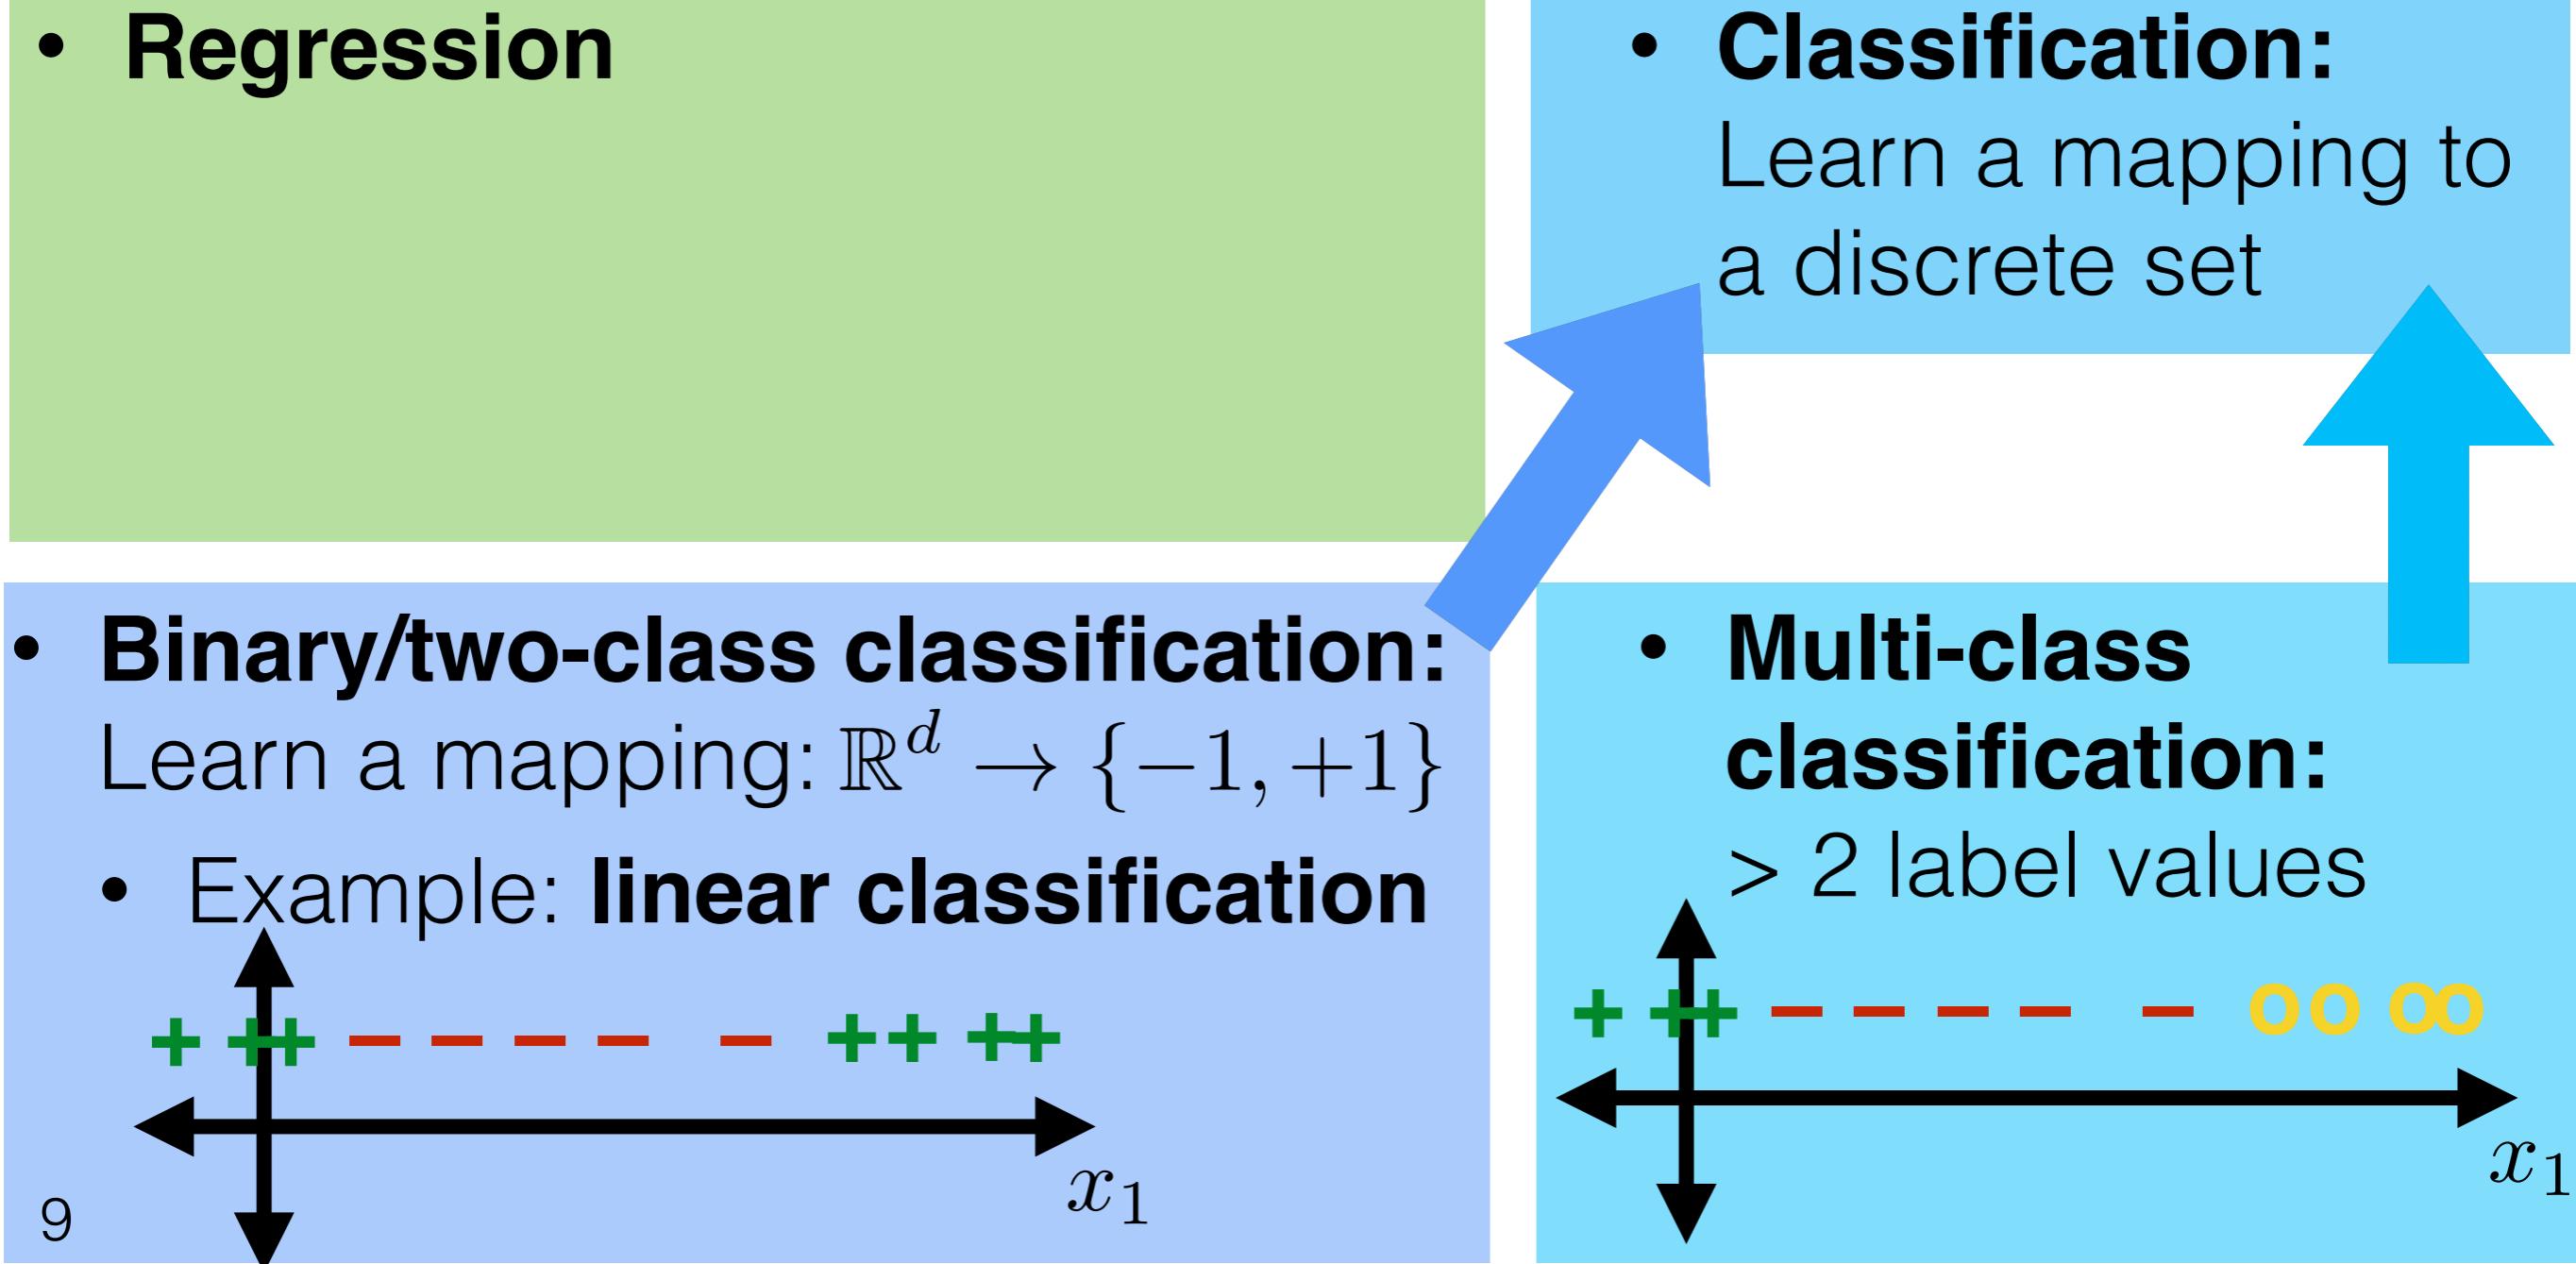


- **Classification:**
Learn a mapping to a discrete set

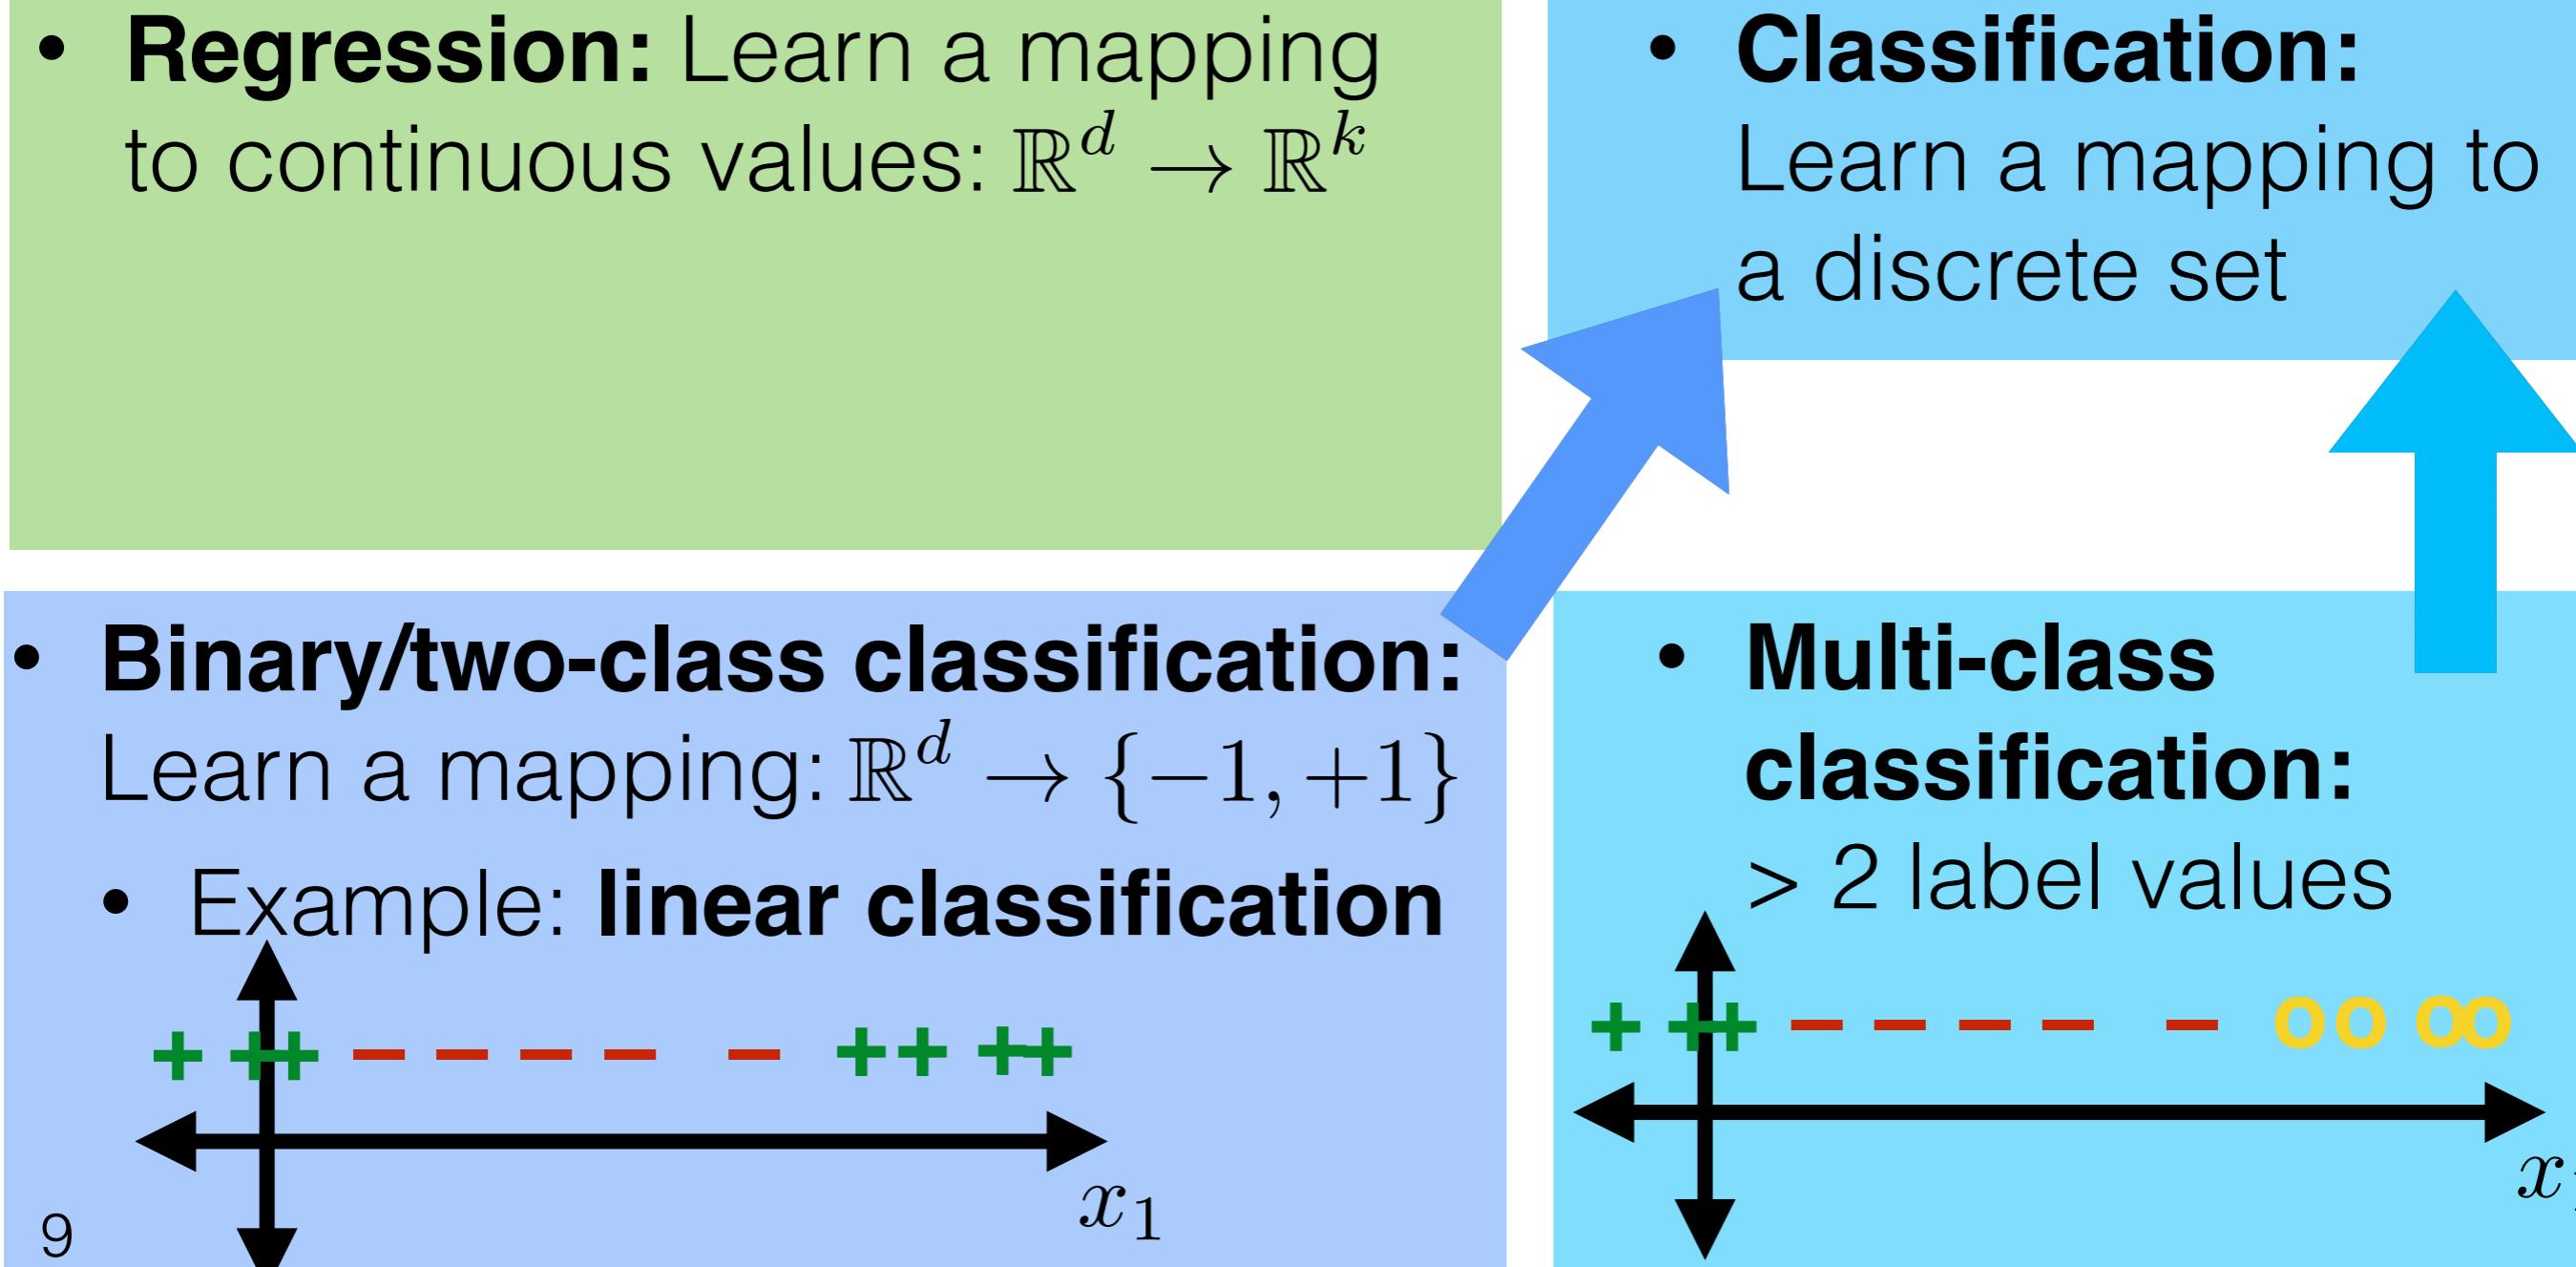
- **Multi-class classification:**
> 2 label values



Machine Learning Tasks

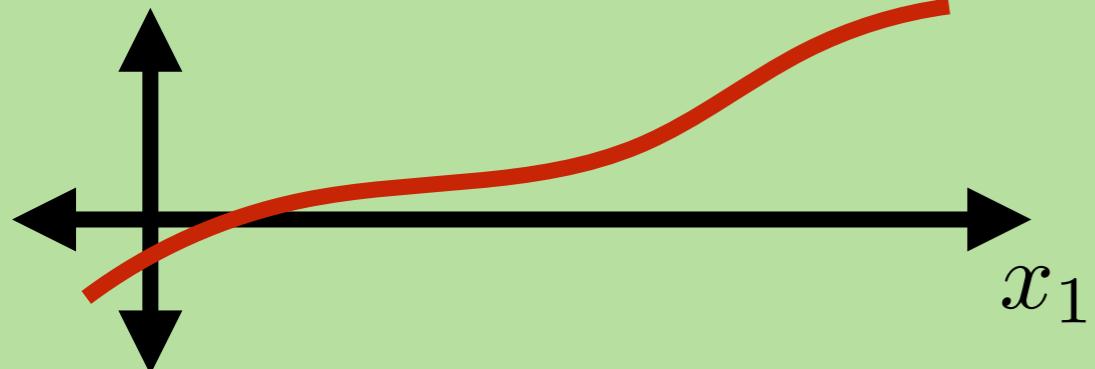


Machine Learning Tasks

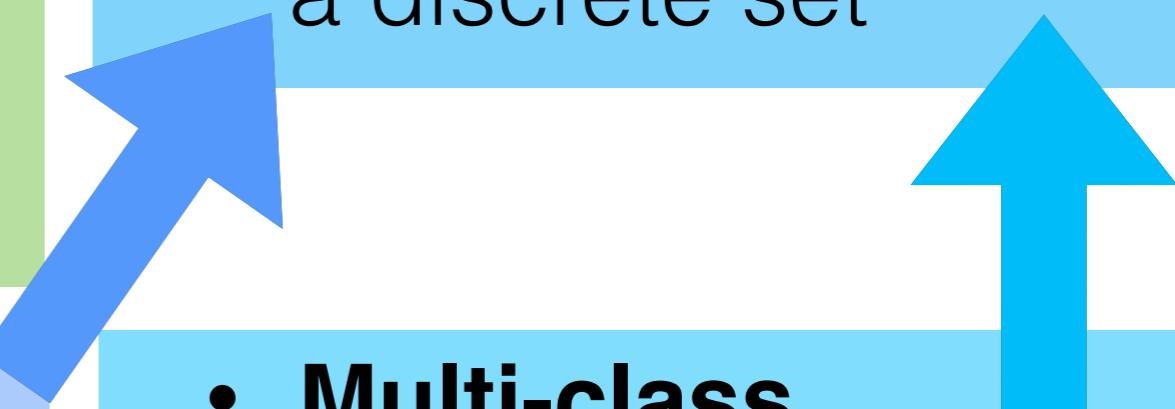


Machine Learning Tasks

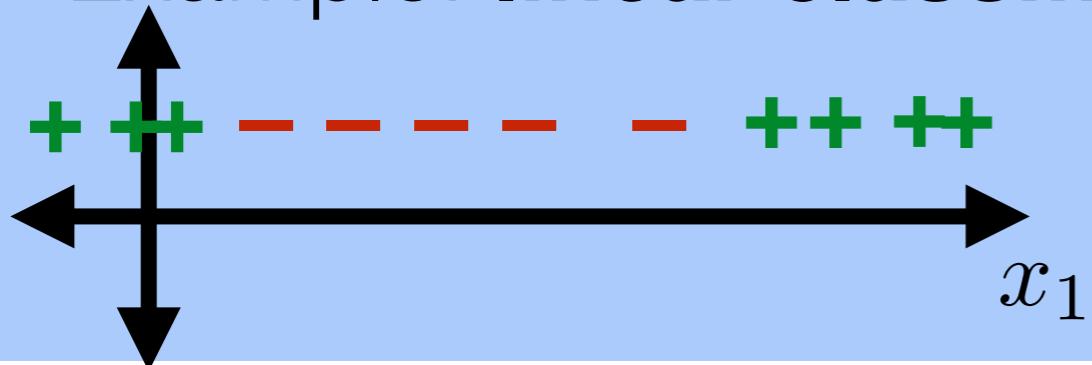
- **Regression:** Learn a mapping to continuous values: $\mathbb{R}^d \rightarrow \mathbb{R}^k$



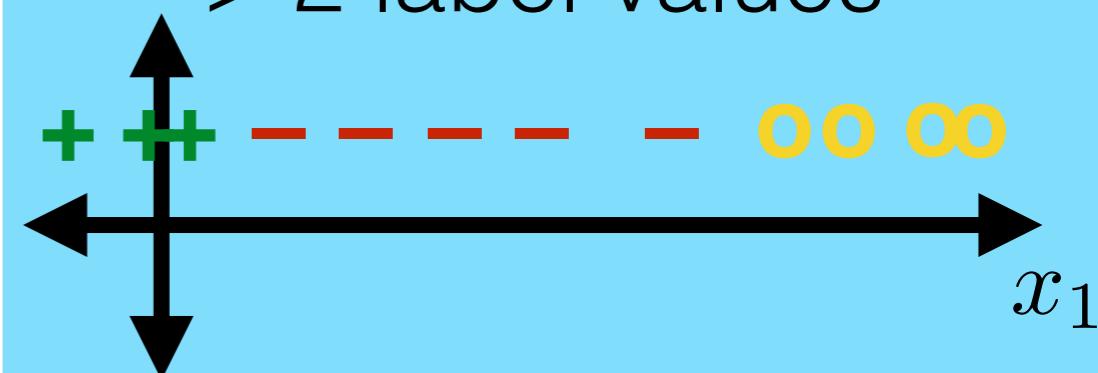
- **Classification:** Learn a mapping to a discrete set



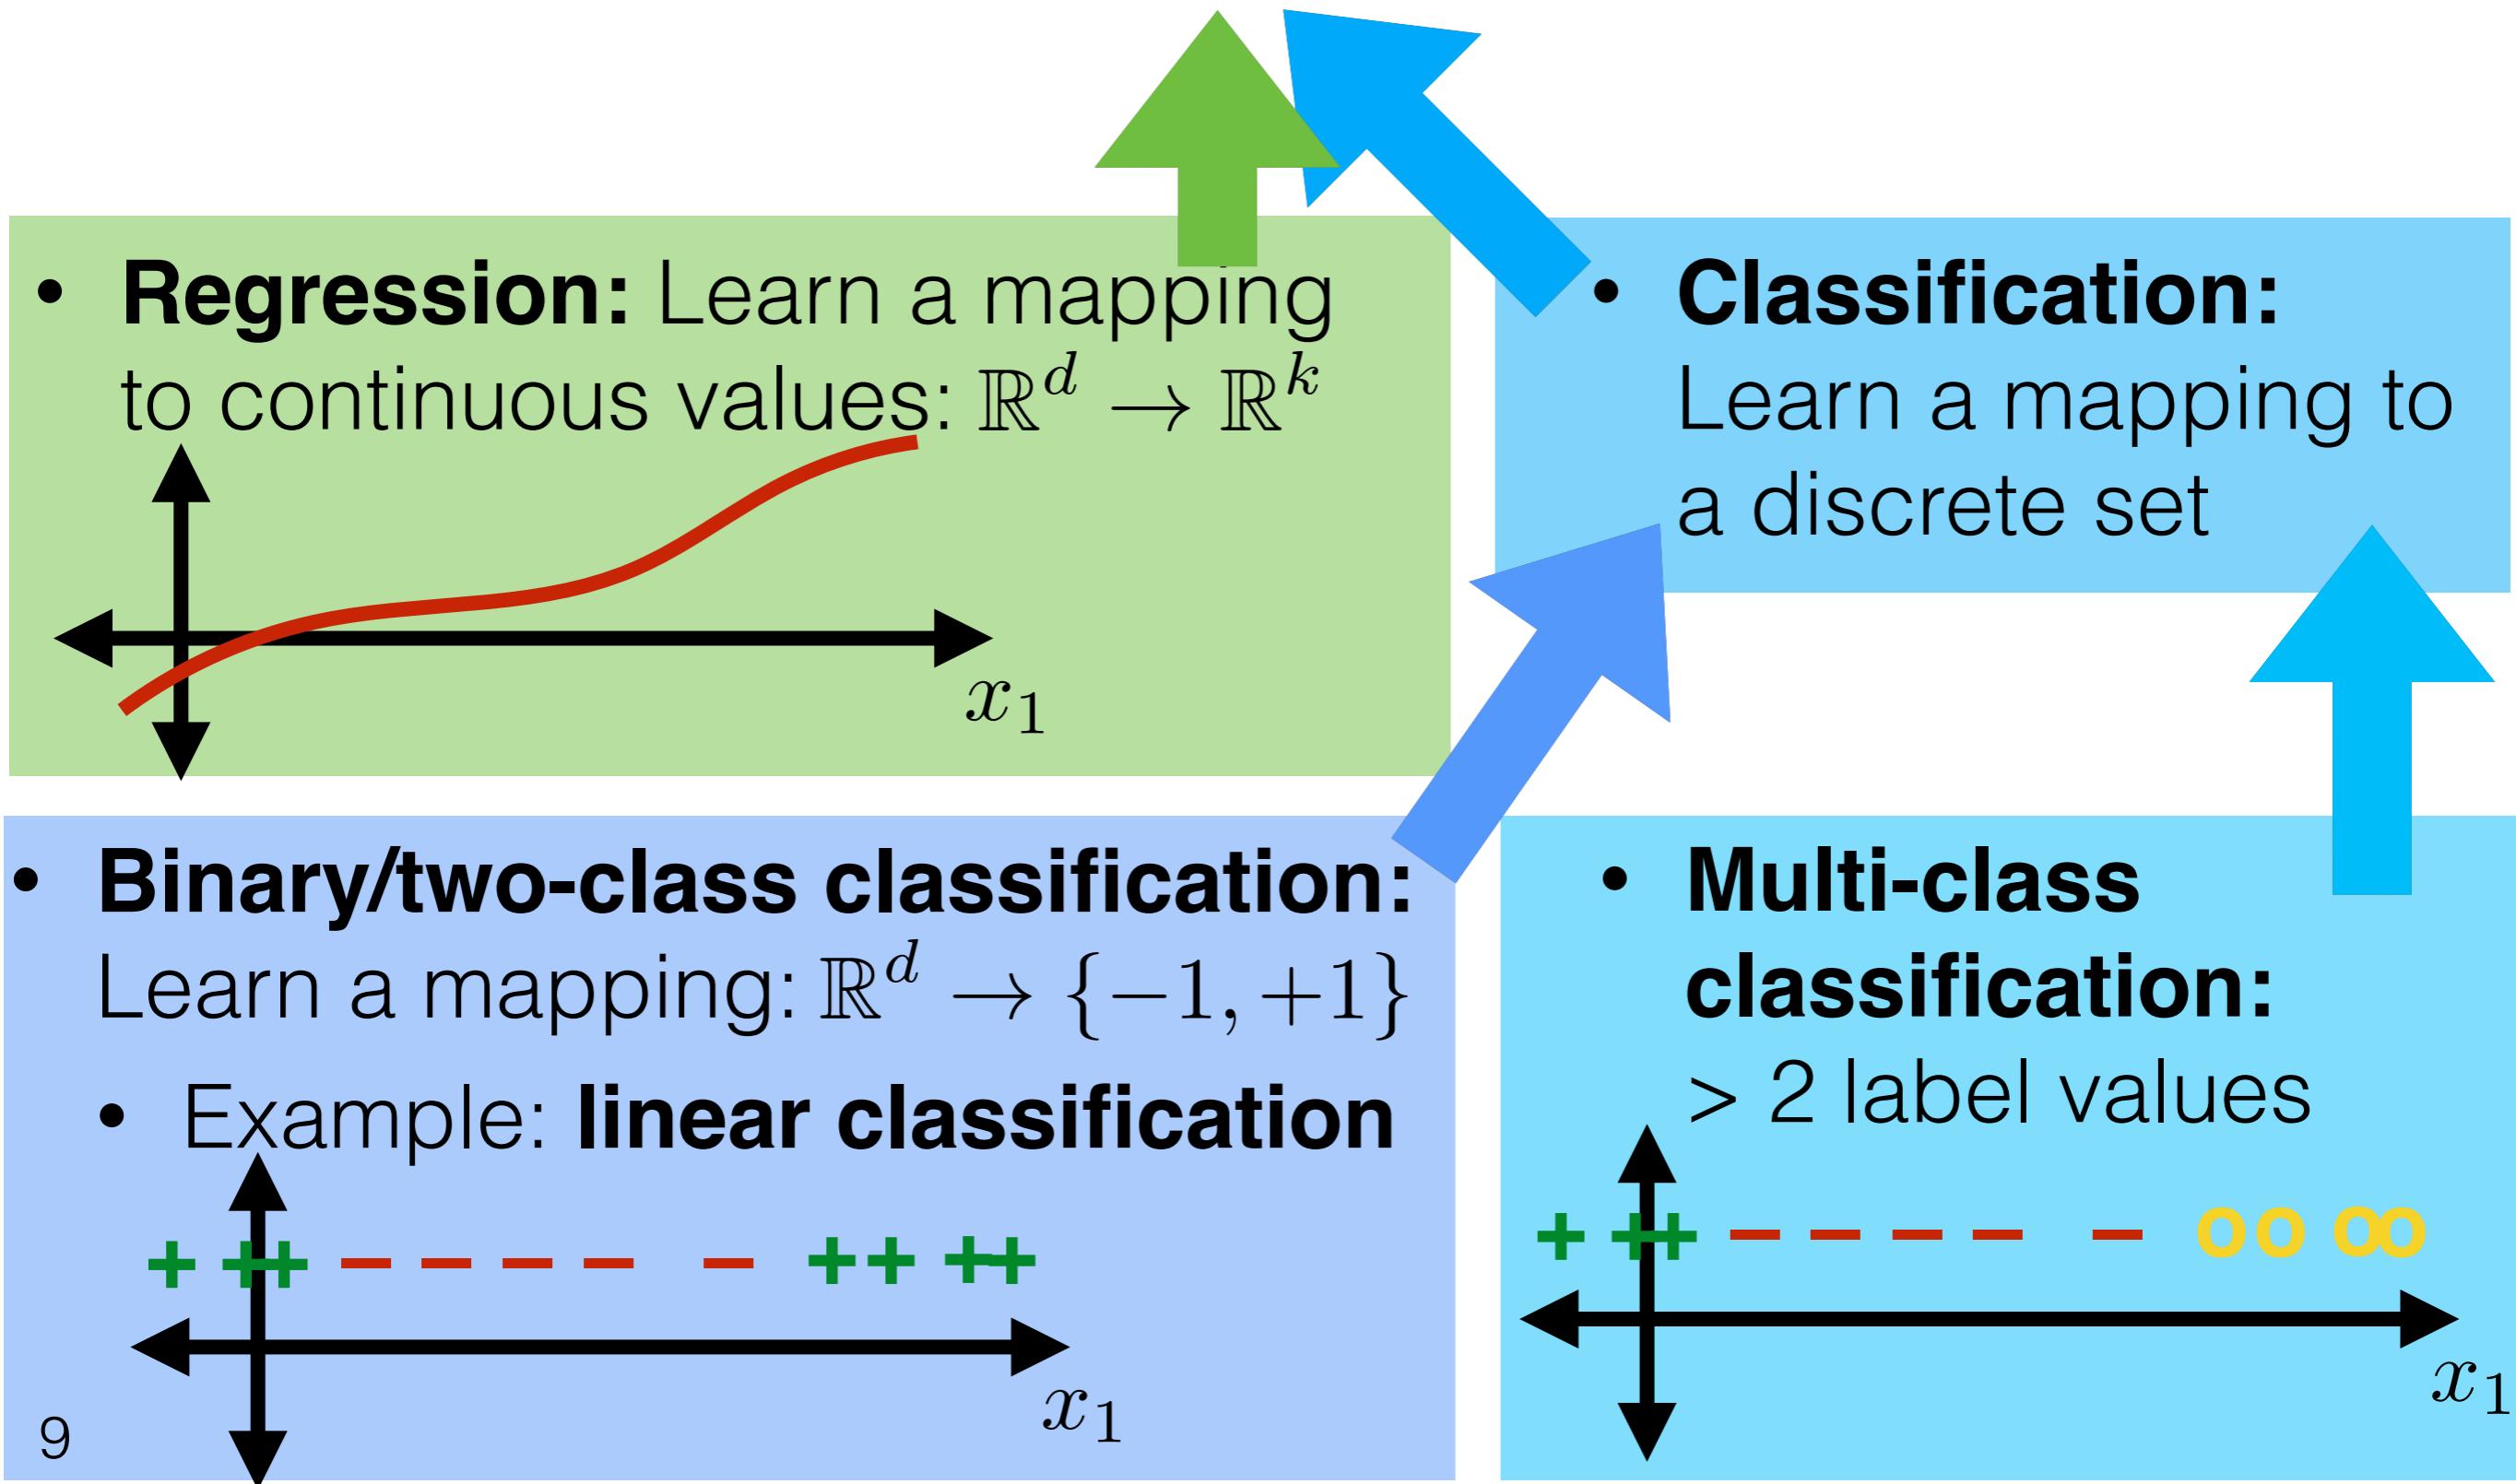
- **Binary/two-class classification:** Learn a mapping: $\mathbb{R}^d \rightarrow \{-1, +1\}$
 - Example: **linear classification**



- **Multi-class classification:** > 2 label values



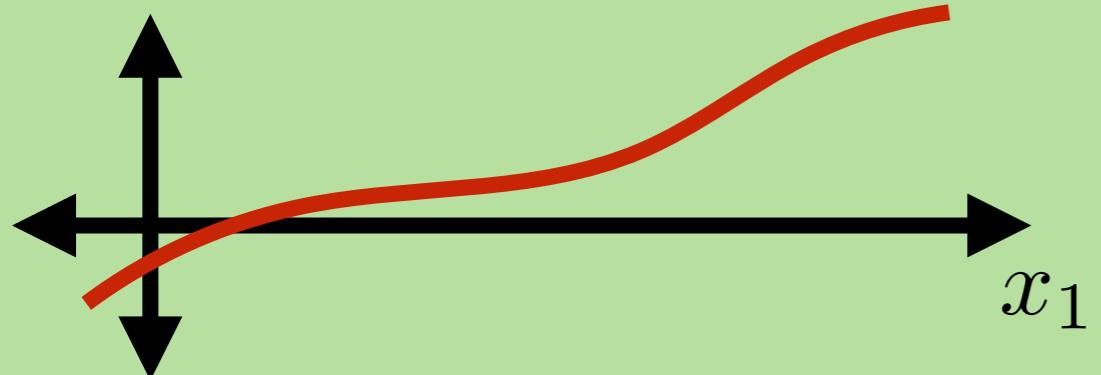
Machine Learning Tasks



Machine Learning Tasks

- **Supervised learning**

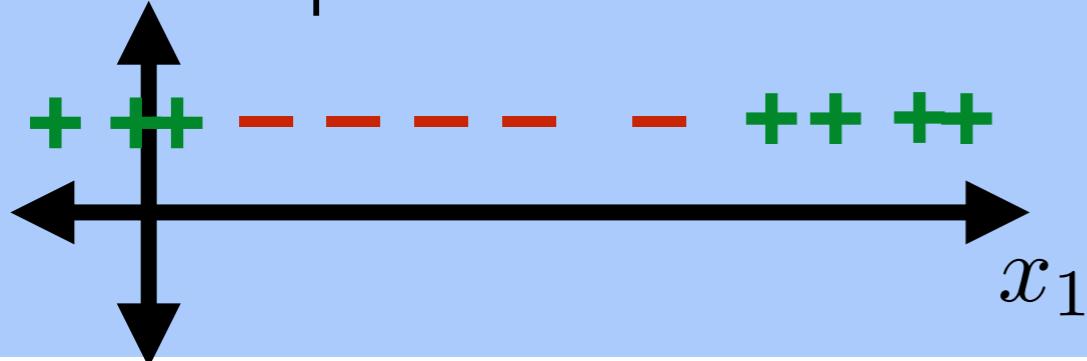
- **Regression:** Learn a mapping to continuous values: $\mathbb{R}^d \rightarrow \mathbb{R}^k$



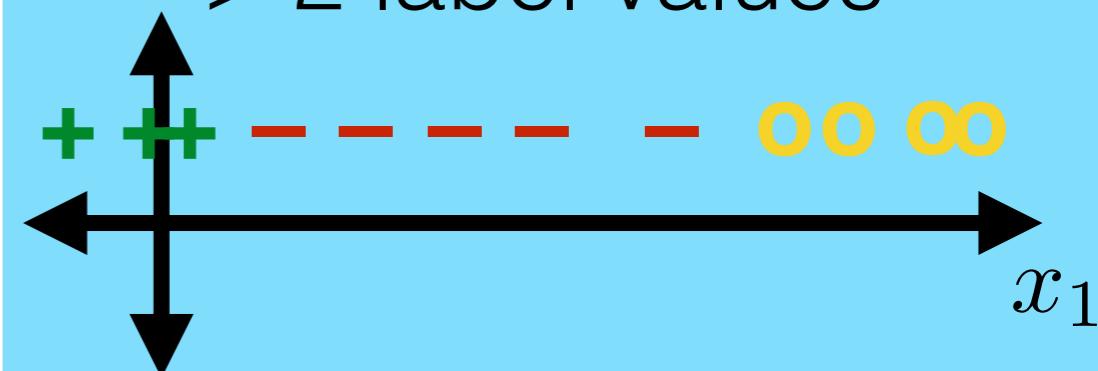
- **Classification:** Learn a mapping to a discrete set

- **Binary/two-class classification:** Learn a mapping: $\mathbb{R}^d \rightarrow \{-1, +1\}$

- Example: **linear classification**

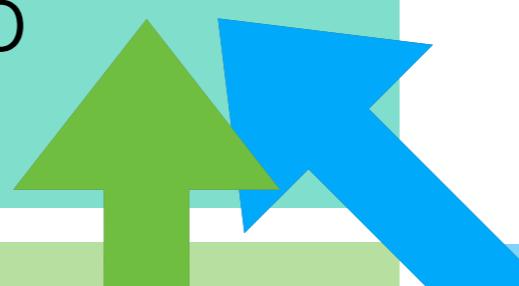


- **Multi-class classification:** > 2 label values



Machine Learning Tasks

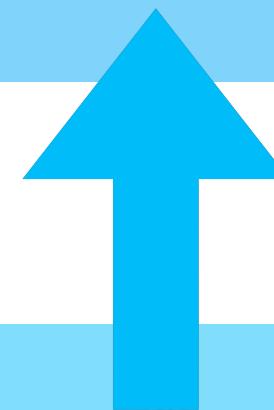
- **Supervised learning:** Learn a mapping from features to labels



- **Regression:** Learn a mapping to continuous values: $\mathbb{R}^d \rightarrow \mathbb{R}^k$

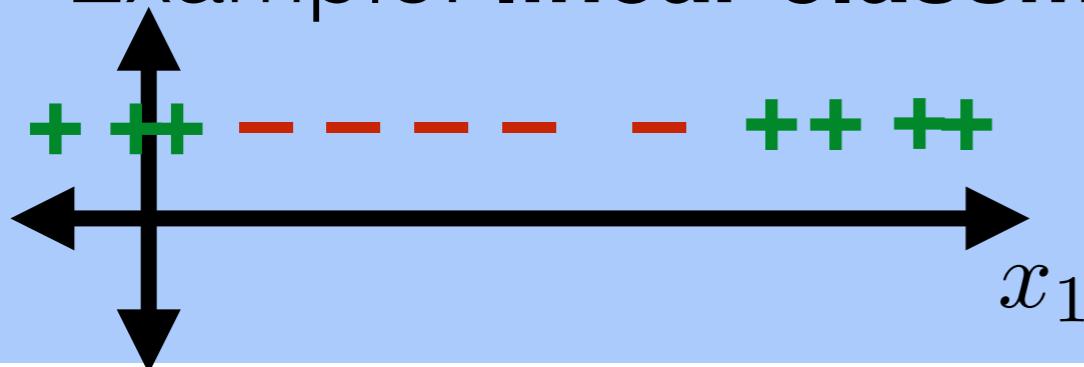


- **Classification:** Learn a mapping to a discrete set

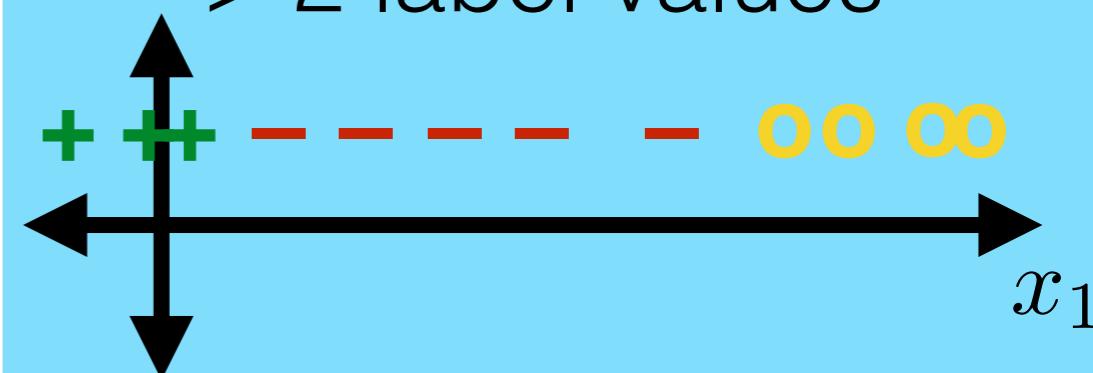


- **Binary/two-class classification:** Learn a mapping: $\mathbb{R}^d \rightarrow \{-1, +1\}$

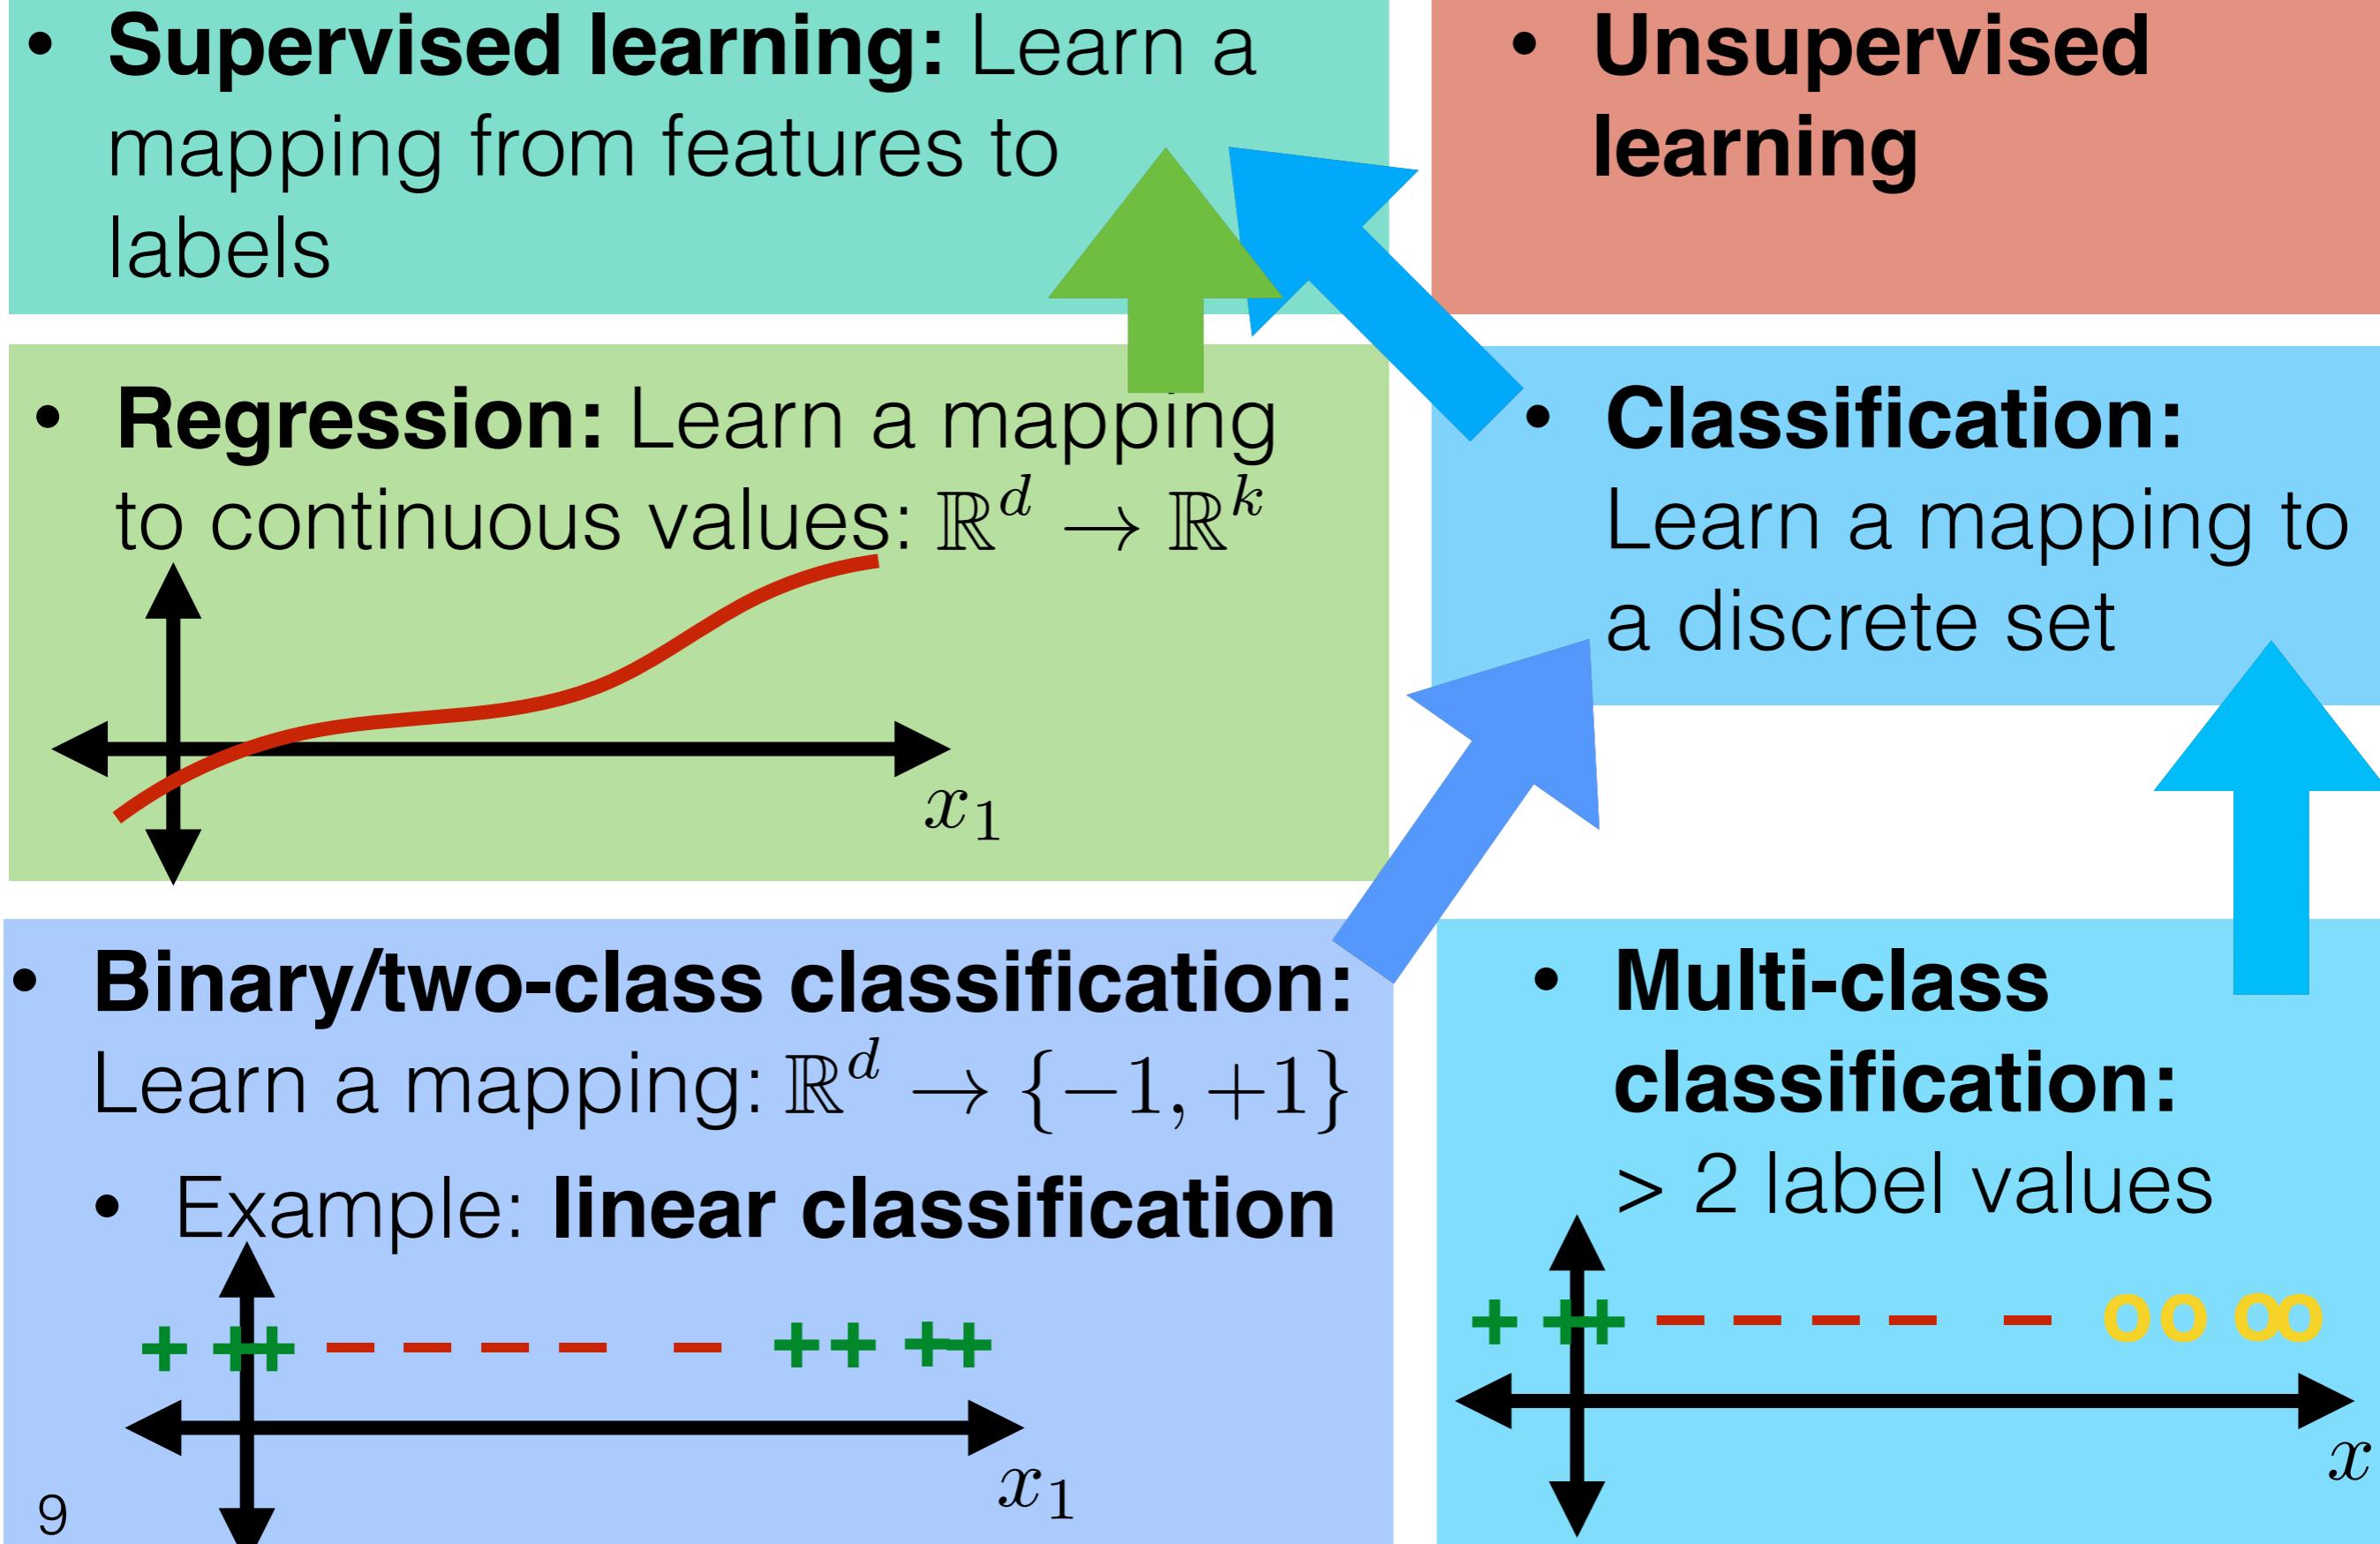
- Example: **linear classification**



- **Multi-class classification:** > 2 label values



Machine Learning Tasks



Machine Learning Tasks

