

MACHINE LEARNING AND BAYESIAN NONPARAMETRICS

Tamara Broderick · Peter Orbanz

MIT and Columbia University

OVERVIEW

OVERVIEW

Morning	Machine Learning
Afternoon	Bayesian Nonparametrics



Machine Learning

- Many latent variable models (finite mixtures, HMMs, ...)
- Linear classifiers can be regarded as boundary cases of models with hidden variables



Bayesian nonparametrics in Machine Learning

- Useful way to impose structure on latent variables
- BNP can be used as building blocks in hierarchies

Neural networks



Graphical models




Grouping dependent variables into
layers is *a good thing*.

NEURAL NETWORKS

“NEURAL NETWORKS”

A neural network represents a function $f : \mathbb{R}^{d_1} \rightarrow \mathbb{R}^{d_2}$.


NEURAL NETWORKS



$$y_i = \phi\left(\sum_j v_{ij}x_j\right)$$

Typically: $\phi = \sigma$ or $\phi(x) = \mathbb{I}\{\pm x > \tau\}$ or $\phi(x) = x$


NEURAL NETWORKS



$$y_i = \phi\left(\sum_j v_{ij}x_j\right)$$

Typically:
$$\phi(x) = \begin{cases} \sigma(x) \text{ (sigmoid)} \\ \mathbb{I}\{\pm x > \tau\} \\ c \text{ (constant)} \\ x \end{cases}$$

NEURAL NETWORKS




$$y_i = \phi\left(\sum_j v_{ij}x_j - c\right)$$

Typically:
$$\phi(x) = \begin{cases} \sigma(x) \text{ (sigmoid)} \\ \mathbb{I}\{\pm x > \tau\} \\ c \text{ (constant)} \\ x \end{cases}$$

LOGISTIC REGRESSION

Sigmoid function

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$



- $\sigma(\theta x)$ is smooth approximation to indicator function $\mathbb{I}\{x \geq 0\}$.
- θ larger \rightarrow closer approximation

Recall: Logistic regression

Covariates $\mathbf{x} \in \mathbb{R}^d$, binary observations y , regressor is logistic function.

$$p(y|\mathbf{x}; \theta) = \text{Bernoulli}(\sigma(\theta^t \mathbf{x}))$$

$\theta \in \mathbb{R}^d$ is parameter to estimate.


NNs AND GRAPHICAL MODELS

Neural networks

- Representation of function using a graph
- Layers:




“ f depends on x only through g ”




Graphical models



- Representation of a distribution using a graph
- Layers:





“ Z is conditionally independent of X given Y ”




Neural networks



Graphical models



LINEAR CLASSIFICATION



HISTORICAL PERSPECTIVE: PERCEPTRON

McCulloch-Pitts neuron model (1943)

$$y = \mathbb{I}\{\mathbf{v}^t \mathbf{x} > c\} \quad \text{for some } c \in \mathbb{R} .$$


Perceptron (Rosenblatt, 1957)

“Train” McCulloch-Pitts model (that is: estimate (c, \mathbf{v})) by applying gradient descent to the function

$$C_p(c, \mathbf{v}) := \sum_{i=1}^n \mathbb{I}\{\text{sgn}(\langle \mathbf{v}, \tilde{\mathbf{x}}_i \rangle - c) \neq \tilde{y}_i\} \left\| \begin{pmatrix} c \\ \mathbf{v} \end{pmatrix}, \begin{pmatrix} 1 \\ \tilde{\mathbf{x}}_i \end{pmatrix} \right\| ,$$

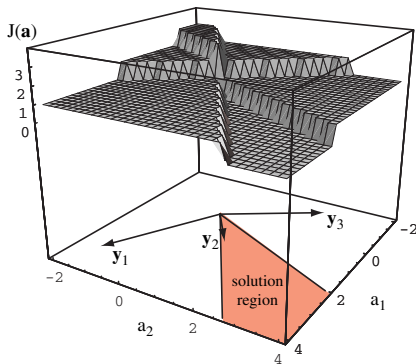
called the **Perceptron cost function**.






PERCEPTRON TRAINING


- The neuron model $\mathbb{I}\{\mathbf{v}'\mathbf{x} > c\}$ is a linear classifier.
- For parameter estimation by optimization, we need an optimization target.
- Choose 0-1 loss as simplest loss for classification.
- Minimize empirical risk (on training data) under this loss.



THE PERCEPTRON CRITERION

- Piece-wise constant function not suitable for numerical optimization.
- Approximate by piece-wise linear function \rightarrow perceptron cost function






NNs AND LINEAR CLASSIFICATION

Linear classifier:

$$f(x) = \text{sgn}(\mathbf{v}^t \mathbf{x} - c)$$

Equivalently:

$$\mathbb{I}\{\mathbf{v}^t \mathbf{x} - c > 0\}$$




NNs AND LINEAR CLASSIFICATION

Linear classifier:


$$f(x) = \text{sgn}(\mathbf{v}^t \mathbf{x} - c)$$

Equivalently:

$$\mathbb{I}\{\mathbf{v}^t \mathbf{x} - c > 0\}$$




REMARKS



- The “neural network” only represents linear classifier, for fixed weights.
- It does not specify the training method.
- “Neural networks” as a field refers to large collection of algorithms and engineering tricks for function estimation.
- These algorithms typically exploit the layered structure.
- Personally, I would recommend to beware of all claims of “neurological plausibility”.

ADDING LAYERS



$$y_i = \psi \left(\sum_j w_{ij} \phi \left(\sum_k v_{jk} x_k \right) \right)$$

$$y_i = \psi \left(\sum_j w_{ij} \phi \left(\sum_k v_{jk} x_k \right) \right)$$

Given: y_i, x_k


Estimate: w_{ij}, v_{jk}

Error Backpropagation

- Use chain rule for differentiation to compute dependence of ψ on w, v
- Differential approximates the correction Δv and Δw we have to apply to produce a correction $\Delta \psi$.

Algorithm

1. For current weights, run network on input x
2. Compare output $\psi(\dots)$ to training values y
3. Make corrections $\Delta v, \Delta w$
4. Repeat



SAMPLE THEORETICAL RESULT

Theorem (Barron, 1993)

If $f : \mathbb{R}^d \rightarrow \mathbb{R}$ satisfies a smoothness conditions, namely

$$C_f = \int \|\omega\| \hat{f}(\omega) d\omega < \infty \quad \text{for Fourier transform } \hat{f},$$

the restriction of f to a disk of radius r can be approximated using N hidden units of simple functions (sigmoids, threshold or ramp functions) to precision $(2rC_f/\sqrt{N})$.


- Note the error rate $2rC_f/\sqrt{N}$ does not seem to depend on dimension d .
- In the wake of this result, a hype declared neural networks can beat the curse of dimensionality.
- Not so: The class of functions satisfying the smoothness condition shrinks with d .
- Much like a filter bank, a NN represents a function basis.

Feature interpretation

Think of applying a 3-layer NN to \mathbf{x} as:

1. Applying a 2-layer NN to \mathbf{x} .
2. Applying another 2-layer NN to the output.

The output of the first NN can be regarded as an intermediate representation.




Convolutional NNs


- Suppose NN is applied to an image.
- Each input vertex is a location in the image (a pixel, say).
- Now constrain the weights applied at different locations to be identical at all locations (“tied weights”).
- Result: Output is translation invariant.
- One can also achieve rotation invariance.
- These invariance properties are very valuable in computer vision.

MAXIMUM MARGIN CLASSIFIERS


GENERALIZATION ERROR



Possible Perceptron solution



Good generalization under a specific distribution (here: Gaussian)



Maximum margin solution


Maximum margin idea

To achieve good generalization (low prediction error), place the hyperplane “in the middle” between the two classes.


Without distributional assumption: Max-margin classifier

- Philosophy: Without distribution assumptions, best guess is symmetric.
- In the Gaussian example, the max-margin solution would *not* be optimal.

LINEAR SEPARATION AND CONVEXITY



Sets can be separated by a hyperplane if and only if their convex hulls are disjoint.



The distance of the plain to the closest set is called the **margin**.

LINEAR CLASSIFIER WITH MARGIN

Recall: Specifying affine plane

Normal vector \mathbf{v}_H .

$$\langle \mathbf{v}_H, \mathbf{x} \rangle - c \begin{cases} > 0 & \mathbf{x} \text{ on positive side} \\ < 0 & \mathbf{x} \text{ on negative side} \end{cases}$$

Scalar $c \in \mathbb{R}$ specifies shift (plane through origin if $c = 0$).

Plane with margin

Demand

$$\langle \mathbf{v}_H, \mathbf{x} \rangle - c > 1 \text{ or } < -1$$

$\{-1, 1\}$ on the right works for any margin: Size of margin determined by $\|\mathbf{v}_H\|$. To increase margin, scale down \mathbf{v}_H .

Classification

Concept of margin applies only to training, not to classification. Classification works as for any linear classifier. For a test point \mathbf{x} :

$$y = \text{sign}(\langle \mathbf{v}_H, \mathbf{x} \rangle - c)$$

SUPPORT VECTOR MACHINE

Finding the hyperplane

For n training points $(\tilde{\mathbf{x}}_i, \tilde{y}_i)$ with labels $\tilde{y}_i \in \{-1, 1\}$, solve optimization problem:

$$\begin{array}{ll} \min_{\mathbf{v}_H, c} & \|\mathbf{v}_H\| \\ \text{s.t.} & \tilde{y}_i(\langle \mathbf{v}_H, \tilde{\mathbf{x}}_i \rangle - c) \geq 1 \quad \text{for } i = 1, \dots, n \end{array}$$

Definition

The classifier obtained by solving this optimization problem is called a **support vector machine**.

DUAL OPTIMIZATION PROBLEM

Solving the SVM optimization problem

$$\begin{aligned} \min_{\mathbf{v}_H, c} \quad & \|\mathbf{v}_H\| \\ \text{s.t.} \quad & \tilde{y}_i (\langle \mathbf{v}_H, \tilde{\mathbf{x}}_i \rangle - c) \geq 1 \quad \text{for } i = 1, \dots, n \end{aligned}$$

is difficult, because the constraint is a function. It is possible to transform this problem into a problem which seems more complicated, but has simpler constraints:

$$\begin{aligned} \max_{\boldsymbol{\alpha} \in \mathbb{R}^n} \quad & W(\boldsymbol{\alpha}) := \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i,j=1}^n \alpha_i \alpha_j \tilde{y}_i \tilde{y}_j \langle \tilde{\mathbf{x}}_i, \tilde{\mathbf{x}}_j \rangle \\ \text{s.t.} \quad & \sum_{i=1}^n \tilde{y}_i \alpha_i = 0 \\ & \alpha_i \geq 0 \quad \text{for } i = 1, \dots, n \end{aligned}$$

This is called the optimization problem **dual** to the minimization problem above. It is usually derived using Lagrange multipliers. We will use a more geometric argument.

DERIVING THE DUAL PROBLEM

Idea

As a consequence of duality on previous slide, we can find the maximum-margin plane as follows:


1. Find shortest line connecting the convex hulls.
2. Place classifier orthogonal to line in the middle.

Convexity of sets ensures that this classifier has correct orientation.

As optimization problem

$$\min_{\substack{\mathbf{u} \in \text{conv}(\mathcal{X}_{\ominus}) \\ \mathbf{v} \in \text{conv}(\mathcal{X}_{\oplus})}}$$

$$\|\mathbf{u} - \mathbf{v}\|^2$$



DUAL OPTIMIZATION PROBLEM

Dual problem

$$\begin{aligned}\left\| \sum_{i \in \mathcal{X}_{\ominus}} \alpha_i \tilde{\mathbf{x}}_i - \sum_{i \in \mathcal{X}_{\oplus}} \alpha_i \tilde{\mathbf{x}}_i \right\|_2^2 &= \left\| \sum_{i \in \mathcal{X}_{\ominus}} \tilde{y}_i \alpha_i \tilde{\mathbf{x}}_i + \sum_{i \in \mathcal{X}_{\oplus}} \tilde{y}_i \alpha_i \tilde{\mathbf{x}}_i \right\|_2^2 \\ &= \left\langle \sum_{i=1}^n \tilde{y}_i \alpha_i \tilde{\mathbf{x}}_i, \sum_{i=1}^n \tilde{y}_i \alpha_i \tilde{\mathbf{x}}_i \right\rangle = \sum_{i,j} \tilde{y}_i \tilde{y}_j \alpha_i \alpha_j \langle \tilde{\mathbf{x}}_i, \tilde{\mathbf{x}}_j \rangle\end{aligned}$$

Note: Minimizing this term under the constraints is equivalent to *maximizing*

$$\sum_i \alpha_i - \frac{1}{2} \sum_{i,j} \tilde{y}_i \tilde{y}_j \alpha_i \alpha_j \langle \tilde{\mathbf{x}}_i, \tilde{\mathbf{x}}_j \rangle$$

under the same constraints, since $\sum_i \alpha_i = 2$ is constant. That is just the dual problem defined four slides back.


CONVEX DUALITY

Sets and Planes

Many dual relations in convex optimization can be traced back to the following fact:

The closest distance between a point \mathbf{x} and a convex set A is the maximum over the distances between \mathbf{x} and all hyperplanes which separate \mathbf{x} and A .


$$d(\mathbf{x}, A) = \sup_{H \text{ separating}} d(\mathbf{x}, H)$$



ENSEMBLE CLASSIFIERS


ENSEMBLES

A *randomly* chosen hyperplane classifier has an *expected* error of 0.5 (i.e. 50%).




ENSEMBLES

A *randomly* chosen hyperplane classifier has an *expected* error of 0.5 (i.e. 50%).



ENSEMBLES

A *randomly* chosen hyperplane classifier has an *expected* error of 0.5 (i.e. 50%).



- Many random hyperplanes combined by majority vote: Still 0.5.
- A single classifier slightly better than random: $0.5 + \epsilon$.
- What if we use m such classifiers and take a majority vote?

Decision by majority vote

- m individuals (or classifiers) take a vote. m is an odd number.
- They decide between two choices; one is correct, one is wrong.
- After everyone has voted, a decision is made by simple majority.

Note: For two-class classifiers f_1, \dots, f_m (with output ± 1):

$$\text{majority vote} = \text{sgn}\left(\sum_{j=1}^m f_j\right)$$

Assumptions

Before we discuss ensembles, we try to convince ourselves that voting can be beneficial. We make some simplifying assumptions:

- Each individual makes the right choice with probability $p \in [0, 1]$.
- The votes are *independent*, i.e. stochastically independent when regarded as random outcomes.

DOES THE MAJORITY MAKE THE RIGHT CHOICE?


Condorcet's rule

If the individual votes are independent, the answer is


$$\Pr\{\text{majority makes correct decision}\} = \sum_{j=\frac{m+1}{2}}^m \frac{m!}{j!(m-j)!} p^j (1-p)^{m-j}$$

This formula is known as **Condorcet's jury theorem**.


Probability as function of the number of votes



$p = 0.55$



$p = 0.45$



$p = 0.85$

Terminology

- An **ensemble method** makes a prediction by combining the predictions of many classifiers into a single vote.
- The individual classifiers are usually required to perform only slightly better than random. For two classes, this means slightly more than 50% of the data are classified correctly. Such a classifier is called a **weak learner**.

Strategy

- We have seen above that if the weak learners are random and independent, the prediction accuracy of the majority vote will increase with the number of weak learners.
- Since the weak learners all have to be trained on the training data, producing random, independent weak learners is difficult.
- Different ensemble methods (e.g. Boosting, Bagging, etc) use different strategies to train and combine weak learners that behave relatively independently.

Input

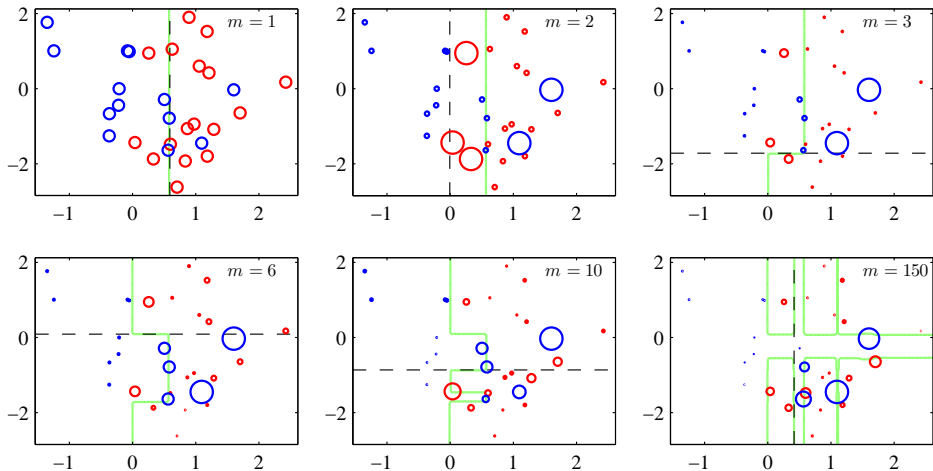
- Training data $(\tilde{\mathbf{x}}_1, \tilde{y}_1), \dots, (\tilde{\mathbf{x}}_n, \tilde{y}_n)$
- Algorithm parameter: Number M of weak learners

Training algorithm

1. Initialize the observation weights $w_i = \frac{1}{n}$ for $i = 1, 2, \dots, n$.
2. For $m = 1$ to M :
 - 2.1 Fit a classifier $g_m(x)$ to the training data using weights w_i .
 - 2.2 Compute
$$\text{err}_m := \frac{\sum_{i=1}^n w_i \mathbb{I}\{y_i \neq g_m(x_i)\}}{\sum_i w_i}$$
 - 2.3 Compute $\alpha_m = \log\left(\frac{1-\text{err}_m}{\text{err}_m}\right)$
 - 2.4 Set $w_i \leftarrow w_i \cdot \exp(\alpha_m \cdot \mathbb{I}(y_i \neq g_m(x_i)))$ for $i = 1, 2, \dots, n$.
3. Output

$$f(x) := \text{sign} \left(\sum_{m=1}^M \alpha_m g_m(x) \right)$$

ILLUSTRATION



Circle = data points, circle size = weight.

Dashed line: Current weak learner. Green line: Aggregate decision boundary.

BOOTSTRAPPING WEAK LEARNERS


Idea

- Recall Boosting: Weak learners are deterministic, but selected to exhibit high variance.
- Strategy now: Randomly distort data set by resampling.
- Train weak learners on resampled training sets.
- Resulting algorithm: **Bagging** (= **B**ootstrap **a**ggregation)

Bagging tree classifiers

- Draw a bootstrap resample of the training data
- Train a tree classifier on the resample
- Repeat

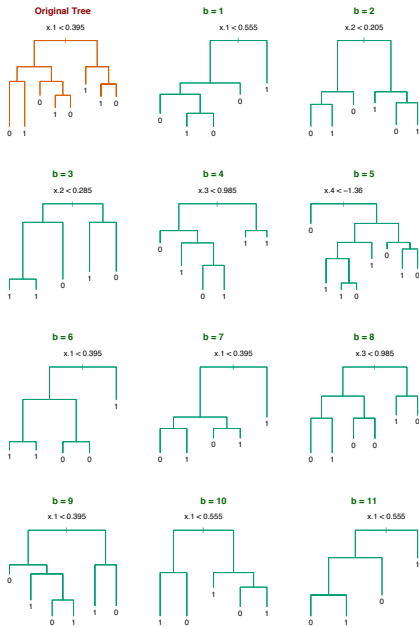
TREES



- Each leaf of the tree corresponds to a region R_m of \mathbb{R}^d .
- Classes $k \in \{1, \dots, K\}$ (not restricted to two classes).
- Training: Greedy splits

BAGGED TREES

- Two classes, each with Gaussian distribution in \mathbb{R}^5 .
- Note the variance between bootstrapped trees.



Bagging vs. Boosting

- Bagging works particularly well for trees, since trees have high variance.
- Boosting typically outperforms bagging with trees.
- The main culprit is usually dependence: Boosting is better at reducing correlation between the trees than bagging is.

Random Forests

Modification of bagging with trees designed to further reduce correlation.

- Tree training optimizes each split over all dimensions.
- Random forests choose a different subset of dimensions *at each split*.
- Optimal split is chosen within the subset.
- The subset is chosen at random out of all dimensions $\{1, \dots, d\}$.

RANDOM FORESTS: ALGORITHM

Training

Input parameter: m (positive integer with $m < d$)

For $b = 1, \dots, B$:

1. Draw a bootstrap sample \mathcal{B}_b of size n from training data.
2. Train a tree classifier f_b on \mathcal{B}_b , where each split is computed as follows:
 - Select m axes in \mathbb{R}_d at random.
 - Find the best split (j^*, t^*) on this subset of dimensions.
 - Split current node along axis j^* at t^* .

Classification

Exactly as for bagging: Classify by majority vote among the B trees. More precisely:

- Compute $f_{\text{avg}}(\mathbf{x}) := (p_1(\mathbf{x}), \dots, p_k(\mathbf{x})) := \frac{1}{B} \sum_{b=1}^B f_b(\mathbf{x})$
- The Random Forest classification rule is

$$f_{\text{Bagging}}(\mathbf{x}) := \arg \max_k \{p_1(\mathbf{x}), \dots, p_k(\mathbf{x})\}$$


RANDOM FORESTS

Remarks


- Recommended value for m is $m = \lfloor \sqrt{d} \rfloor$ or smaller.
- RF typically achieve similar results as boosting. Implemented in most packages, often as standard classifier.

Example: Synthetic Data


- This is the RF classification boundary on the synthetic data we have already seen a few times.
- Note the bias towards axis-parallel alignment.



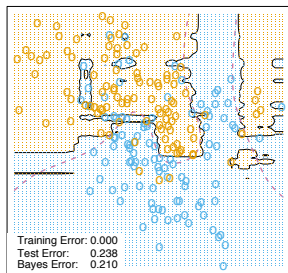
COMPARISON: CLASSIFIERS



Linear SVM



RBF SVM



Random forest

COMPARISON: CLASSIFIERS


Perhaps best off-the-shelf choices

- SVM (with RBF kernel)
- Random Forest

Every method has its idiosyncrasies

- SVM: Treats feature space *geometrically*
- RF:
 - Approximates correlation by piece-wise constant decision surface
 - Applicable to almost any type of data
 - Works naturally with multiple classes
- Boosting (with decision stumps = AdaBoost): Weights provide built-in feature selection

	Decision boundary	Feature space	Tuning parameters
Perceptron	Linear	Euclidean	“Learning rate”
SVM (linear)	Linear	Euclidean	Slack
SVM (RBF)	Nonlinear	Euclidean	Slack Bandwidth
Random forest	Nonlinear	Flexible	Tree depth # resamples # dimensions



MARKOV MODELS


MARKOV MODELS


Markov models

The sequence $(X_n)_n$ is called a (first-order) **Markov chain** if

$$P(X_n | X_{n-1}, \dots, X_1) = P(X_n | X_{n-1}) .$$

Example: Random binary sequence





Parametrization

$$P_{\text{init}} := \begin{pmatrix} \Pr\{X_0 = 1\} \\ \dots \\ \Pr\{X_0 = d\} \end{pmatrix} \quad \mathbf{p} := (p_{i \rightarrow j})_{j, i \leq d} = \begin{pmatrix} p_{1 \rightarrow 1} & \dots & p_{d \rightarrow 1} \\ \vdots & & \vdots \\ p_{1 \rightarrow d} & \dots & p_{d \rightarrow d} \end{pmatrix}$$

The chain $(X_n)_{n \in \mathbb{N}}$ is **stationary** if $p_{s \rightarrow t}$ does not depend on n .

General case

If X_n takes values in infinite set, \mathbf{p} becomes a probability kernel.

STATE PROBABILITIES

Probability after $n = 1$ steps

$$P_1(s_1) = \sum_{s_0 \in \mathbf{X}} p_{s_0 \rightarrow s_1} P_{\text{init}}(s_0) \quad \text{hence} \quad P_1 = \mathbf{p} \cdot P_{\text{init}} .$$

Probability after n steps

$$P_n = \mathbf{p}^n P_{\text{init}}$$

Limiting distribution

$$P_\infty := \lim_{n \rightarrow \infty} P_n = \lim_{n \rightarrow \infty} \mathbf{p}^n P_{\text{init}}$$

(need not exist!)

INVARIANT DISTRIBUTION

Observation

If the limit P_∞ exists, then

$$\mathbf{p} \cdot P_\infty = \mathbf{p} \cdot \lim_{n \rightarrow \infty} \mathbf{p}^n P_{\text{init}} = \lim_{n \rightarrow \infty} \mathbf{p}^n P_{\text{init}} = P_\infty .$$

Therefore, P_∞ is called an **equilibrium** or **invariant distribution** of the chain.

Theorem

Suppose a Markov chain $(\mathbf{p}, P_{\text{init}})$ is stationary, and for each state $s \in \mathbf{X}$:

1. There is a path (with non-zero probability) from s to every other state (i.e. the chain is irreducible).
2. $p_{s \rightarrow s} > 0$ (i.e. the chain is aperiodic).

Then:

- The limit distribution P_∞ exists.
- The limit distribution is also the equilibrium distribution.
- The equilibrium distribution is unique.

WHAT CAN GO WRONG?


Problem 1: The limit may not exist



P_n oscillates between $P_{\text{even}} = \begin{pmatrix} 1 \\ 0 \end{pmatrix}$ and $P_{\text{odd}} = \begin{pmatrix} 0 \\ 1 \end{pmatrix}$.

Calculus analogy: $\lim_n (-1)^n$ does not exist.

Problem 2: The equilibrium distribution may not be unique



Both $P = (0, 1, 0)$ and $P' = (0, 0, 1)$ are valid equilibria.

APPLICATION: INTERNET SEARCH


Random walks as Markov chains

- Simple random walk on a connected graph is a Markov chain.
- Its state space is the vertex set.
- Its equilibrium is the degree-biased distribution.

Internet search

- Idea: Popularity of a web page is a good proxy for whether it is interesting.
- A web crawler can determine which web page links which.
- It cannot determine how often a link is followed.
- “Random surfer model”: Popularity score = probability that a random web surfer would land on the page.
- Random surfer is modeled as simple random walk.
- Google’s PageRank algorithm approximates the equilibrium distribution of simple random walk on the web graph.

APPLICATION: HMMs



- In a HMM, the *latent* variables (Z_n) form a Markov chain.
- Since these are not observed, the sequence (X_n) can exhibit long-range dependencies.
- This is another example of how multiple layers of simple structures express more complicated structures.

Reversible Markov chains

A Markov chain (X_1, \dots, X_n) is **reversible** if

$$(X_1, \dots, X_n) \stackrel{d}{=} (X_n, \dots, X_1)$$


Detailed balance condition

$$P_\infty(dx)P(X_n \in dx' | X_{n-1} = x) = P_\infty(dx')P(X_n \in dx | X_{n-1} = x')$$

If a stationary Markov chain satisfied detailed balance, it is reversible.


MANY HIDDEN LAYERS

DIRECTED GRAPHICAL MODEL



Graphical model distribution on $\{0, 1\}^n$ with joint law

$$P(x_1, \dots, x_n) = \frac{e^{\mathbf{x}'W\mathbf{x} + \mathbf{c}'\mathbf{x}}}{Z(W, \mathbf{c})}$$




- This is a Markov random field.
- The Markov blanket of X_i are those X_j with $W_{ij} \neq 0$.
- For $\mathbf{x} \in \{-1, 1\}^n$: “Potts model with external magnetic field”.
- Statisticians might call this an exponential family with sufficient statistics $\mathbb{E}[X_i X_j]$ and $\mathbb{E}[X_i]$.

BOLTZMANN MACHINE

Graphical model distribution on $\{0, 1\}^n$ with joint law

$$P(x_1, \dots, x_n) = \frac{e^{\mathbf{x}'W\mathbf{x} + \mathbf{c}'\mathbf{x}}}{Z(W, \mathbf{c})}$$



- This is a Markov random field.
- The Markov blanket of X_i are those X_j with $W_{ij} \neq 0$.
- For $\mathbf{x} \in \{-1, 1\}^n$: “Potts model with external magnetic field”.
- Statisticians might call this an exponential family with sufficient statistics $\mathbb{E}[X_i X_j]$ and $\mathbb{E}[X_i]$.


We dispense with tedious accuracy and call P a **Boltzmann machine**.

RESTRICTED BOLTZMANN MACHINE

With observations


If some vertices represent observation variables Y_i :

$$P(x_1, \dots, x_n, y_1, \dots, y_m) = \frac{e^{(\mathbf{x}, \mathbf{y})^t \mathbf{W}(\mathbf{x}, \mathbf{y}) + \mathbf{c}^t \mathbf{x} + \tilde{\mathbf{c}}^t \mathbf{y}}}{Z(\mathbf{W}, \mathbf{c}, \tilde{\mathbf{c}})}$$



Recall our hierarchical design approach

- Only permit layered structure.
- Obvious grouping: One layer for X , one for Y .
- As before: No connections *within* layers.
- Since the graph is undirected, that makes it bipartite.




A **restricted Boltzmann machine (RBM)** is a BM that is bipartite.

GIBBS SAMPLING BOLTZMANN MACHINES

Full conditionals: General case

$$P(\mathbf{X} = \mathbf{x}) = \frac{e^{\mathbf{x}^t \mathbf{W} \mathbf{x} + \mathbf{c}^t \mathbf{x}}}{Z(\mathbf{W}, \mathbf{c})}$$

$$P(X_i = 1 | x_{(i)}) = \sigma(\mathbf{W}_i^t \mathbf{x} + c_i)$$




Full conditionals: RBM


- X's are conditional independent given Y's and vice versa
- Two groups of conditionals: $X|Y$ and $Y|X$
- Blocked Gibbs samplers

$$P(\mathbf{X} = \mathbf{x} | \mathbf{Y} = \mathbf{y}) = \sigma(\mathbf{W}^t \mathbf{y} + \mathbf{c}')$$


$$P(\mathbf{Y} = \mathbf{y} | \mathbf{X} = \mathbf{x}) = \sigma(\mathbf{W}^t \mathbf{x} + \mathbf{c})$$



DIRECTED GRAPHICAL MODEL



RECALL: EXPLAINING AWAY



Conditioning on Z makes X and Y dependent.


COMPLEMENTARY PRIOR

Observation

$X^{(1)}, X^{(2)}, \dots, X^{(N)}$ is a Markov chain.

Complementary prior idea

- Suppose Markov chain is reversible.
- Then all arrows can be reversed.
- Now: Inference easy.




COMPLEMENTARY PRIOR

Observation

$X^{(1)}, X^{(2)}, \dots, X^{(N)}$ is a Markov chain.

Complementary prior idea

- Suppose Markov chain is reversible.
- Then all arrows can be reversed.
- Now: Inference easy.



BUILDING A COMPLEMENTARY PRIOR

- Find reversible Markov chain with

$$P^{(1)} = P_{\infty}$$


- Let \mathbf{p}_T be its transition kernel
- Choose

$$P^{(n+1)}(\bullet | X^{(n)} = \mathbf{x}) = \mathbf{p}_T(\bullet | \mathbf{x})$$

- Then $P^{(2)} = \dots = P^{(n)} = P_{\infty}$
- Since chain is reversible,

$$P^{(n)}(\bullet | X^{(n+1)} = \mathbf{x}) = \mathbf{p}_T(\bullet | \mathbf{x})$$

and edges flip.



BUILDING A COMPLEMENTARY PRIOR

- Find reversible Markov chain with

$$P^{(1)} = P_{\infty}$$


- Let \mathbf{p}_T be its transition kernel
- Choose

$$P^{(n+1)}(\bullet | X^{(n)} = \mathbf{x}) = \mathbf{p}_T(\bullet | \mathbf{x})$$

- Then $P^{(2)} = \dots = P^{(n)} = P_{\infty}$
- Since chain is reversible,


$$P^{(n)}(\bullet | X^{(n+1)} = \mathbf{x}) = \mathbf{p}_T(\bullet | \mathbf{x})$$

and edges flip.



WHERE DO WE GET THE MARKOV CHAIN?


Start with an RBM



Blocked Gibbs sampling alternates between X and Y

$$X^{(1)} \rightarrow Y^{(1)} \rightarrow X^{(2)} \rightarrow Y^{(2)} \rightarrow \dots \rightarrow X^{(N)} \rightarrow Y^{(N)}$$

“Roll off” this chain into a graphical model




The *Gibbs sampler* for the RBM becomes the *model* for the directed network.

- “Deep learning” refers to hierarchies with multiple hidden layers.
- There are other methods to build/train “deep” hierarchies.
- Many training methods are of the form: Initialize with method (a), fine-tune with method (b), etc
- “Deep” now refers to everything that is hierarchical and has more than two layers.
- Some deep networks are indeed deep: ImageNet now has >100 layers

How deep is deep?


- Hinton/Osindero/Teh, 2006: Input + output + 3 hidden layers.
- Models currently used for computer vision problems: >100 layers.

AUTOENCODERS



$$\text{Prediction} = f(\text{Input})$$

AUTOENCODERS



$$\text{Prediction} = f(\text{Input})$$

AUTOENCODERS

